



# SD for SPD Systems – Behaviour

Steepest descent can suffer from slow convergence for poorly conditioned  $\mathbf{A}$


- as  $\text{cond}(\mathbf{A}) = \sigma_{\max}/\sigma_{\min}$  becomes large,  $\mathbf{d}_k = -\nabla f(\mathbf{x}_k)$  may not point in the direction of a (global) minimum of  $f(\mathbf{x})$



# SD for SPD Systems – Behaviour




Well conditioned  $A$



Poorly conditioned  $A$

# Motivating Conjugate Gradients

Instead of descending down the gradient direction, the **conjugate gradient** scheme takes directions designed to avoid zig-zagging





# Conjugate Gradient – Better Directions

We arrived at SD from GD by choosing the optimal step size and *leaving the direction as the gradient*

- we observed that each pair of iterative descent directions are *perpendicular to each other*
  - zig zag
- CG seeks to find **descent directions** that avoid zig-zagging

Starting from the general iterative update rule, we now leave **both** the descent direction **and** step size variable:

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$$

# CG – General Step Size Optimum

$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k \longrightarrow f(\mathbf{x}_{k+1}) \longrightarrow$  and solve for  $\alpha_k$  in  $df(\mathbf{x}_{k+1}) / d\alpha_k = 0$

$$\begin{aligned} f(\mathbf{x}_{k+1}) &= \frac{1}{2} (\mathbf{x}_k + \alpha_k \mathbf{d}_k)^T \mathbf{A} (\mathbf{x}_k + \alpha_k \mathbf{d}_k) - \mathbf{b}^T (\mathbf{x}_k + \alpha_k \mathbf{d}_k) + c \\ &= \frac{1}{2} \mathbf{x}_k^T \mathbf{A} \mathbf{x}_k + \alpha_k \mathbf{d}_k^T \mathbf{A} \mathbf{x}_k + \frac{1}{2} \alpha_k^2 \mathbf{d}_k^T \mathbf{A} \mathbf{d}_k - \mathbf{b}^T \mathbf{x}_k - \alpha_k \mathbf{b}^T \mathbf{d}_k + c \\ &= \left[ \frac{1}{2} \mathbf{x}_k^T \mathbf{A} \mathbf{x}_k - \mathbf{b}^T \mathbf{x}_k + c \right] + [\alpha_k \mathbf{d}_k^T (\mathbf{A} \mathbf{x}_k - \mathbf{b})] + \frac{1}{2} \alpha_k^2 \mathbf{d}_k^T \mathbf{A} \mathbf{d}_k \\ &= f(\mathbf{x}_k) + \alpha_k \mathbf{d}_k^T \nabla f(\mathbf{x}_k) + \frac{1}{2} \alpha_k^2 \mathbf{d}_k^T \mathbf{A} \mathbf{d}_k \end{aligned}$$

$$\frac{df(\mathbf{x}_{k+1})}{d\alpha_k} = 0 = \mathbf{d}_k^T \nabla f(\mathbf{x}_k) + \alpha_k \mathbf{d}_k^T \mathbf{A} \mathbf{d}_k \longrightarrow \alpha_k = - \frac{\mathbf{d}_k^T \nabla f(\mathbf{x}_k)}{\mathbf{d}_k^T \mathbf{A} \mathbf{d}_k}$$

# CG – Building Conjugate Directions

To find better search directions  $\mathbf{d}_k$  we need a few definitions:

Two vectors  $\mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^n$  are  $\mathbf{A}$ -conjugate (or  $\mathbf{A}$ -orthogonal) if

$$\mathbf{x}_1^T \mathbf{A} \mathbf{x}_2 = 0$$

- geometric interpretation?

A set of vectors  $\mathcal{S} = \{\mathbf{x}_1, \dots, \mathbf{x}_k\} \in \mathbb{R}^n$  are an  $\mathbf{A}$ -conjugate set if

$$\mathbf{x}_i^T \mathbf{A} \mathbf{x}_j = 0, \forall i \neq j$$

# CG – Building Conjugate Directions

Note: if  $\mathbf{A} > 0$  and  $\mathbf{A} = \mathbf{A}^T$  then  $\mathcal{S}$  are also linearly independent

- recall that linear independence means that

$$\alpha_1 \mathbf{x}_1 + \alpha_2 \mathbf{x}_2 + \cdots + \alpha_{k-1} \mathbf{x}_{k-1} + \alpha_k \mathbf{x}_k = \mathbf{0}$$

can only hold if  $\alpha_i = 0, i = 1, \dots, k$ .

- pre-multiplying by  $\mathbf{A}$  and then by  $\mathbf{x}_i^T$  we arrive at

$$\alpha_1 \mathbf{A} \mathbf{x}_1 + \alpha_2 \mathbf{A} \mathbf{x}_2 + \cdots + \alpha_{k-1} \mathbf{A} \mathbf{x}_{k-1} + \alpha_k \mathbf{A} \mathbf{x}_k = \mathbf{0},$$

$$\alpha_1 \mathbf{x}_i^T \mathbf{A} \mathbf{x}_1 + \alpha_2 \mathbf{x}_i^T \mathbf{A} \mathbf{x}_2 + \cdots + \alpha_{k-1} \mathbf{x}_i^T \mathbf{A} \mathbf{x}_{k-1} + \alpha_k \mathbf{x}_i^T \mathbf{A} \mathbf{x}_k = 0,$$

$$\alpha_i \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i = 0.$$

- since  $\alpha_i \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i > 0, \forall i$  it stands that  $\alpha_i = 0, \forall i$ .

# CG – Building Conjugate Directions

Given these definitions and relationships, we can now outline our strategy for choosing descent directions  $\mathbf{d}_k$

- the descent directions  $\mathbf{d}_k \in \mathbb{R}^n$  will form a finite  $\mathbf{A}$ -conjugate set
  - as such, they will also form a basis for  $\mathbb{R}^n$
- we can thus express solutions  $\mathbf{x}$  of  $\mathbf{Ax} = \mathbf{b}$  as  $\mathbf{x} = \sum_{i=1}^n \alpha_i \mathbf{d}_i$

We will iteratively build the  $\mathbf{A}$ -conjugate set (and, so too, the basis) such that at each step  $j$  we choose a new descent direction  $\mathbf{d}_j$  to be  $\mathbf{A}$ -conjugate to all preceding descent directions, i.e.,  $\mathbf{d}_j^T \mathbf{A} \mathbf{d}_i = 0, \forall i < j$



# CG – Building Conjugate Directions

Drawing from steepest descent, we will seek  $\mathbf{d}_k$ s of the form

$$\mathbf{d}_{k+1} = -\nabla f(\mathbf{x}_{k+1}) + \beta_k \mathbf{d}_k$$

where  $\mathbf{d}_k$  is  $\mathbf{A}$ -conjugate to the previous directions  $\mathbf{d}_i, \forall i < k$

- this reduces the problem to finding the appropriate  $\beta_k$ s

- note that with  $\beta_k = 0, \forall k$  we recover the steepest descent directions

Pre-multiplying the equation for  $\mathbf{d}_{k+1}$  by  $\mathbf{A}$  and then  $\mathbf{d}_k^T$  gives:

$$\begin{aligned}\mathbf{A}\mathbf{d}_{k+1} &= -\mathbf{A}\nabla f(\mathbf{x}_{k+1}) + \beta_k \mathbf{A}\mathbf{d}_k \\ \mathbf{d}_k^T \mathbf{A}\mathbf{d}_{k+1} &= -\mathbf{d}_k^T \mathbf{A}\nabla f(\mathbf{x}_{k+1}) + \beta_k \mathbf{d}_k^T \mathbf{A}\mathbf{d}_k\end{aligned}$$

- this reduces the problem to finding the appropriate  $\beta_k$ s
  - note that with  $\beta_k = 0, \forall k$  we recover the steepest descent directions

Pre-multiplying the equation for  $\mathbf{d}_{k+1}$  by  $\mathbf{A}$  and then  $\mathbf{d}_k^T$  gives:

$$\begin{aligned}\mathbf{A}\mathbf{d}_{k+1} &= -\mathbf{A}\nabla f(\mathbf{x}_{k+1}) + \beta_k\mathbf{A}\mathbf{d}_k \\ \mathbf{d}_k^T\mathbf{A}\mathbf{d}_{k+1} &= -\mathbf{d}_k^T\mathbf{A}\nabla f(\mathbf{x}_{k+1}) + \beta_k\mathbf{d}_k^T\mathbf{A}\mathbf{d}_k\end{aligned}$$

We need to force the LHS to be 0 to maintain  $\mathbf{A}$ -conjugacy, which yields

$$\beta_k = \frac{\mathbf{d}_k^T\mathbf{A}\nabla f(\mathbf{x}_{k+1})}{\mathbf{d}_k^T\mathbf{A}\mathbf{d}_k}$$

# Conjugate Gradients – Algorithm

The **conjugate gradient algorithm** is a modification to SD as:

1. begin with any vector  $\mathbf{x}_1 \in \mathbb{R}^n$  and set  $\mathbf{d}_1 = -\nabla f(\mathbf{x}_1)$ , then iteratively solve for
2. the minimizing distance  $\alpha_k = \frac{-\mathbf{d}_k^\top \nabla f(\mathbf{x}_k)}{\mathbf{d}_k^\top \mathbf{A} \mathbf{d}_k}$  and next CG iterate  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$ , before forming the next CG descent direction as
3.  $\mathbf{d}_{k+1} = -\nabla f(\mathbf{x}_{k+1}) + \beta_k \mathbf{d}_k$  with offset distance  $\beta_k = \frac{\mathbf{d}_k^\top \mathbf{A} \nabla f(\mathbf{x}_{k+1})}{\mathbf{d}_k^\top \mathbf{A} \mathbf{d}_k}$


# Conjugate Gradients for Linear Systems

Conjugate gradient has much nicer convergence properties:


- it is **guaranteed\*** to converge to the solution in **at most**  $n$  steps
  - this means  $O(n^2)$  overall cost for sparse systems, and at worst the same  $O(n^3)$  cost as direct methods for dense systems – in practice, often  $n' \ll n$  iterations are needed
- the conjugate gradient algorithm follows the same structure as gradient descent, but requires some additional linear algebraic manipulation when constructing the conjugate descent directions



# Gradient Descent vs. Conjugate Gradient




Well conditioned  $A$




Poorly conditioned  $A$

# Gradient Descent vs. Conjugate Gradient



Well conditioned  $A$



Poorly conditioned  $A$



# Gradient Descent – Summary

---

Gradient descent is a simple and powerful technique

- many extensions and deeply studied area

In its simplest form, it requires “only” the ability to evaluate the gradient of the function we wish to minimize

- can set step size manually, or take several steps with, e.g., adaptive sizes
- *line search* can be worth the additional costs; requires function evaluations for the 1D optimization

Gradient descent can be specialized to a linear solver with an optimal step size (that doesn't require a line search)

- conjugate gradient has better numerics and stronger guarantees