

Week 3

Synchronization Primitives

Oana Balmau
January 19, 2023

Overview of Assignment 1

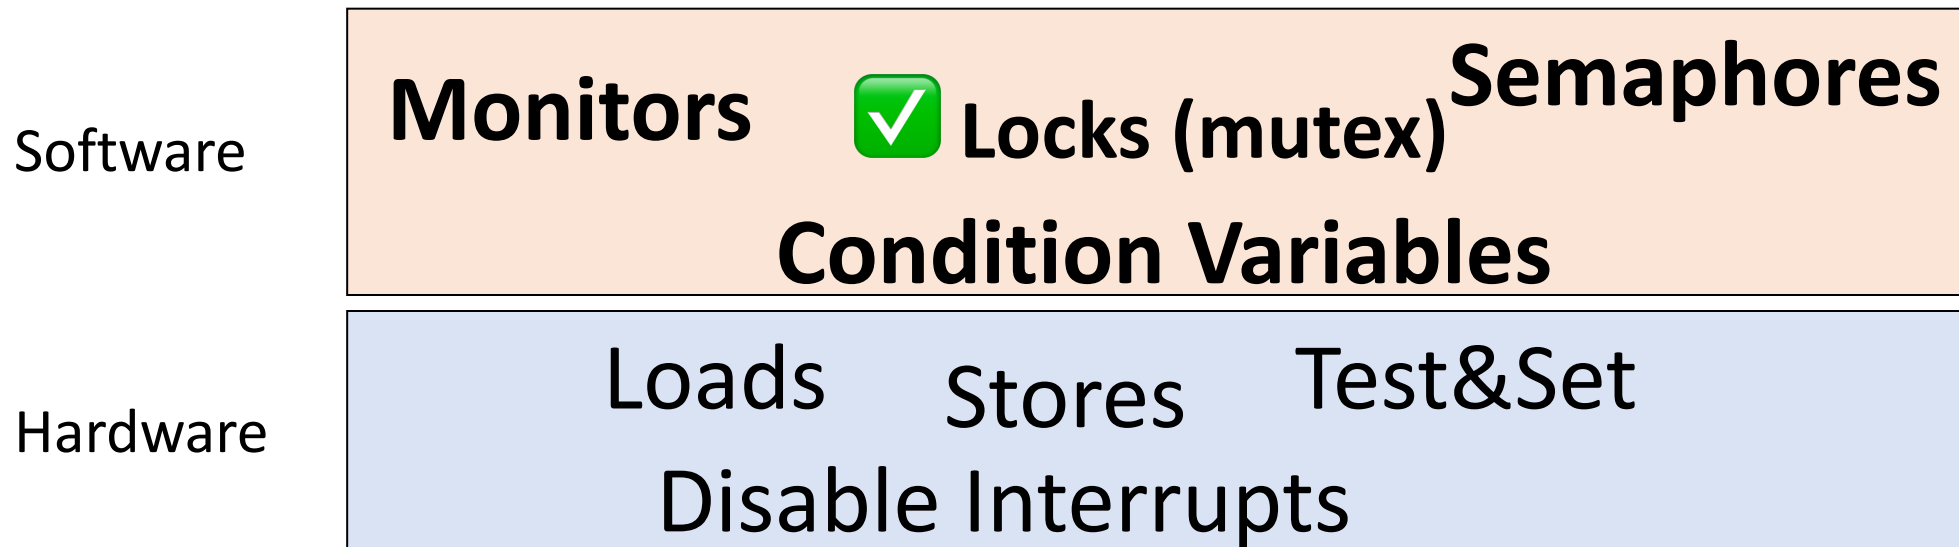
- Presented by Jiaxuan.
- Recording will be posted on MyCourses.

Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



Condition Variables

- Used when thread A needs to wait for an event done by thread B
- A waits until a certain condition is true
 - First test condition,
 - If condition not true, call `pthread_cond_wait()`
 - A blocks until condition is true.
- At some point B makes the condition true
 - Then B calls `pthread_cond_signal()`, which unblocks A.

Advantage of Condition Variable of Locks

- A waits until a certain condition is true
 - A goes to sleep here.
 - Does not keep CPU busy.

Condition variables use (incorrect)

Thread A

```
x = f (a , b) ;  
if (x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Thread B

```
//change a and b;  
x = f (a , b) ;  
if (x >= 0 && x <= 9)  
    pthread_cond_signal (&cv);
```

Find the data race.

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal (&cv);
```

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f ( a , b ) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal (&cv);
```


Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal(&cv);
```

:(Broadcast missed by A

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

A waits forever...

Thread B

```
//change a and b;  
x = f ( a , b ) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv);
```

Condition variables use (still incorrect)

Thread A

```
pthread_mutex_lock(&mutex );  
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv, &mutex);  
pthread_mutex_unlock(&mutex) ;
```

Thread B

```
pthread_mutex_lock(&mutex );  
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv , &mutex);  
pthread_mutex_unlock(&mutex);
```

Remember: Condition variable is a shared resource between A and B

→ Every time you use a condition variable you must also use a mutex to prevent the race condition.

One more issue...

Sometimes, the wait function might return even though the condition variable has not actually been signaled.

Thread A

```
pthread_mutex_lock(&mutex );  
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv, &mutex);  
pthread_mutex_unlock(&mutex) ;
```

Thread B

```
pthread_mutex_lock(&mutex );  
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv , &mutex);  
pthread_mutex_unlock(&mutex);
```

Example:

If process P running A and B receives an OS signal

- Any thread in P can be chosen to process the signal.
- → A might be chosen to process the signal handling function
- → wait returns with an error code → A runs even if condition is not true...

How can we fix this?

Condition variables use (correct)

- Retest the condition after `pthread_cond_wait()` returns.
 - This is most easily done using a loop.

Thread A

```
pthread_mutex lock(&mutex ) ;  
while (1){  
    x = f ( a , b ) ;  
    if ( x < 0 || x > 9)  
        pthread_cond_wait (&cv, &mutex);  
    else break;  
}  
pthread_mutex_unlock(&mutex) ;
```

Thread B

```
pthread_mutex lock(&mutex ) ;  
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal (&cv, &mutex);  
pthread_mutex_unlock(&mutex);
```

Conditional Variables Interface

- `pthread_cond_init(pthread_cond_t *cv, pthread_condattr_t *cattr)`
 - Initialize the conditional variable, cattr can be NULL
- `pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex)`
 - Block thread until condition is true, and atomically unblock mutex.
- `pthread_cond_signal(pthread_cond_t *cv)`
 - Unblock one thread at random that is blocked by the condition variable
- `pthread_cond_broadcast(pthread_cond_t *cv)`
 - Unblock all threads that are blocked on the condition variable pointed to by cv.

Condition Variable Example

```
pthread_cond_t is_zero;
pthread_mutex_t mutex;
int shared_data = 100;

void *thread_func(void *arg){
    while(shared_data > 0) {
        pthread_mutex_lock(&mutex);
        shared_data--;
        printf("%d ", shared_data);
        pthread_mutex_unlock(&mutex);
    }

    printf("Signaling main\n");
    pthread_cond_signal(&is_zero);
    return NULL;
}
```

```
int main (void){

    pthread_t tid;
    void * exit_status;
    int i;

    pthread_cond_init(&is_zero, NULL);
    pthread_mutex_init(&mutex, NULL);

    pthread_create(&tid, NULL, thread_func, NULL);

    pthread_mutex_lock(&mutex);
    printf("Start waiting in main\n");
    while(shared_data!=0)
        pthread_cond_wait(&is_zero, &mutex);
    pthread_mutex_unlock(&mutex);

    printf("Done waiting in main!\n");

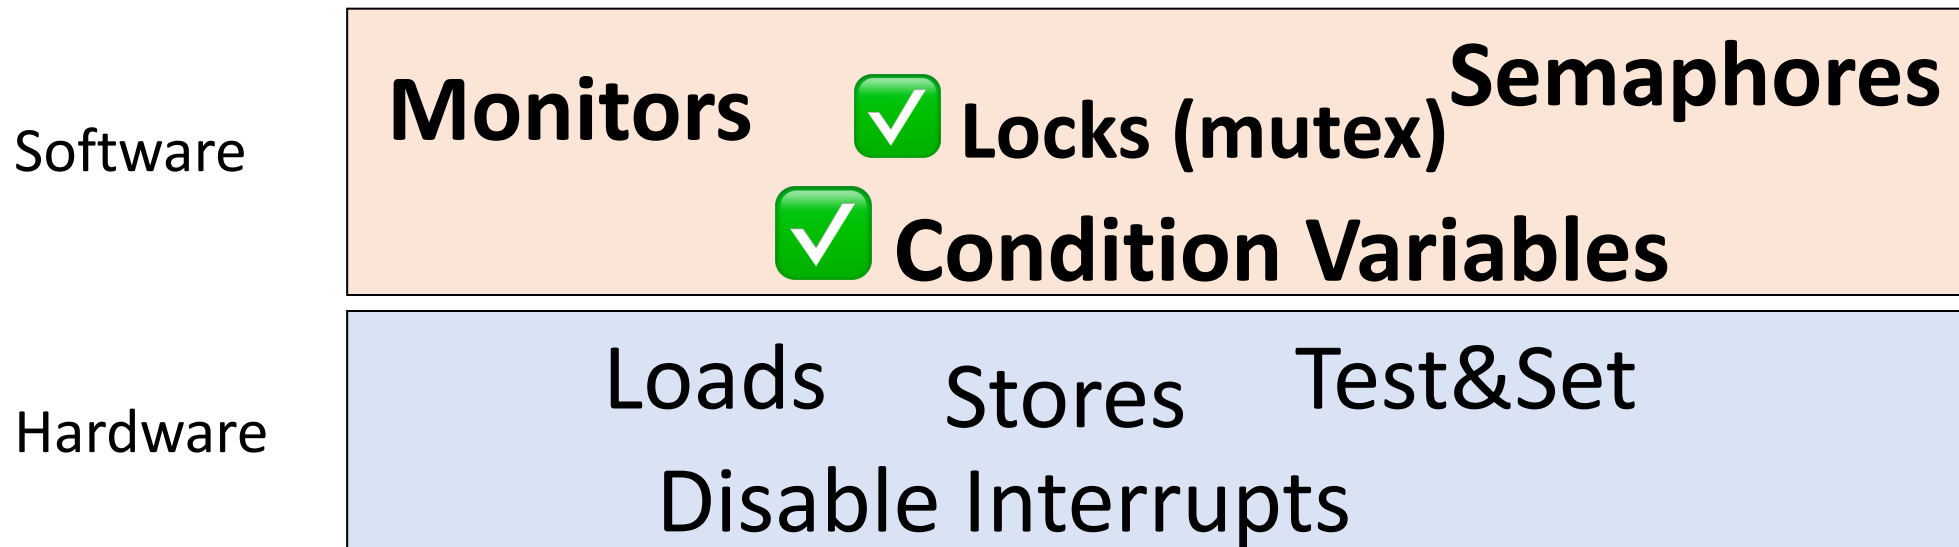
    pthread_join(tid, &exit_status);
    return 0;
}
```

Synchronization


Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



Semaphores



What are Semaphores?

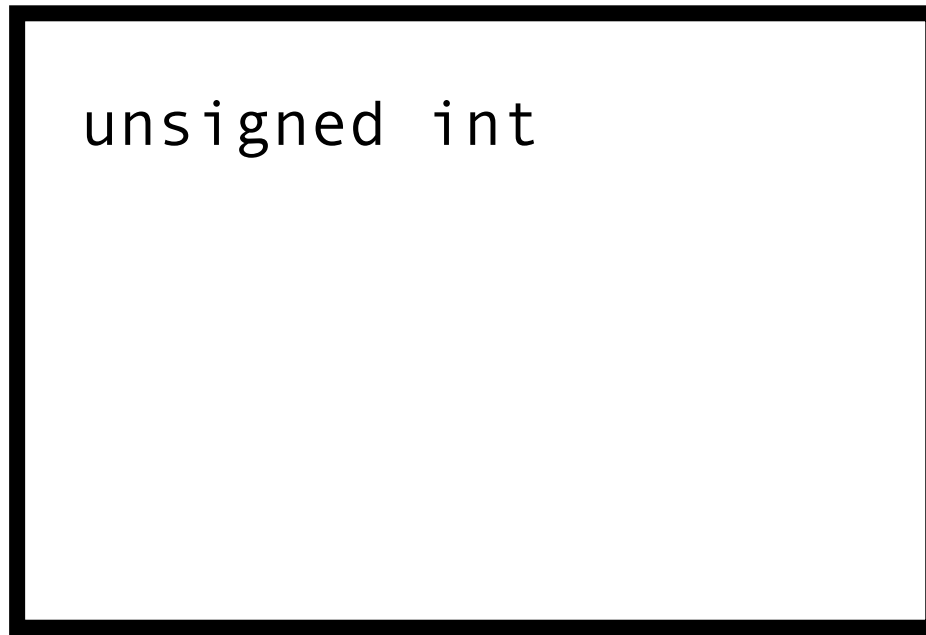
- A shared, non-negative counter.
- Two primary operations:
 - Wait → attempts to decrement the counter; blocks when counter is 0.
 - Post (or Signal) → attempts to increment the counter.
- `#include<semaphore.h>`

Changes are atomic



unsigned int

Semaphore



Semaphore

Changes are atomic
2 operations

`wait()`

`post()`



unsigned int

Semaphore

Changes are atomic


2 operations

wait()

try to decrement semaphore value
block if value = 0

post()

increment semaphore value



Semaphore

```
wait() {
```

```
while(1){
```


```
    if (v>0){  
        v--;  
        return;  
    }
```

```
}
```

```
}
```

Atomic execution

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

```
wait() {
```

```
while(1){
```

```
if (v>0){
```

```
    v--;
```

```
    return;
```


```
}
```

```
}
```

```
}
```

Atomic execution

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

```
wait() {
```

```
  while(1){
```

```
    if (v>0){
```

```
      v--;
```

```
      return;
```

```
    }
```


```
  }
```

```
}
```

Atomic execution

Wait until semaphore value becomes positive again

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

```
wait() {
```

```
while(1){
```


```
    if (v>0){  
        v--;  
        return;  
    }
```

Atomic execution

```
}
```

Once the value is positive, the thread that
Was waiting is able to decrement the value.

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore


```
post() {
```

```
    v++;  
    return;
```

```
}
```

Atomic execution

Note: post() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

```
post() {
```

```
    v++;
```

```
    return;
```

```
}
```

Atomic execution

Note: post() is not actually implemented like this. This is how the behavior looks like to the programmer



unsigned int

Semaphore

One last thing:

- Semaphore value can be initialized upon creation to a positive value
- or 0

Semaphore Uses

Mutual exclusion

- A semaphore with its counter initialized to 1 is equivalent to a lock.

Bound the concurrency

- Only allow X threads out of N to proceed.

Producer-consumer problem

- More complex use of semaphores. Will see in 2 weeks.

Semaphores interface

- `int sem_init(sem_t *sem, int pshared, unsigned value);`
- `int sem_post(sem_t *sem);`
- `int sem_wait(sem_t *sem);`

For more details: <https://man7.org/linux/man-pages/man0/semaphore.h.0p.html>

```
int sem_init(sem_t *sem, int pshared, unsigned value);
```

- Initializes the semaphore * *sem*. The initial value of the semaphore is *value*.
- If *pshared* is 0, the semaphore is shared among all threads of a process.
- If *pshared* is not zero, the semaphore is shared but should be in shared memory.
- Return 0 on success, -1 on failure.

```
int sem_wait(sem_t *sem);
```

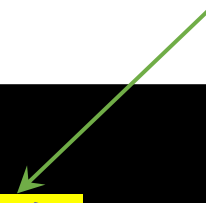
- If the *sem* has a value > 0 , decrement the value by 1.
- If *sem* has value 0, the caller will be blocked (busy-waiting or more likely on a queue) until *sem* has a value larger than 0.
- Return 0 on success, -1 on failure.


```
int sem_post(sem_t *sem);
```

- Increment the value of *sem* by 1.
- If threads are blocked waiting for the semaphore, one of them (at random) will return successfully from its call to *sem_wait()*; the semaphore value is immediately decremented.
- Return 0 on success, -1 on failure.

Semaphores example

Note: If 2 is changed to 1, we have a lock behavior



```
#include <pthread.h>
#include <semaphore.h>

pthread_t threads[5];
int tid[5];
sem_t sem;

void * thread_func(void *arg){
    int tid_ = tid[* (int *) arg];
    printf("Thread %d created\n", tid_);
    int j;

    sem_wait(&sem);
    for (j=0; j<3; j++){
        printf("T%d run %d\n", tid_, j);
        sleep(2);
    }
    sem_post(&sem);
}
```

```
int main(){
    sem_init(&sem, 0, 2);
    //sem initialized for all threads in the
    process; allow only 2 threads in the critical
    section at a time;

    int i;
    for (i=0; i<5; i++){
        tid[i]=i;
        pthread_create(&threads[i], NULL,
        thread_func, &tid[i]);
    }


    for (i=0; i<5; i++){
        pthread_join(threads[i], NULL);
    }
    sem_destroy(&sem);
    return 0;
}
```

Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



Monitors

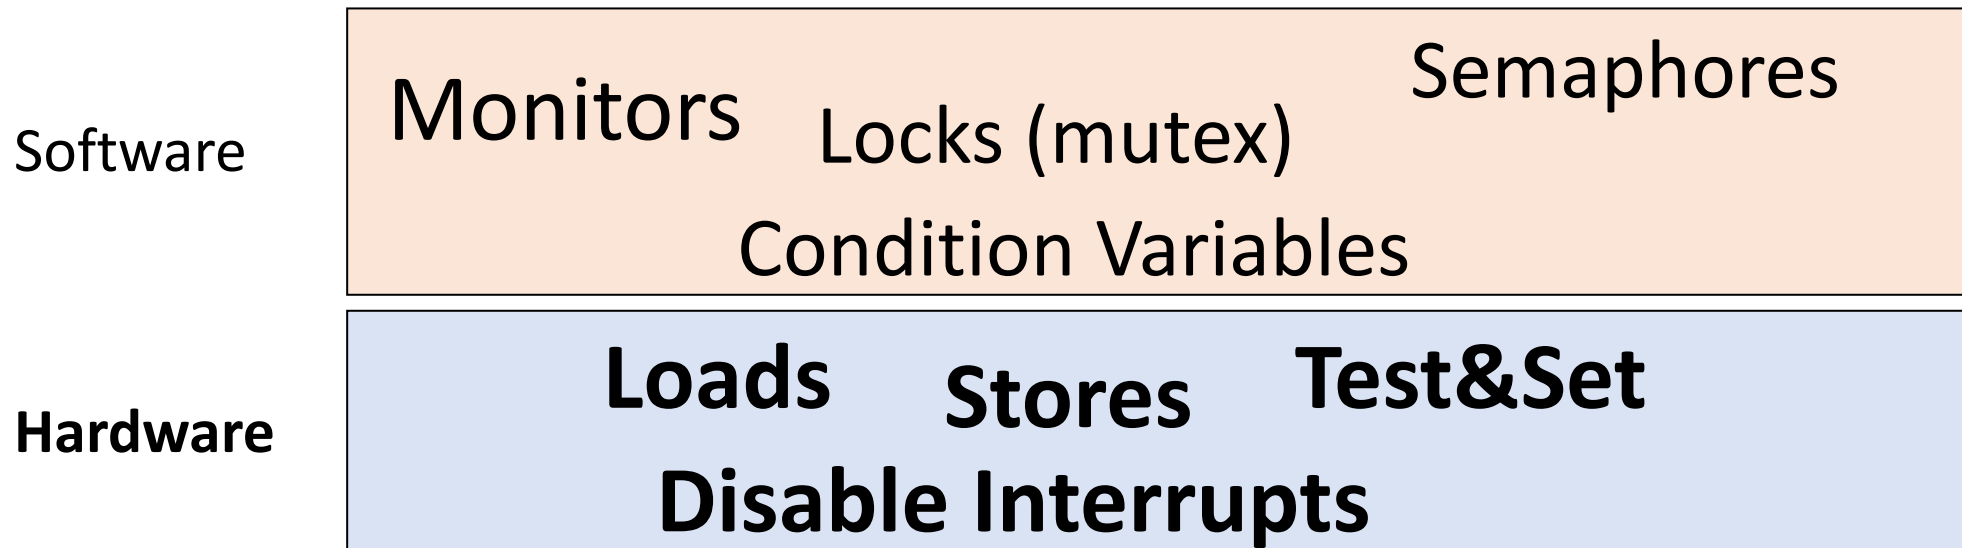
- Collection of variables and functions
- Threads can only access monitor functions
 - Variables are private to the monitor
- Only one process at a time can execute code inside the monitor.

Monitors

- Pthreads does not offer a monitor primitive 😞
- ... but possible to implement monitor semantics using mutex and condition variables 😊
- We will see how in 2 weeks.

Synchronization

How are locks, semaphores, cond var, etc **implemented?**



Lock implementation motivating example

“Too much milk”

- Alice and Bob are roommates
- They want to coordinate grocery shopping
 - Need to be careful to not buy too much of perishable items, like milk.



Schedule that leads to too much milk

Time	Alice	Bob
3:00	Look in Fridge. Out of milk.	
3:05	Leave for store.	
3:10	Arrive at store.	Look in Fridge. Out of milk.
3:15	Buy milk.	Leave for store.
3:20	Arrive home, put milk away.	Arrive at store.
3:25		Buy milk.
3:30		Arrive home, put milk away.

Problem specifications

Safety

- Never more than one person buys

Liveness

- Someone buys if needed

Lock implementation: first attempt

Restrict ourselves to use only atomic load and store operations as building blocks.

Idea: Use **a note** to avoid buying too much milk:

- Leave a note on fridge before buying (kind of “lock”)
- Remove note after buying (kind of “unlock”)
- Don’t buy if note (wait)

Lock implementation: first attempt (incorrect)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note; }  
}
```



Lock implementation: first attempt (incorrect)

```
if (noMilk) { load (atomic)
  if (noNote) {
    leave Note;
    buy milk; store (atomic)
    remove note; }
}
```



Lock implementation: first attempt (incorrect)

Thread A

```
if (noMilk) {  
  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

Thread B

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove note; }  
}
```

Lock implementation: first attempt (incorrect)

Thread A

```
if (noMilk) {  
  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

Remember: scheduler can create any interleaving
We are assuming a malicious scheduler

Thread B

```
if (noMilk) {  
    if (noNote) {  
  
  
        leave Note;  
        buy milk;  
        remove note; }  
}
```

First attempt result

- Still too much milk but only occasionally!
 - This is worse than a consistent error, because it is harder to catch.
- Thread can get context switched after checking milk and note but before buying milk!

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```

Problem solved?

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

Proc switch leave note A;
if (noNote B) {
 if (noMilk) {
 buy milk;
 }
}
remove note A;

Thread B

leave note B;
if (noNote A) {
 if (noMilk) {
 buy milk;
 }
}
remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

```
if (noNote B) {  
    if (noMilk) {  
        buy milk;  
    }  
}
```

remove note A;

Proc switch

Thread B

leave note B;

```
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}
```

remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

if (noNote B) {
 if (noMilk) {
 buy milk;
 }

}

Proc switch

remove note A;

Thread B

leave note B;

if (noNote A) {
 if (noMilk) {
 buy milk;
 }

}

remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```

Proc switch

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```

Not quite...

Nobody buys milk 😞

Lock implementation: third attempt (correct)

- How about labeled notes?

Thread A

```
leave note A;  
while (note B)  
do nothing;  
if (noMilk)  
    buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

This works!
B has priority to buy

Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
    do nothing;  
if (noMilk)  
    buy milk;  
remove note A;
```

Happened before



Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```


Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
    do nothing;
```

```
if (noMilk)  
    buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
do nothing;
```

Wait for note B
to be removed

```
if (noMilk)  
    buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
do nothing;
```



Wait for note B
to be removed

```
if (noMilk)  
buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 2

Thread A

```
leave note A;  
while (note B)  
    do nothing;  
if (noMilk)  
    buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 2

Thread A

```
leave note A;  
while (note B)  
    do nothing;  
if (noMilk)  
    buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 2

Thread A

```
leave note A;  
while (note B)  
do nothing;
```

```
if (noMilk)  
buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before

Wait for note B
to be removed

Attempt 3 discussion

Solution 3 works, but it's really unsatisfactory.

- Complex, even for this simple example.
 - Hard to convince yourself that this really works.
- A's code is different from B's –what if lots of threads?
 - Code would have to be slightly different for each thread.
- While A is waiting, it is consuming CPU time.
 - This is called “busy-waiting”.

There must be a better way!

- Have higher-level hardware primitives than atomic load & store
- Build higher-level abstractions on this hardware support

Disabling interrupts

- Lock implementation code executed in kernel mode.
- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

Disabling interrupts

```
int value = FREE;
```

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts? }  
    else {  
        value = BUSY;  
        enable interrupts;  
    }
```

```
Unlock() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue;  
        place on ready queue;}  
    else {  
        value = FREE;  
        enable interrupts; }  
}
```

Disabling interrupts discussion

Why do we need to disable interrupts?

- Avoid interruption between checking and setting lock value
- Otherwise two threads could think that they both have lock

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts? }  
    else {  
        value = BUSY;  
        enable interrupts;  
    }
```

Critical section is short
in kernel mode

What about enabling interrupts when going to sleep?

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts? }  
    else {  
        value = BUSY;}  
    enable interrupts;}  
}
```

Enable interrupts here?

☹ Release can check the queue and not wake up thread

What about enabling interrupts when going to sleep?

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue:  
        Go to sleep();  
        // Enable interrupts? }  
    else {  
        value = BUSY;  
        enable interrupts;  
    }
```

Enable interrupts here?

- Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
- Misses wakeup and still holds lock (deadlock!) 😞😞

What about enabling interrupts when going to sleep?

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts? }  
    else {  
        value = BUSY;  
        enable interrupts;  
    }
```

Want to enable interrupts after sleep()

- But how?

What about enabling interrupts when going to sleep?

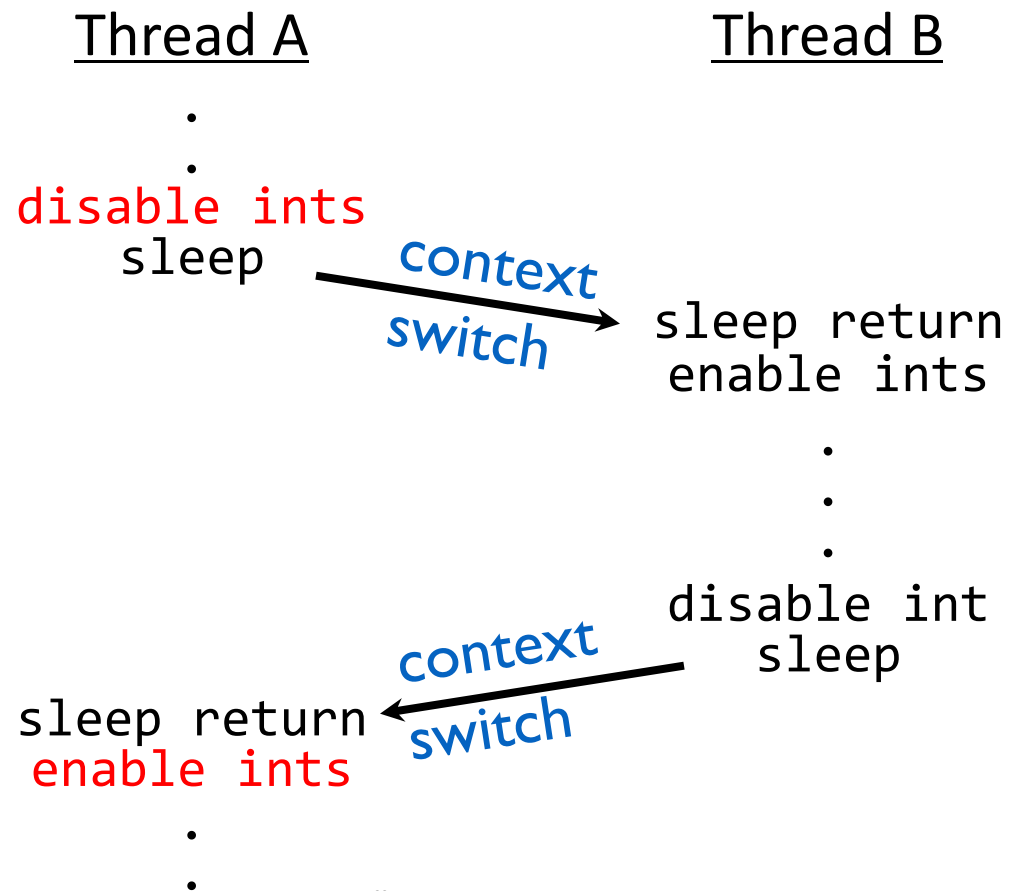
```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts? }  
    else {  
        value = BUSY;  
        enable interrupts;  
    }
```

Want to enable interrupts after sleep()

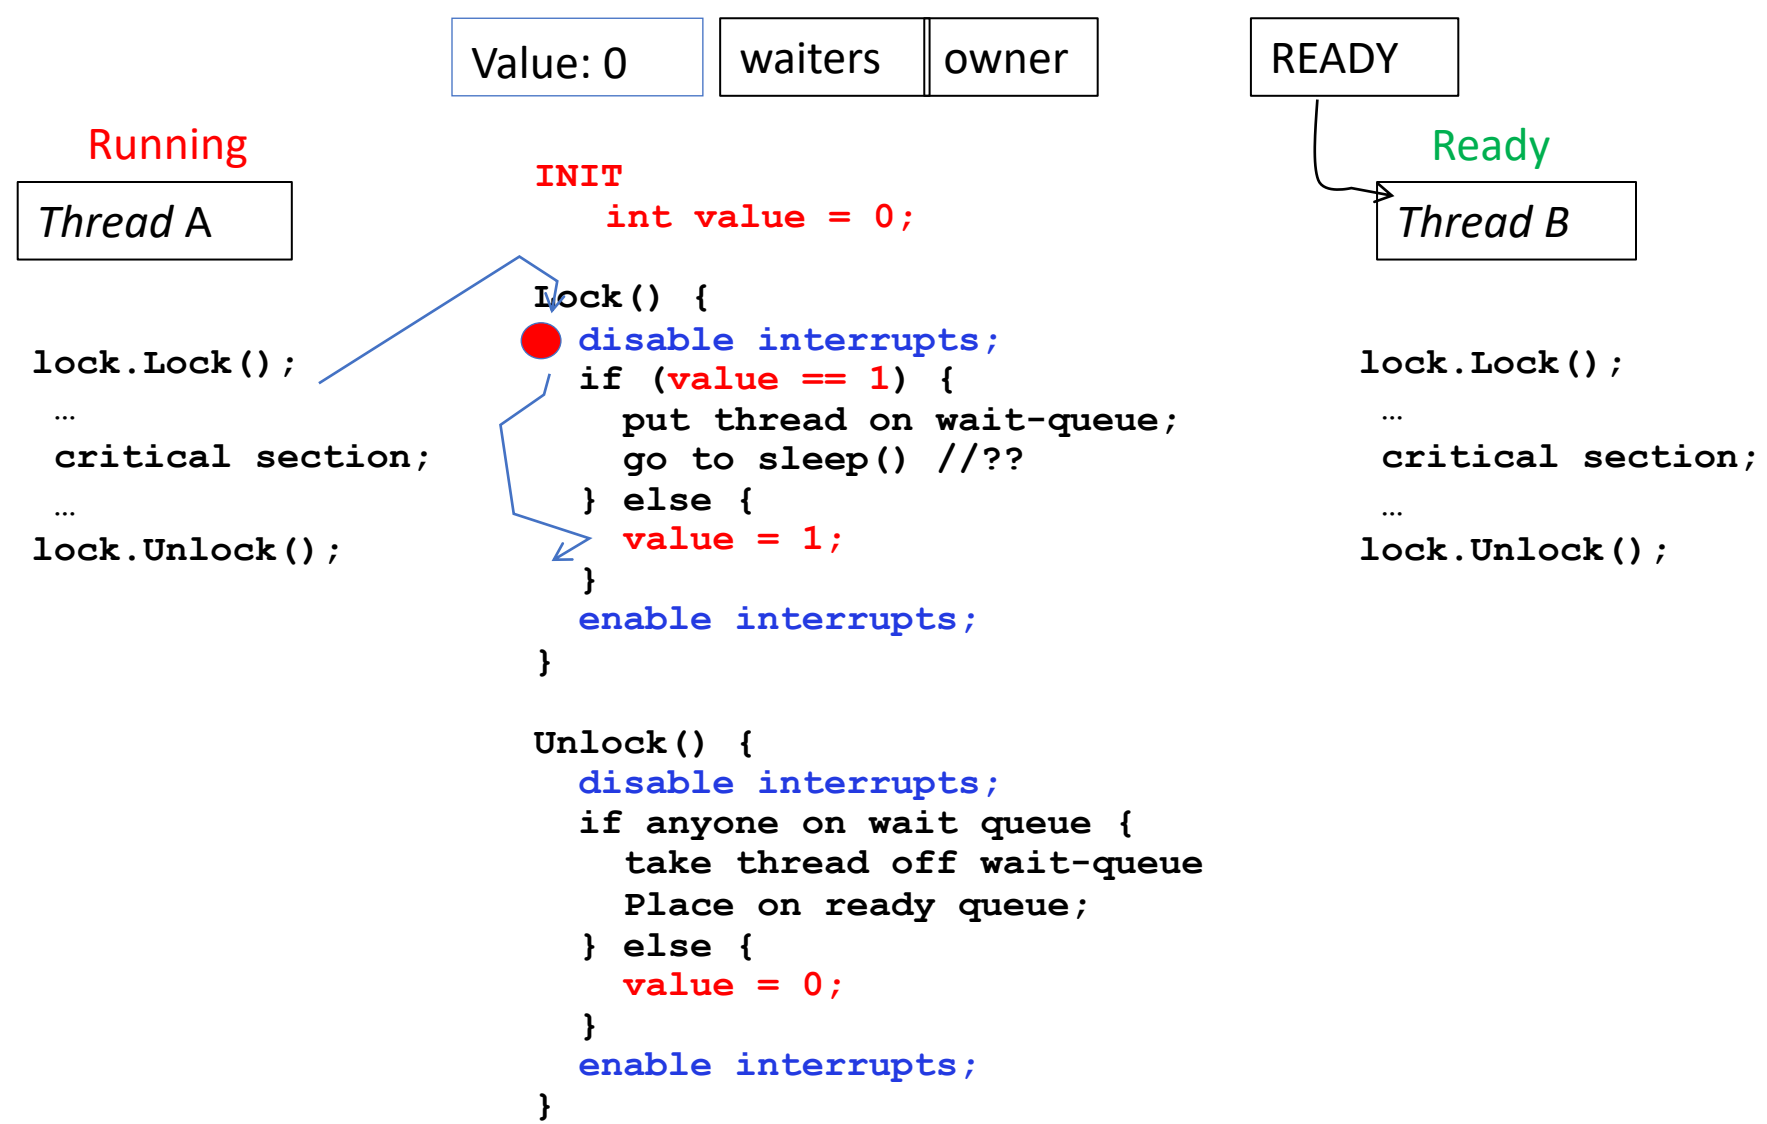
- But how? In scheduler.

In scheduler, since interrupts are disabled when you call sleep():

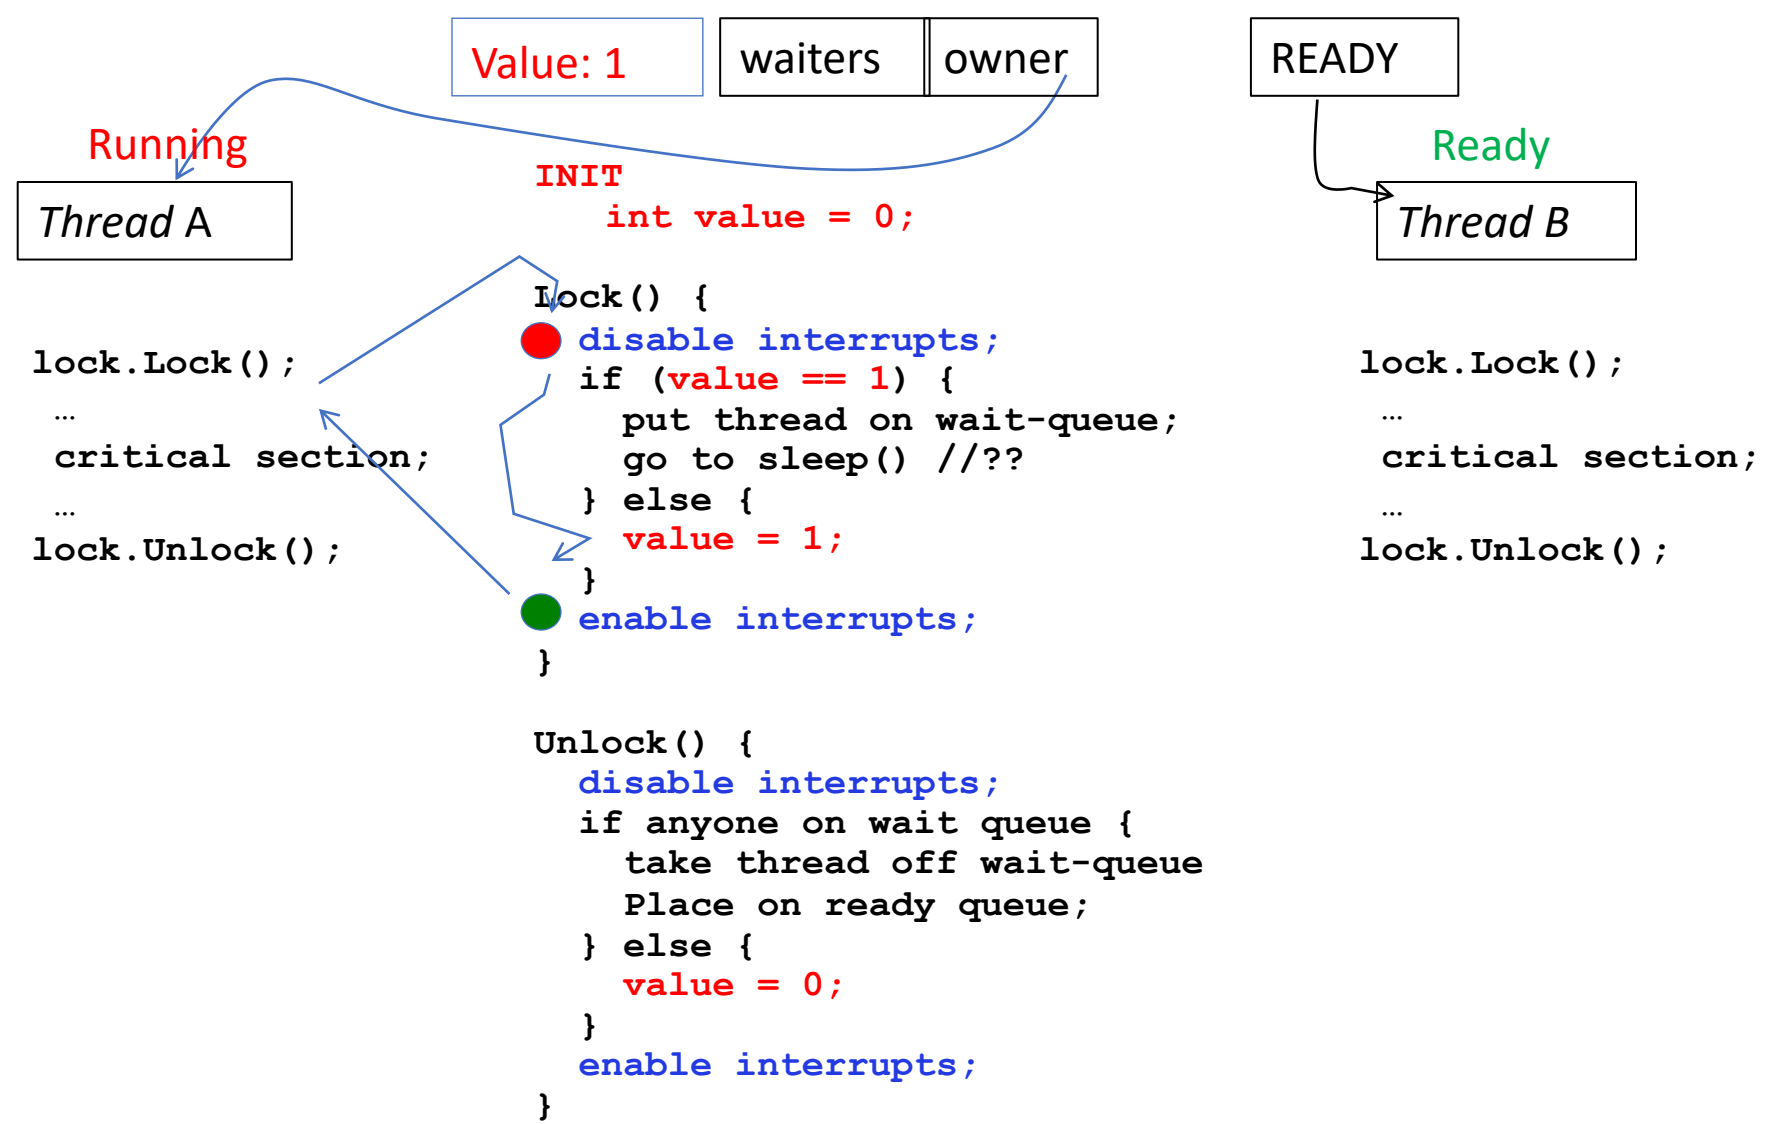
- Responsibility of the next thread to re-enable interrupts
- When the sleeping thread wakes up, returns to lock() and re-enables interrupts



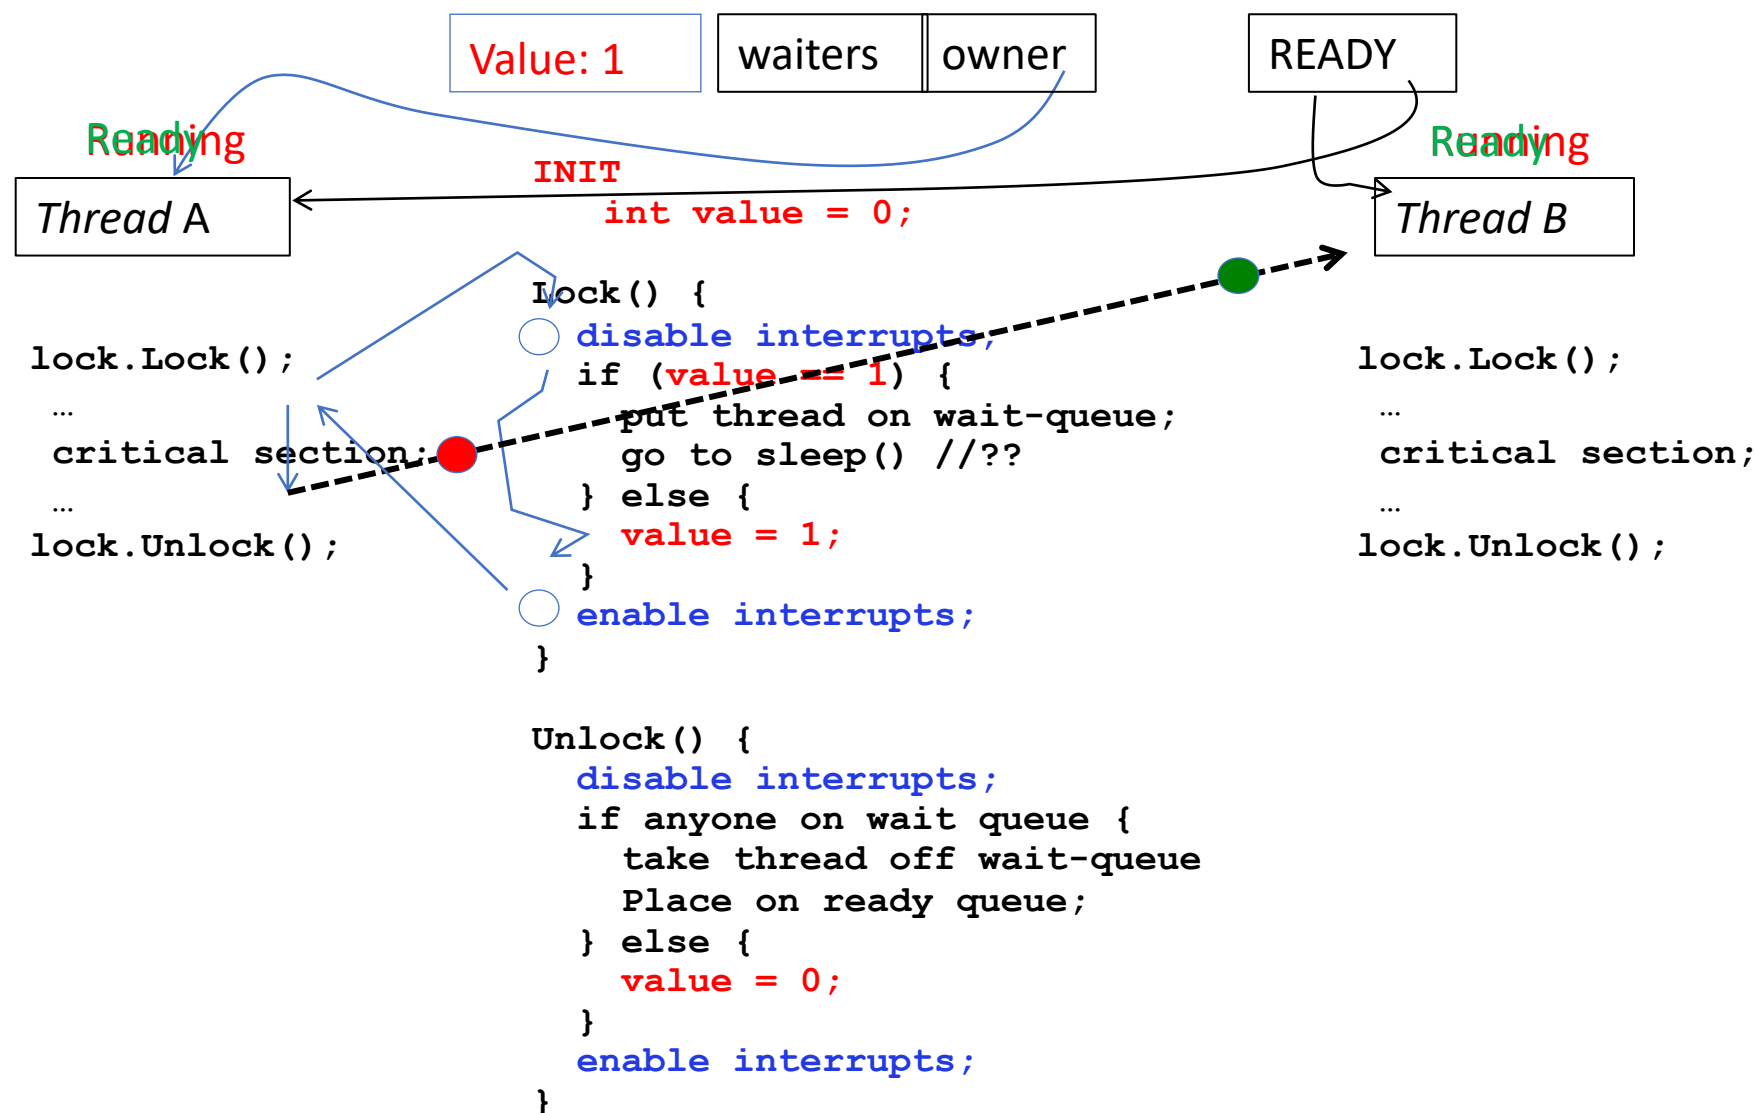
Simulation of locks with disabled interrupts



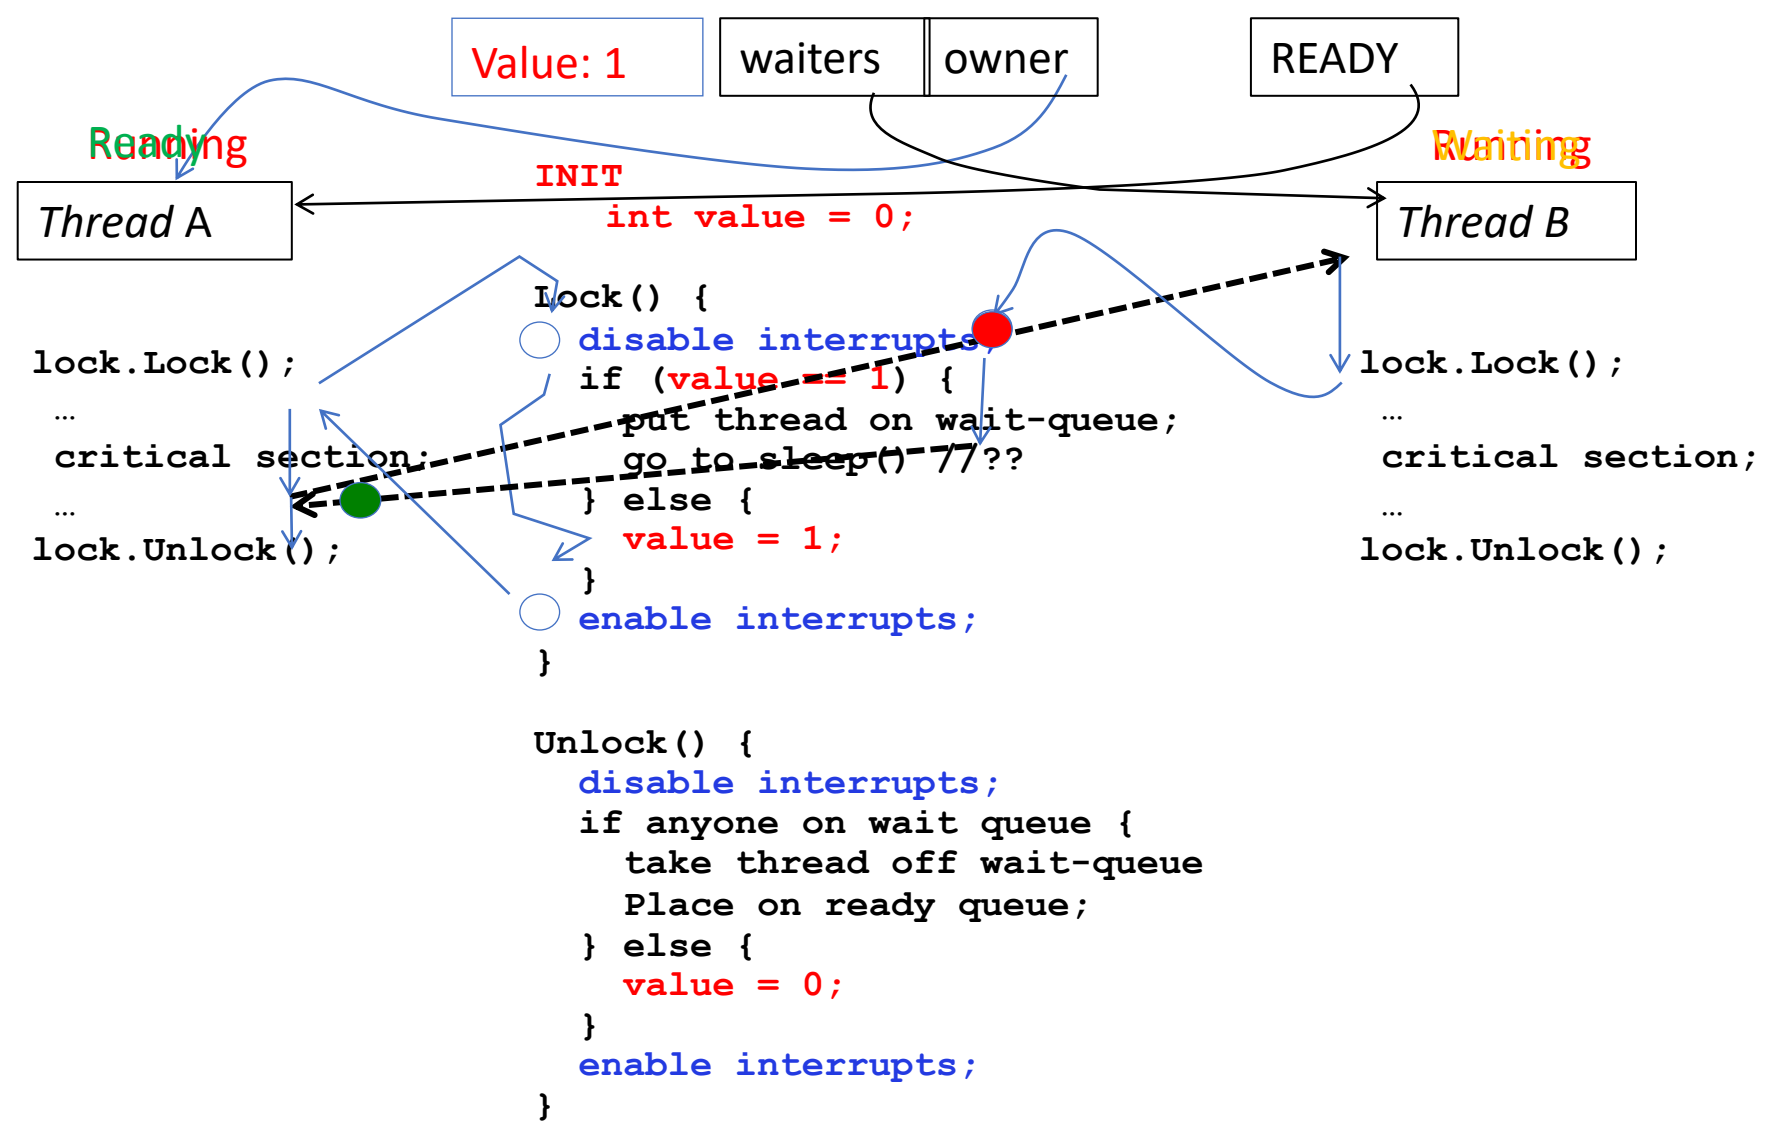
Simulation of locks with disabled interrupts



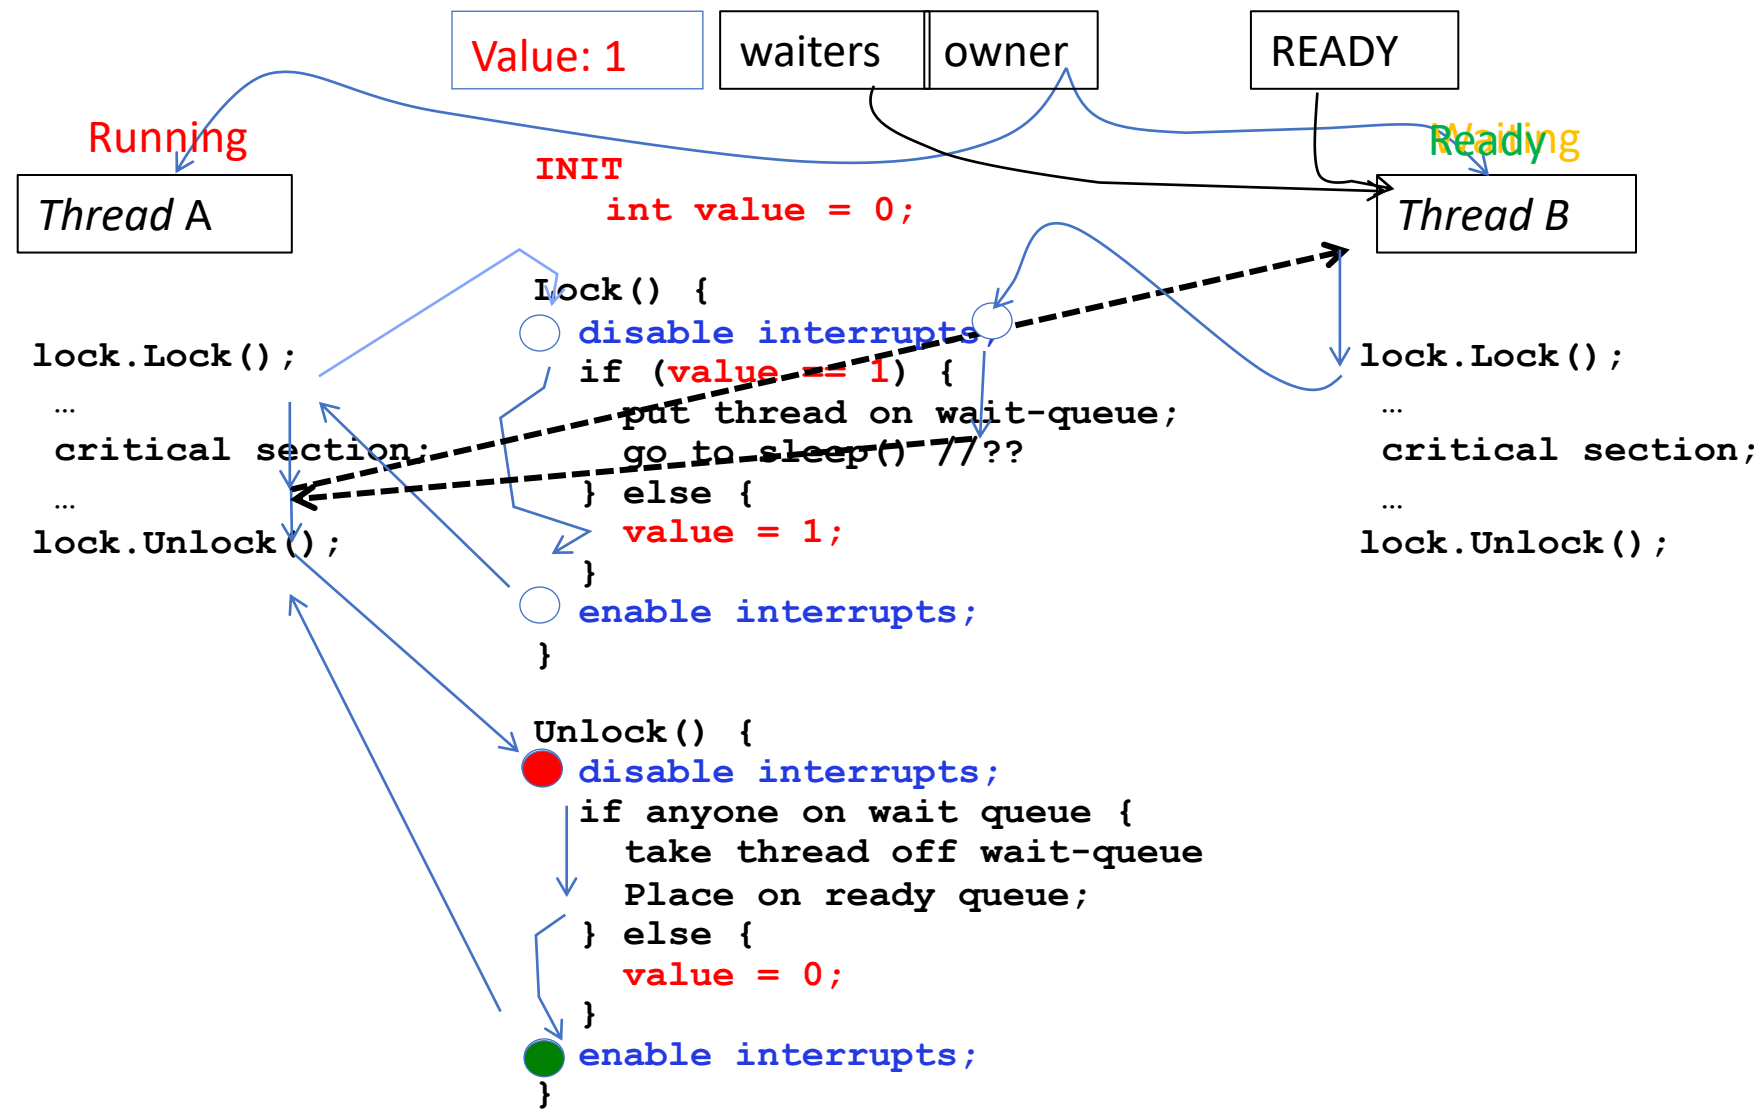
Simulation of locks with disabled interrupts



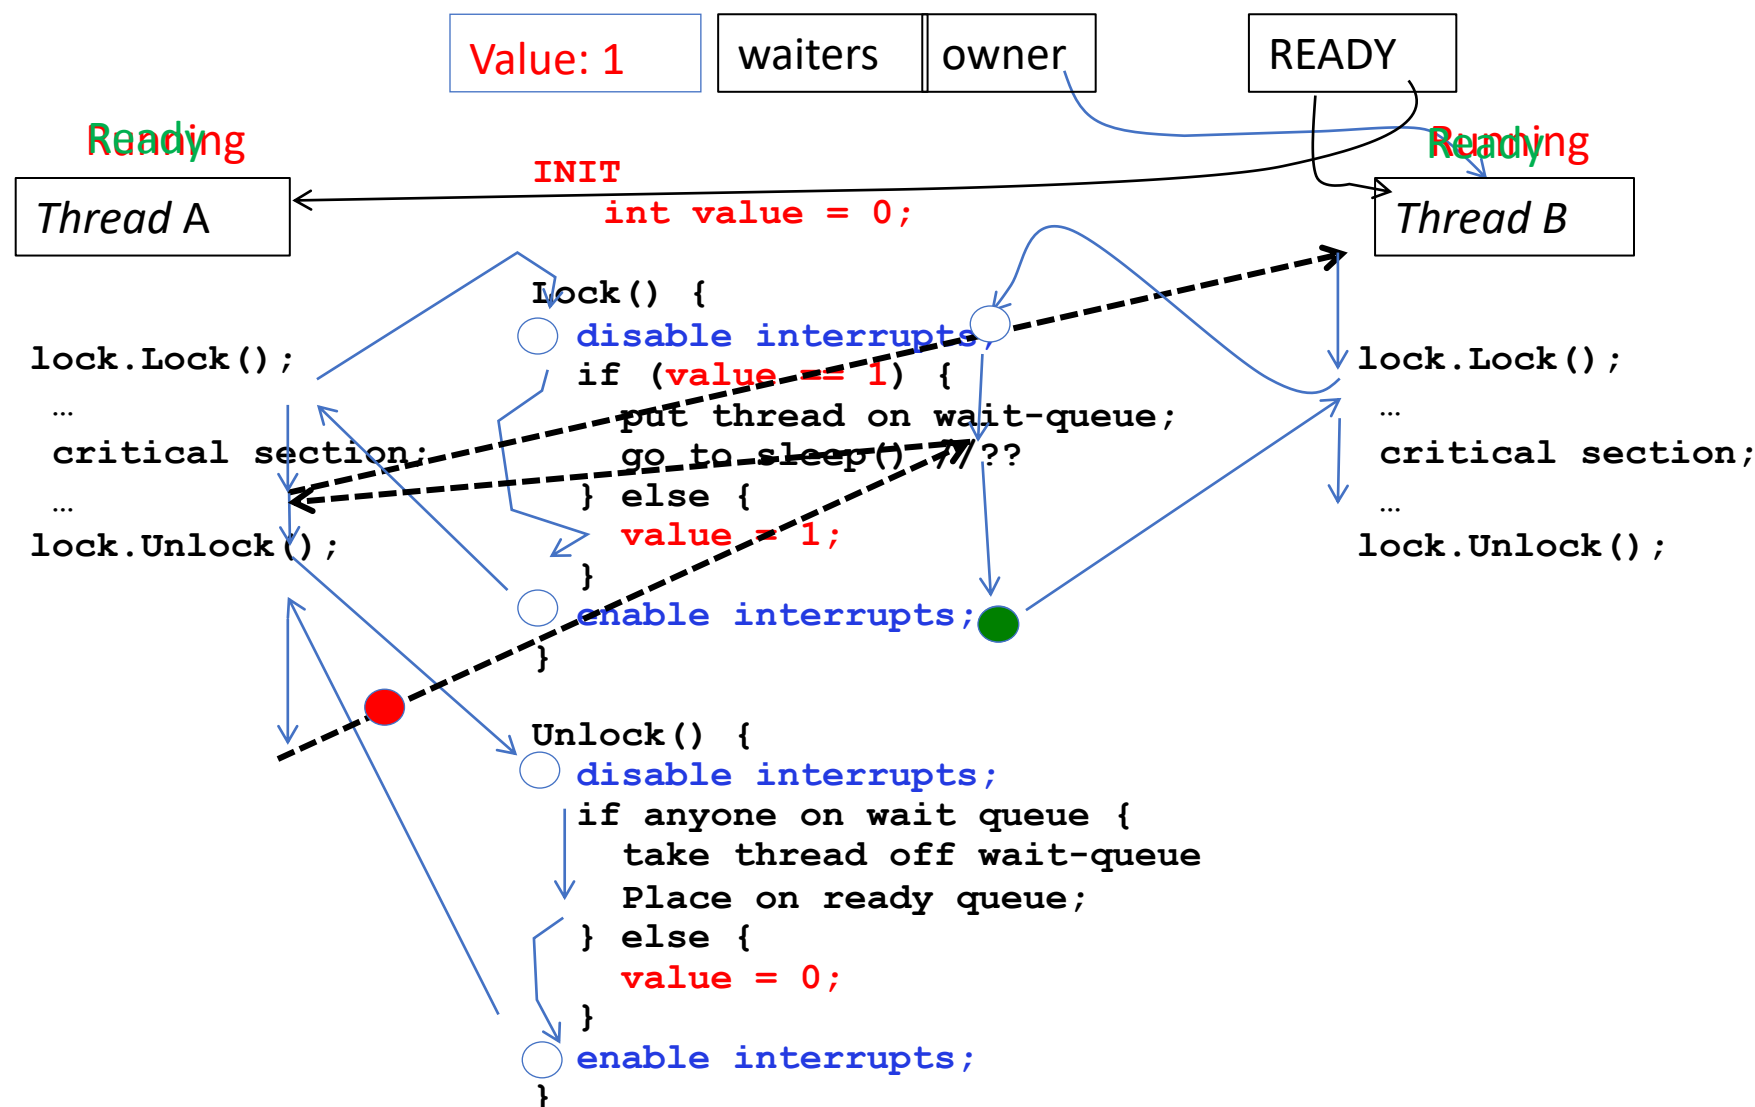
Simulation of locks with disabled interrupts



Simulation of locks with disabled interrupts



Simulation of locks with disabled interrupts



Further Optional Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 25 – 31 (inclusive)

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

Reading on concurrency: Herlihy & Shavit: The Art of Multiprocessor Programming, 2nd edition.

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Maurice Herlihy (Brown University), Prof. Natacha Crooks (UC Berkeley).

