

# Week 4

# **Multiprocess Communication**

Oana Balmau  
January 24, 2023

# Schedule Highlights

Week 4 Process Management	jan 23 C Tools: GDB basics	jan 24 Multi-process Structuring (1/2) Team registration deadline	jan 25	jan 26 Multi-process Structuring (2/2)	jan 27
Week 5 Process Management	jan 30 C Review: Pointers & Memory Allocation I	jan 31 Multithreading (1/2) Practice Exercises Sheet: Process Management	feb 1	feb 2 Multithreading (2/2)	feb 3
Week 6 Memory Management	feb 6 C Review: C files	feb 7 Virtual Memory (1/2) Optional reading: <a href="#">OSTEP</a> Chapters 12 – 18	feb 8 Scheduling Assignment Released	feb 9 Virtual Memory (2/2) Scheduling Assignment Overview — with Jiaxuan	feb 10

Please sign up teams of 1 as well!

# Recap Week 3

## Concurrency – Option 1

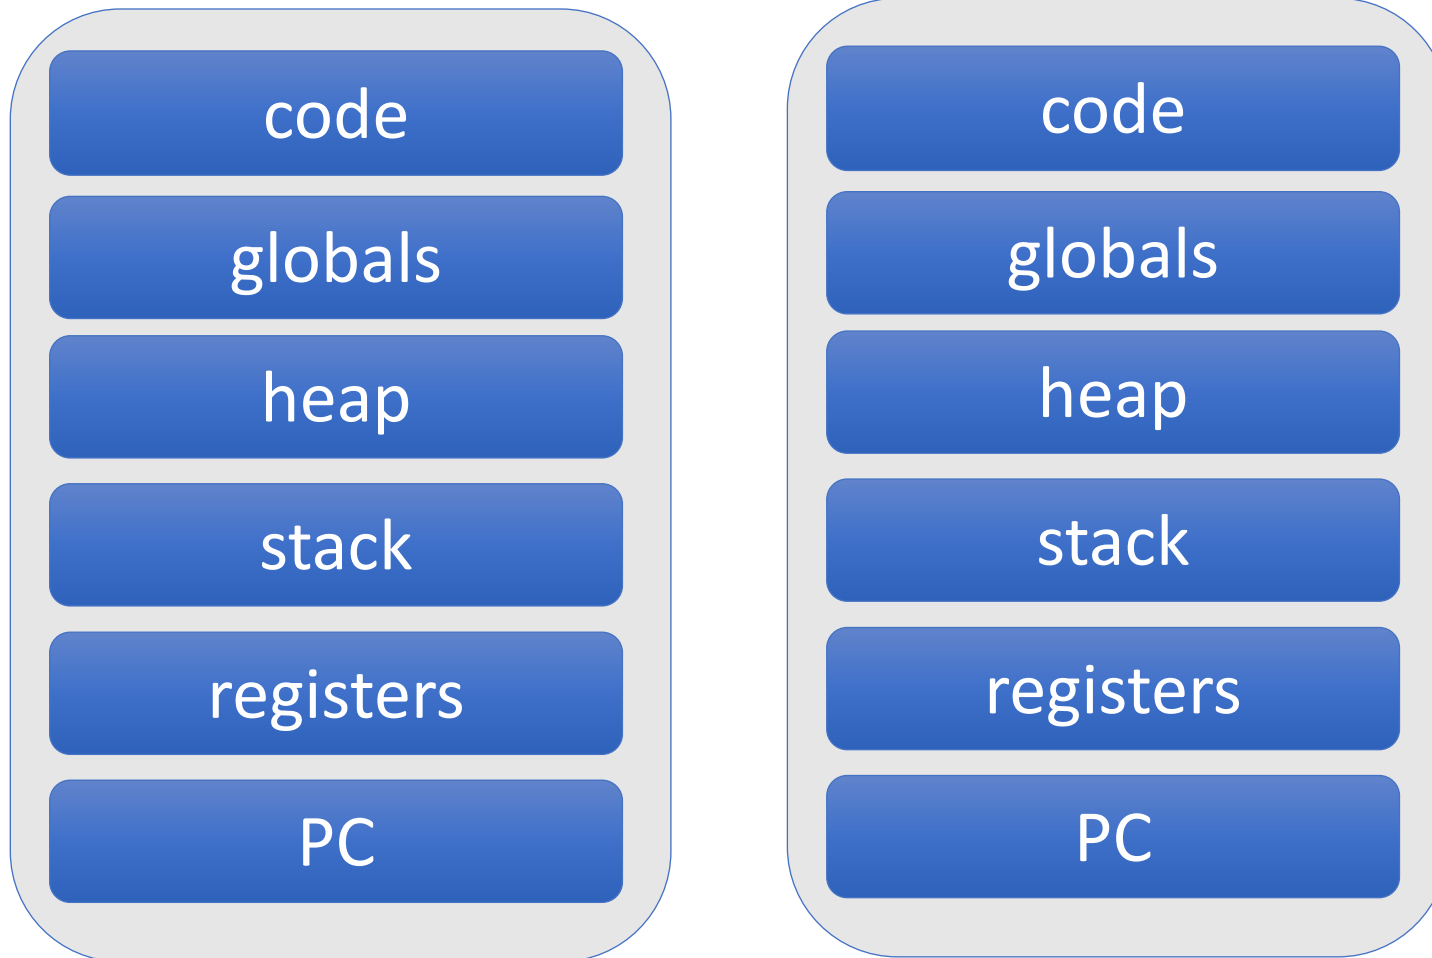
- Build apps from many communicating **processes**
- Communicate through message passing
  - No shared memory
- Pros
  - If one process crashes, other processes unaffected
- Cons
  - High communication overheads
  - Expensive context switching



This week's  
focus

# Recap Week 3

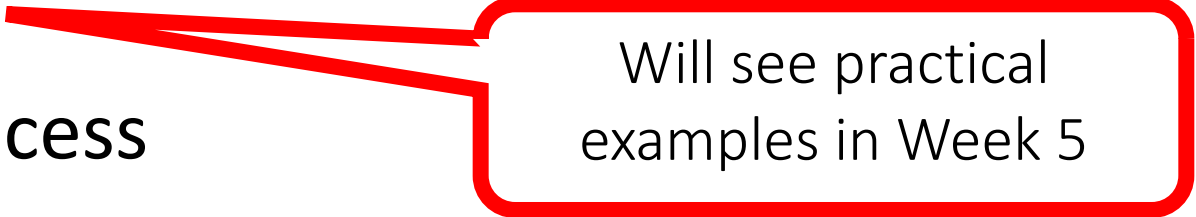
## Two Processes



# Recap Week 3

## Concurrency – Option 2

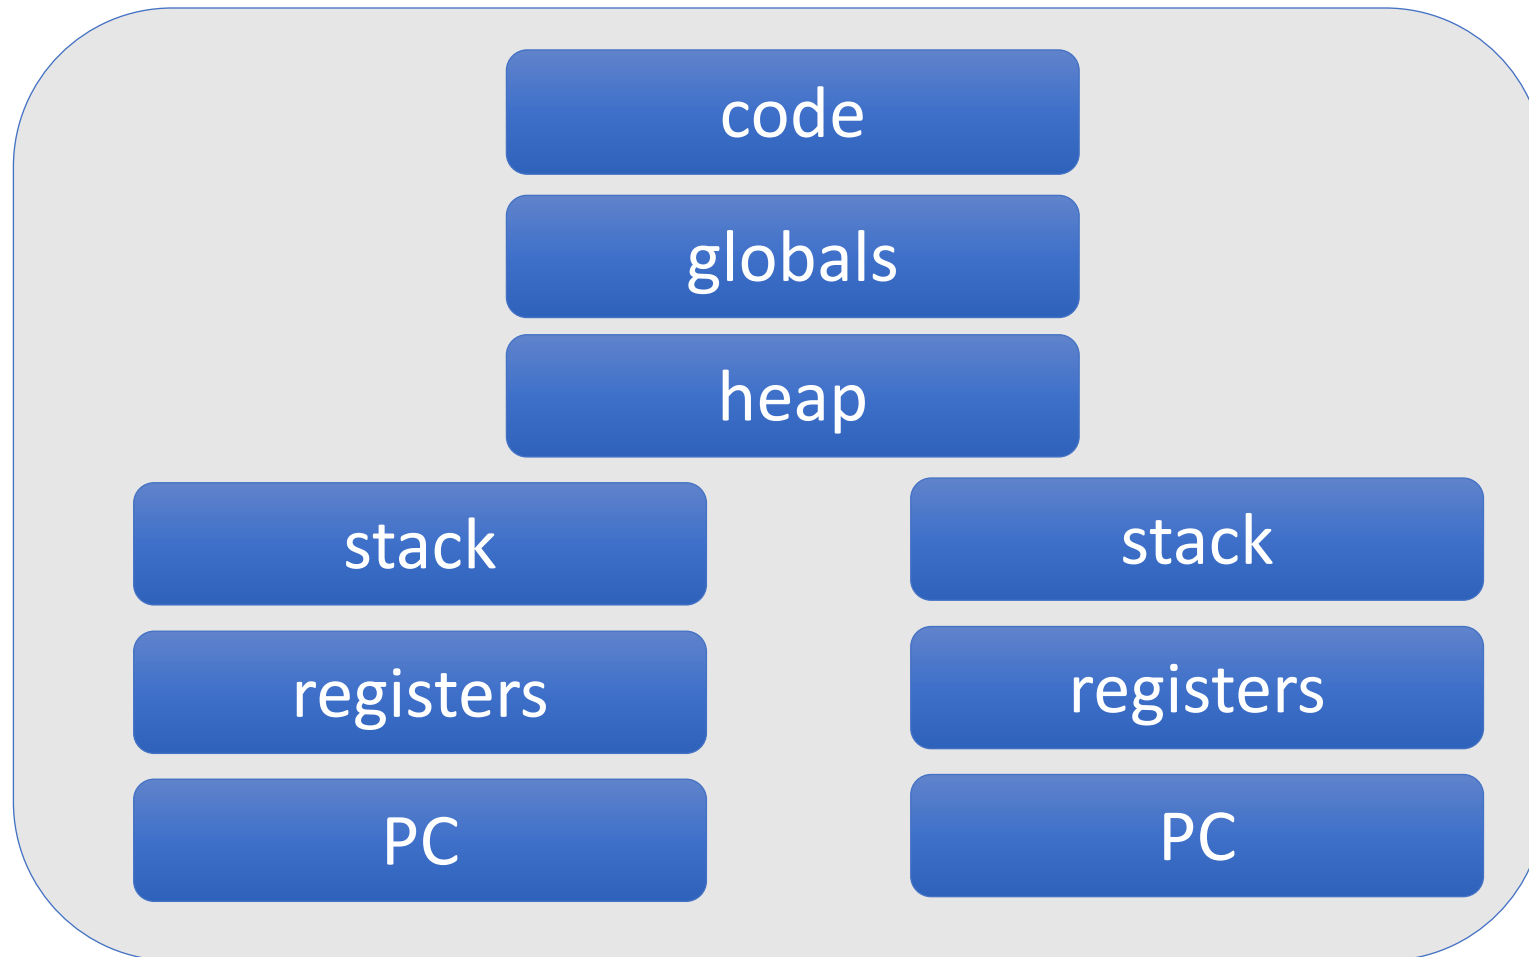
- New abstraction: **thread**
- Multiple threads in a process
- Threads are like processes except
  - Multiple threads in the same process share an address space
    - Communicate through shared address space
  - If one thread crashes,
    - the entire process, including other threads, crashes



Will see practical examples in Week 5

# Recap Week 3

## Two Threads in a Process

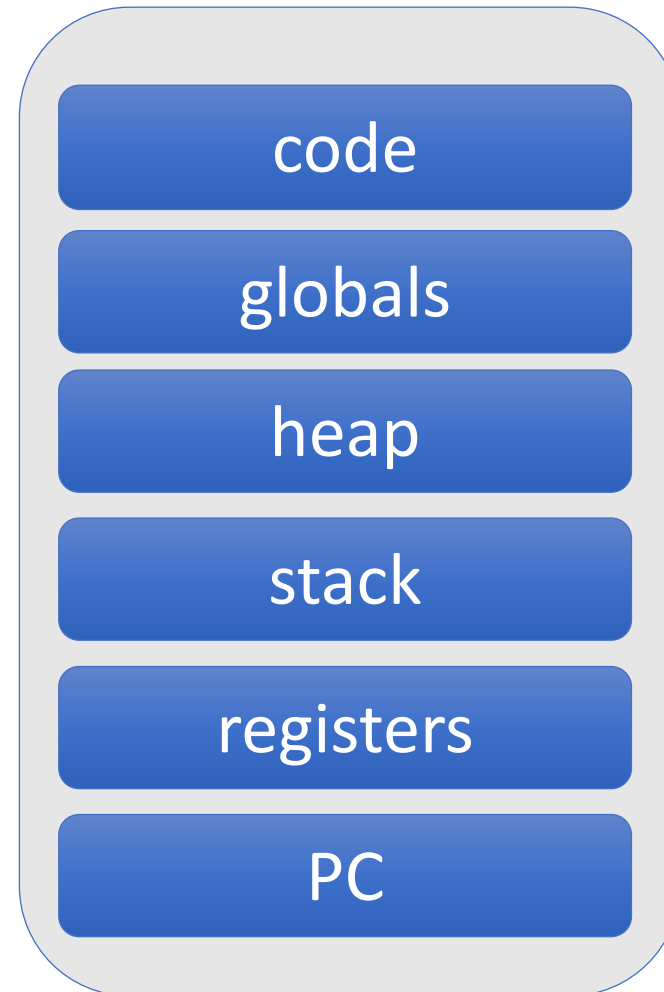


# Key Concepts for Today

- Interprocess communication
- Message passing
- Remote procedure call (RPC)

# So far

- One program  
= one process
- Examples:
  - Shell
  - Compiler
  - ...





# This is not always the case

- One program  
= multiple processes
- Example:
  - Web server



# (Very Simple) Web Server

```
WebServerProcess {  
    forever {  
        wait for an incoming request  
        read file from disk  
        send file back in response  
    }  
}
```

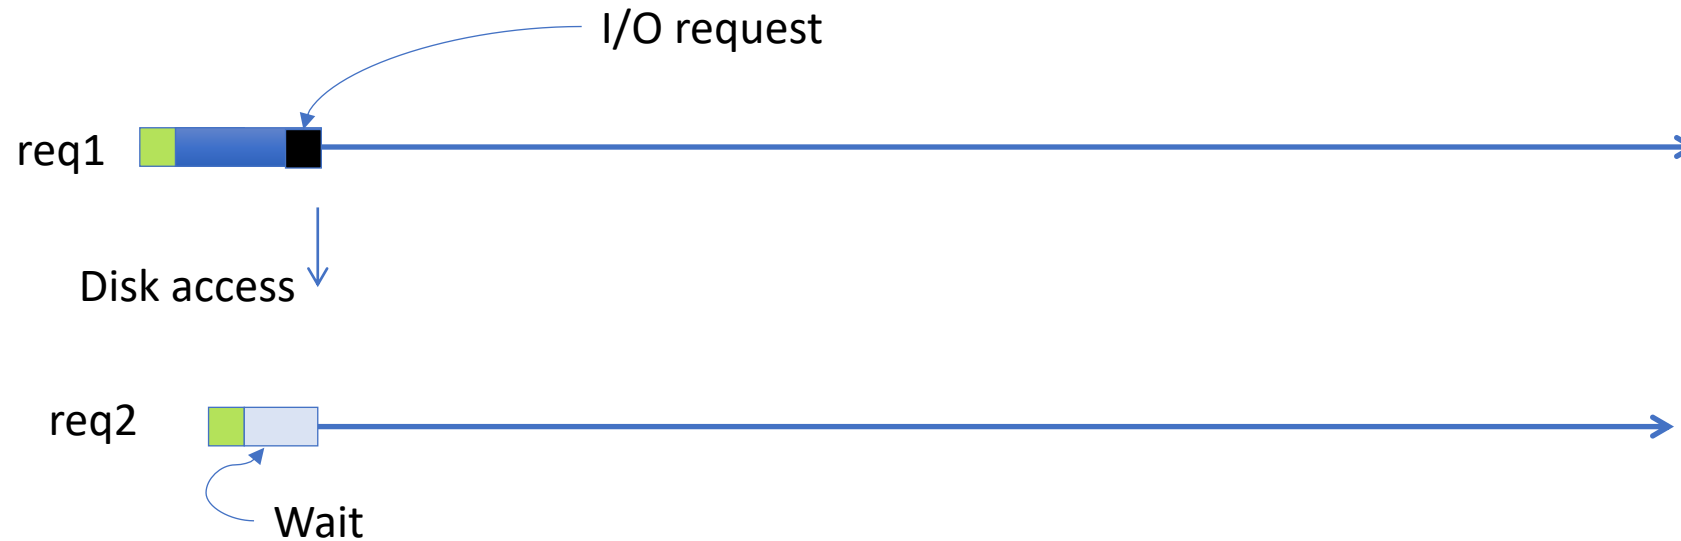
# Single-Process Web Server

Example: Web server receives two requests in quick succession




# Single-Process Web Server

Example: Web server receives two requests in quick succession




# Single-Process Web Server

Example: Web server receives two requests in quick succession



# Single-Process Web Server

Example: Web server receives two requests in quick succession




# Multiprocess Web Server

```
ListenerProcess {  
    forever {  
        wait for incoming request  
        CreateProcess( worker, request )  
    }  
}
```

```
WorkerProcess(request) {  
    read file from disk  
    send response  
    exit  
}
```

# Multi vs. Single-process Web Server


Example: Web server receives two requests in quick succession






# Multi vs. Single-process Web Server

Example: Web server receives two requests in quick succession




# Multi vs. Single-process Web Server

Example: Web server receives two requests in quick succession




# Multi vs. Single-process Web Server

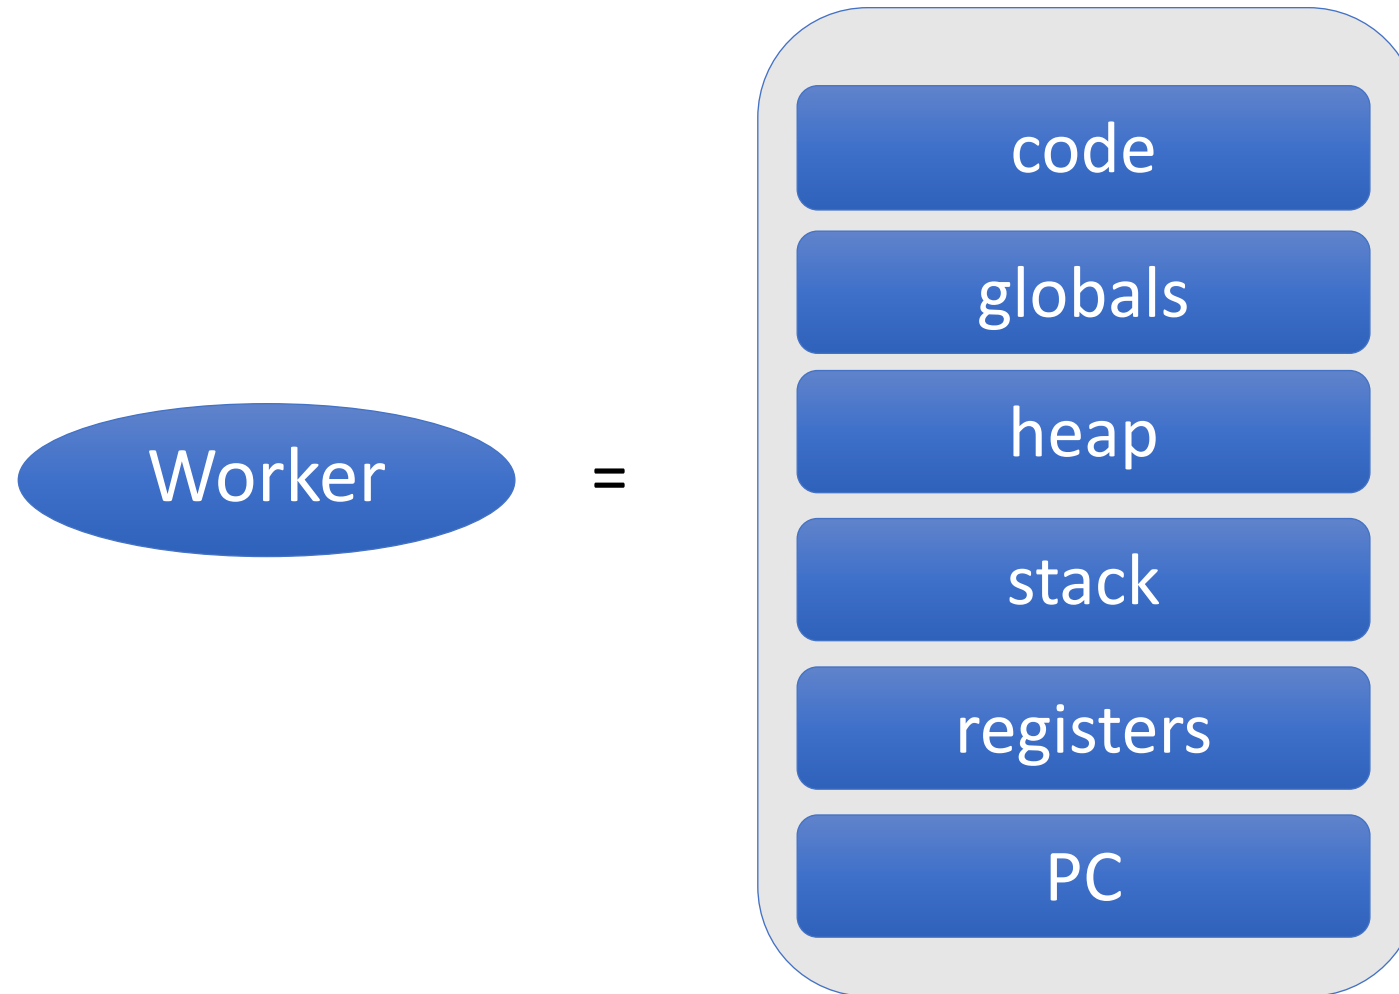
Example: Web server receives two requests in quick succession



# Multiprocess Web Server



# Each Worker is a Process



# Amount of work on server per request

- Receive network packet
- Run listener process
- Create worker process
- Read file from disk
- Send network packet

# Amount of work on server per request

- Receive network packet
- Run listener process
- *Create worker process is expensive*
- Read file from disk
- Send network packet

# Multiprocess Web Server

```
ListenerProcess {  
    forever {  
        wait for incoming request  
        CreateProcess( worker, request )  
    }  
}
```



How can we avoid this?

```
WorkerProcess(request) {  
    read file from disk  
    send response  
    exit  
}
```



# Process Pool


- Create worker processes during initialization
- Hand incoming request to them

# Multiprocess Web Server with Process Pool

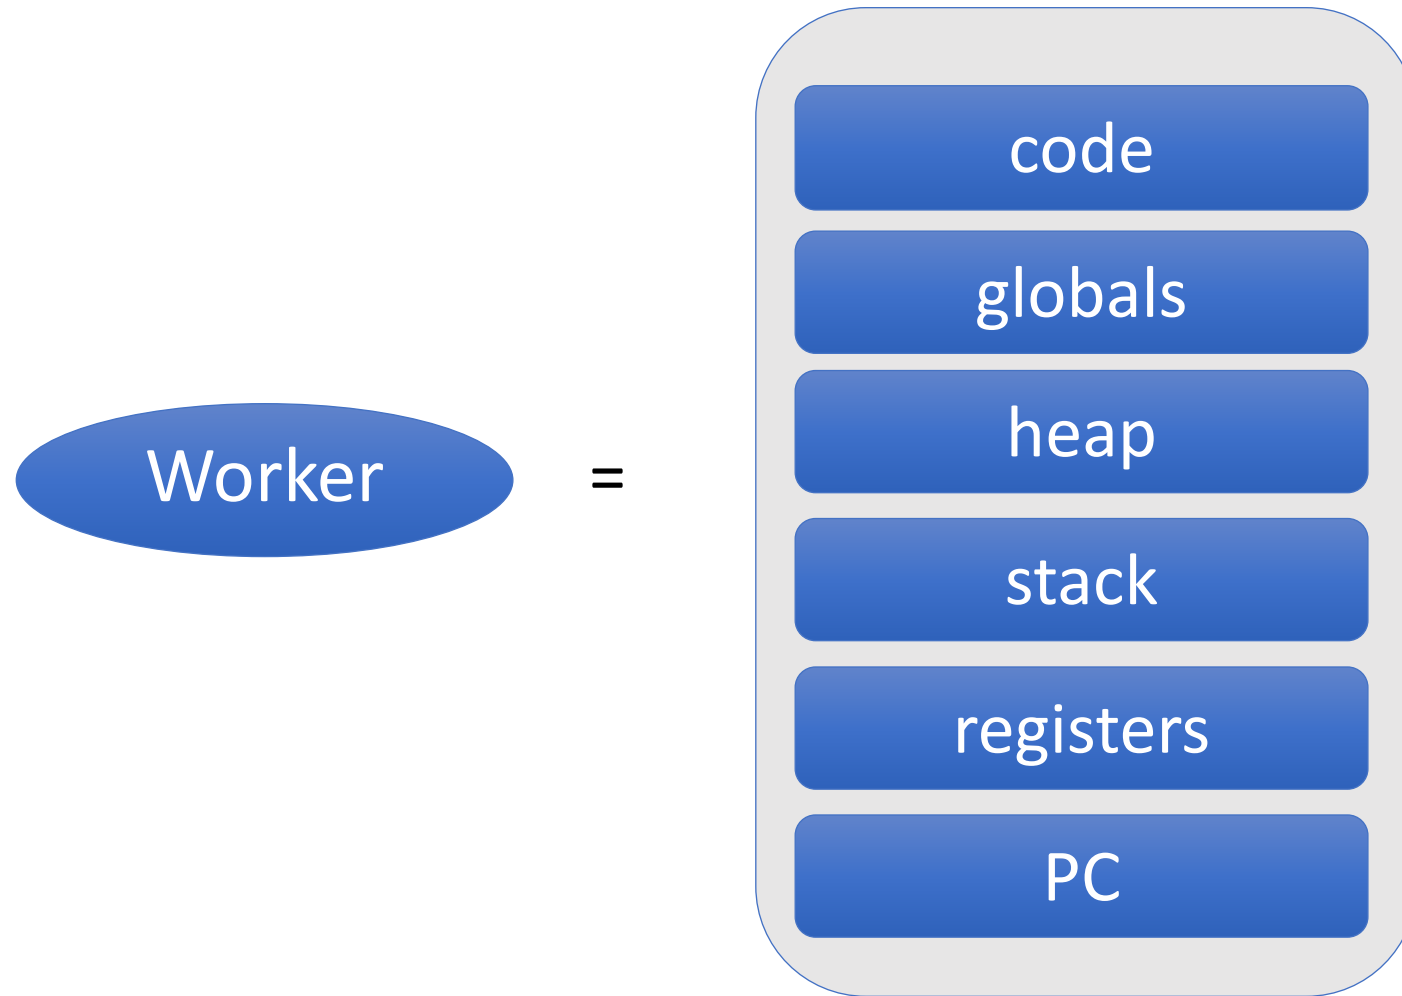
```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        wait for incoming request  
        send(request, process[?])  
    }  
}
```

```
WorkerProcess[?] {  
    forever {  
        wait for message(&request)  
        read file from disk  
        send response  
    }  
}
```

# Pictures remain the same



# Pictures remain the same



# What changed:

## Amount of work on server per request

- Receive network packet
- Run listener process
- *Send message to worker process (cheaper)*
- Read file from disk
- Send network packet

# Interprocess Communication (IPC)

# Interprocess Communication (IPC)

- OS support to allow the processes to manage shared data
  - Through message passing
  - Through remote procedure calls (RPC)

# Where do you need IPC?



# Multiprocess Web Server with Process Pool

```
ListenerProcess {  
    for( i=0; i<MAX_PROCESSES; i++ ) process[i] = CreateProcess(worker)  
    forever {  
        receive incoming request  
        send( request, process[?] )  
    }  
}
```

```
WorkerProcess[?] {  
    forever {  
        wait for message( &request )  
        read file from disk  
        send response  
    }  
}
```

Need IPC here

**For client-server communication**

# Multiprocess Web Server with Process Pool

```
ListenerProcess {  
    for( i=0; i<MAX_PROCESSES; i++ ) process[i] = CreateProcess(worker)  
    forever {  
        receive incoming request  
        send( request, process[?] )  
    }  
}
```

```
WorkerProcess[?] {  
    forever {  
        wait for message( &request )  
        read file from disk  
        send response  
    }  
}
```

Need IPC here

For communication between  
cooperating processes  
(e.g., between listener and workers)

# Where do you need IPC?

- Between client and server
- Between cooperating processes

# How to do IPC?

- Message passing
- Remote procedure calls (RPC)

# Message Passing Primitives

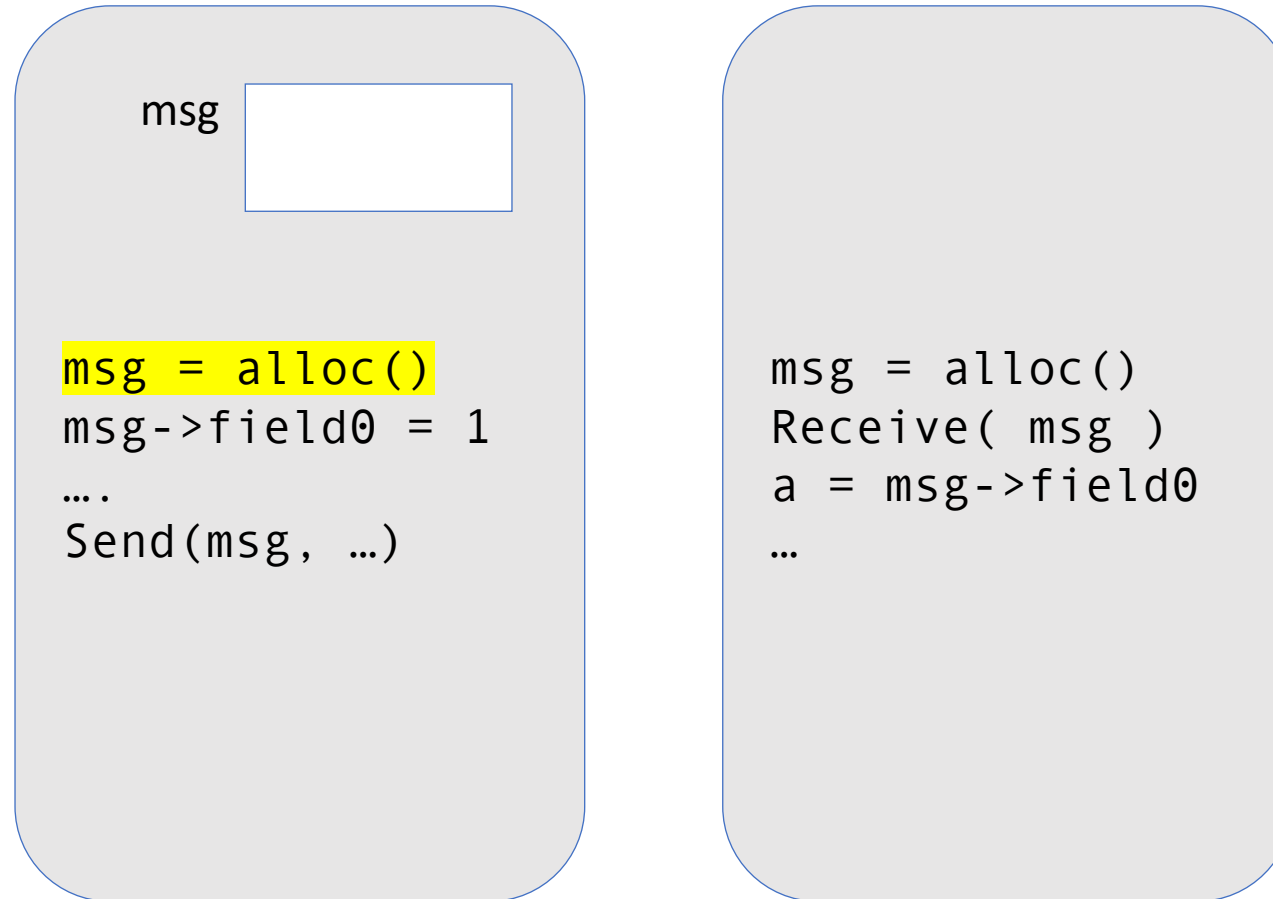
- Send message
- Receive message

# Message Passing Send / Receive

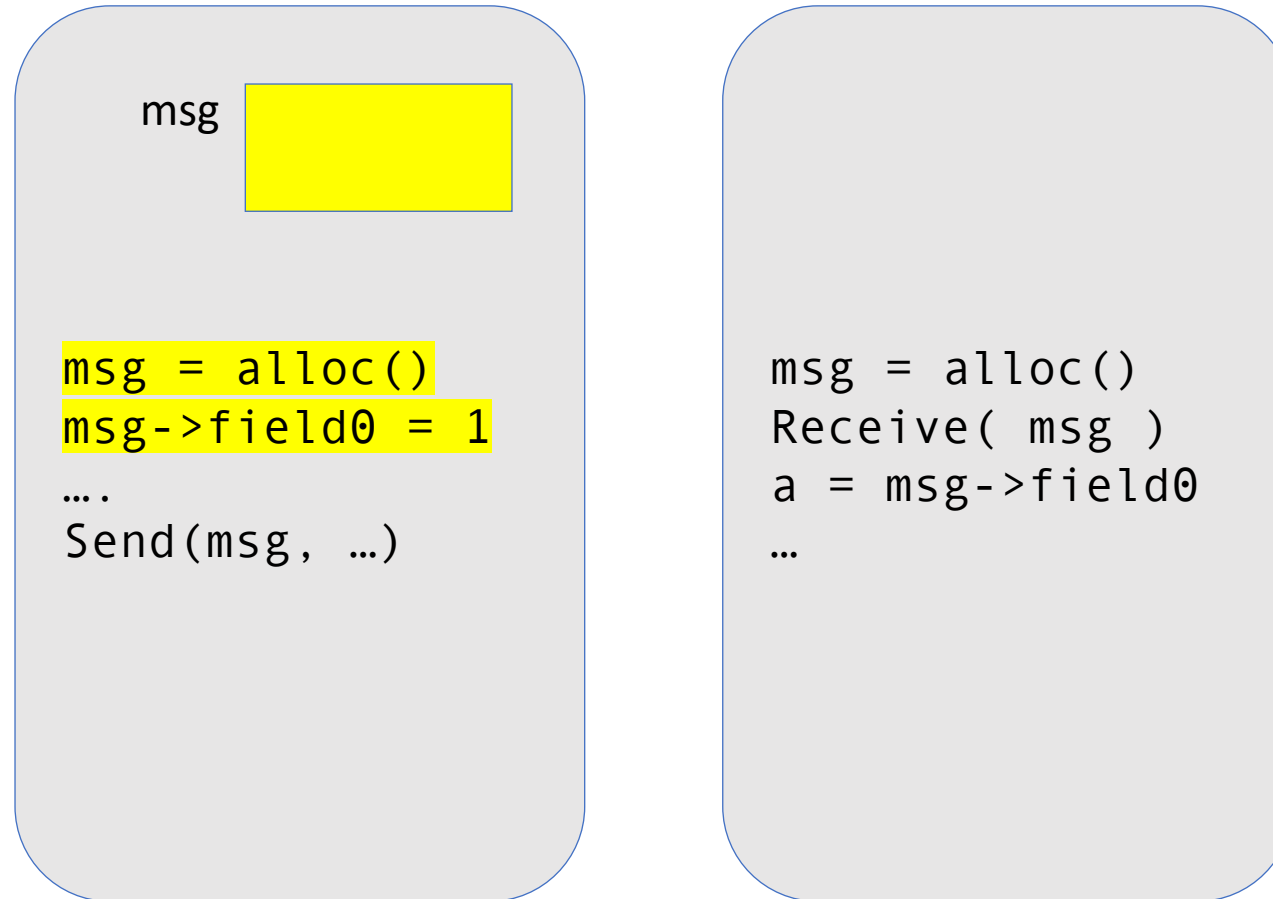
```
msg = alloc()  
msg->field0 = 1  
...  
Send(msg, ...)
```

```
msg = alloc()  
Receive( msg )  
a = msg->field0  
...
```

# Message Passing Send / Receive

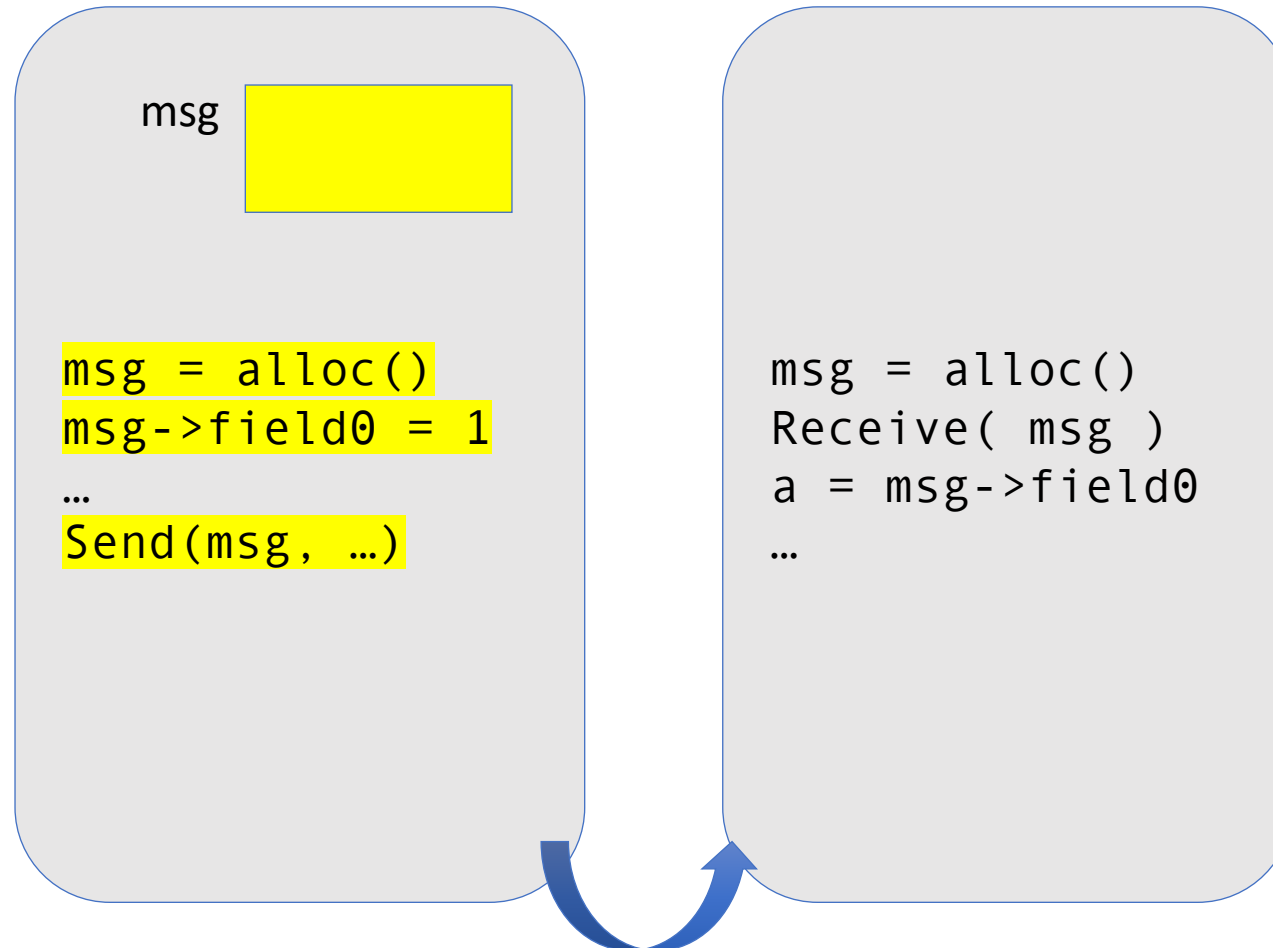


# Message Passing Send / Receive

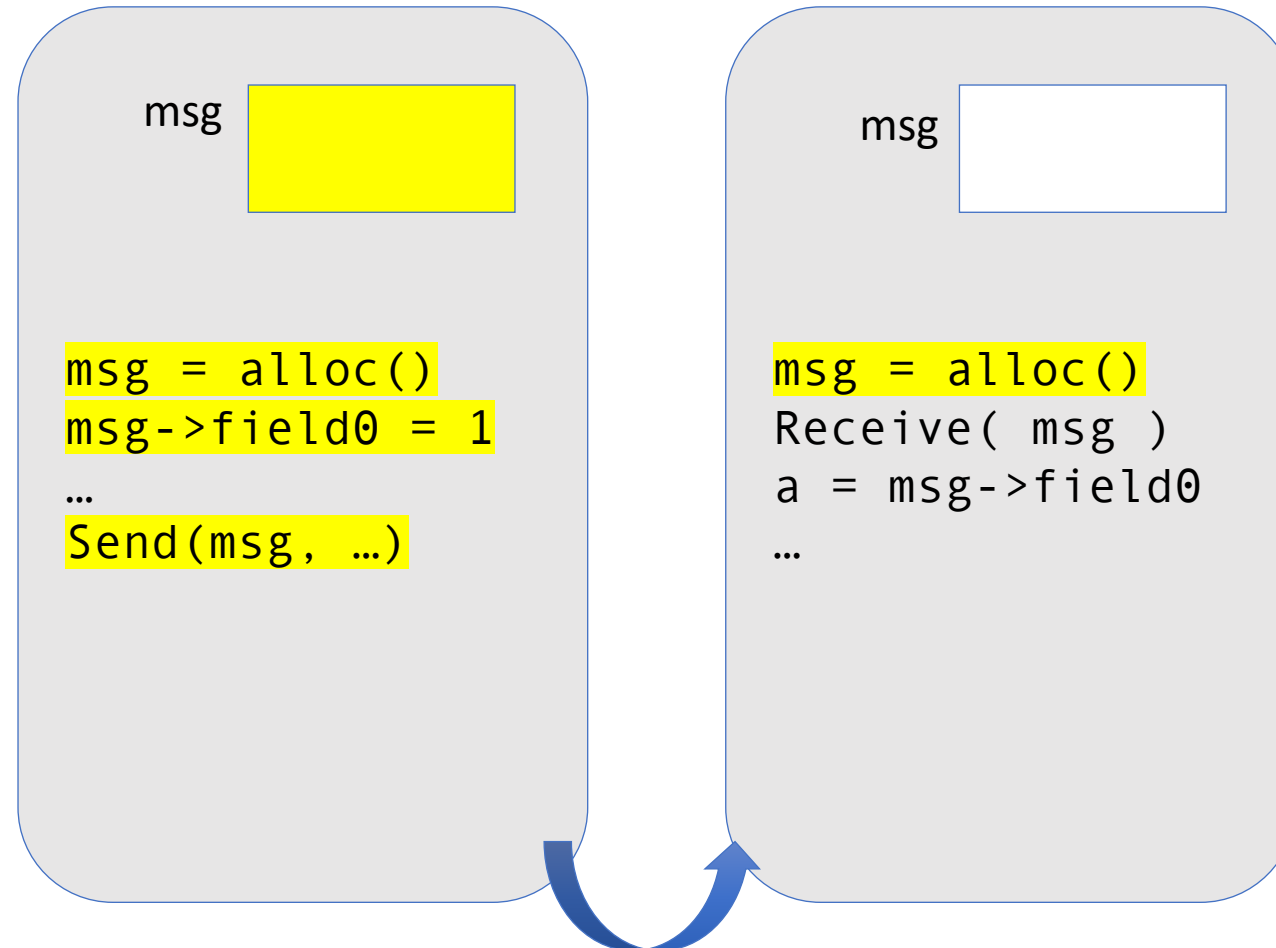




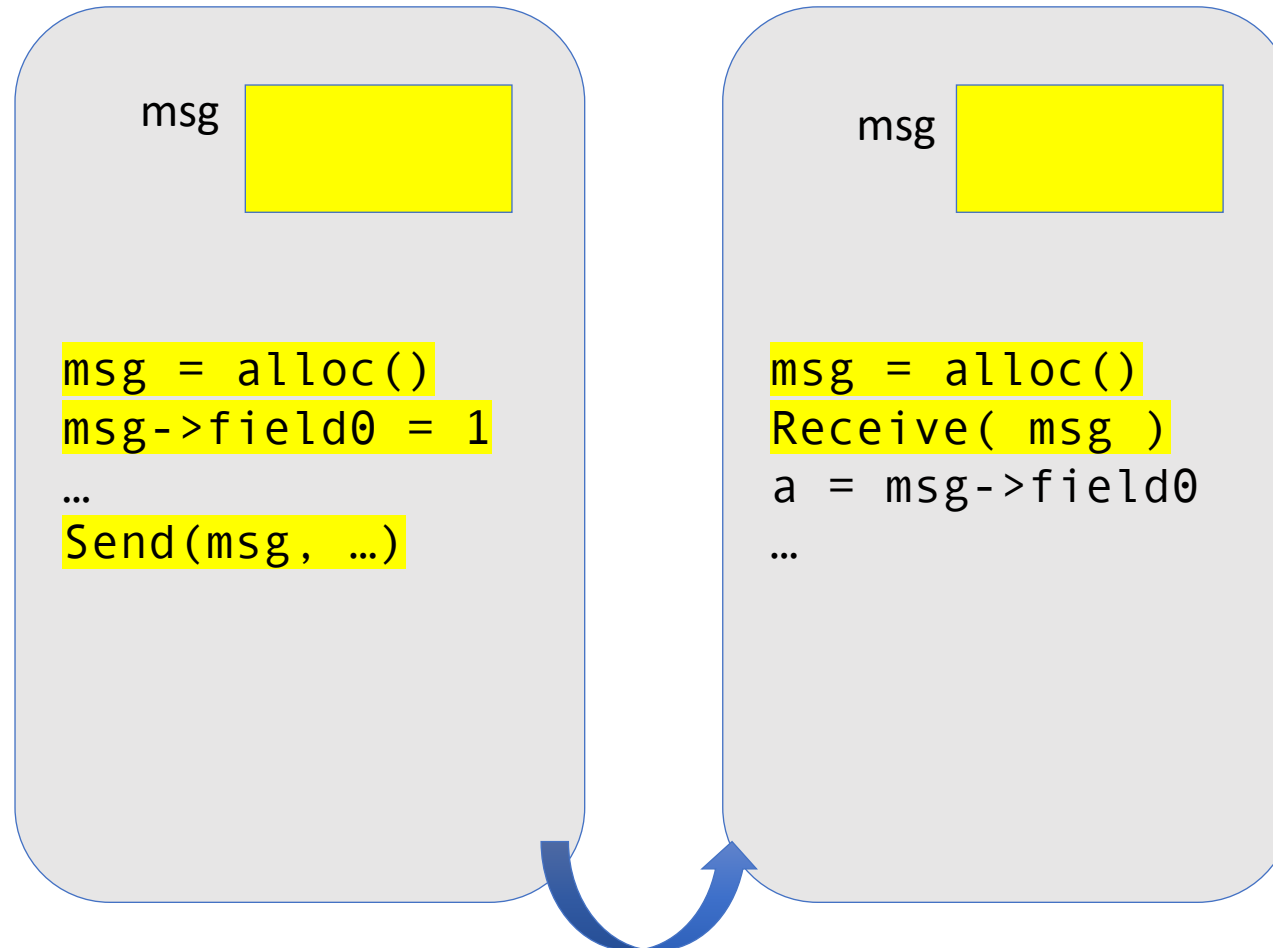
# Message Passing Send / Receive



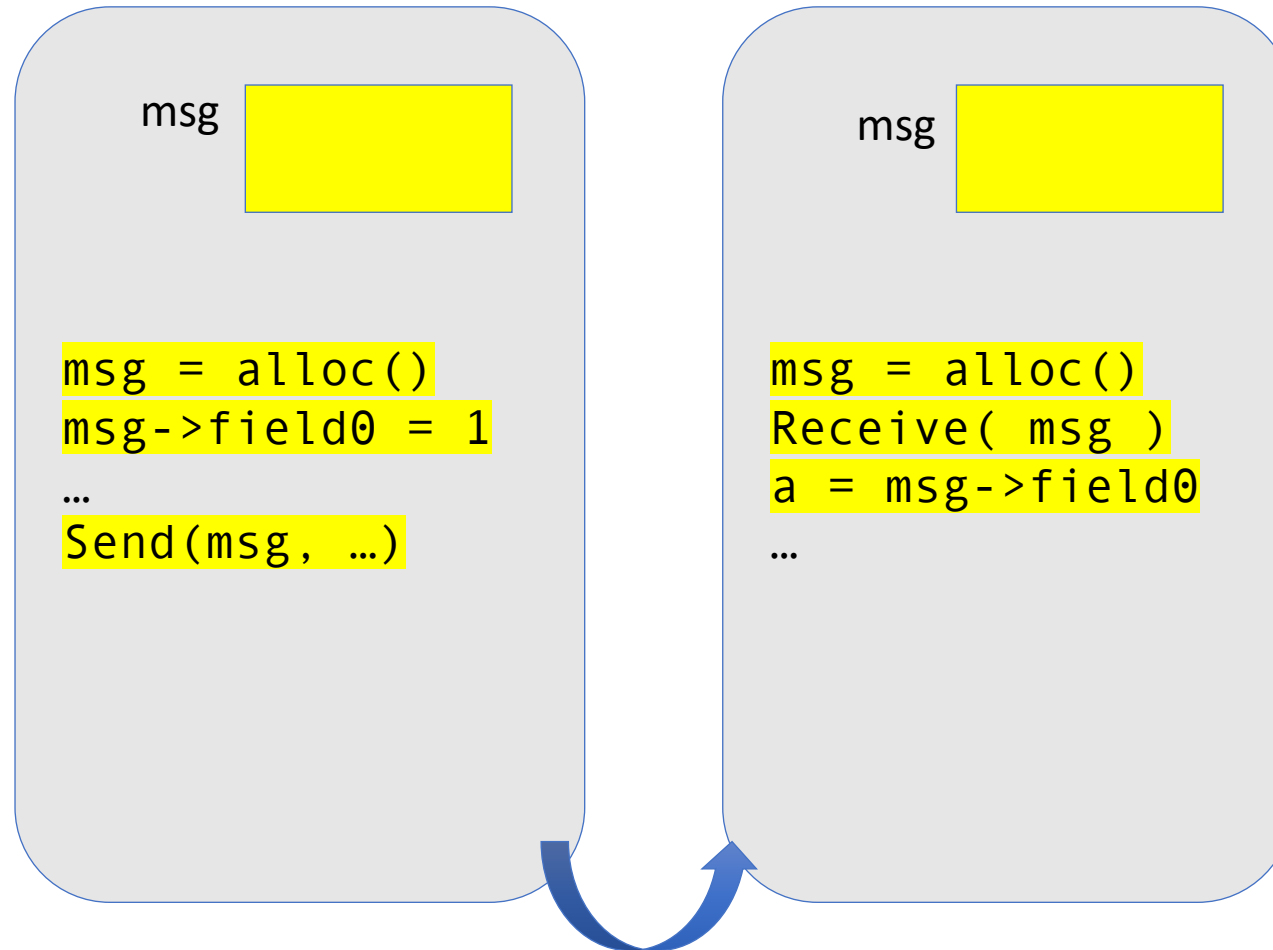
# Message Passing Send / Receive



# Message Passing Send / Receive



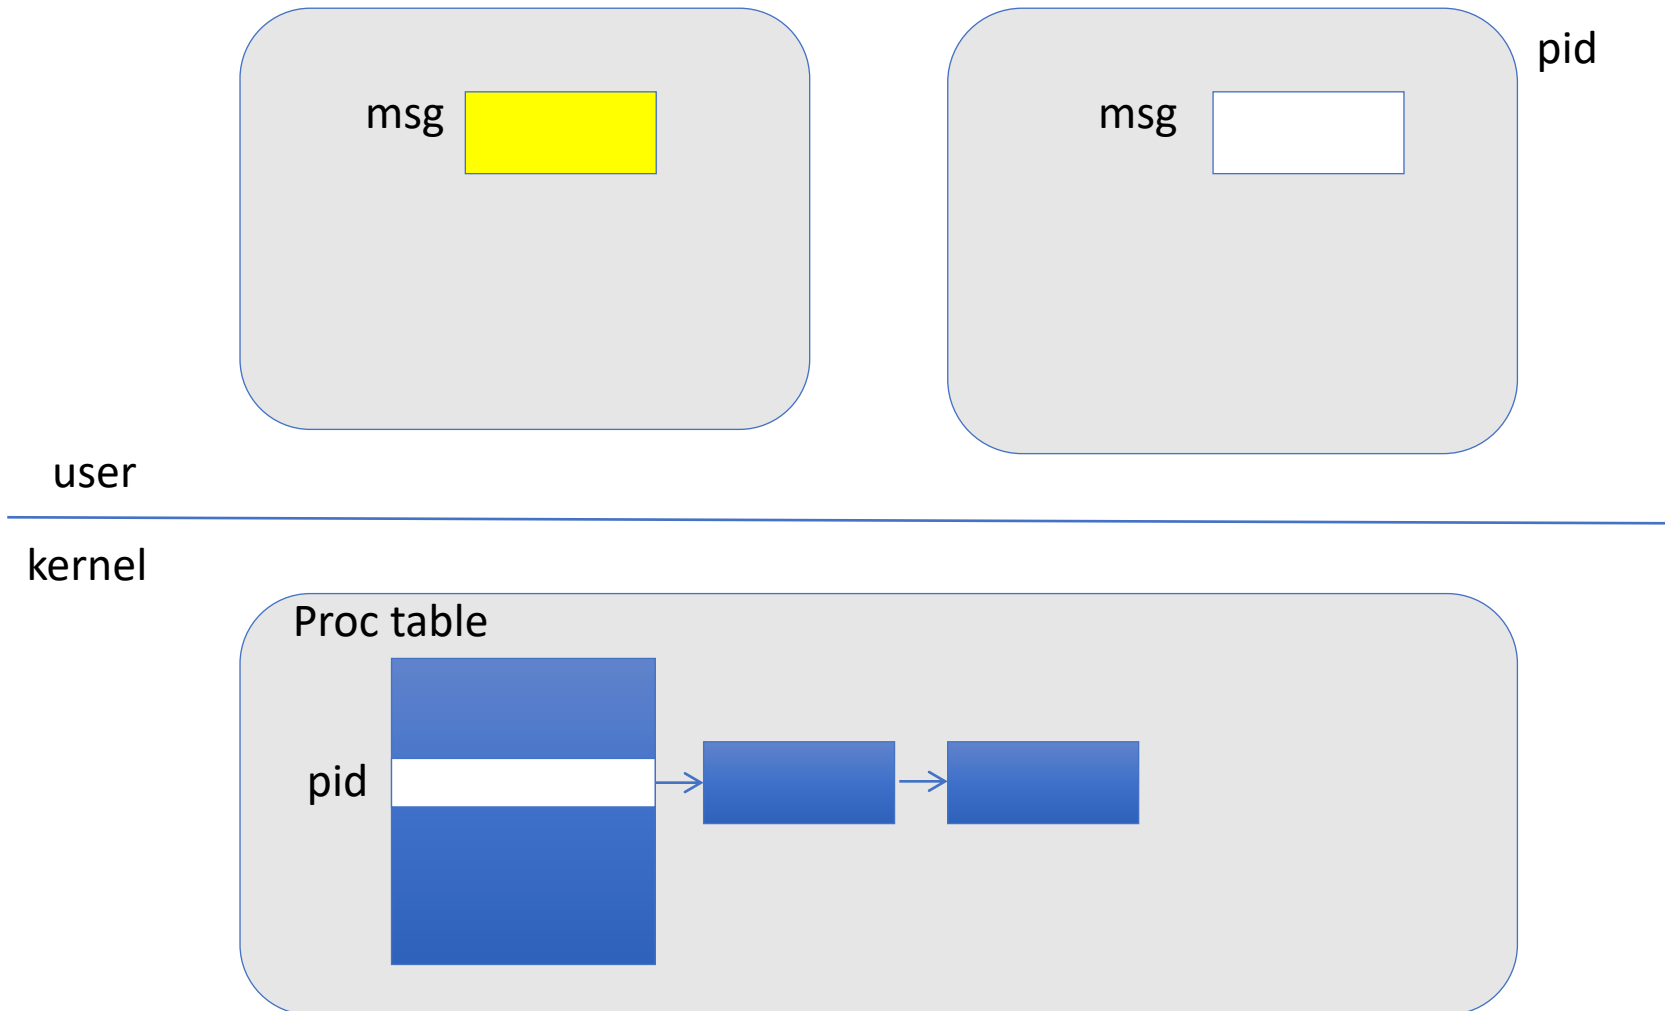
# Message Passing Send / Receive



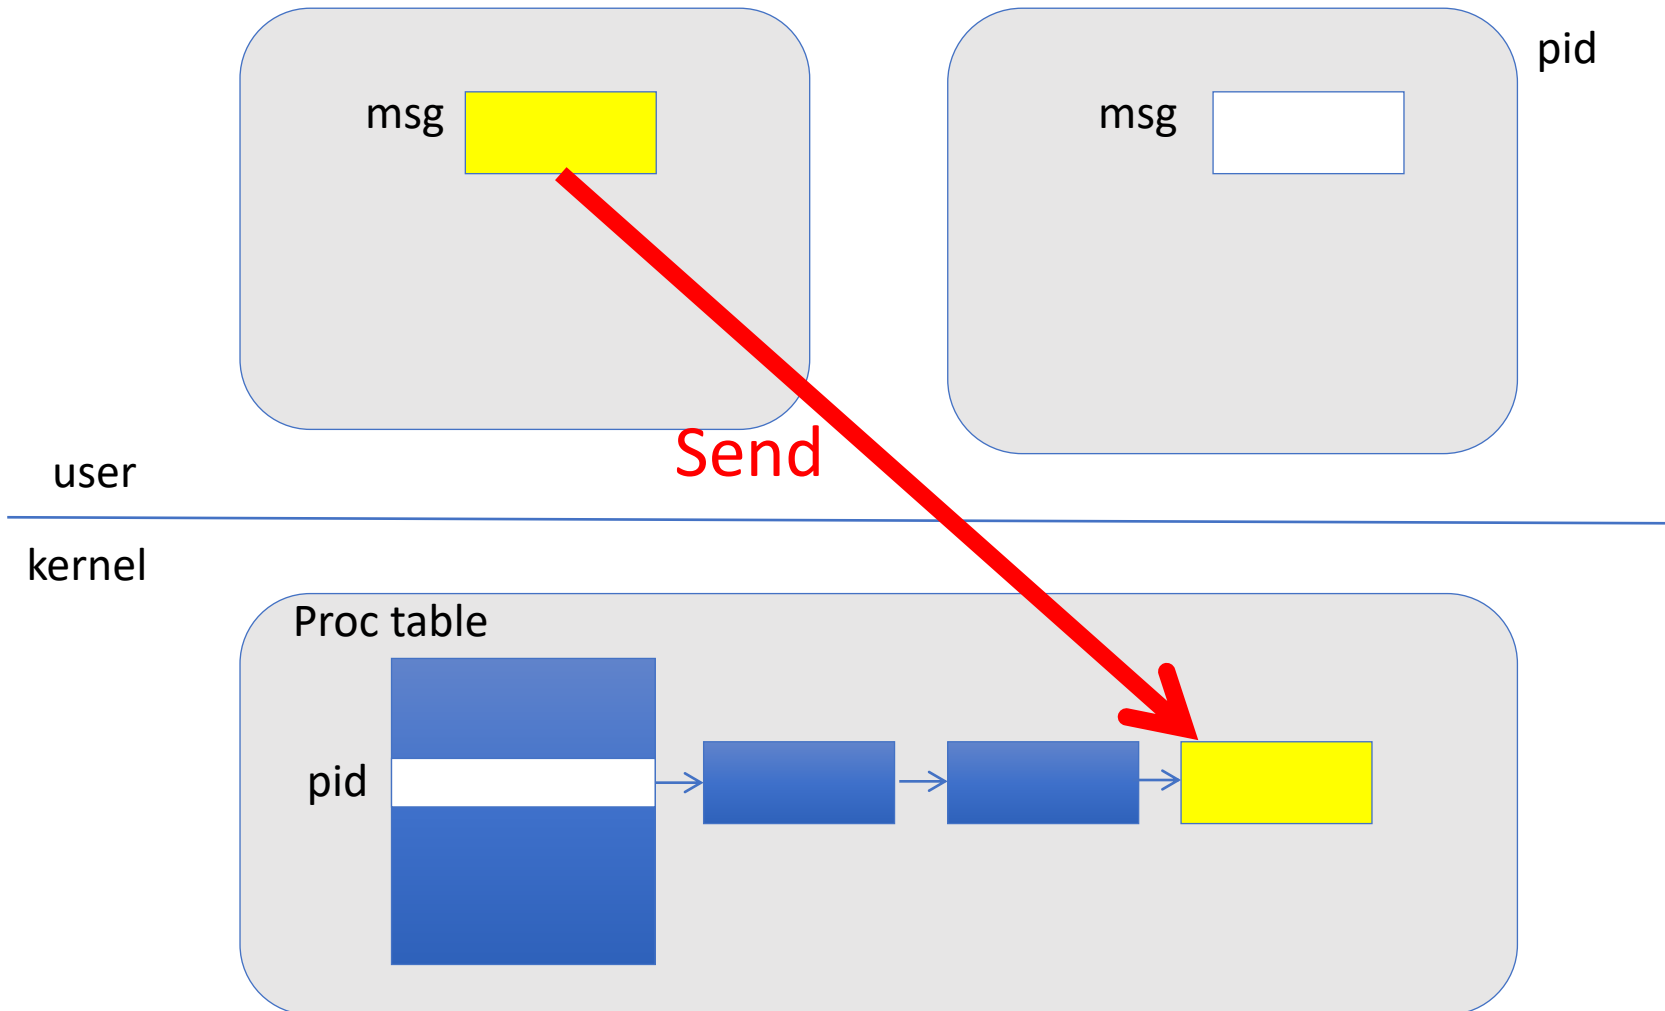
# Message Passing

- By value communication
- **Never by reference**
- Receiver cannot affect message in sender

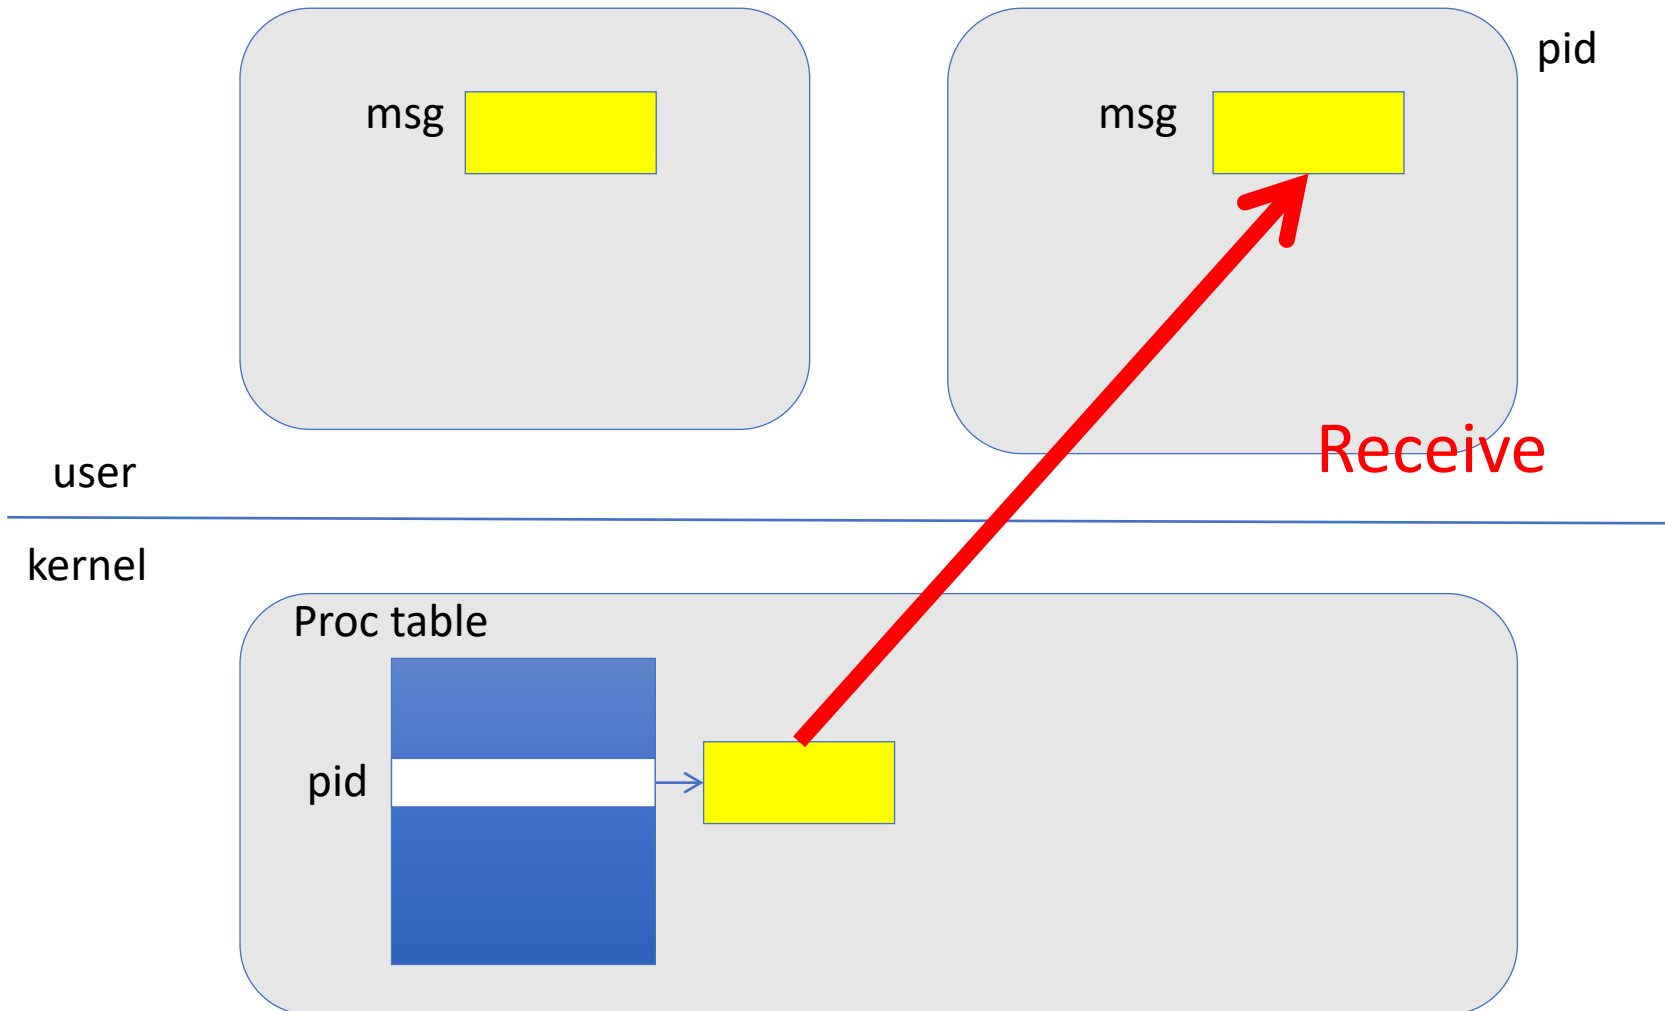
# Message Passing Implementation



# Message Passing Implementation



# Message Passing Implementation





# Message Passing Alternatives

- Symmetric / asymmetric addressing
- Blocking / nonblocking

# Symmetric Addressing

- `Send(msg, to_pid)`
- `Receive(msg, from_pid)`
- Message is (typically) a struct
- `to_pid`, `from_pid` are process identifiers
- Symmetric addressing seldom used

# Asymmetric Addressing

- **Send(msg, pid)**
  - Send msg to process pid
- **pid = Receive(msg)**
  - Receive msg from *any* process
  - Return the pid of sending process
- More common and useful form of addressing

# Blocking or Nonblocking Send

- Nonblocking:
  - Send returns immediately after message is sent
- Blocking
  - Sender blocks until message is delivered
- Nonblocking is the more common form

# Blocking or Nonblocking Receive

- Nonblocking
  - Receive returns immediately
  - Regardless of message present or not
- Blocking
  - Receive blocks until message is present
- Blocking is the more common form

# (Slightly Rewritten) Example: Multiprocess Web Server with Process Pool

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg)  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg)  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```

# Asymmetric Addressing: Send

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg)  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg)  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```

# Asymmetric Addressing: Receive

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg) /* receive msg from any client */  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg) /* receive msg' from listener; could be symmetric */  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```



# Blocking Receive

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg) /* nothing else to do*/  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg) /* nothing else to do*/  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```

# Nonblocking Send

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg)  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i]) /* must not block */  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg)  
        read file from disk  
        send(resp, client_pid) /* must not block */  
    }  
}
```

# Client-Server Communication

# (Server Side) Client-Server Communication

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        receive incoming request  
        send( request, process[?] )  
    }  
}
```

```
WorkerProcess[?] {  
    forever {  
        wait for message( &request )  
        read file from disk  
        send response  
    }  
}
```

# (Client-Side) Client-Server Communication

`send(msg to server)`

`receive(reply msg from server)`

# A Very Common Pattern

- Client:

- Send `/* send request to server */`
- Blocking receive `/* wait for reply */`

- Server

- Blocking receive `/* wait for request */`
- Send `/* send reply */`


# This looks like ...

- Client:
  - Send
  - Blocking receive
- Server
  - Blocking receive
  - Send

calling site  
call procedure  
return  
callee site  
invoke procedure  
return

# Remote Procedure Call (RPC)

- Client:
  - Send
  - Blocking receive
- Server
  - Blocking receive
  - Send



RPC: when client wants to call a function  
that belongs to server code



# RPC Interface

- Interface
  - List of remotely callable procedures
  - With their arguments and return values
- Example: file system interface
  - `Open(string filename)`
  - returns `int fd`
    - `fd` = file descriptor; will see later in course
  - ...

# RPC Client Code

- Import file system interface
- `fd = open("/a/b/c")`
- `nbytes = read(fd, buffer, size)`

# RPC Server Code

- Export file system interface
- `int open(stringname) { ... }`
- `int read(fd, buffer, nbytes) { ... }`
- ...

# Problem

- Want a procedure call interface
- Have only message passing between processes
- How to bridge the gap?

# Solution: Stub Library

- Client stub and server stub
- Client stub linked with client process
- Server stub linked with server process

# Two Message Types

- Call message
  - From client to server
  - Contains arguments
- Return message
  - From server to client
  - Contains return values

# Client Stub

- Sends arguments in call message
- Receives return values in return message

# Server Stub

- Receives arguments in call message
- Invokes procedure
- Sends return values in return message



# RPC Implementation

client  
process


client  
code

server  
process


server  
code

# Client and Server Stubs


client  
process




server  
process




# RPC Implementation: Call



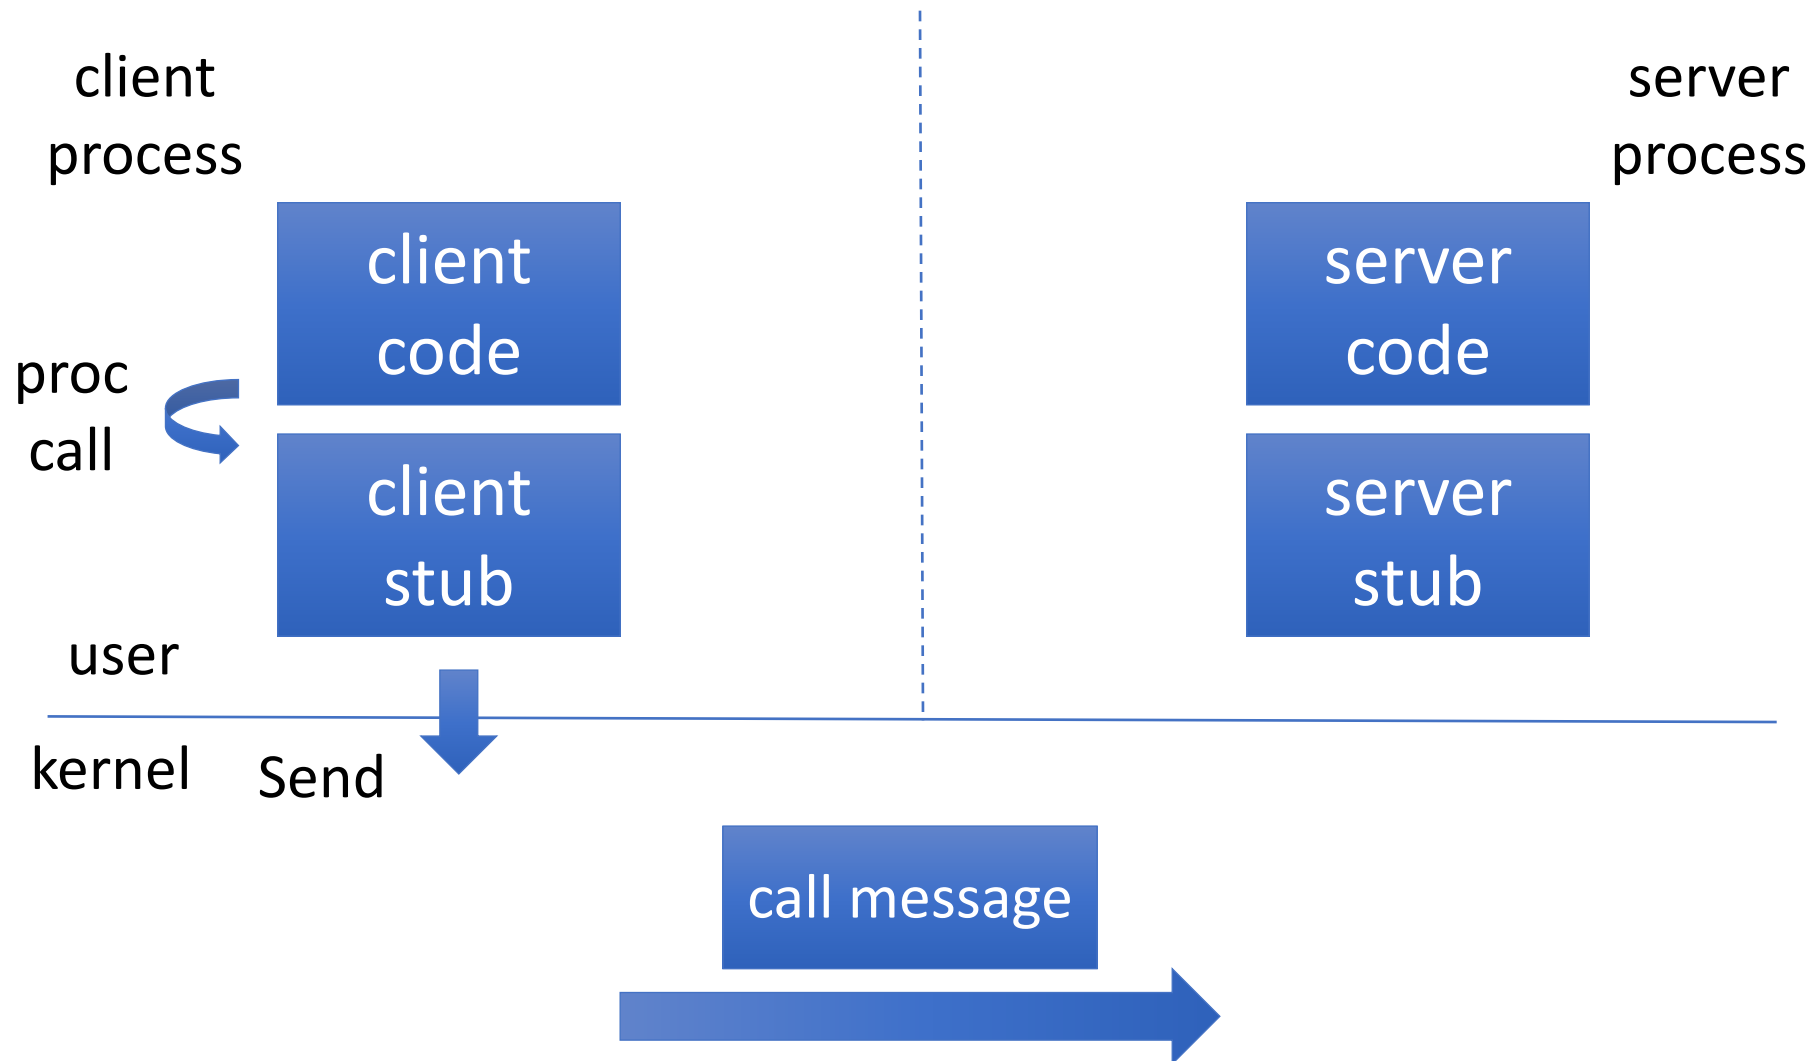
# RPC Implementation: Call



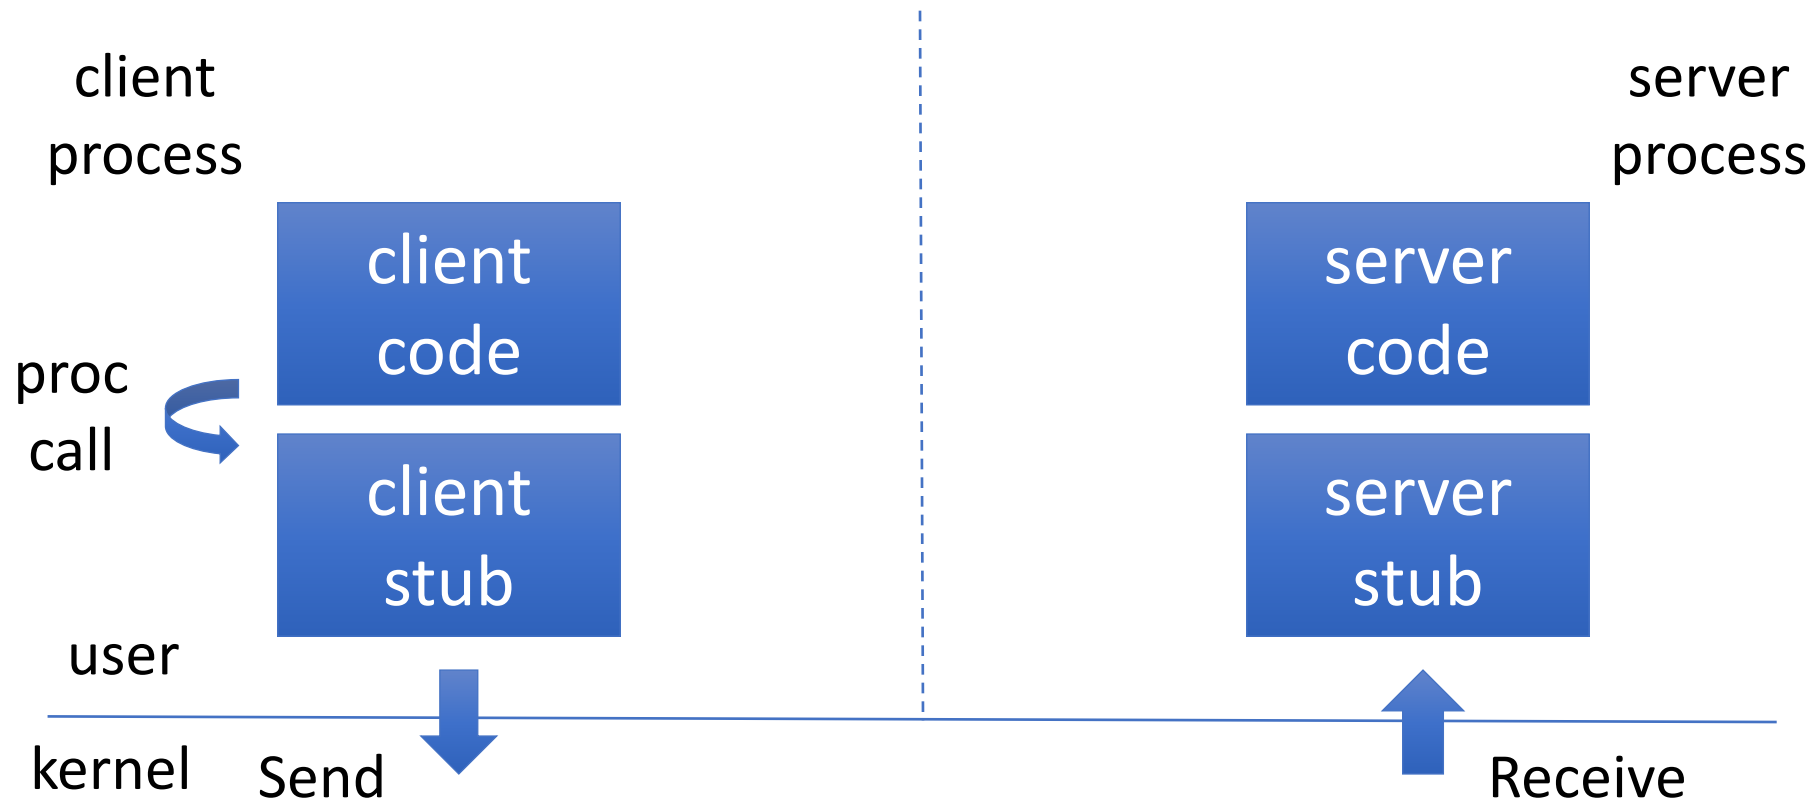
# RPC Implementation: Call




# RPC Implementation: Call



# RPC Implementation: Call




# RPC Implementation: Call







# RPC Implementation: Return



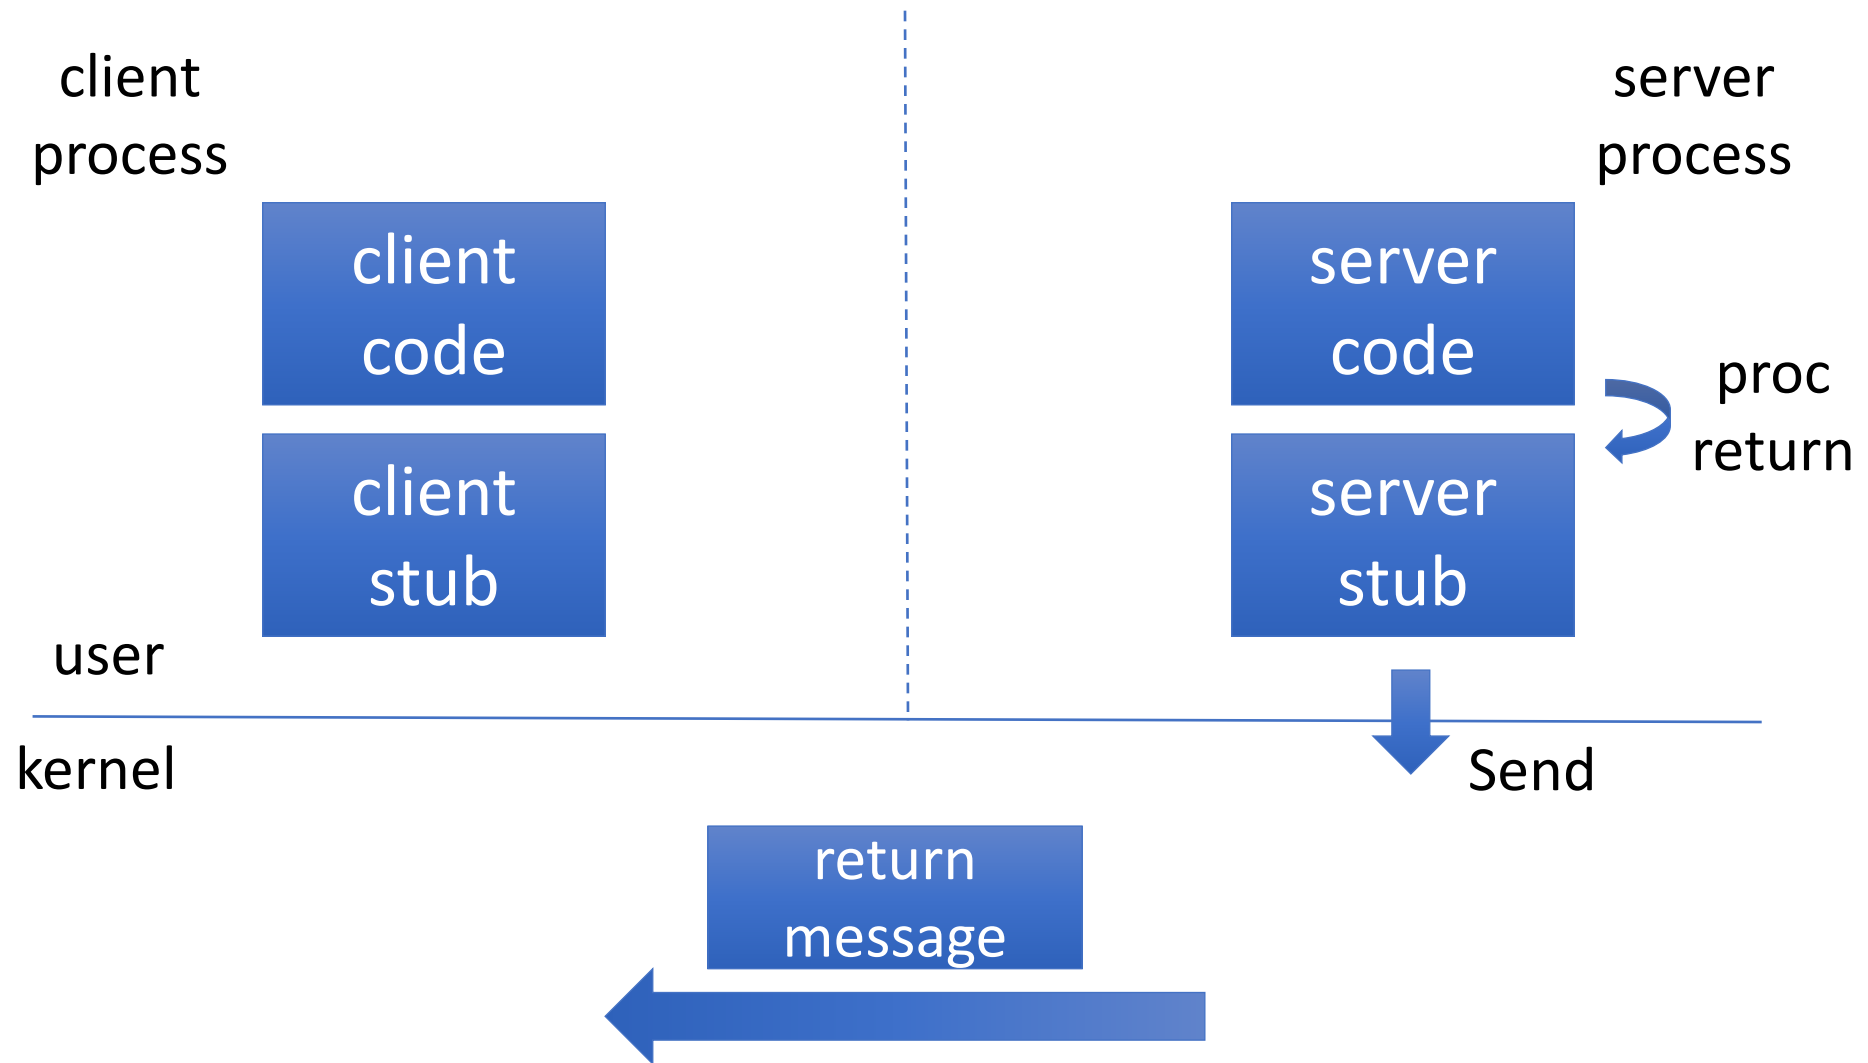
# RPC Implementation: Return



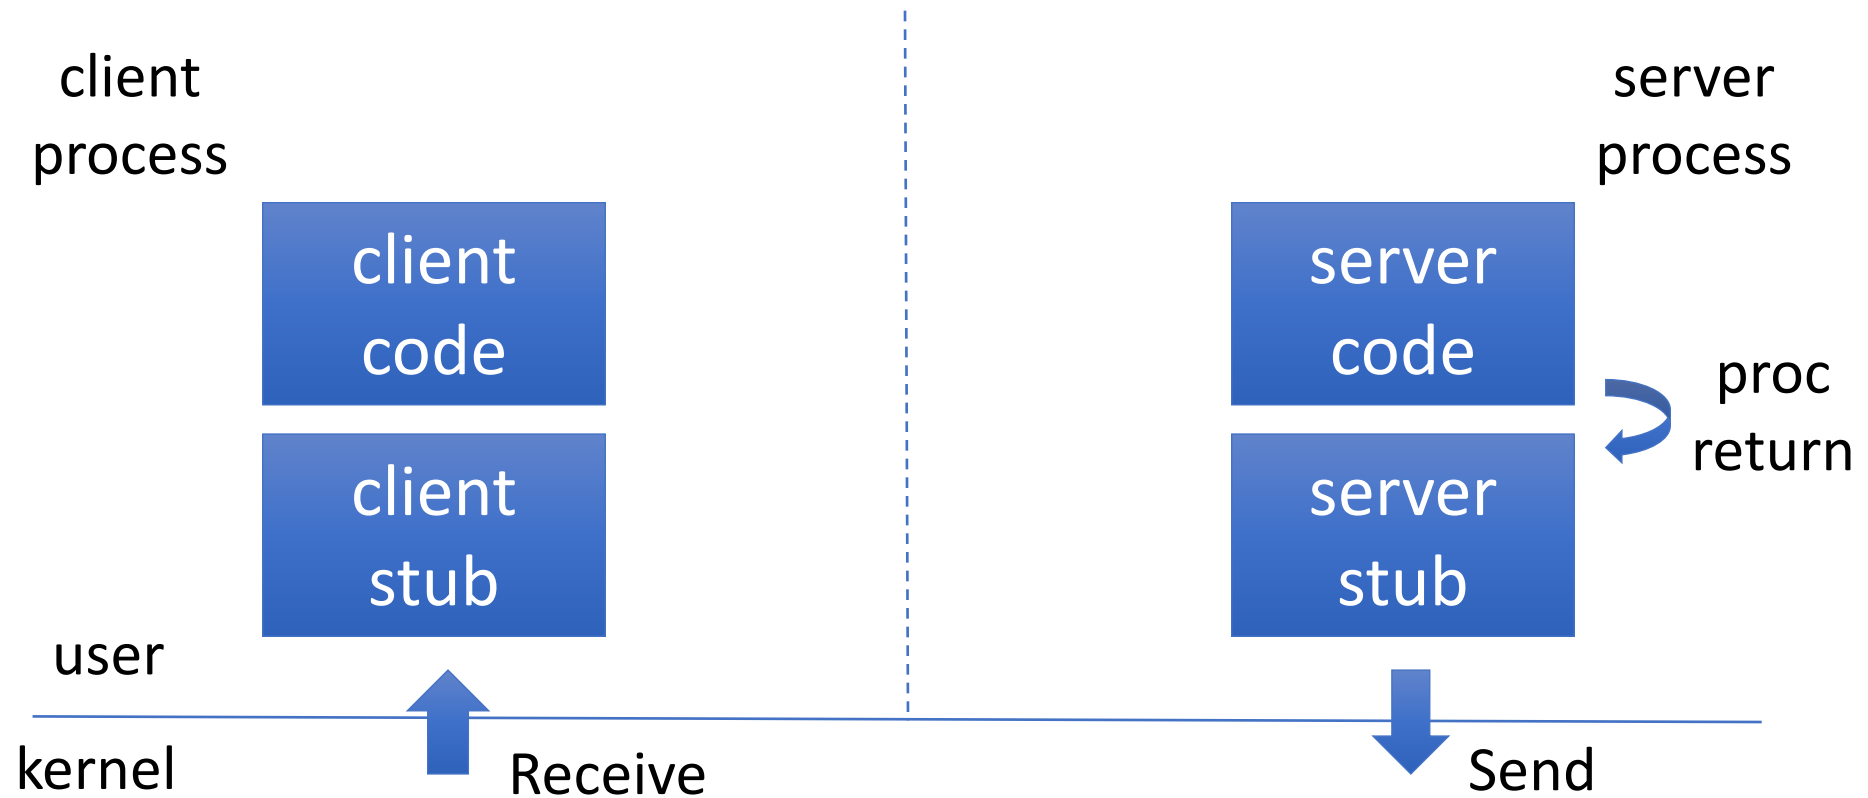
# RPC Implementation: Return




# RPC Implementation: Return



# RPC Implementation: Return



# RPC Implementation: Return



# Further Optional Reading

## **Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau**

Chapters 25 – 32 (inclusive) <https://pages.cs.wisc.edu/~remzi/OSTEP/>

### **Credits:**

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Maurice Herlihy (Brown University)