

Week 4

Multiprocess Communication

Oana Balmau
January 26, 2023

Message Passing Alternatives

- Symmetric / asymmetric addressing
- Blocking / nonblocking

Symmetric Addressing

- `Send(msg, to_pid)`
- `Receive(msg, from_pid)`
- Message is (typically) a struct
- `to_pid`, `from_pid` are process identifiers
- Symmetric addressing seldom used

Asymmetric Addressing

- **Send(msg, pid)**
 - Send msg to process pid
- **pid = Receive(msg)**
 - Receive msg from *any* process
 - Return the pid of sending process
- More common and useful form of addressing

Blocking or Nonblocking Send

- Nonblocking:
 - Send returns immediately after message is sent
- Blocking
 - Sender blocks until message is delivered
- Nonblocking is the more common form

Blocking or Nonblocking Receive

- Nonblocking
 - Receive returns immediately
 - Regardless of message present or not
- Blocking
 - Receive blocks until message is present
- Blocking is the more common form

(Slightly Rewritten) Example: Multiprocess Web Server with Process Pool

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg)  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg)  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```

Asymmetric Addressing: Send

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg)  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg)  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```


Asymmetric Addressing: Receive

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg) /* receive msg from any client */  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg) /* receive msg' from listener; could be symmetric */  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```

Blocking Receive

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg) /* nothing else to do*/  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i])  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg) /* nothing else to do*/  
        read file from disk  
        send(resp, client_pid)  
    }  
}
```

Nonblocking Send

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        client_pid = receive(msg)  
        msg' = slightly modify msg to include client_pid  
        send(msg', worker_process[i]) /* must not block */  
    }  
}
```

```
WorkerProcess[i] {  
    forever {  
        receive(msg)  
        read file from disk  
        send(resp, client_pid) /* must not block */  
    }  
}
```

Client-Server Communication

(Server Side) Client-Server Communication

```
ListenerProcess {  
    for(i=0; i<MAX_PROCESSES; i++) process[i] = CreateProcess(worker)  
    forever {  
        receive incoming request  
        send( request, process[?] )  
    }  
}
```

```
WorkerProcess[?] {  
    forever {  
        wait for message( &request )  
        read file from disk  
        send response  
    }  
}
```

(Client-Side) Client-Server Communication

`send(msg to server)`

`receive(reply msg from server)`

A Very Common Pattern

- Client:

- Send `/* send request to server */`
- Blocking receive `/* wait for reply */`

- Server

- Blocking receive `/* wait for request */`
- Send `/* send reply */`

This looks like ...

- Client:

- Send
- Blocking receive

- Server

- Blocking receive
- Send

calling site

call procedure

return

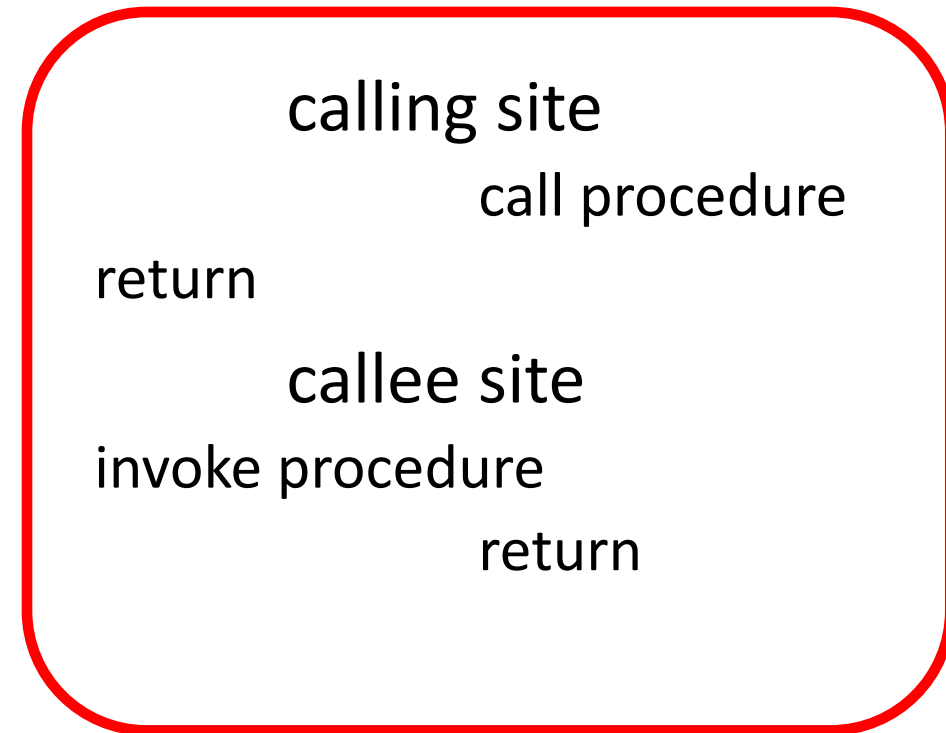
callee site

invoke procedure

return

Remote Procedure Call (RPC)

- Client:
 - Send
 - Blocking receive
- Server
 - Blocking receive
 - Send



RPC: when client wants to call a function
that belongs to server code

RPC Interface

- Interface
 - List of remotely callable procedures
 - With their arguments and return values
- Example: file system interface
 - `Open(string filename)`
 - returns `int fd`
 - `fd` = file descriptor; will see later in course
 - ...

RPC Client Code

- Import file system interface
- `fd = open("/a/b/c")`
- `nbytes = read(fd, buffer, size)`

RPC Server Code

- Export file system interface
- `int open(stringname) { ... }`
- `int read(fd, buffer, nbytes) { ... }`
- ...

Problem

- Want a procedure call interface
- Have only message passing between processes
- How to bridge the gap?

Solution: Stub Library

- Client stub and server stub
- Client stub linked with client process
- Server stub linked with server process

Two Message Types

- Call message
 - From client to server
 - Contains arguments
- Return message
 - From server to client
 - Contains return values

Client Stub

- Sends arguments in call message
- Receives return values in return message

Server Stub

- Receives arguments in call message
- Invokes procedure
- Sends return values in return message

RPC Implementation

client
process


client
code

server
process


server
code

Client and Server Stubs


client
process



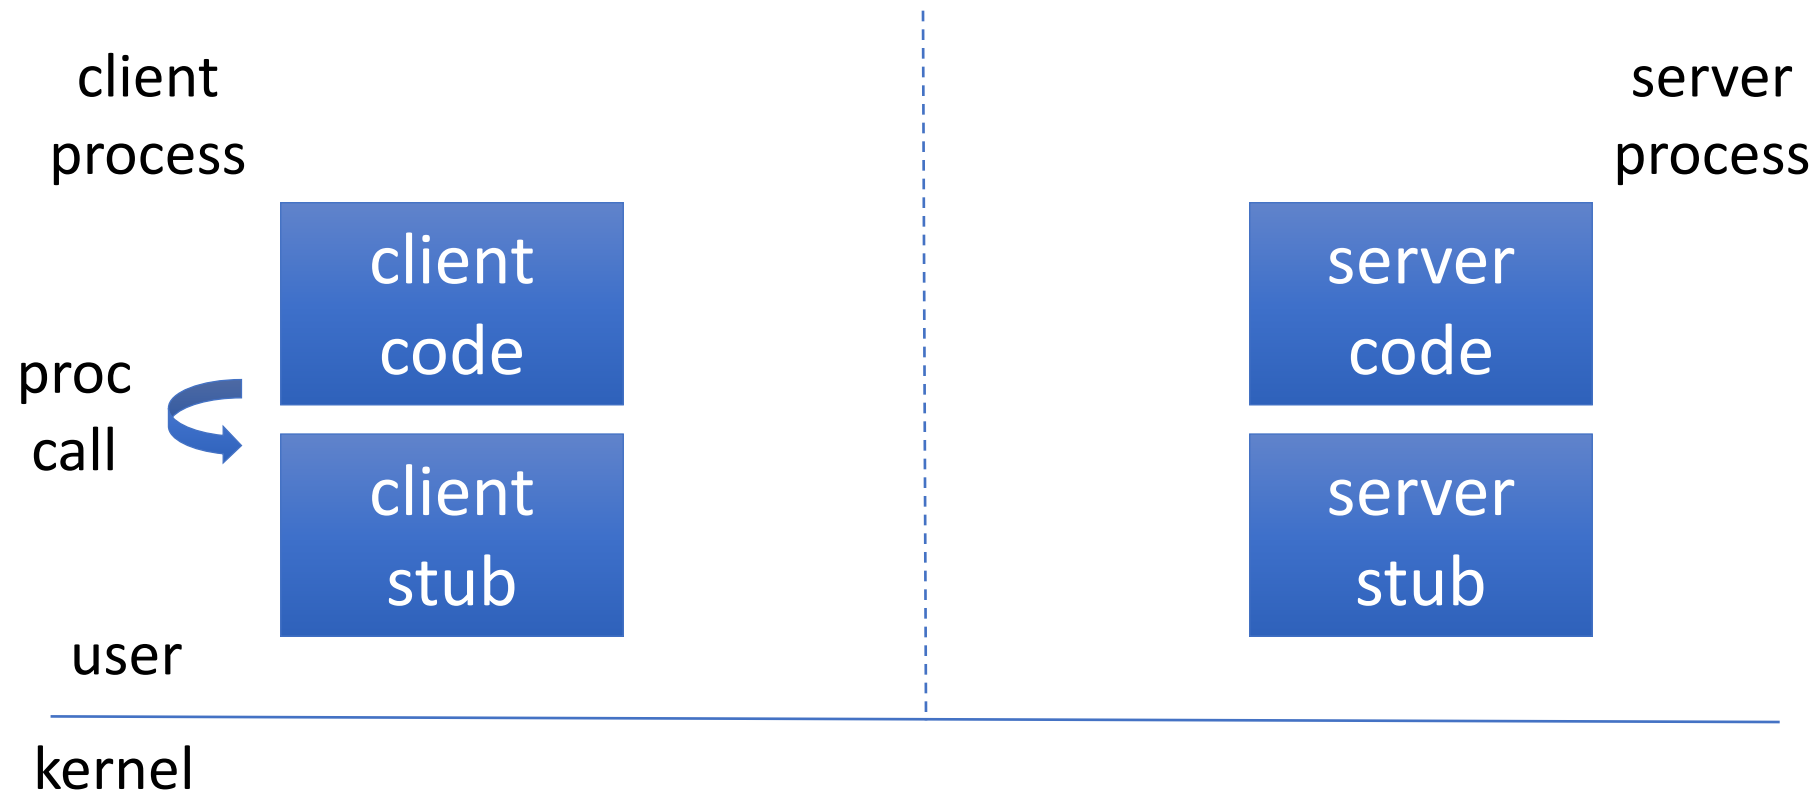
server
process




RPC Implementation: Call



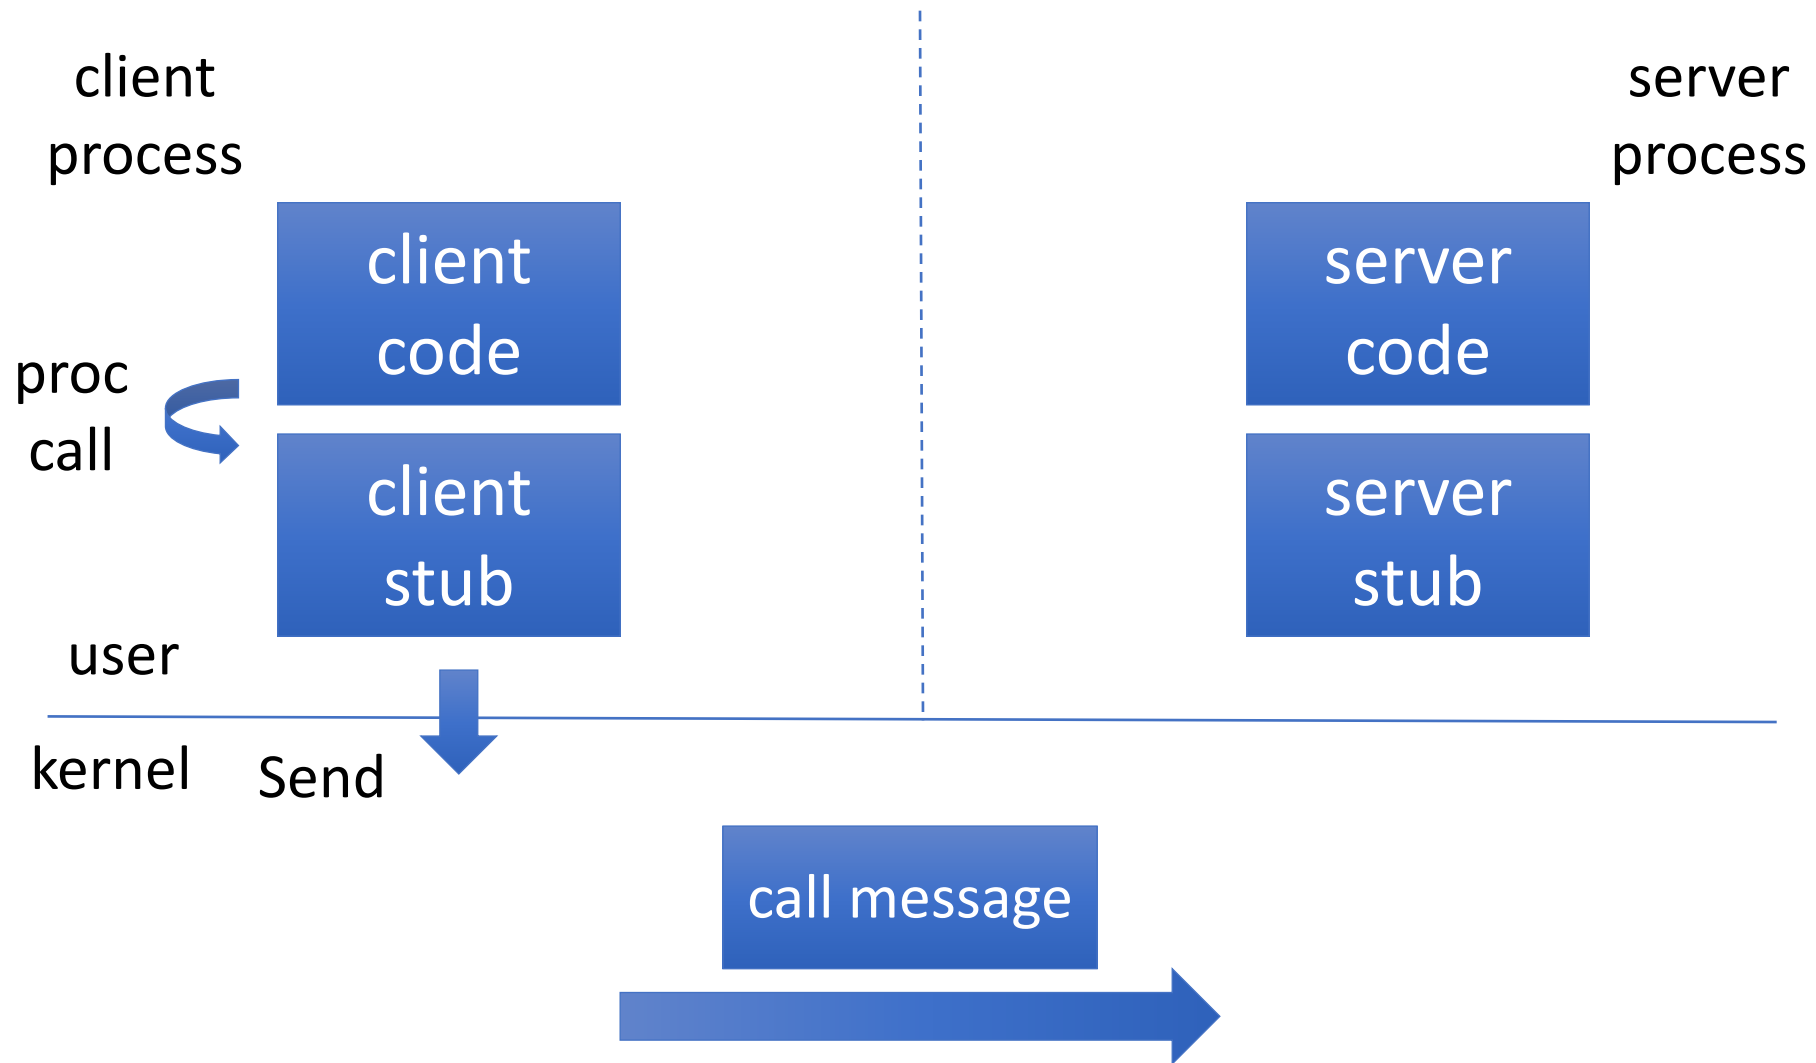
RPC Implementation: Call




RPC Implementation: Call



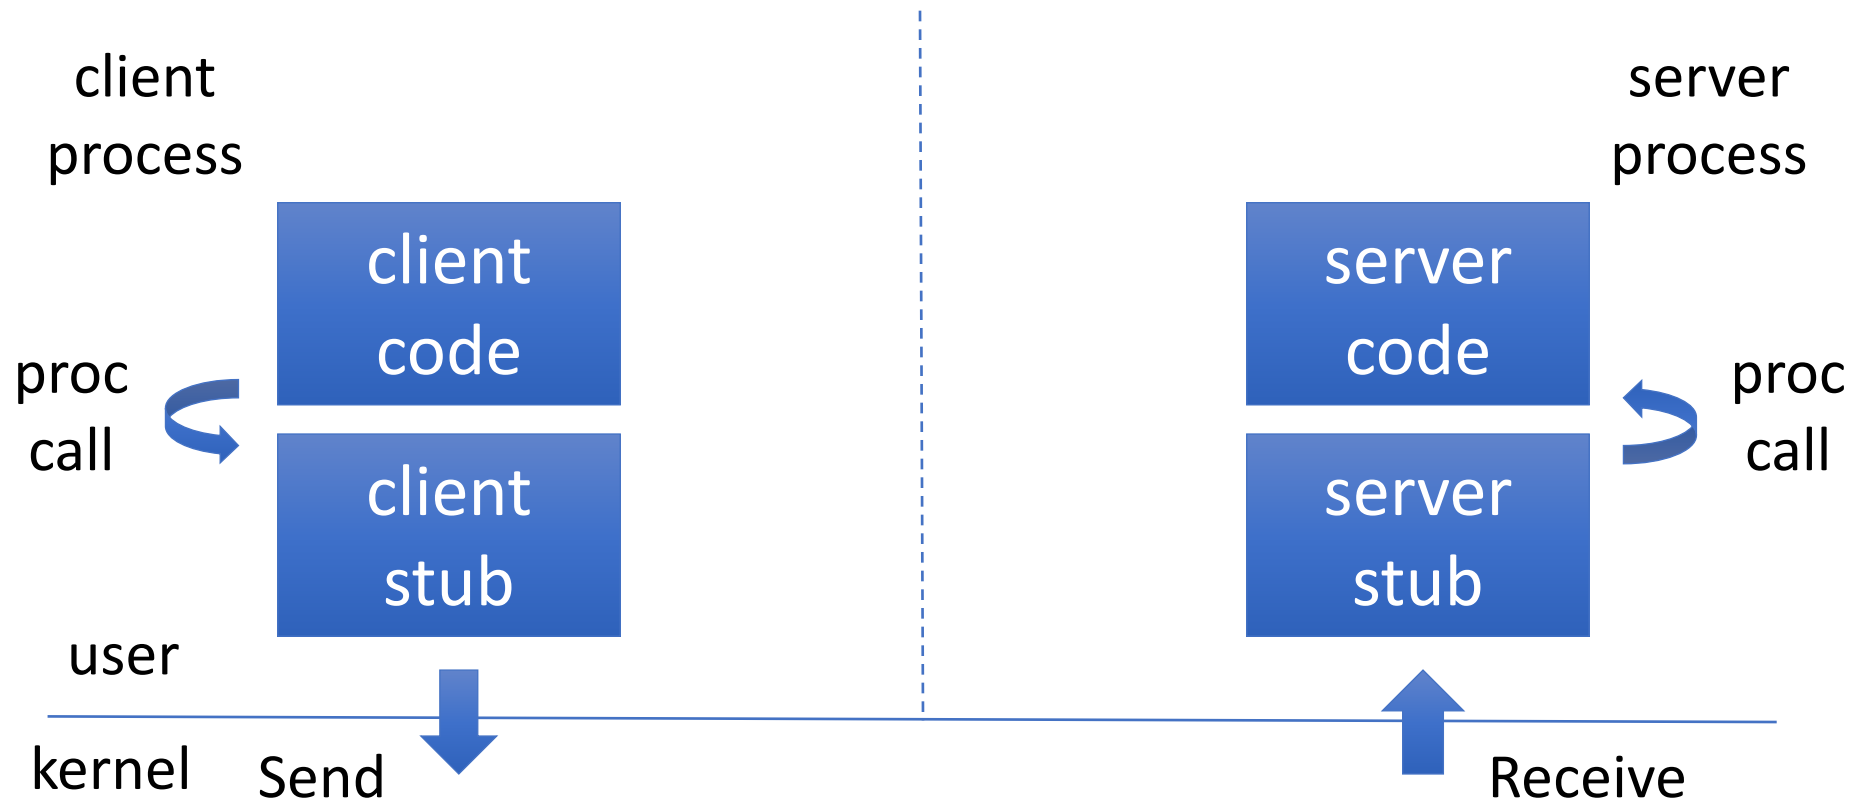
RPC Implementation: Call




RPC Implementation: Call




RPC Implementation: Call




RPC Implementation: Return



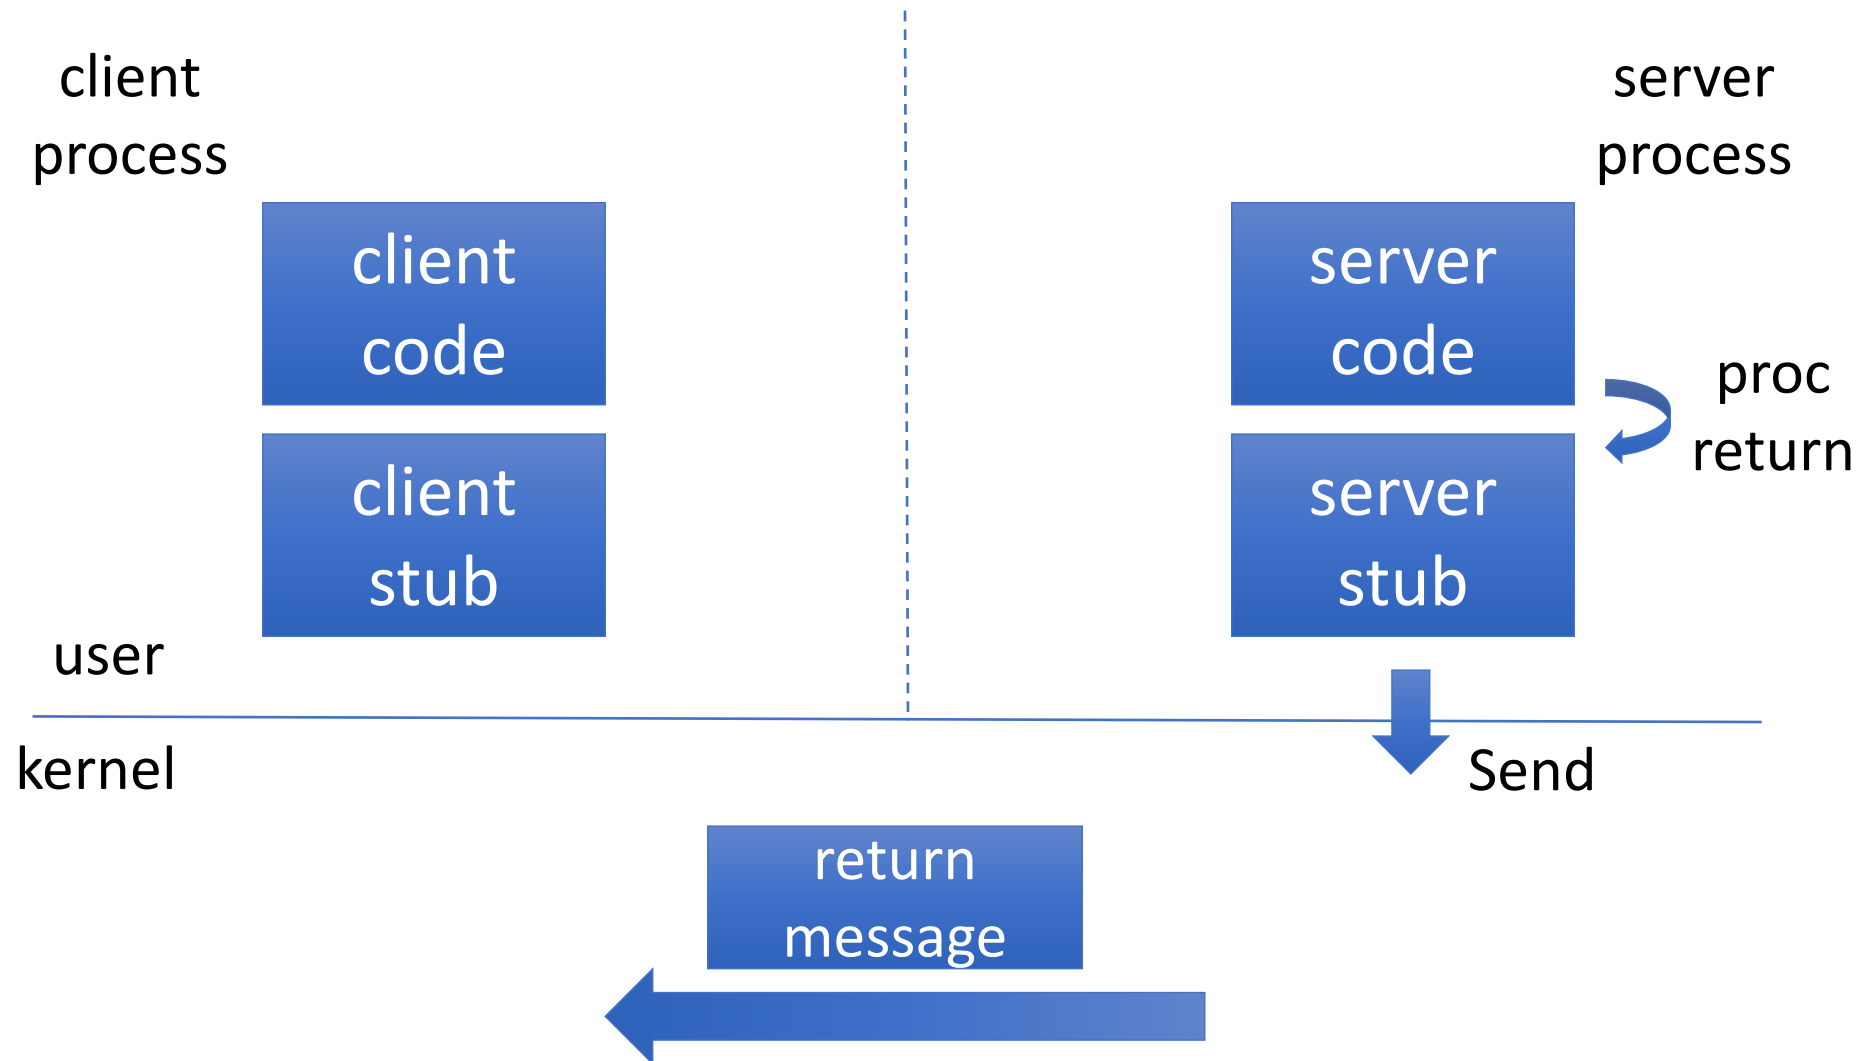
RPC Implementation: Return




RPC Implementation: Return




RPC Implementation: Return



RPC Implementation: Return



RPC Implementation: Return



Example 1: Timeserver

- Supports `GetTime()` and `SetTime()`

Interface

```
long GetTime()  
boolean SetTime(long time)
```

Server Code

```
GetTime()  
{  
    return(ReadHardwareClock())  
}  
SetTime(time)  
{  
    WriteHardwareClock(time)  
    return(1)  
}
```

Client Code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

Message Format

- We already saw:
 - Call message contains arguments
- Must also include which procedure is called

Message Format

Call Message

procno
arg0

Return Message

retval0

Client Stub

```
GetTime()  
{  
    msg->procno = 1  
    Send(msg)  
    Receive(msg)  
    return(msg->retval0)  
}  
SetTime(long time)  
{  
    msg->procno = 2  
    msg->arg0 = time  
    Send(msg)  
    Receive(msg)  
    return(msg->retval0)  
}
```

Server Stub

```
while(true) do
{
    Receive(msg)
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }

        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```


client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code

```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

client code

```
main()
{
    time = GetTime()
    SetTime(time + 100)
}
```

client stub

```
GetTime()
{
    msg->procno = 1
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}

SetTime(long time)
{
    msg->procno = 2
    msg->arg0 = time
    Send(msg)
    Receive(msg)
    return(msg->retval0)
}
```

server code


```
GetTime()
{
    return(ReadHardwareClock())
}

SetTime(time)
{
    WriteHardwareClock(time)
    return(1)
}
```

server stub

```
while(true) do
{
    Receive( msg )
    switch msg->procno {
        case 1: { time = GetTime()
                  msg->retval0 = time
                  Send(msg) }
        case 2: { ret = SetTime(msg->arg0)
                  msg->retval0 = ret
                  Send(msg) }
    }
}
```

Note: Stubs Generated Automatically



Let's see how this works in Linux

RPC library

- Make procedure calls on other machines across the network
 - Can also be used for inter-process communication on localhost
 - <https://man7.org/linux/man-pages/man3/rpc.3.html>
- Beyond the scope of the class to be fluent in RPC. We will just see a simple example.
 - Can continue exploring this in Networking classes (e.g., Distributed systems COMP512, Computer Networks COMP535).

Automatic stub generation

- **rpcbind** — universal addresses to RPC program number mapper
 - A server that converts RPC program numbers into universal addresses.
 - It must be running on the host to be able to make RPC calls on a server on that machine.
 - **sudo apt-get install rpcbind**

Automatic stub generation cont'd

rpcgen – tool that generates remote program interface modules.

- compiles source code written in the RPC language.
- RPC language is similar in syntax and structure to C.
- produces one or more C language source modules, which are then compiled by a C compiler.
- ☺ rpcgen reduces the development time that would otherwise be spent developing low-level routines.
 - Have a look at a [tutorial](#)

Another example: Sum Server

- Client
 - sends 2 integers: a and b
- Server
 - replies with $a+b$

Step 1: Defining the data structures with XDR routines, in add.x file

```
//add.x

struct numbers{
    int a;
    int b;
};

program ADD_PROG{                                /* program number */
    version ADD_VERS{                             /* version number */
        int add(numbers)=1;                       /* procedure */
    }=1;
}=0x12345;
```

Step 2: Compile XDR routines into C code

```
//add.x
```

```
struct numbers{  
    int a;  
    int b;  
};
```

```
program ADD_PROG{  
    version ADD_VERS{  
        int add(numbers)=1;  
    }=1;  
}=0x12345;
```

```
~/COMP310-2023/RPC$ rpcgen -a -C add.x
```

```
/* program number */  
/* version number */  
/* procedure */
```

Step 3: Modify the server/client code.

```
balmau@disclab-server1:~/COMP310-2023/RPC/RPC-class$ ls
add_client.c  add.h        add_svc.c    add_xdr.c
add_clnt.c    add_server.c add.x         Makefile.add
```

- `add_client.c` : client code that we need to implement
- `add_clnt.c` : automatically generated stub code – No need to change
- `add_server.c` : server code that we need to implement
- `add_svc.c` : automatically generated stub code – No need to change

Step 3.1: Modify the server code add_server.c

```
#include "add.h"

int * add_1_svc(numbers *argp, struct svc_req *rqstp)
{
    static int result;
    /*
     * insert server code here
     */
    return &result;
}
```


Step 3.1: Modify the server code add_server.c

```
#include "add.h"

int * add_1_svc(numbers *argp, struct svc_req *rqstp)
{
    static int result;
    /*
     * insert server code here
     */
    printf("add(%d,%d) is called\n", argp->a, argp->b);
    result = argp->a + argp->b;
    return &result;
}
```

Step 3.2: Modify the client code add_client.c

```
int main (int argc, char *argv[]) {
    char *host;
    if (argc < 2) {
        printf ("usage: %s server_host\n",
argv[0]);
        exit (1);
    }
    host = argv[1];
    add_prog_1 (host);
    exit (0);
}
```

```
void add_prog_1(char *host) {
    CLIENT *clnt;
    int *result_1;
    numbers add_1_arg;

#ifdef DEBUG
    clnt = clnt_create (host, ADD_PROG,
ADD_VERS, "udp");
    if (clnt == NULL) {
        clnt_pcreateerror (host);
        exit (1);
    }
#endif /* DEBUG */

    result_1 = add_1(&add_1_arg, clnt);
    if (result_1 == (int *) NULL) {
        clnt_perror (clnt, "call failed");
    }
#ifdef DEBUG
    clnt_destroy (clnt);
#endif /* DEBUG */
}
```

Step 3.2: Modify the client code add_client.c

```
int main (int argc, char *argv[]) {
    char *host;
    if (argc < 4) {
        printf ("usage: %s server_host\n",
argv[0]);
        exit (1);
    }
    host = argv[1];
    add_prog_1 (host);
    add_prog_1 (host, atoi(argv[2]),
atoi(argv[3]));

    exit (0);
}
```

```
void add_prog_1(char *host, int x, int y) {
    CLIENT *clnt;
    int *result_1;
    numbers add_1_arg;

    //...

    add_1_arg.a=x;
    add_1_arg.b=y;

    result_1 = add_1(&add_1_arg, clnt);
    if (result_1 == (int *) NULL) {
        clnt_perror (clnt, "call failed");
    } else {
        printf("Result:%d\n", *result_1);
    }
    //...}
}
```

Step 4: Compile C code

```
balmau@discslab-server1:~/COMP310-2023/RPC/RPC-class$ ls  
add_client.c  add.h      add_svc.c  add_xdr.c  
add_clnt.c   add_server.c  add.x      Makefile.add
```

```
~/COMP310-2023/RPC$ make -f Makefile.add
```

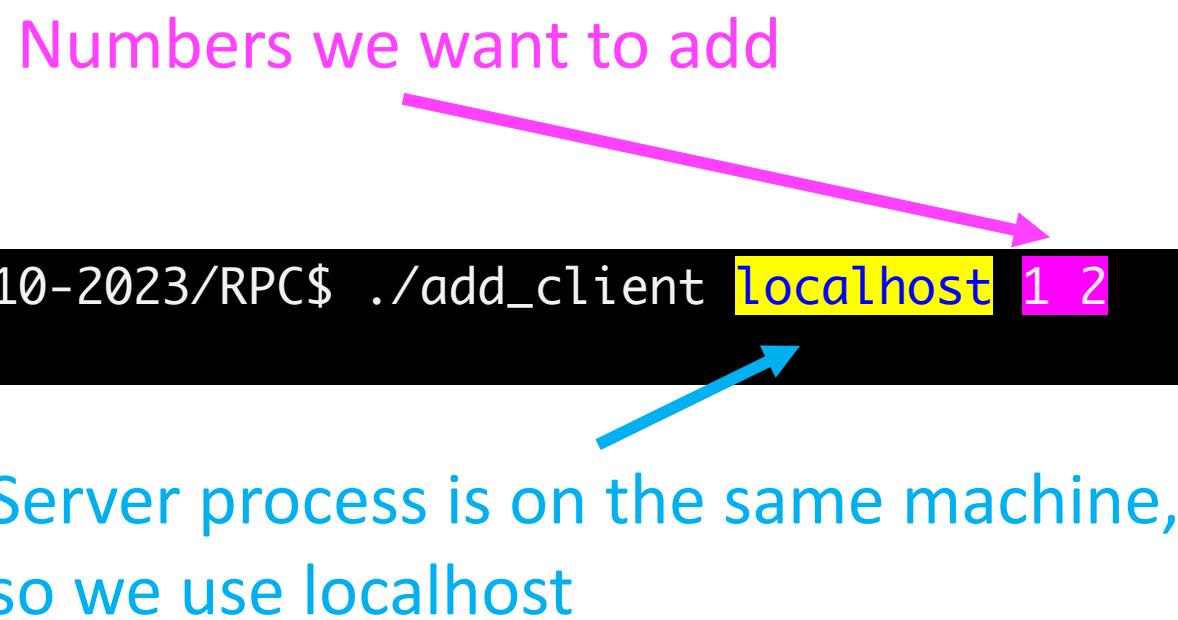
Step 5: Launch the server

```
balmau@discslab-server1:~/COMP310-2023/RPC$ ./add_server
```

Step 6: Use the client

Open new terminal window

```
balmau@discslab-server1:~/COMP310-2023/RPC$ ./add_client localhost 1 2  
Result:3
```



Numbers we want to add

Server process is on the same machine,
so we use localhost

Meanwhile, in the server window:

```
balmau@discslab-server1:~/COMP310-2023/RPC$ ./add_server  
add(1,2) is called
```

Optional Homework

- Use the same process to create the Timeserver in example 1
- *Hint: The linked rpcgen [tutorial](#) will help*

Key Concepts for Today

- Interprocess communication
- Message passing
- Remote procedure call

Further Optional Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 25 – 32 (inclusive) <https://pages.cs.wisc.edu/~remzi/OSTEP/>

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Maurice Herlihy (Brown University)