

Week 5

Multithreading

Oana Balmau
February 2, 2023

Goals for today

- Practice multi-threading synchronization
 - Join Implementation
 - Dining Philosophers
 - Alice and Bob are back with pet dragons!
- Q&A for first module

Problem 1: Join Implementation

- Implement the equivalent of `pthread_join` seen in class.

Reminder: pthread_join(threadid, &status)

- Wait for thread threadid to exit
- Receive status, if any

Ordering Example: Join

```
pthread_t p1, p2;  
pthread_create(&p1, NULL, mythread, "A");  
pthread_create(&p2, NULL, mythread, "B");
```

```
// join waits for the threads to finish
```

```
pthread_join(p1, NULL);  
pthread_join(p2, NULL);  
printf("Done!\n");  
return 0;
```

how to implement join?

Reminder: Condition Variables

Condition Variable \sim queue of waiting threads

B waits for a signal on CV before running

- wait(CV, ...)

A sends signal to CV when time for **B** to run

- signal(CV, ...)

Reminder: Condition Variables

wait(cond_t *cv, mutex_t *lock)

- assumes the lock is held when wait() is called
- puts caller to sleep + releases the lock (atomically)
- when awoken, reacquires lock *immediately* before returning

signal(cond_t *cv)

- wake a single waiting thread (if ≥ 1 thread is waiting)
- if there is no waiting thread, just return, doing nothing

Thinking time; How to implement join?

```
pthread_t p1, p2;  
pthread_create(&p1, NULL, mythread, "A");  
pthread_create(&p2, NULL, mythread, "B");  
  
// join waits for the threads to finish  
pthread_join(p1, NULL);  
pthread_join(p2, NULL);  
printf("Done!\n");  
return 0;
```

wait(cond_t *cv, mutex_t *lock)

- assumes the lock is held when wait() is called
- puts caller to sleep + releases the lock (atomically)
- when awoken, reacquires lock *immediately* before returning

signal(cond_t *cv)

- wake a single waiting thread (if ≥ 1 thread is waiting)
- if there is no waiting thread, just return, doing nothing

Hint: Use mutex + condition variables

Join Implementation: Attempt 1

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // x  
    Cond_wait(&c, &m);        // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    Cond_signal(&c);          // b  
    Mutex_unlock(&m);         // c  
}
```

Join Implementation: Attempt 1

Parent:

```
void thread_join() {
    Mutex_lock(&m);           // x
    Cond_wait(&c, &m);        // y
    Mutex_unlock(&m);         // z
}
```

Child:

```
void thread_exit() {
    Mutex_lock(&m);           // a
    Cond_signal(&c);          // b
    Mutex_unlock(&m);         // c
}
```

Example schedule:

Parent: x y z

Child:

Works!

Join Implementation: Attempt 1

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // x  
    Cond_wait(&c, &m);        // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    Cond_signal(&c);          // b  
    Mutex_unlock(&m);         // c  
}
```

Example schedule:

Parent:	x	y				z
Child:			a	b	c	

Can you construct a schedule that doesn't work?

Join Implementation: Attempt 1 (incorrect)

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // x  
    Cond_wait(&c, &m);        // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    Cond_signal(&c);          // b  
    Mutex_unlock(&m);         // c  
}
```

Example broken schedule:

Parent:

x y

Child:

a b c

Parent waits forever!

Idea

- **Keep state** in addition to CV's!
- CV's are used to signal threads when state changes
- If state is already as needed, thread doesn't wait for a signal!

Join Implementation: Attempt 2

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;                 // a  
    Cond_signal(&c);          // b  
}
```

Join Implementation: Attempt 2

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;                 // a  
    Cond_signal(&c);          // b  
}
```

Fixes previous broken ordering:

Parent:

w x ~~y~~ z

Child:

a b

Join Implementation: Attempt 2

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;                // a  
    Cond_signal(&c);          // b  
}
```

Fixes previous broken ordering:

Can you construct ordering that does not work?

Parent:

w x y z

Child:

a b

Join Implementation: Attempt 2 (incorrect)

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;                 // a  
    Cond_signal(&c);          // b  
}
```

Parent: w x

y

... sleep forever ...

Child: a b

Join Implementation: Correct

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1;                 // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```


Parent: w x y z


Child: a b c

Use mutex to ensure no race between interacting with state and wait/signal

Problem 2: Dining Philosophers

- Classic problem in synchronization
 - Dijkstra '71






Each philosopher has 2 states:


- **Eating**
 - Need 2 chopsticks to eat
- **Thinking**

When they're not eating,
they're thinking (and vice-versa)




Dinner rules

- Philosophers can't speak to each other.
- Philosophers can only pick up one chopstick at a time.
- Philosophers can only pick up adjacent chopsticks.
- Infinite food supply.



Our task for today

Design a behavior such that no philosopher will starve, i.e. each can forever continue to alternate between eating and thinking.



Each philosopher has 2 states:

- **Eating**
 - Need 2 chopsticks to eat
- **Thinking**

When they're not eating, they're thinking (and vice-versa)


Dinner rules


- Philosophers can't speak to each other.
- Philosophers can only pick up one chopstick at a time.
- Philosophers can only pick up adjacent chopsticks.
- Infinite food supply.

Our task

Design a behavior (i.e., an algorithm) such that no philosopher will starve, i.e. each can forever continue to alternate between eating and thinking.


Brainstorming Slide





A simple solution (incorrect)


```
do forever{  
  think()  
  wait(chopstick_R)  
  grab(chopstick_R)  
  wait(chopstick_L)  
  grab(chopstick_L)  
  eat()  
  releaseChopsticks()  
}
```



A simple solution (incorrect)

```
do forever{  
  think()  
  wait(chopstick_R)  
  grab(chopstick_R)  
  wait(chopstick_L)  
  grab(chopstick_L)  
  eat()  
  releaseChopsticks()  
}
```

Deadlock Danger!



A simple solution (incorrect)

```
do forever{  
  think()  
  wait(chopstick_R)  
  grab(chopstick_R)  
  wait(chopstick_L)  
  grab(chopstick_L)  
  eat()  
  releaseChopsticks()  
}
```

Deadlock Danger!




A simple solution (incorrect)

```
do forever{  
  think()  
  wait(chopstick_R)  
  grab(chopstick_R)  
  wait(chopstick_L)  
  grab(chopstick_L)  
  eat()  
  releaseChopsticks()  
}
```

Deadlock Danger!


How can we fix this?



Solution 1

Asymmetric algorithm


- Assign odd and even IDs to philosophers.
- Odd philosophers pick up left chopstick, then right.
- Even philosophers pick up right chopstick, then left.



Solution 1

Asymmetric algorithm


- Assign odd and even IDs to philosophers.
- Odd philosophers pick up left chopstick, then right.
- Even philosophers pick up right chopstick, then left.



Solution 1

Asymmetric algorithm


- Assign odd and even IDs to philosophers.
- Odd philosophers pick up left chopstick, then right.
- Even philosophers pick up right chopstick, then left.



Solution 1

Asymmetric algorithm


- Assign odd and even IDs to philosophers.
- Odd philosophers pick up left chopstick, then right.
- Even philosophers pick up right chopstick, then left.





Solution 1

Asymmetric algorithm



- Assign odd and even IDs to philosophers.
- Odd philosophers pick up left chopstick, then right.
- Even philosophers pick up right chopstick, then left.





Any other approaches?





**Only allow 4 philosophers
At the table**




**Only allow 4 philosophers
At the table**



**Only allow 4 philosophers
At the table**



**Only allow 4 philosophers
At the table**



In line to eat

**Only allow 4 philosophers
At the table**


Dining philosophers – more solutions

- Other solutions are possible as well:
 - Different asymmetry: make one philosopher grab the left fork first and then the right fork; all others grab the right fork first and then the left fork.
 - Use an arbiter who determines the order in which the philosophers can eat. Arbiter allows philosophers to pick up 2 chopsticks at once.
 - Use backoffs and randomness to break deadlock.

Problem 3:

Alice and Bob Share a Pond and Pet Dragons


- The following story was told by a famous multiprocessing pioneer
- Leslie Lamport.




Mutual Exclusion, or “Alice & Bob share a pond”



Alice has a pet



Bob has a pet



The Problem



Formalizing the Problem

- Mutual exclusion
 - Both pets never in pond simultaneously
- No deadlock
 - If one wants in, it gets in
 - If both want in, one gets in

Formalizing the Problem

- Mutual exclusion
 - Both pets never in pond simultaneously
 - Safety property!
- No deadlock
 - If one wants in, it gets in
 - If both want in, one gets in
 - Liveness property!

Simple Protocol

- Idea
 - Just look at the pond
- Gotcha
 - Not atomic
 - “Trees obscure the view”




Interpretation

- Threads can't “see” what other threads are doing
- Explicit communication required for coordination

Cell-phone protocol


- Idea
 - Bob calls Alice (or vice-versa)
- Gotcha
 - Bob takes shower
 - Alice recharging battery
 - Bob out shopping for pet food




Interpretation

- Message-passing with shared memory doesn't work
- Recipient might not be
 - Listening
 - There at all
- Communication must be
 - Persistent (like writing)
 - Not transient (like speaking)


Flag Protocol



Alice's Protocol (sort of)



Bob's Protocol (sort of)



Alice's Protocol

- Raise flag
- Wait until Bob's flag is down
- Unleash pet
- Lower flag when pet returns

Bob's protocol

- Raise flag
- Wait until Alice's flag is down
- Unleash pet
- Lower flag when pet returns

deadlock
danger!

Bob's protocol (2nd try)

- Raise flag
- While Alice's flag is up...
 - Lower flag
 - Wait for Alice's flag to go down
 - Raise flag
- Unleash pet
- Lower flag when pet returns

Bob's protocol (2nd try)

- Raise flag
- While Alice's flag is up...
 - Lower flag
 - Wait for Alice's flag to go down
 - Raise flag
- Unleash pet
- Lower flag when pet returns

Bob defers to
Alice

The Flag Principle

- Raise Flag
- Look at other's flag
- Flag principle:
 - If each raises and looks, then
 - Last to look must see both flags up

Alice-Bob-Pond: What does it mean? (1/3)

- What do all these story elements mean in OS?
- Alice & Bob = threads in the same process
- Pond = Shared memory region that needs to be accessed in a mutually exclusive way
- Pets = functions in the threads' code that need to access the shared memory region
 - When pets are in the pond, it means that **threads are in the critical section**

Alice-Bob-Pond: What does it mean? (2/3)

- Trees = a metaphor to signify that threads cannot observe the state of a shared memory region with confidence. Why?
 - Because threads can be interrupted by the OS at any time,
 - for an indefinite amount of time.

Thread A

```
1 mem_state = read(shared_mem)
2 switch(mem_state)
3 case(x) { do X}
4 case(y) { do Y}
...
```

Alice-Bob-Pond: What does it mean? (2/3)

- Trees = a metaphor to signify that threads cannot observe the state of a shared memory region with confidence. Why?
 - Because threads can be interrupted by the OS at any time,
 - for an indefinite amount of time.

Thread A

```
1 mem_state = read(shared_mem)
2 switch(mem_state)
3 case(x) { do X}
4 case(y) { do Y}
...
```

Scheduler can interrupt A between lines 1 and 2
Thread B can run and change shared_mem

A's mem_state variable is stale

Alice-Bob-Pond: What does it mean? (3/3)

- Flags = Bits that threads use to coordinate access to the shared memory region.
 - In a system with N threads, generally we define an array of size N .
 - Each thread can write in one entry (i.e., write only their flag)
 - Can read all the other entries (i.e., look at all the others' flags)

Proof of Mutual Exclusion (Sketch)

- Proof by contradiction
- Assume both pets in pond somehow
 - Derive a contradiction
 - By reasoning backwards
- Consider the last time Alice and Bob each looked before letting the pets in.
- Without loss of generality, assume Alice was the last to look...

Proof of Mutual Exclusion

Thread A

```
Raise flag  
Wait until B's flag is down (looking)  
Release pet  
Lower flag when pet returns
```

Thread B

```
Raise flag  
While (A's flag is up)  
  Lower B flag  
  Wait for A flag to go down (looking)  
  Raise flag  
Release pet  
Lower flag when pet returns
```

Protocol: first raise flag, then look
for both A and B




Proof of Mutual Exclusion

Thread A

```
Raise flag  
Wait until B's flag is down (looking)  
Release pet  
Lower flag when pet returns
```

Thread B

```
Raise flag  
While (A's flag is up)  
  Lower B flag  
  Wait for A flag to go down (looking)  
  Raise flag  
Release pet  
Lower flag when pet returns
```




Proof of Mutual Exclusion

Thread A

```
Raise flag  
Wait until B's flag is down (looking)  
Release pet  
Lower flag when pet returns
```

Thread B

```
Raise flag  
While (A's flag is up)  
  Lower B flag  
  Wait for A flag to go down (looking)  
  Raise flag  
Release pet  
Lower flag when pet returns
```




Proof of Mutual Exclusion

Thread A

```
Raise flag  
Wait until B's flag is down (looking)  
Release pet  
Lower flag when pet returns
```

Thread B

```
Raise flag  
While (A's flag is up)  
  Lower B flag  
Wait for A flag to go down (looking)  
Raise flag  
Release pet  
Lower flag when pet returns
```



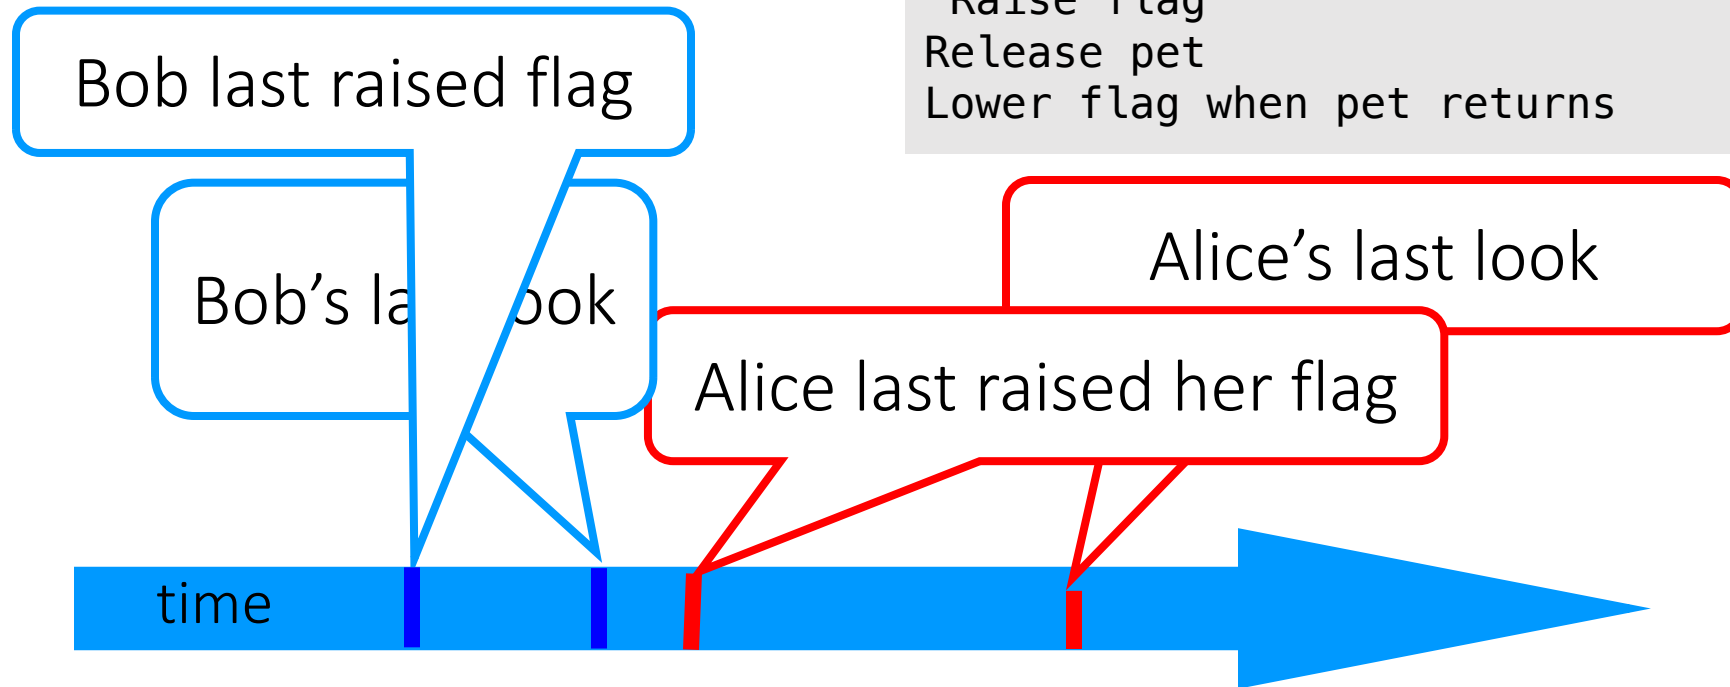
Proof of Mutual Exclusion

Thread A

Raise flag
Wait until B's flag is down (looking)
Release pet
Lower flag when pet returns

Thread B

Raise flag
While (A's flag is up)
Lower B flag
Wait for A flag to go down (looking)
Raise flag
Release pet
Lower flag when pet returns




Proof of Mutual Exclusion

Thread A

Raise flag
Wait until B's flag is down (looking)
Release pet
Lower flag when pet returns

Thread B

Raise flag
While (A's flag is up)
Lower B flag
Wait for A flag to go down (looking)
Raise flag
Release pet
Lower flag when pet returns



Alice must have seen Bob's Flag. A Contradiction

Alice and Bob Mutual Exclusion Proof cont'd

Assume by way contradiction that we do not have mutual exclusion. **We are assuming that both pets are in the pond.**

Therefore, both A and B had a last “looking” action before they let their pet entered the pond.

Consider the one who finished this looking action first.

- Someone must be first (i.e., in a single-CPU, instructions are executed sequentially)
- **When B (or A) looked, B (or A) saw that the other one's flag was down.**

Alice and Bob Mutual Exclusion Proof cont'd

Without loss of generality, assume it was B who finished the looking action first (i.e., A looked last).

- so B had (A's flag == down) as true, otherwise B couldn't have entered the critical section (i.e., send pet in the pond).
- So it follows that A's flag was up **after** B finished the looking action.
- Therefore, **A's looking was completely after** the end of B's raising of their flag (remember, threads first raise flag, then look)
- So A must have seen B's flag up and could not have entered the critical section,
→ Contradiction with the initial assumption that both threads are in the critical section at the same time.

Proof of No Deadlock

- Claim: If only one pet wants in, it gets in
- Deadlock requires both continually trying to get in
- If Bob sees Alice's flag, he backs off, gives her priority
(Alice's lexicographic privilege)

Further Optional Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 25 – 31 (inclusive) <https://pages.cs.wisc.edu/~remzi/OSTEP/>

For a very helpful alternative explanation on the producer/consumer problem, with C code tutorial, check out this [link](#) from the CodeVault YouTube channel.

You are also encouraged to check the other CodeVault tutorials on multi-processing and multi-threading in C/Unix ([link](#)).

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney).