

Week 5

Multithreading

Oana Balmau
January 31, 2023

Schedule Highlights

		drop deadline			
Week 4 Process Management	jan 23 C Tools: GDB basics	jan 24 Multi-process Structuring (1/2) Team registration deadline	jan 25	jan 26 Multi-process Structuring (2/2)	jan 27
Week 5 Process Management	jan 30 C Review: Pointers & Memory Allocation I	jan 31 Multithreading (1/2)	feb 1	feb 2 Multithreading (2/2) Practice Exercises Sheet: Process Management	feb 3
Week 6 Memory Management	feb 6 C Review: C files	feb 7 Virtual Memory (1/2) Optional reading: OSTEP Chapters 12 – 18	feb 8 Scheduling Assignment Released	feb 9 Virtual Memory (2/2) Scheduling Assignment Overview — with Jiaxuan	feb 10
Week 7 Memory Management	feb 13 OS Shell Assignment Due C Review: Working with pthreads I	feb 14 Demand Paging (1/3) Optional reading: OSTEP Chapters 19 – 22	feb 15	feb 16 Demand Paging (2/3)	feb 17


Practice exercises released.

Assignment 2 released next week

Recap Week 4

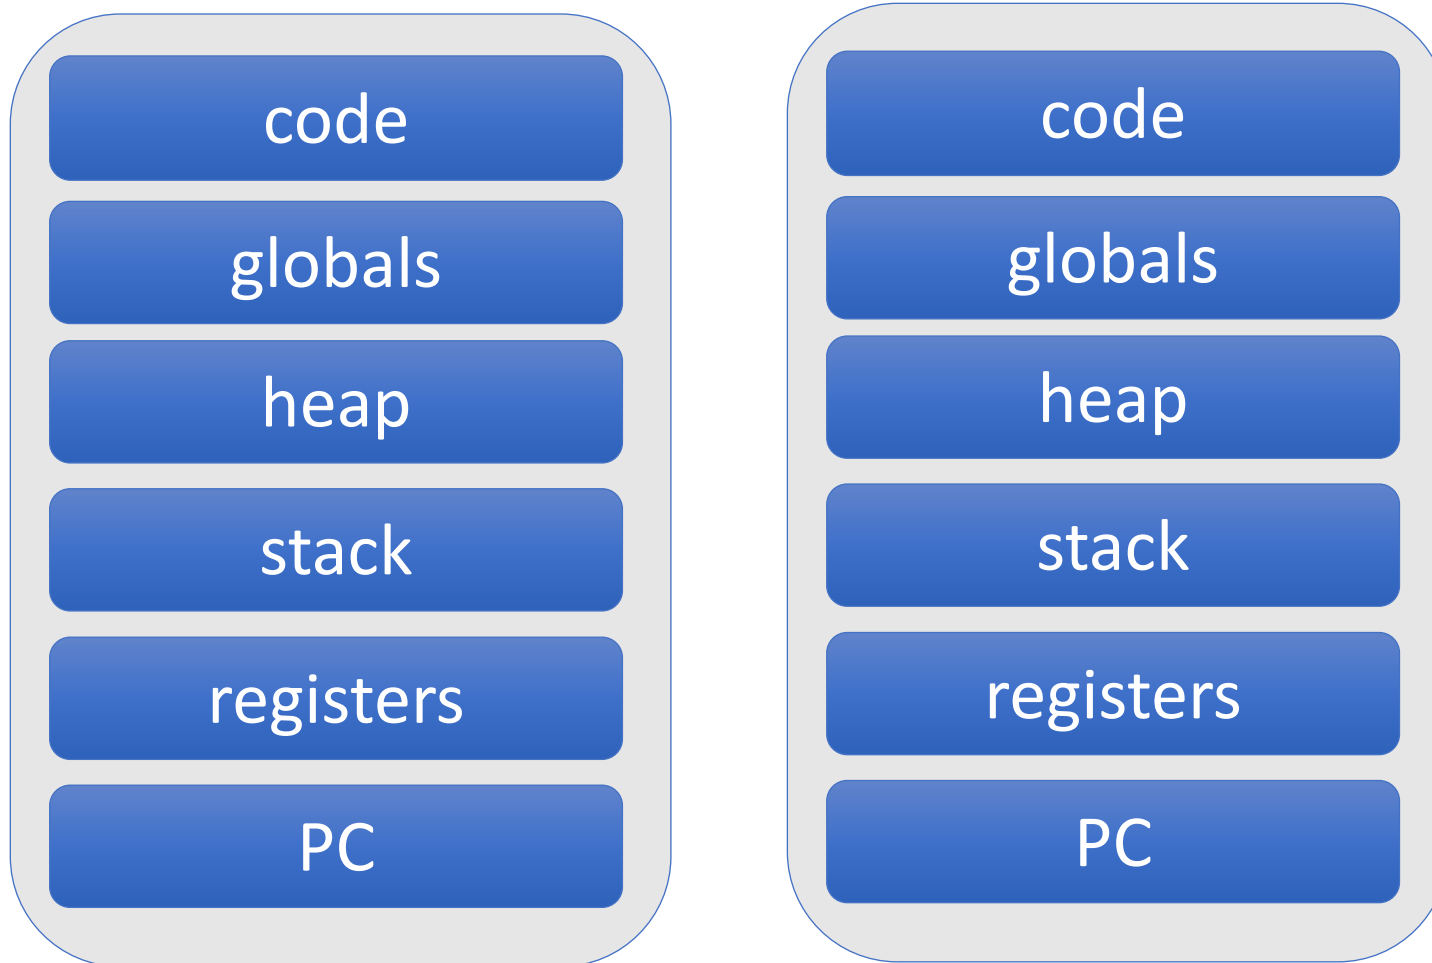
Concurrency – Option 1

- Build apps from many communicating **processes**
- Communicate through **message passing / RPCs**
- Pros
 - If one process crashes, other processes unaffected
- Cons
 - High communication overheads



Last week's
focus

Recap Week 4: Two Processes



Recap Week 4 RPC Implementation

client
process

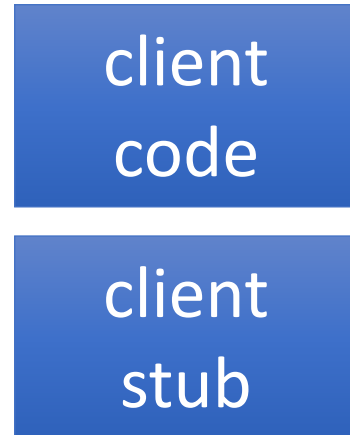
client
code

server
process

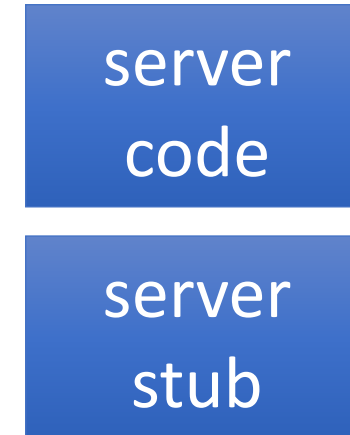
server
code

Recap Week 4 Client and Server Stubs

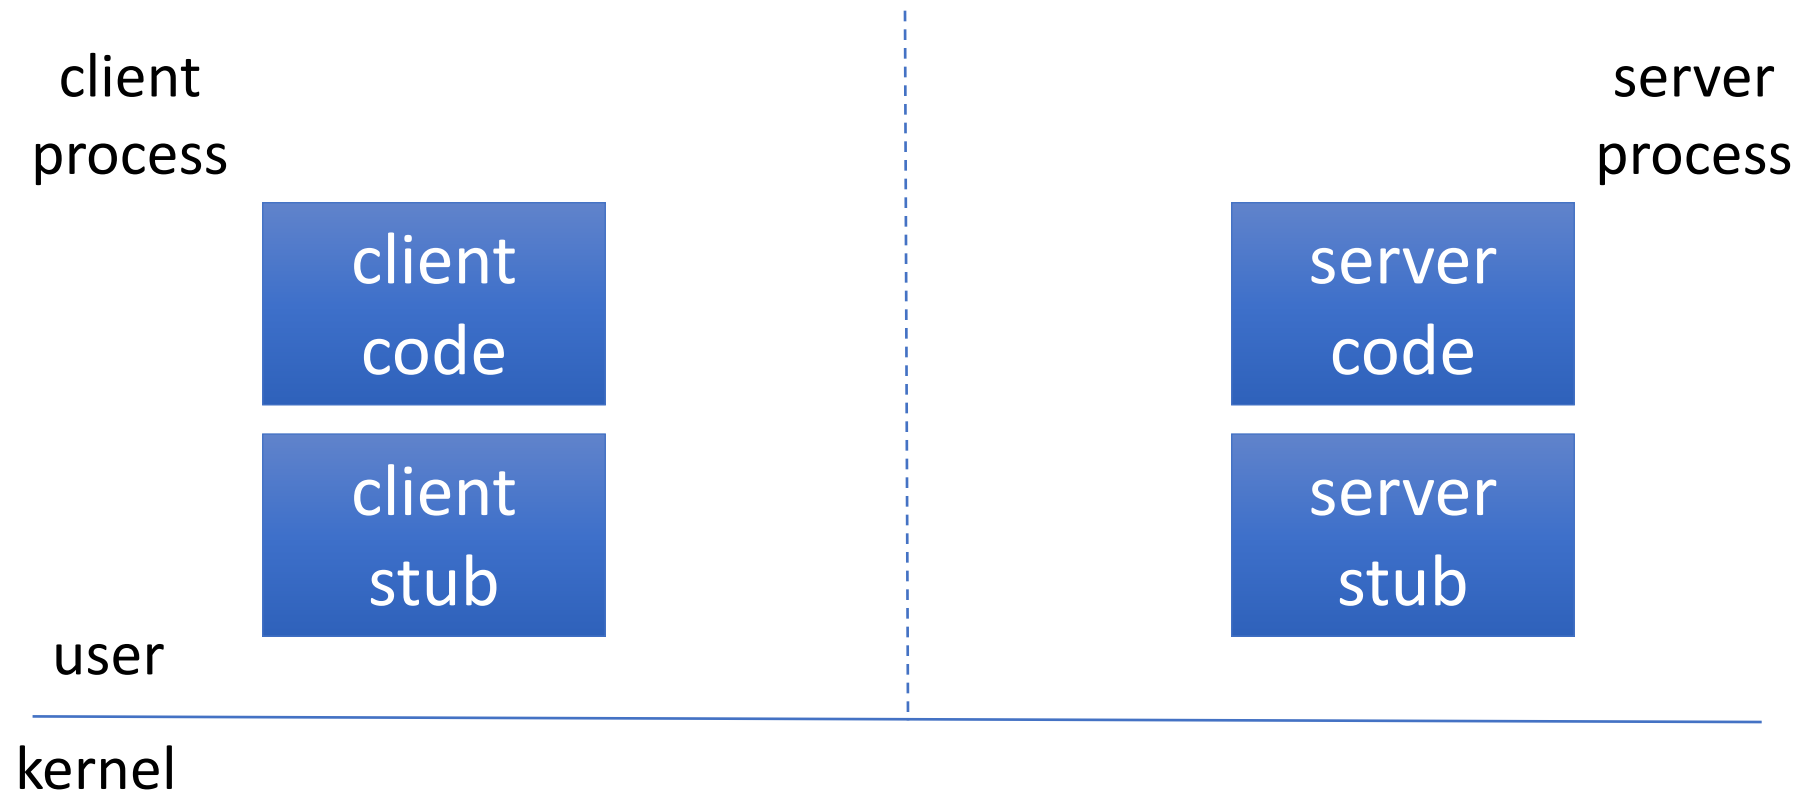
client
process



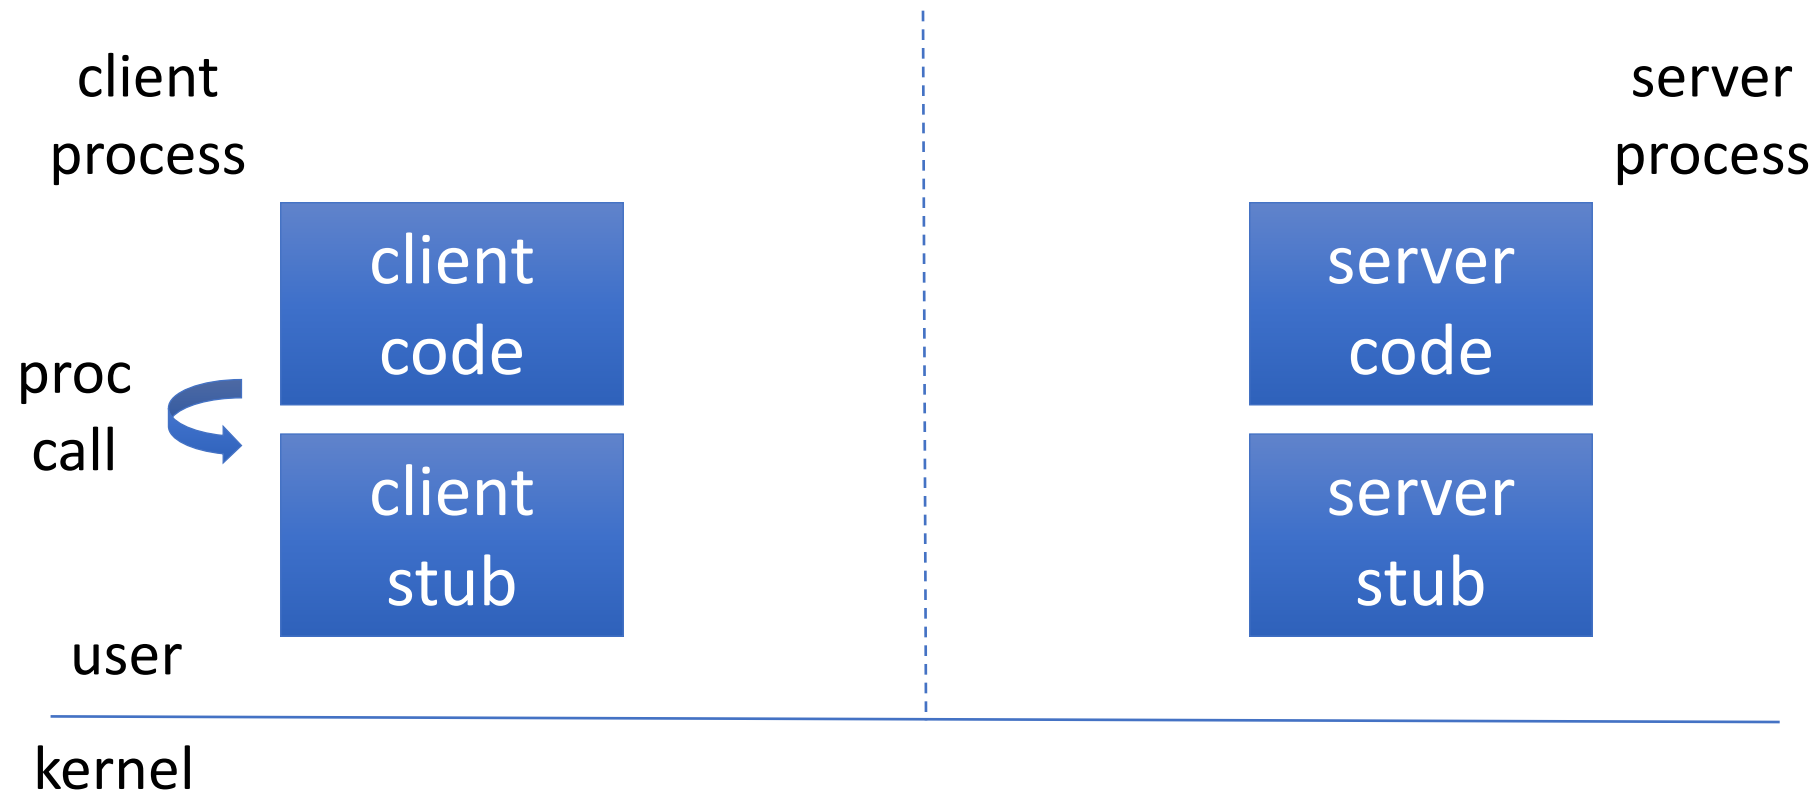
server
process



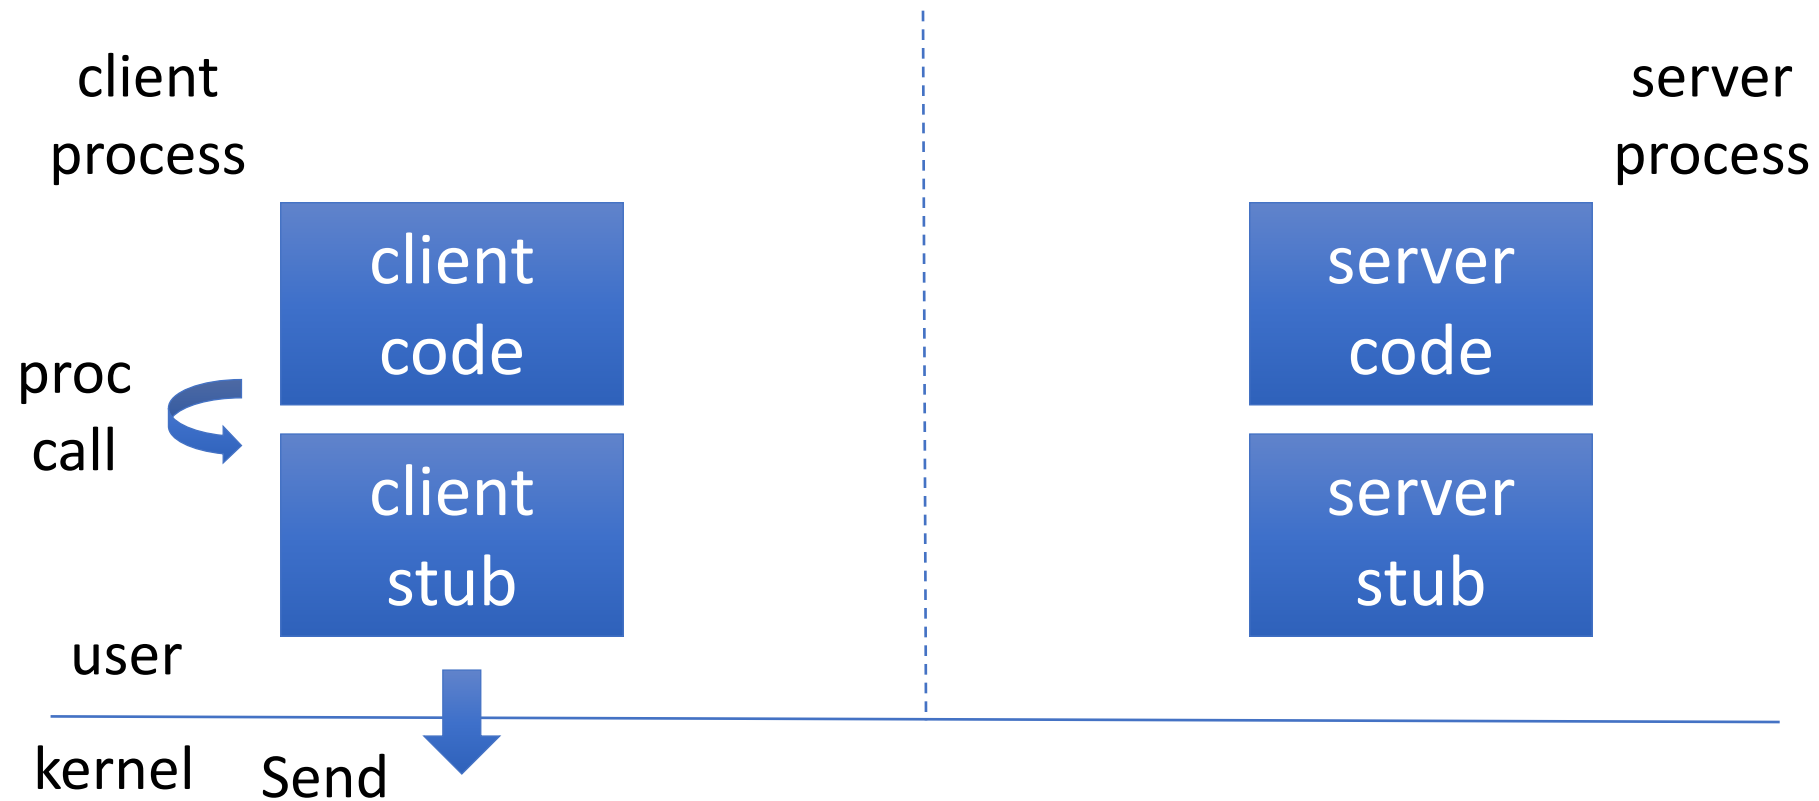
Recap Week 4 RPC Implementation: Call



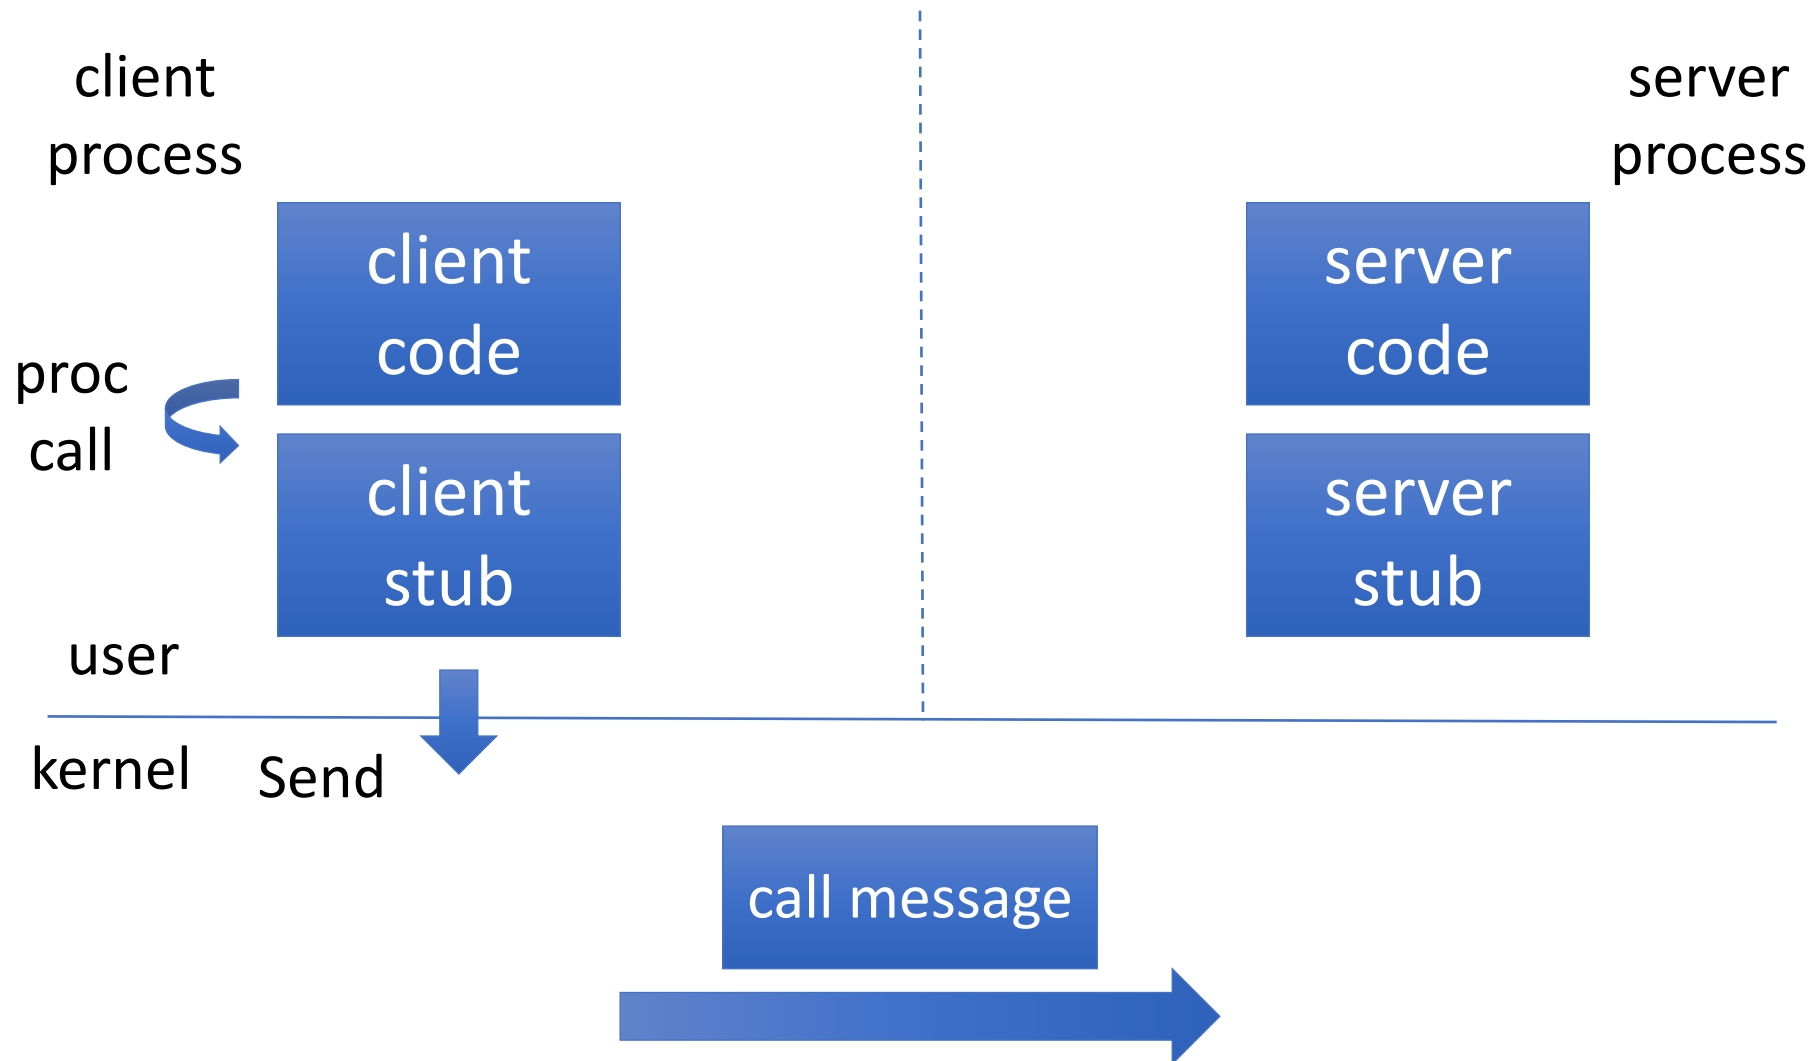
Recap Week 4 RPC Implementation: Call



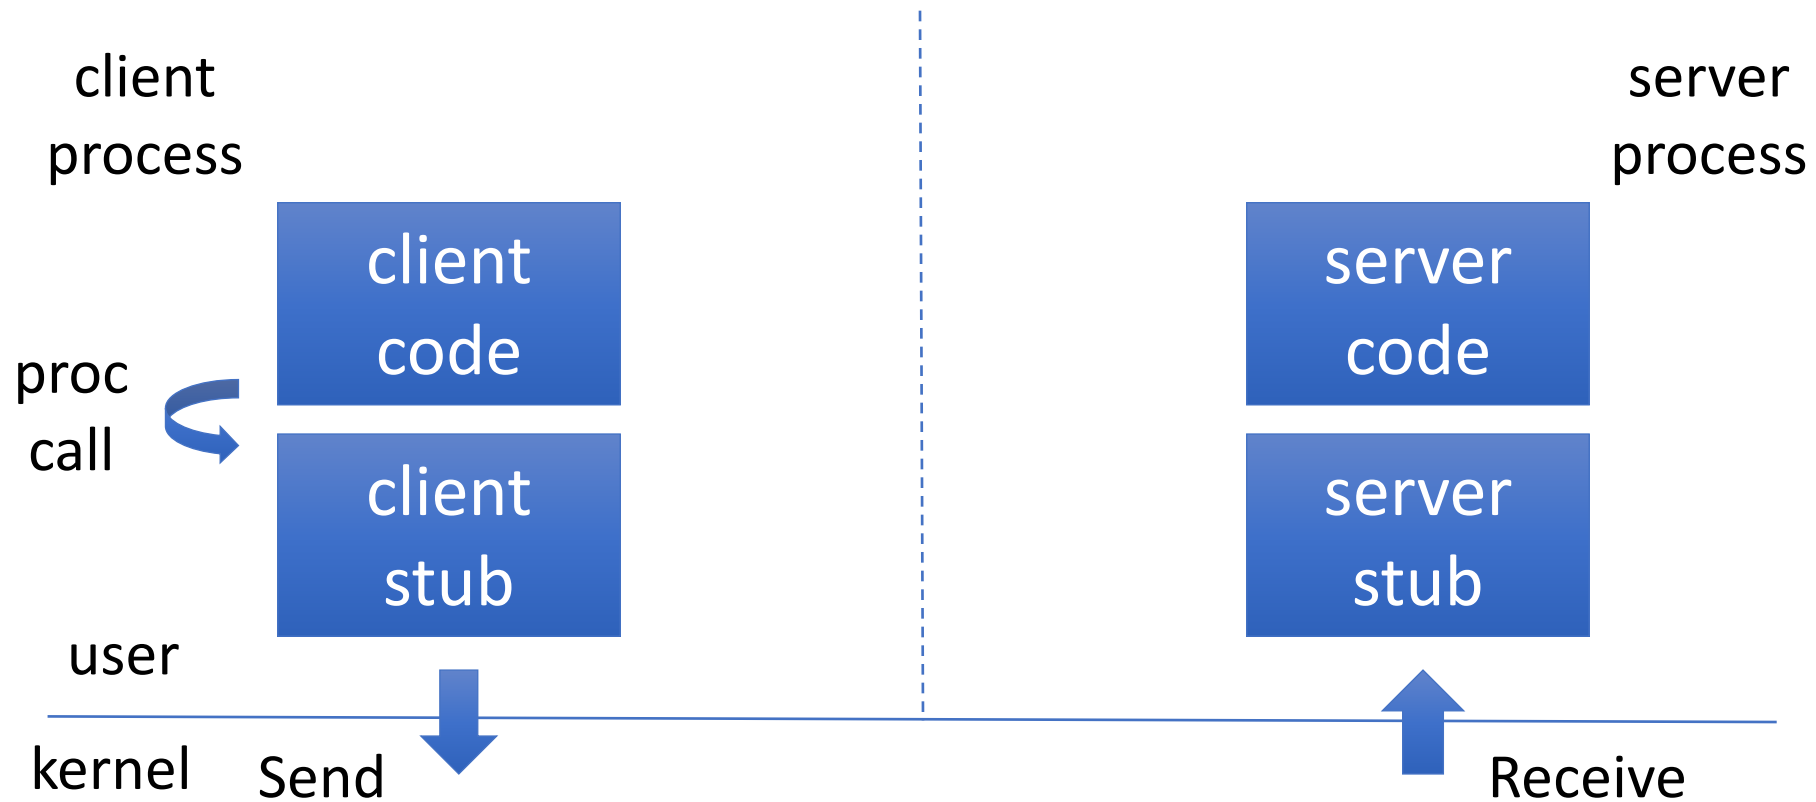
Recap Week 4 RPC Implementation: Call



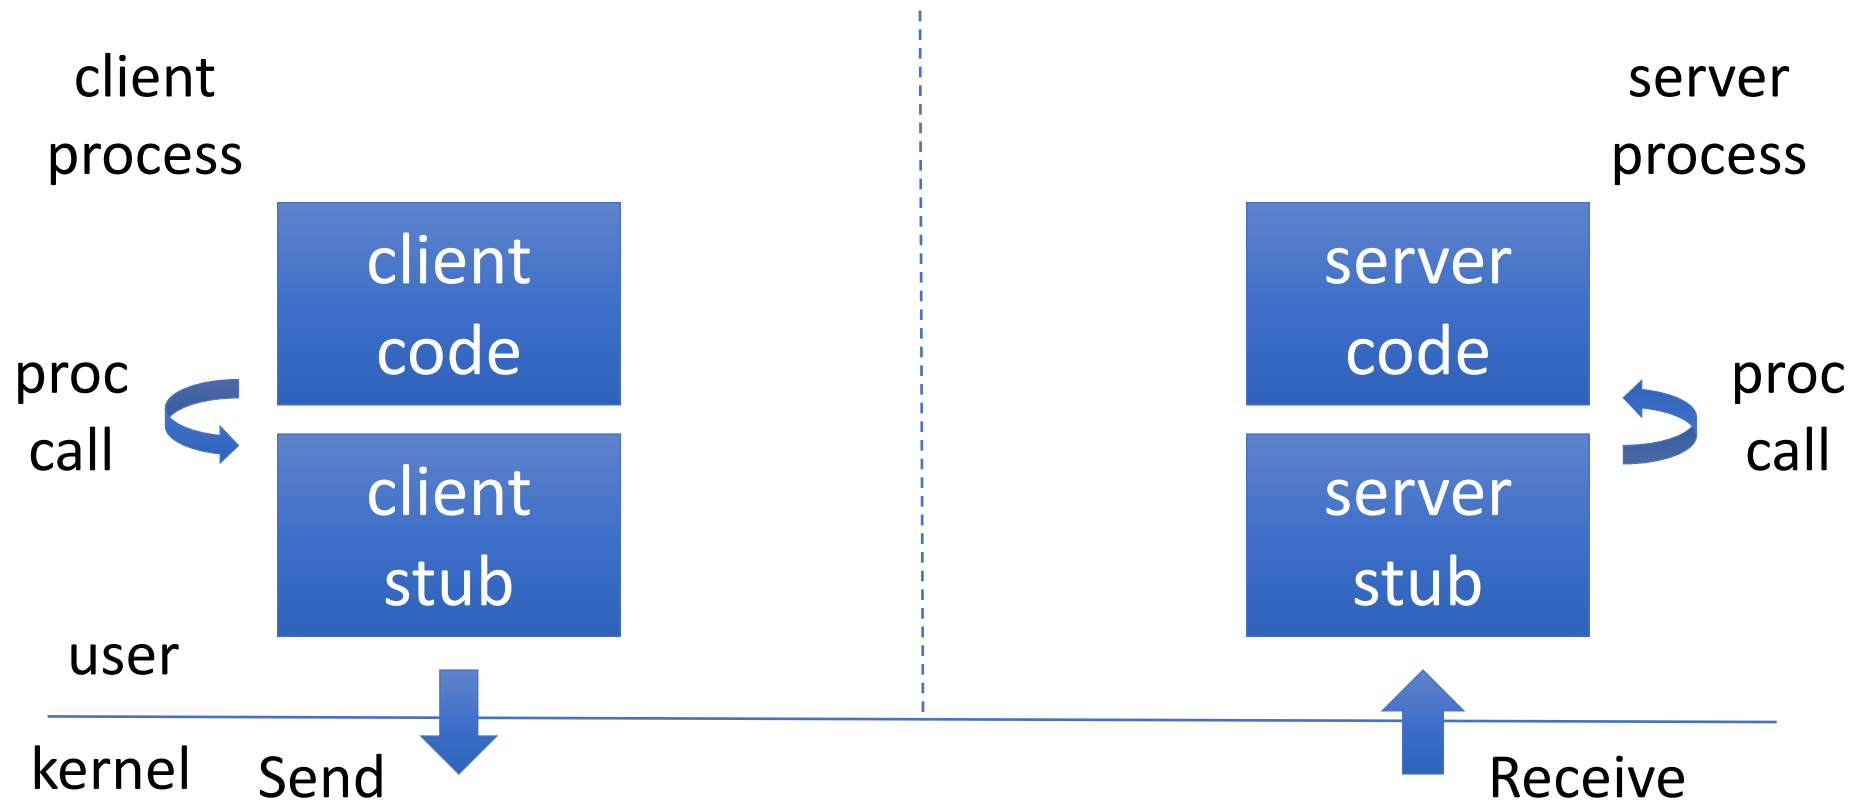
Recap Week 4 RPC Implementation: Call



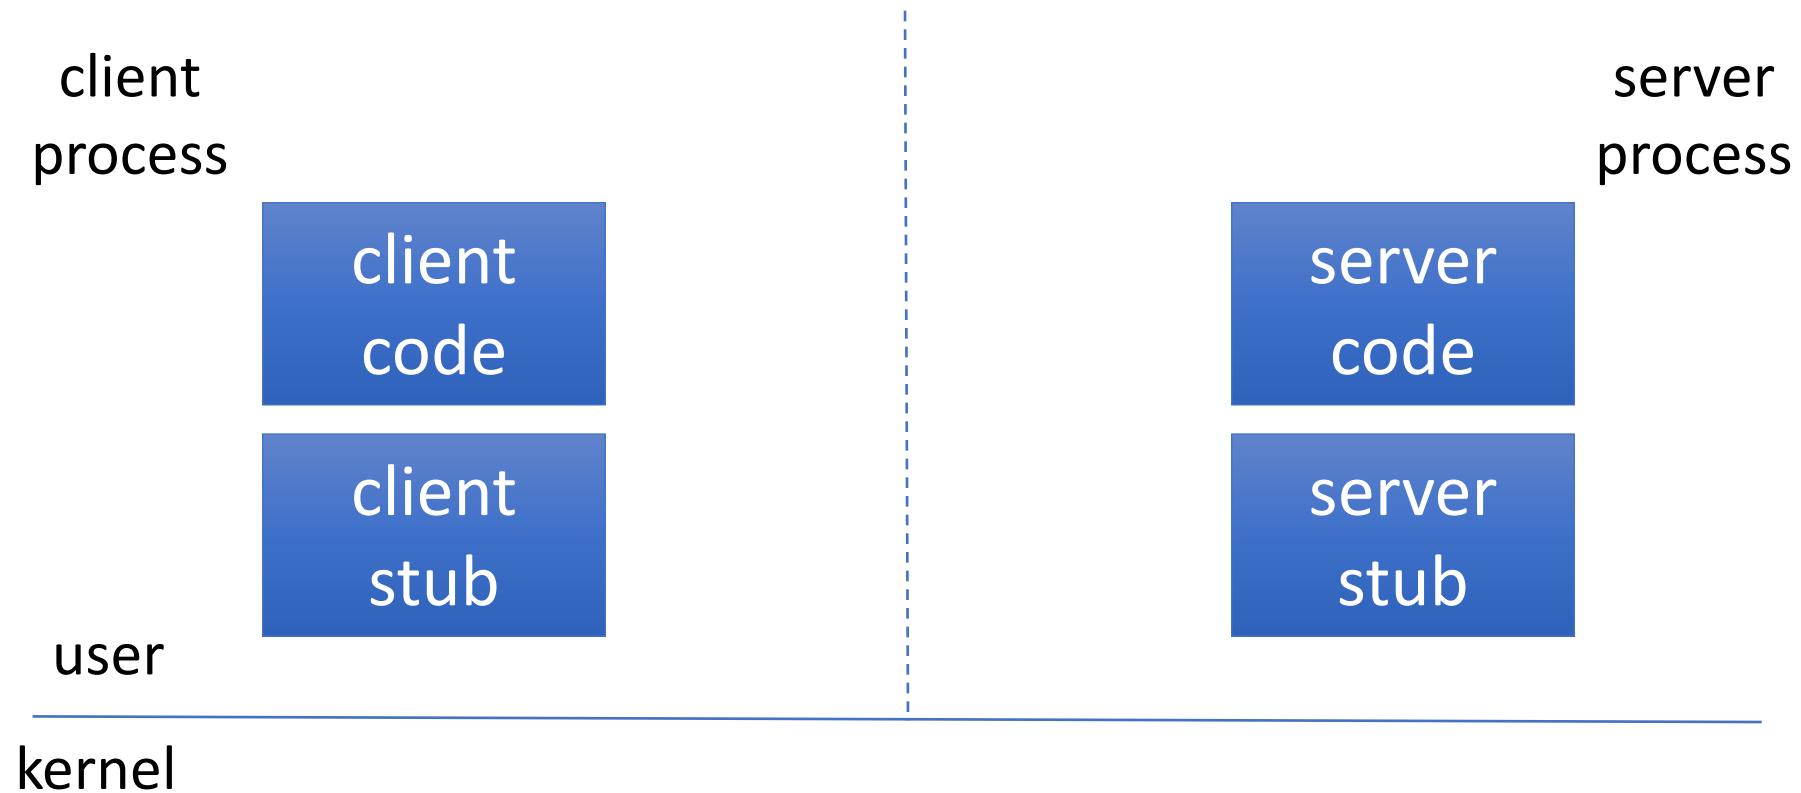
Recap Week 4 RPC Implementation: Call



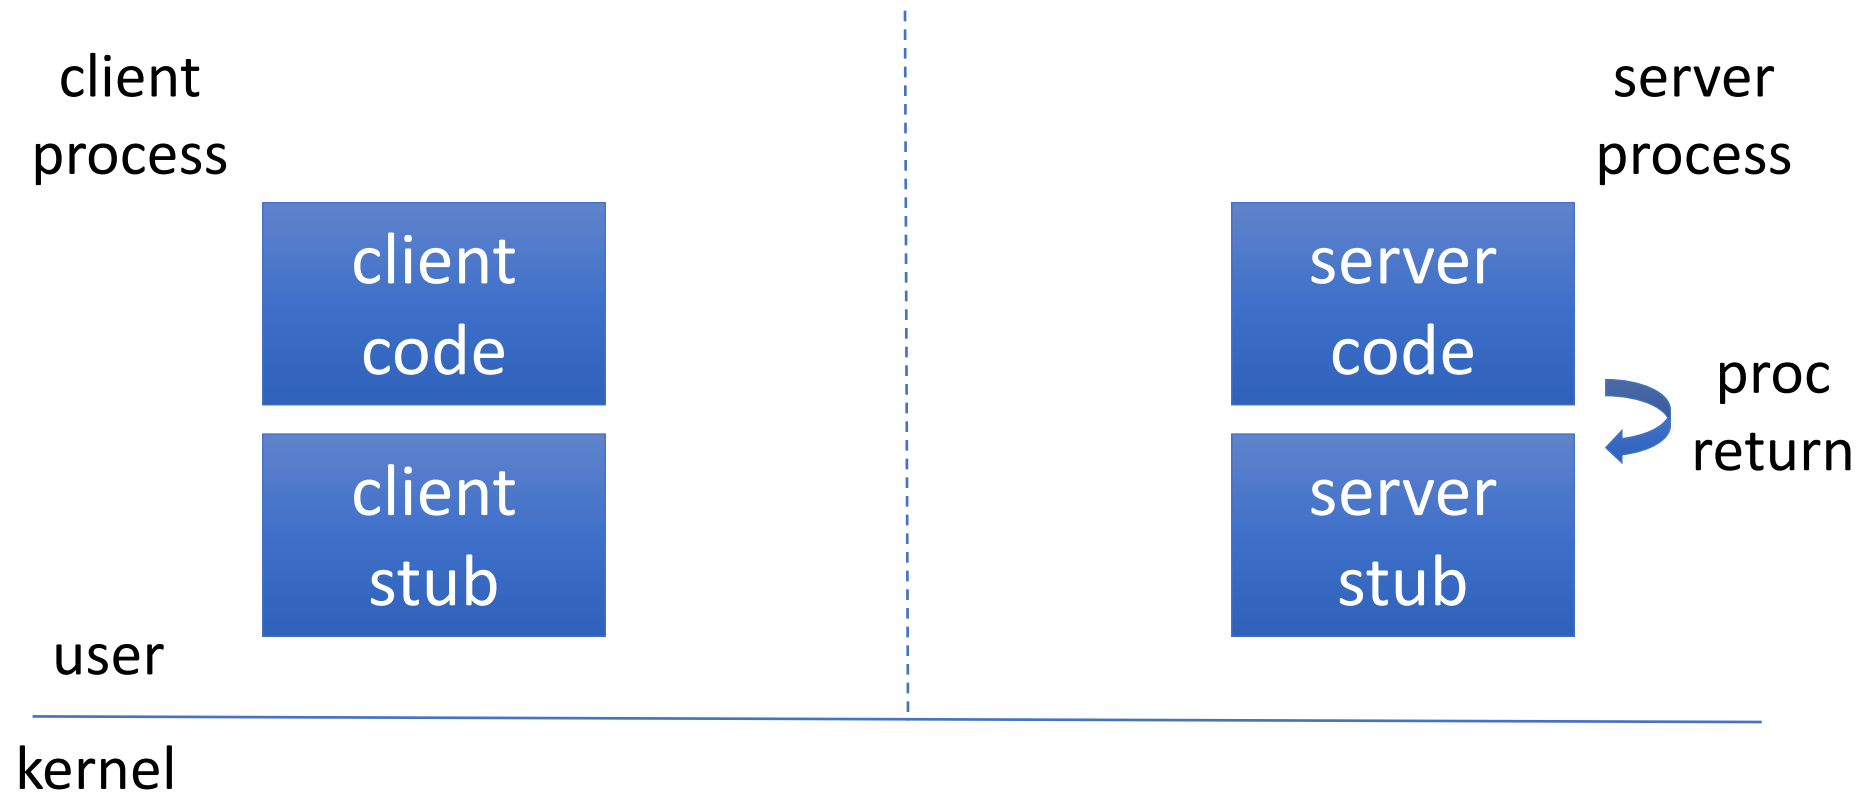
Recap Week 4 RPC Implementation: Call



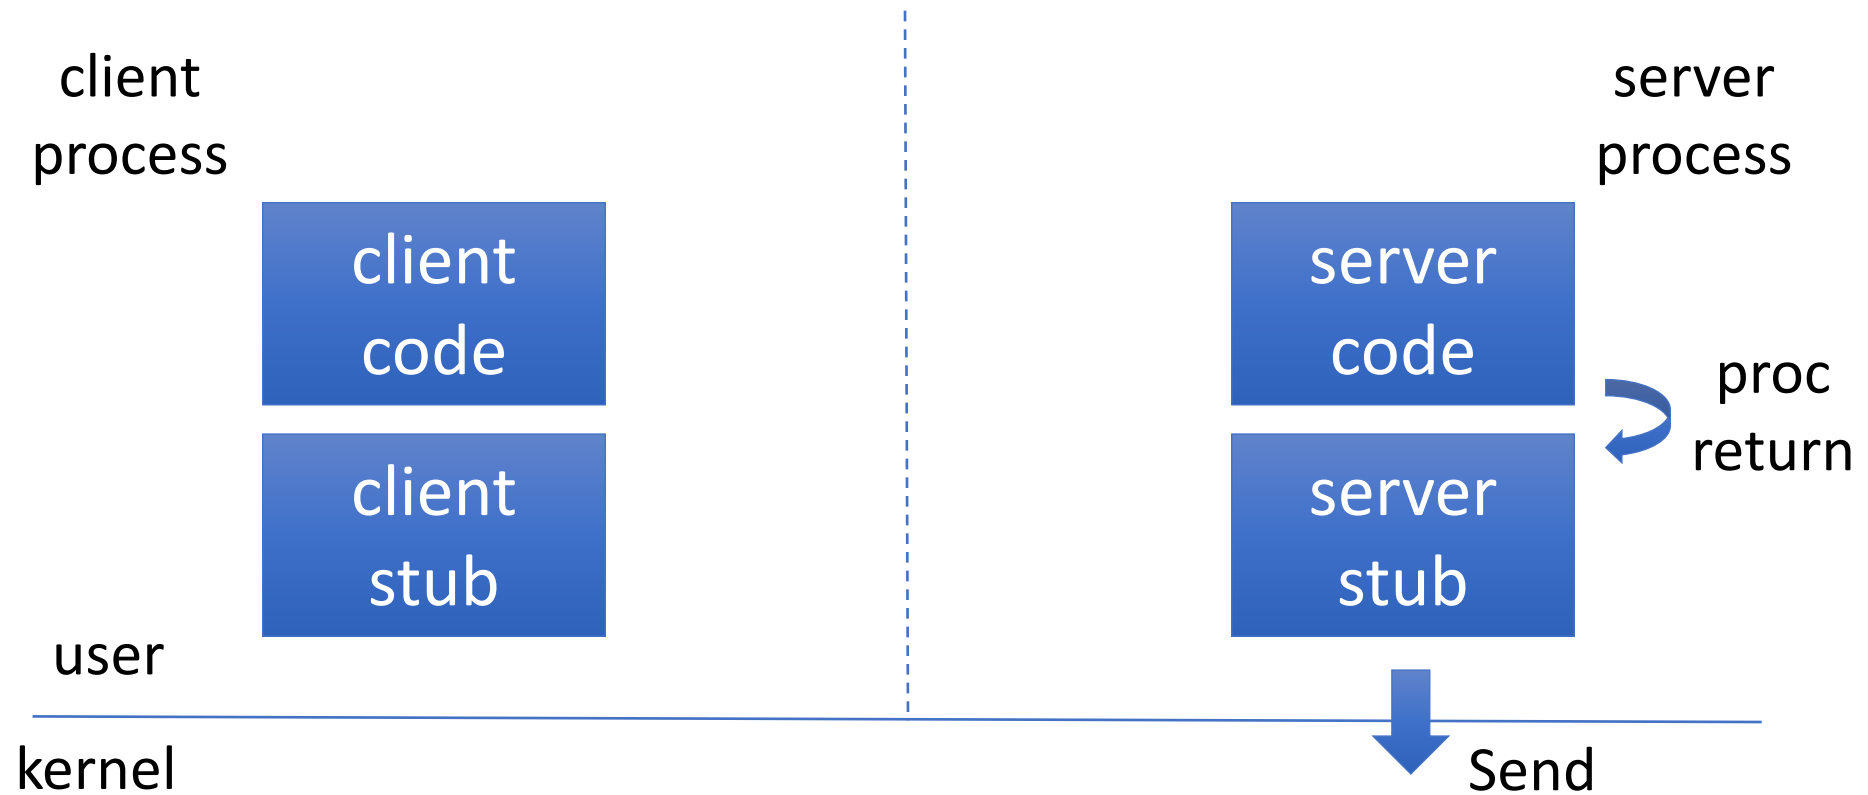
Recap Week 4 RPC Implementation: Return



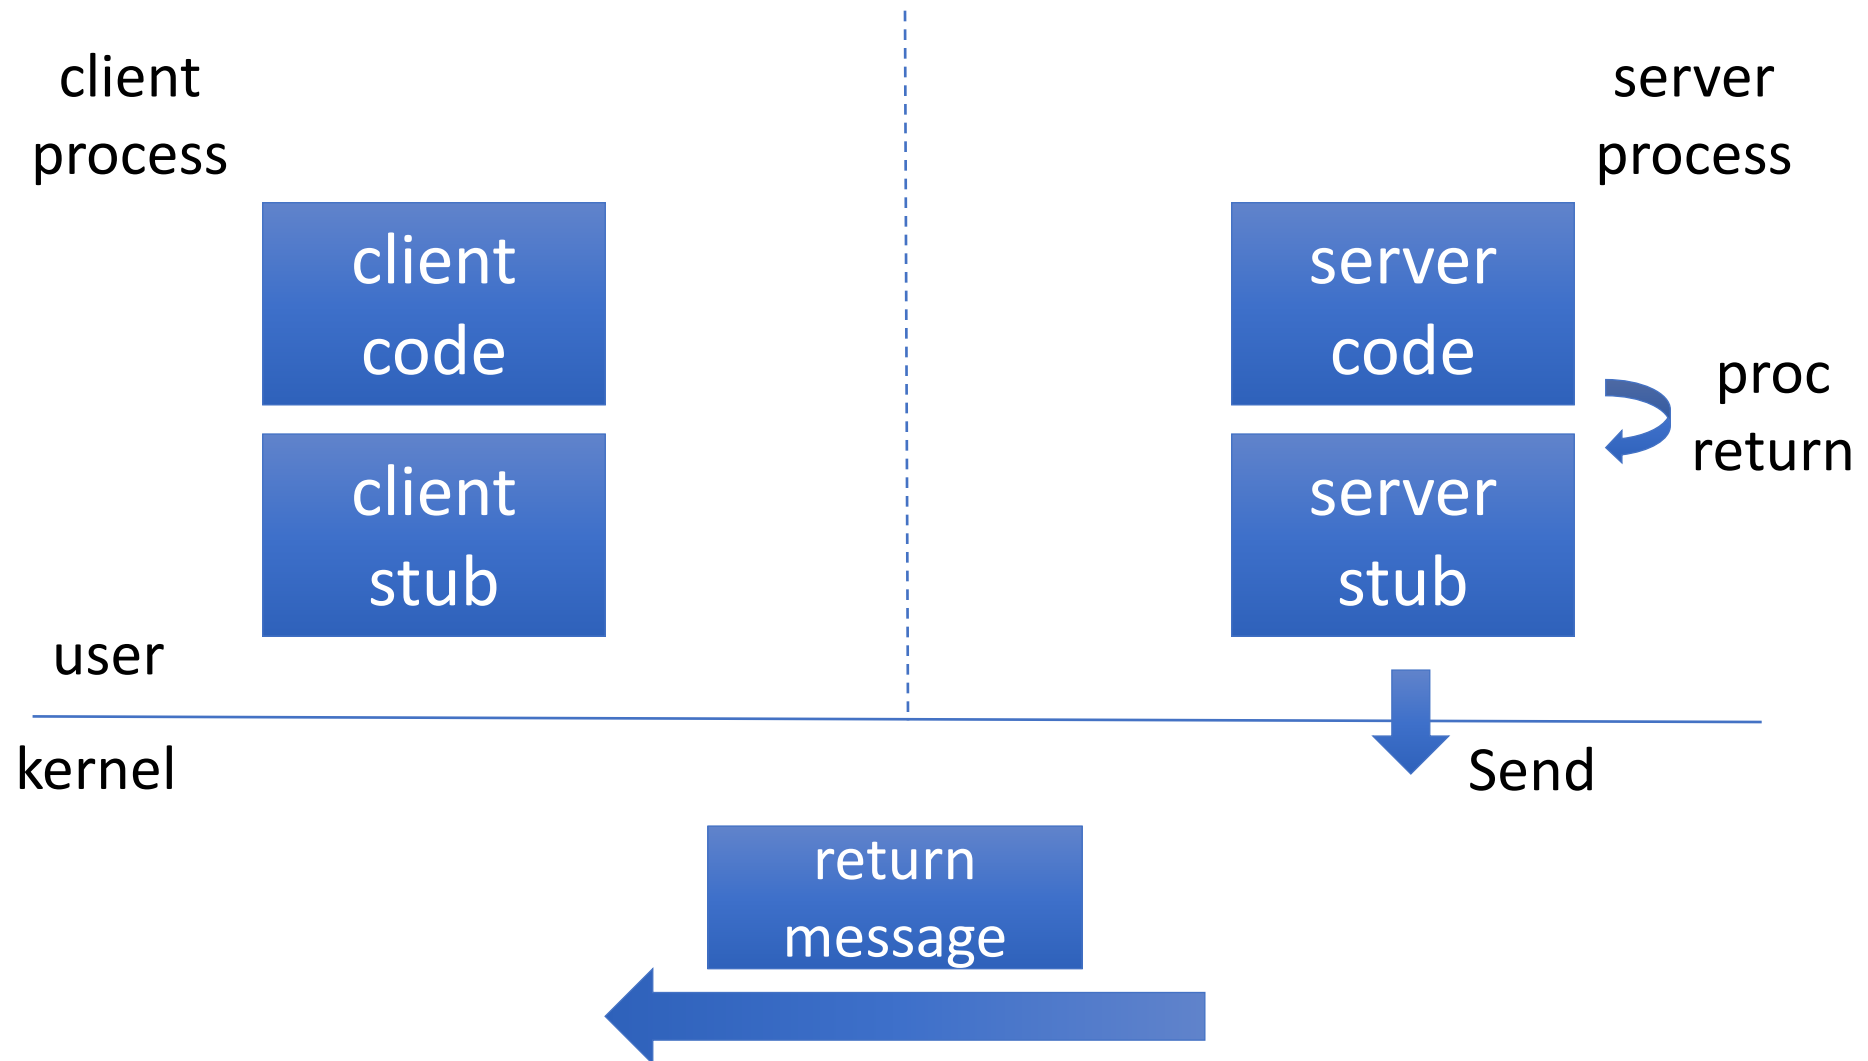
Recap Week 4 RPC Implementation: Return



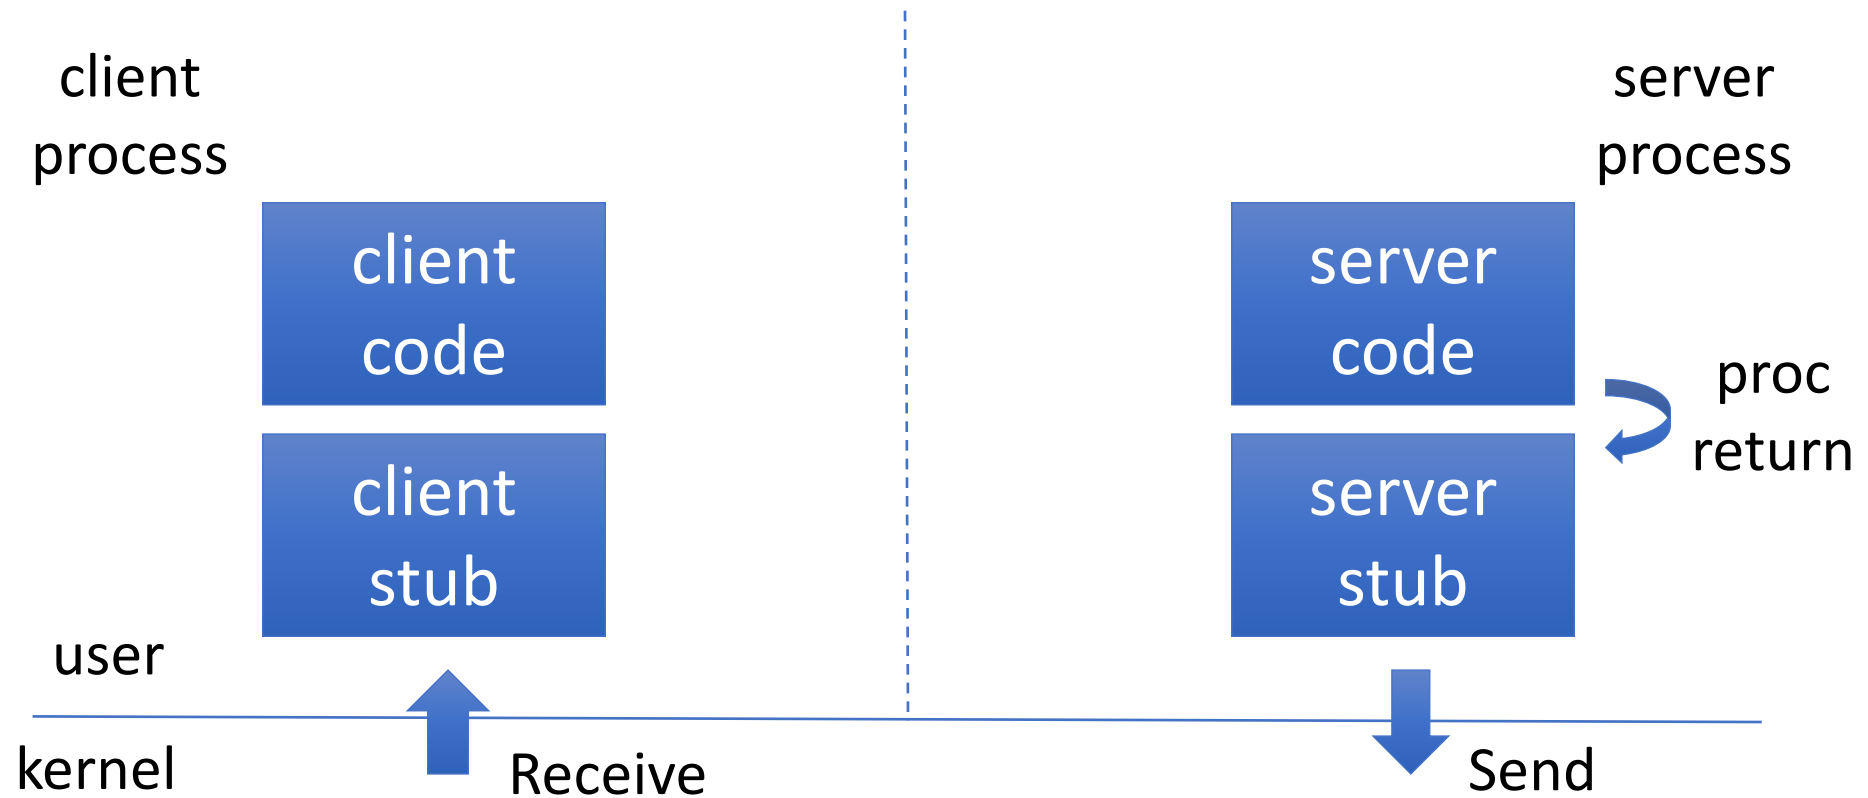
Recap Week 4 RPC Implementation: Return



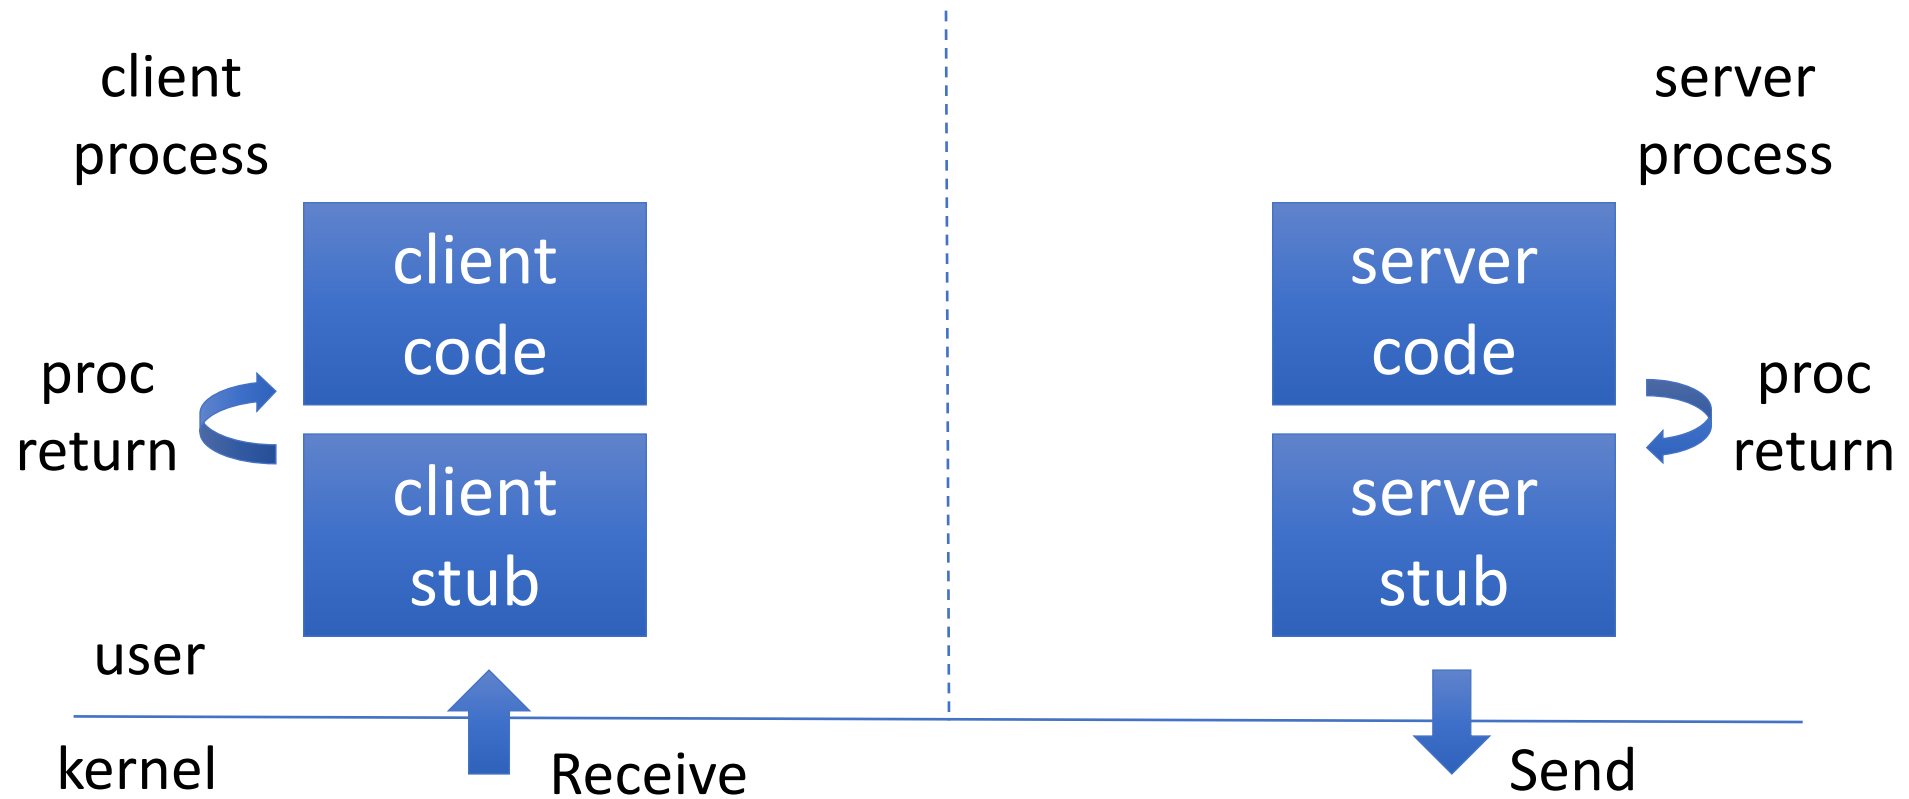
Recap Week 4 RPC Implementation: Return



Recap Week 4 RPC Implementation: Return



Recap Week 4 RPC Implementation: Return

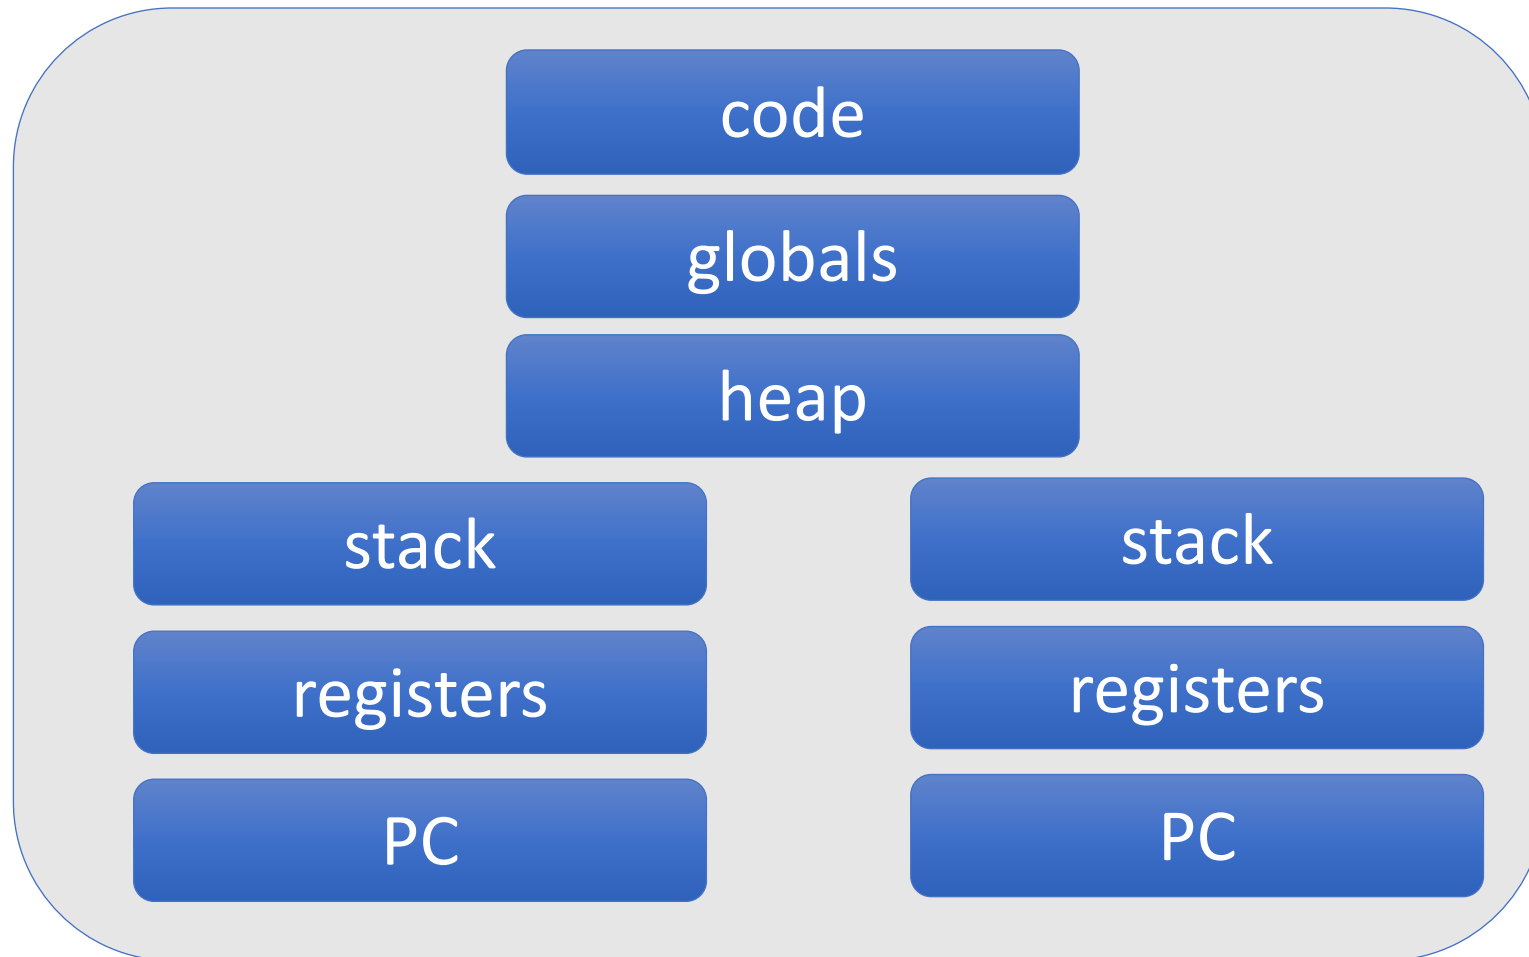


Recap Week 3

Concurrency – Option 2

- New abstraction: **thread**
- Multiple threads in a process
- Threads are like processes except
 - Multiple threads in the same process share an address space
 - Communicate through shared address space
 - If one thread crashes,
 - the entire process, including other threads, crashes

Recap Week 3 Two Threads in a Process



Key Concepts for Today

- Multithreading techniques
 - Division of work
 - Synchronization of shared data
 - Fine-grain locking
 - Privatization
 - Producer/consumer problem

In General

- Processes provide separation
 - In particular, memory separation (no shared data)
 - Suitable for coarse-grain interaction
- Threads do not provide separation
 - In particular, threads share memory (shared data)
 - Suitable for tighter integration

Concrete Example: Web Server

- Serving static content (files)
 - Probably no bugs
 - Can easily be done in a **multithreaded process**
- Serving dynamic (third-party) content
 - No guarantees about bugs
 - Keep in a different process

Shared Data

- Advantage:
 - Many threads can read/write it
- Disadvantage:
 - Many threads can read/write it
 - Can lead to *data races*

Basic Approach to Multithreading

1. Divide “work” among multiple threads &
2. Share data
 - Which data is shared?
 - **Global variables and heap**
 - Not local variables, not read-only variables
 - Where is shared data accessed?
 - Put shared data access in **critical section**

Example: Single-Threaded Code

```
main(){
    int i
    int sum = 0
    int prod = 1
    for( i=0; i<MAX; i++ ){
        c = a[i] * b[i]
        sum += c
        prod *= c
    }
}
```

Approach to Multithreading

- Divide “work” among multiple threads
- Example: give each thread equal number of iterations

Example: Divide Work

```
main() {  
    int i  
    int sum= 0, prod = 1  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_create(...) }  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_join(...) }  
    printf(sum)  
    printf(prod)  
}
```

```
Threadcode() {  
    int i, c  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        sum += c  
        prod *=c  
    }  
}
```

Example: Divide Work

```
main() {  
    int i  
    int sum= 0, prod = 1  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_create(...) }  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_join(...) }  
    printf(sum)  
    printf(prod)  
}
```

```
Threadcode() {  
    int i, c  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        sum += c  
        prod *= c  
    }  
}
```

local data: i, c, my_min, my_max

Shared read-only data: a, b

Shared data: sum, prod

Example: Divide Work

```
main() {  
    int i  
    int sum= 0, prod = 1  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_create(...) }  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_join(...) }  
    printf(sum)  
    printf(prod)  
}
```

```
Threadcode() {  
    int i, c  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        sum += c  
        prod *=c  
    }  
}
```

Protect access to shared data with mutex

Shared data: sum, prod

Example: Divide Work

```
main() {  
    int i  
    int sum= 0, prod = 1  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_create(...) }  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_join(...) }  
    printf(sum)  
    printf(prod)  
}
```

Protection not necessary here because
only the main thread accesses sum, prod

```
Threadcode() {  
    int i, c  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        sum += c  
        prod *= c  
    }  
}
```

Protect access to shared data with mutex

Shared data: sum, prod

Example: Synchronization with 1 mutex

```
Threadcode() {  
    int i  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        pthread_mutex_lock(biglock)  
        sum += c  
        prod *= c  
        pthread_mutex_unlock(biglock)  
    }  
}
```


Why it will not work very well

- Single lock inhibits parallelism
- Two approaches:
 - Fine-grain locking:
 - Multiple locks on individual pieces of shared data
 - Privatization:
 - Make shared data accesses into private data accesses

Fine Grain Locking

- Define separate lock for sum and prod

Example: Finer-Grain Locking

```
Threadcode() {  
    int i, c  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        pthread_mutex_lock(sumlock)  
        sum += c  
        pthread_mutex_unlock(sumlock)  
        pthread_mutex_lock(prodlock)  
        prod *= c  
        pthread_mutex_unlock(prodlock)  
    }  
}
```

Example: Privatization

- Define for each thread
 - A local variable representing its sum
 - A local variable representing its product
- Use those for accesses in the loop
 - Become local accesses
 - No need for lock
- Only access shared data after the loop
 - Use lock there

Example: Privatization

```
Threadcode() {  
    int i, c  
    local_sum = 0  
    local_prod = 1  
  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        local_sum += c  
        local_prod *= c  
    }  
  
    pthread_mutex_lock(sumlock)  
    sum += local_sum  
    pthread_mutex_unlock(sumlock)  
    pthread_mutex_lock(prodlock)  
    prod *= local_prod  
    pthread_mutex_unlock(prodlock)  
}
```

Example: Privatization

```
Threadcode() {  
    int i, c  
    local_sum = 0  
    local_prod = 1  
  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        local_sum += c  
        local_prod *= c  
    }
```

```
    pthread_mutex_lock(sumlock)  
    sum += local_sum  
    pthread_mutex_unlock(sumlock)  
    pthread_mutex_lock(prodlock)  
    prod *= local_prod  
    pthread_mutex_unlock(prodlock)  
}
```

Only **one access** to each lock per thread

Example: Privatization;

Compare to before finer-grained lock accesses

```
Threadcode() {  
    int i, c  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        pthread_mutex_lock(sumlock)  
        sum += c  
        pthread_mutex_unlock(sumlock)  
        pthread_mutex_lock(prodlock)  
        prod *= c  
        pthread_mutex_unlock(prodlock)  
    }  
}
```

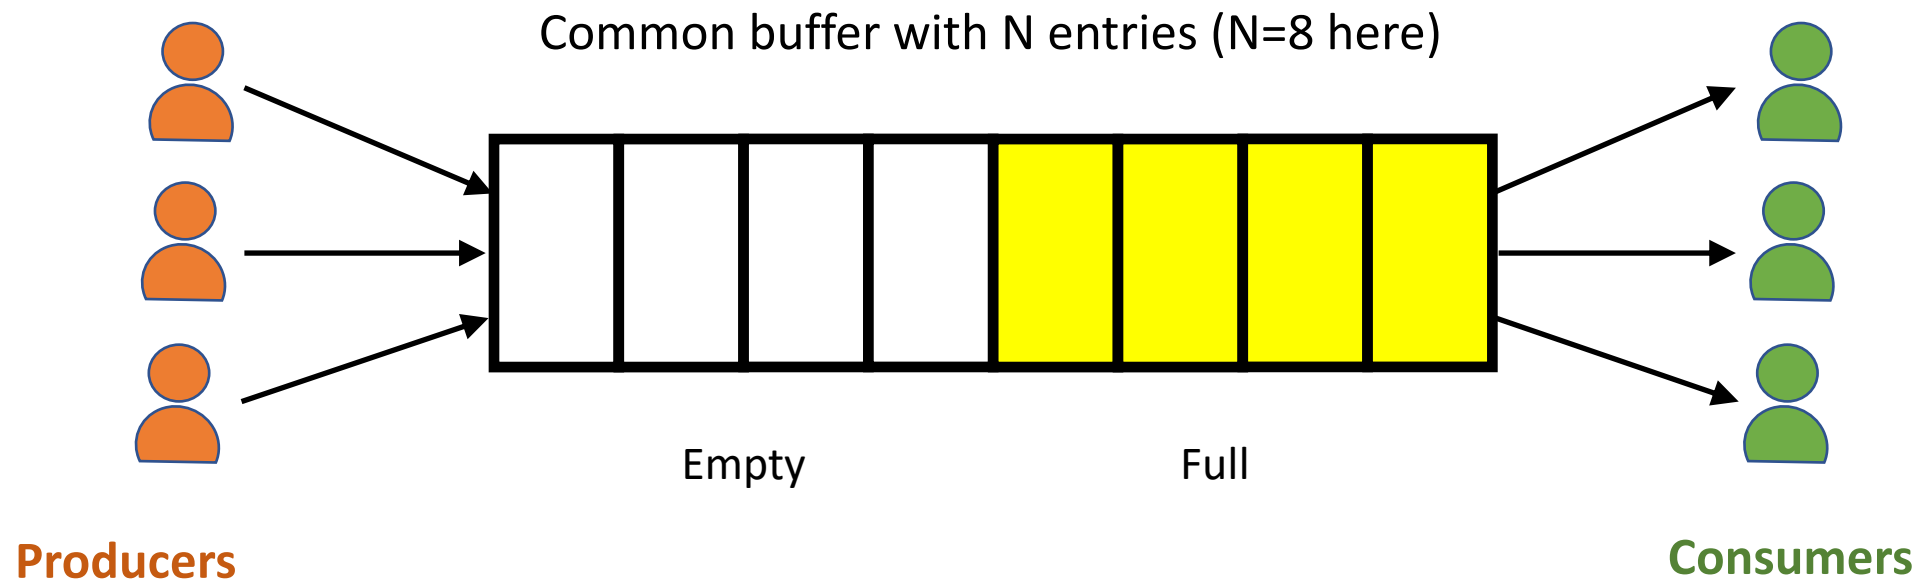
**2 lock accesses
per thread, per iteration**

```
Threadcode() {  
    int i, c  
    local_sum = 0  
    local_prod = 1  
  
    for( i=my_min; i<my_max; i++ ) {  
        c = a[i] * b[i]  
        local_sum += c  
        local_prod *= c  
    }  
  
    pthread_mutex_lock(sumlock)    2 lock accesses  
    sum += local_sum               per thread, in total  
    pthread_mutex_unlock(sumlock)  
    pthread_mutex_lock(prodlock)  
    prod *= local_prod  
    pthread_mutex_unlock(prodlock)  
}
```

Producer/Consumer Problem

Producer/Consumer Problem

- Arises when two or more threads communicate with each other.
 - Some threads “produce” data and other threads “consume” this data.
- Example of producer/consumer in OS: I/O queues

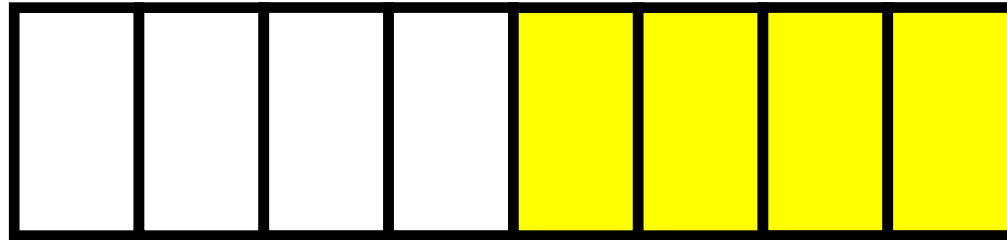
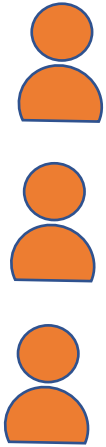


Producer/Consumer Problem

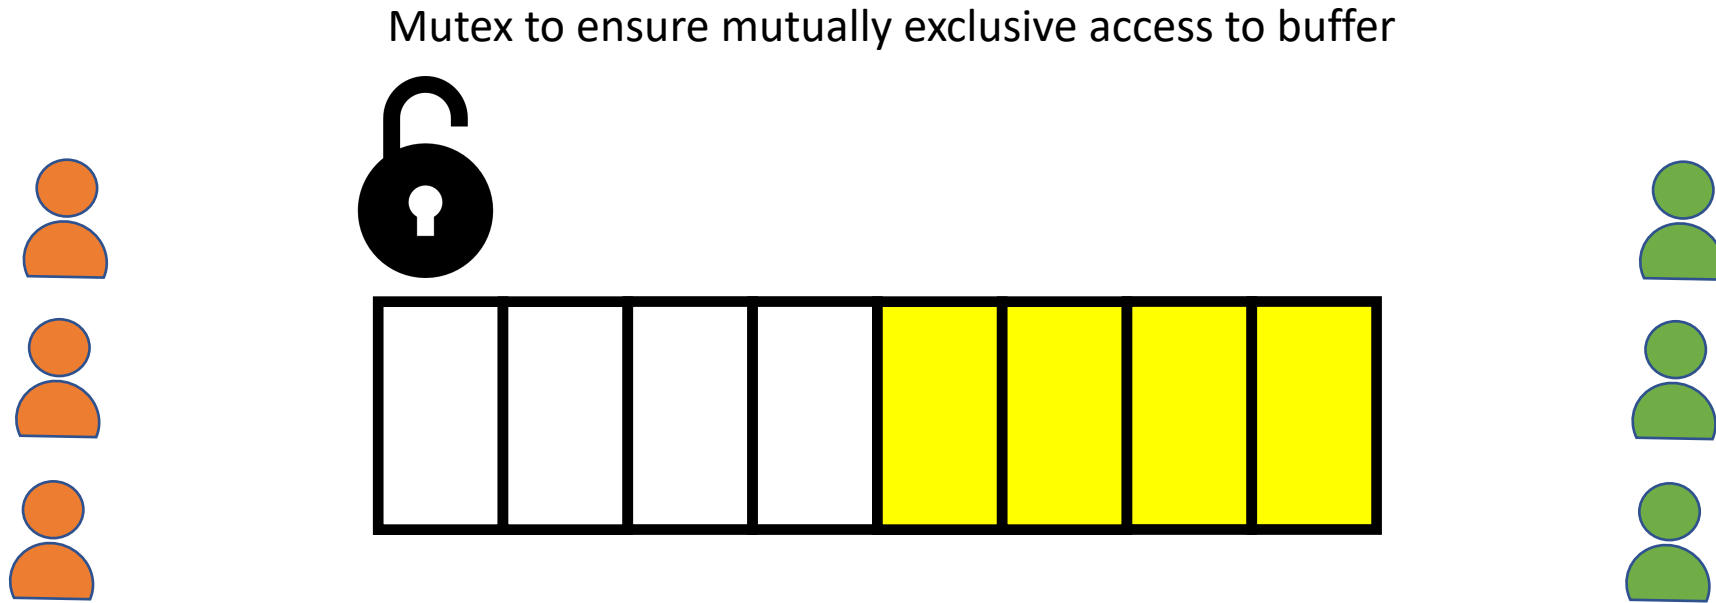
- Multiple producer-threads.
- Multiple consumer-threads.
- One shared bounded buffer with N entries.
- Requirements:
 - No production when all N entries are full.
 - No consumption when no entry is full.
 - Access to the buffer is **mutually exclusive**.

Solve Producer-Consumer with Locks?

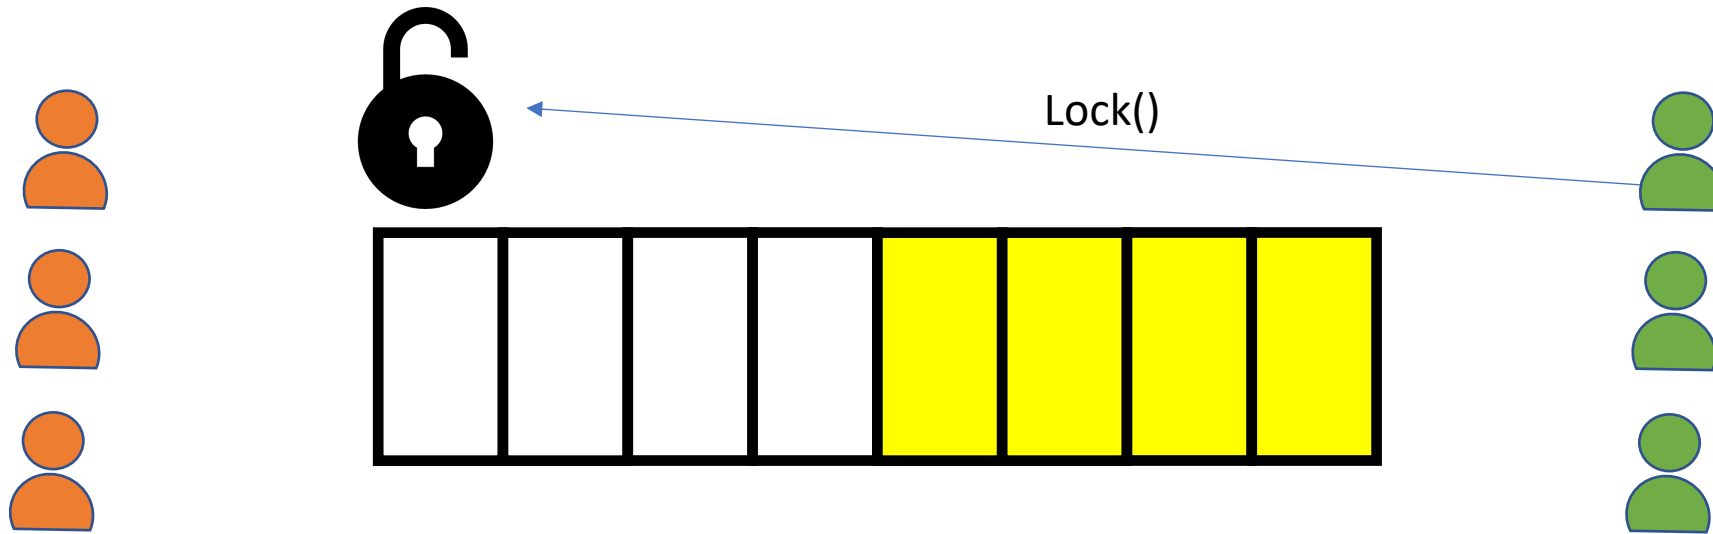
Think!



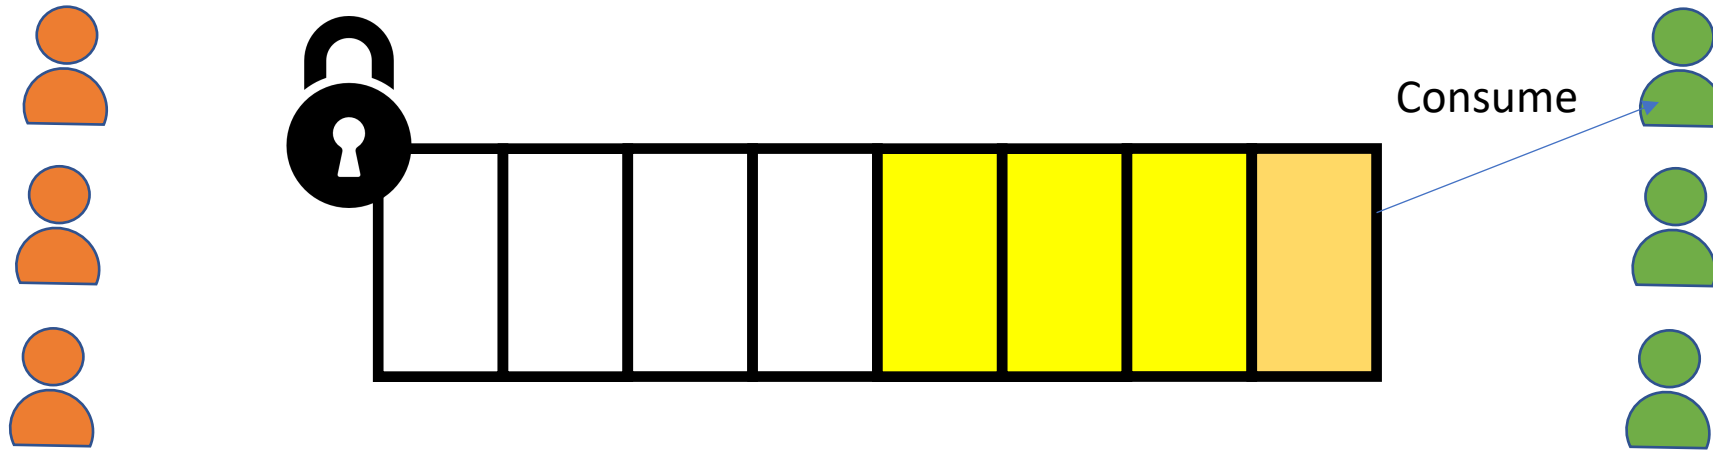
Solve Producer-Consumer with Locks? (Incorrect)



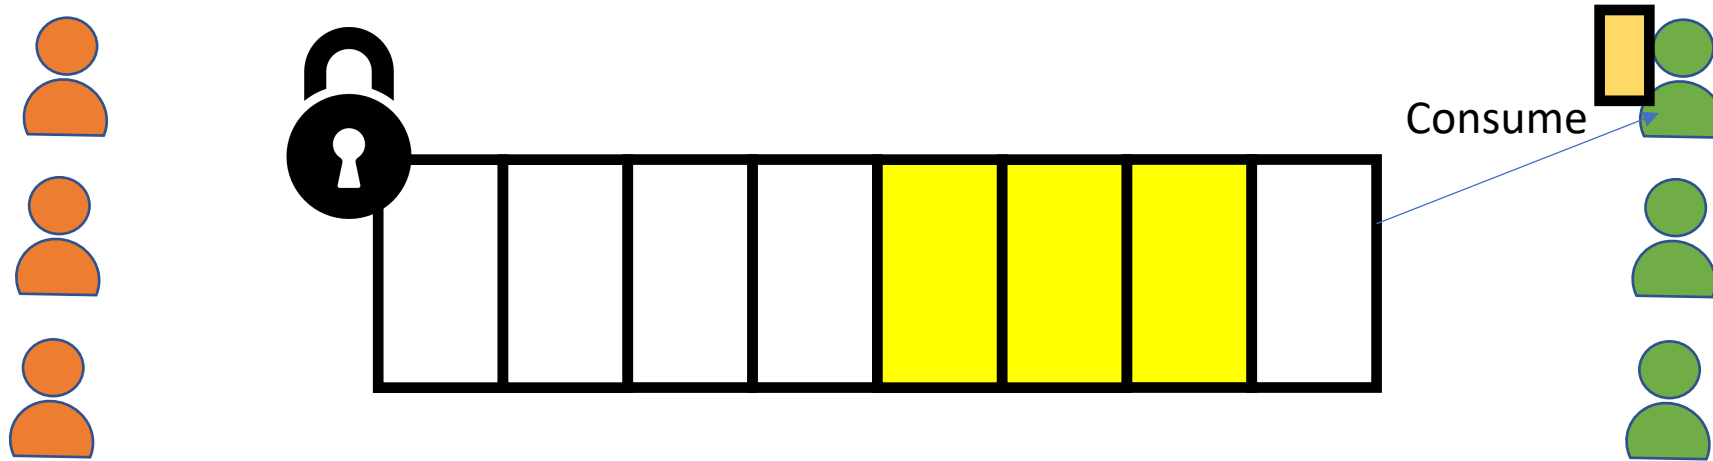
Solve Producer-Consumer with Locks? (Incorrect)



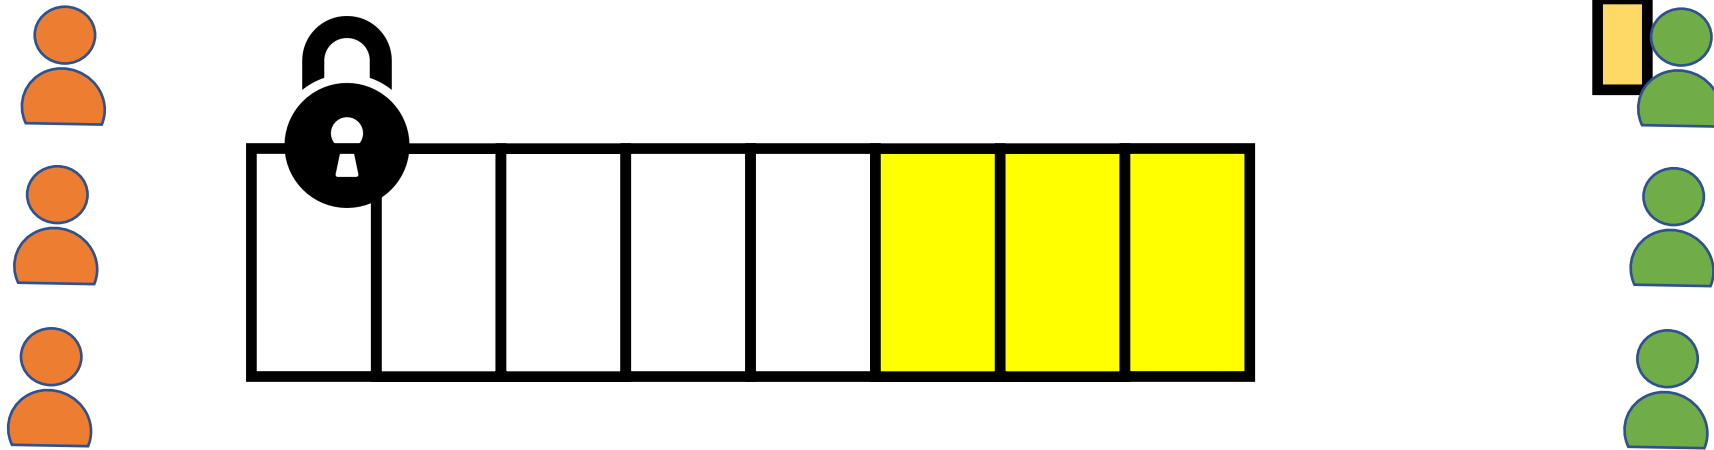
Solve Producer-Consumer with Locks? (Incorrect)



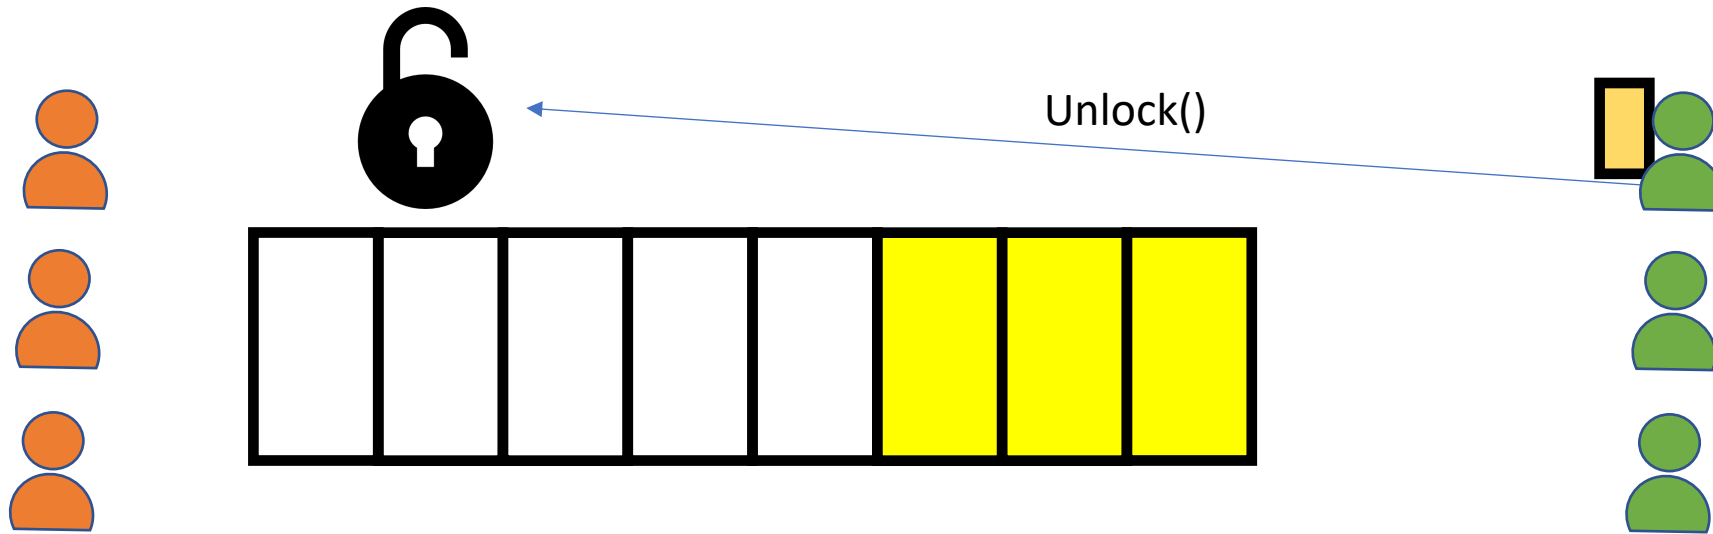
Solve Producer-Consumer with Locks? (Incorrect)



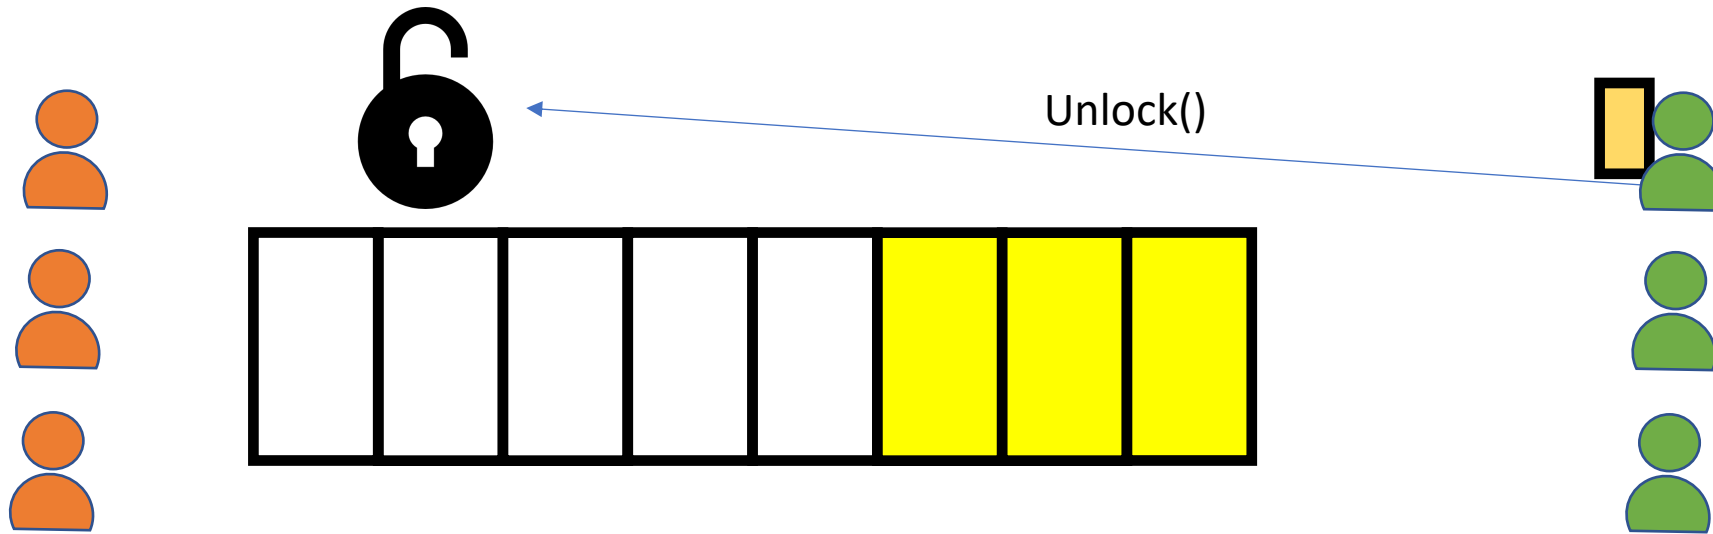
Solve Producer-Consumer with Locks? (Incorrect)



Solve Producer-Consumer with Locks? (Incorrect)



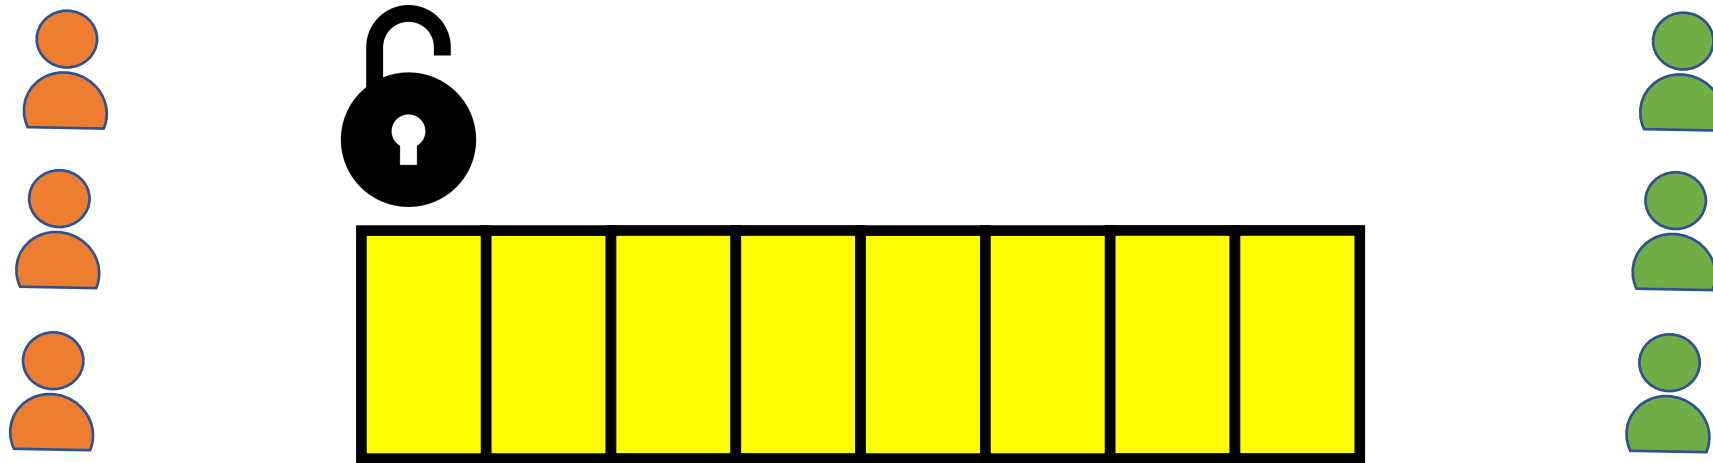
Solve Producer-Consumer with Locks? (Incorrect)



Problem? Think!

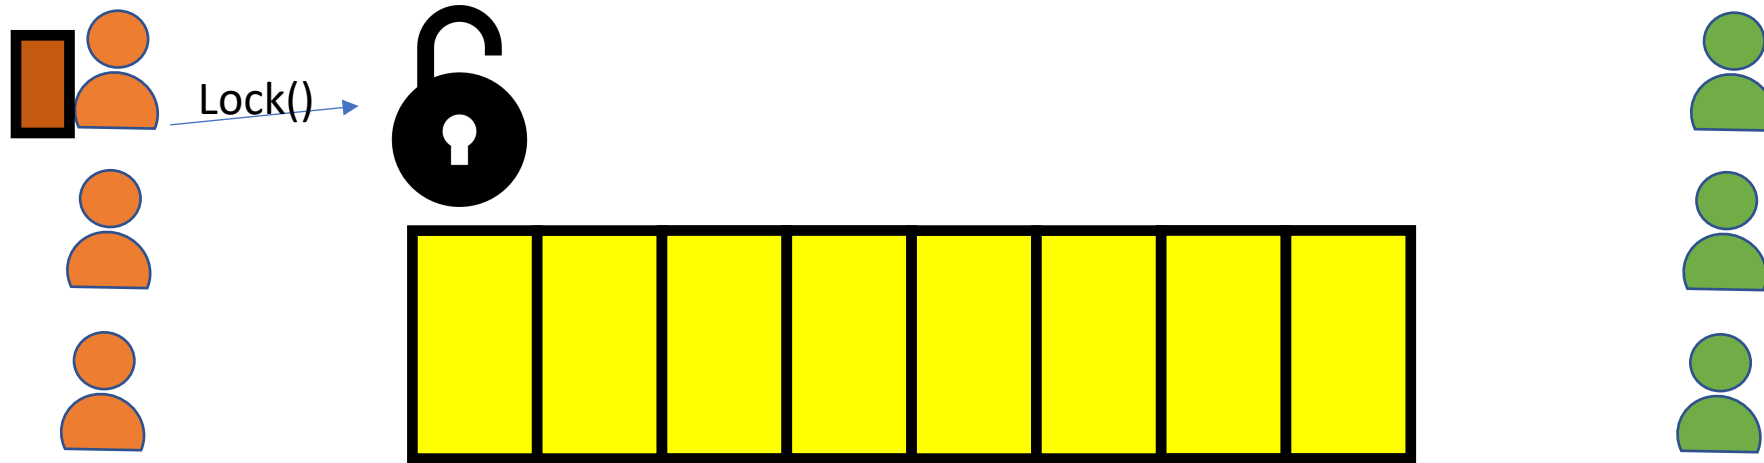
Problem: Threads can't know the state of the buffer before acquiring the lock

- Problem 1: Buffer is full and producer want to add entries



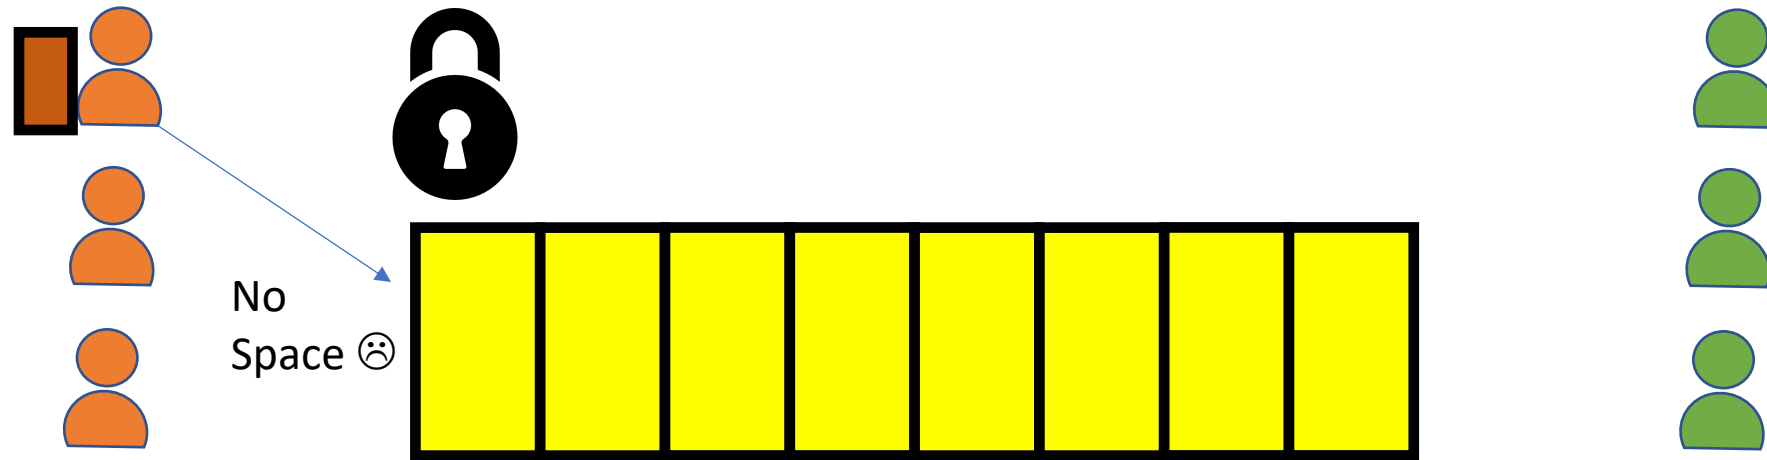
Problem: Threads can't know the state of the buffer before acquiring the lock

- Problem 1: Buffer is full and producer want to add entries



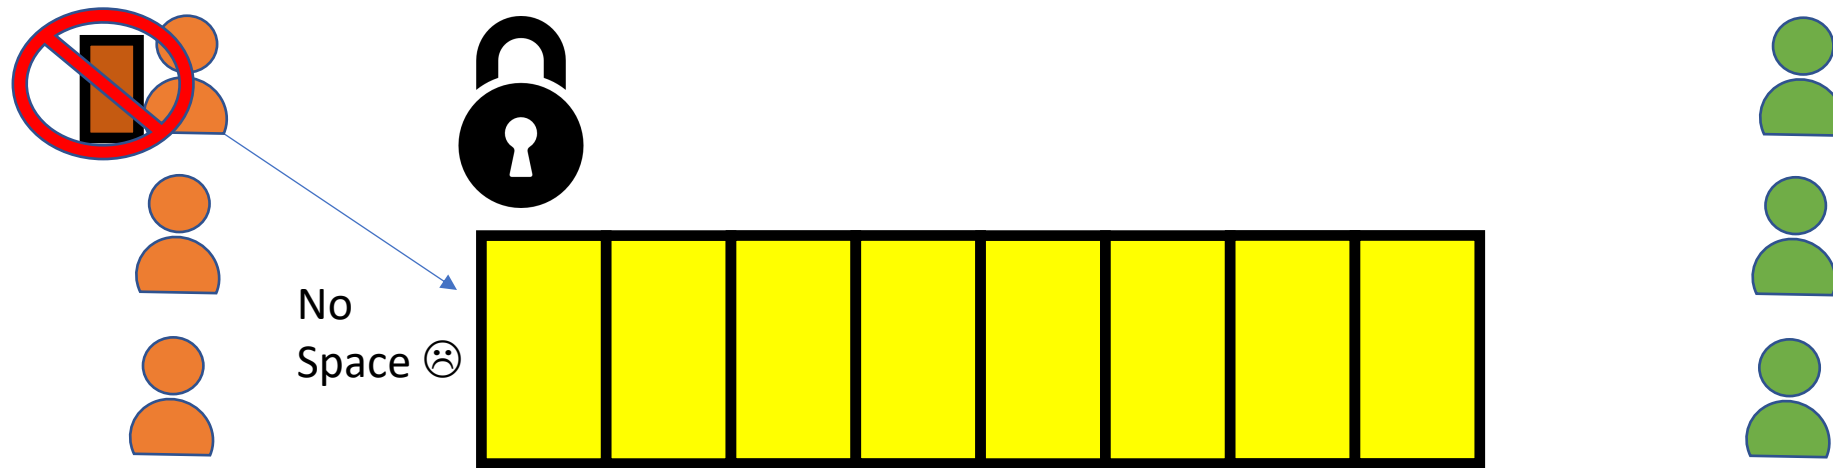
Problem: Threads can't know the state of the buffer before acquiring the lock

- Problem 1: Buffer is full and producer want to add entries



Problem: Threads can't know the state of the buffer before acquiring the lock

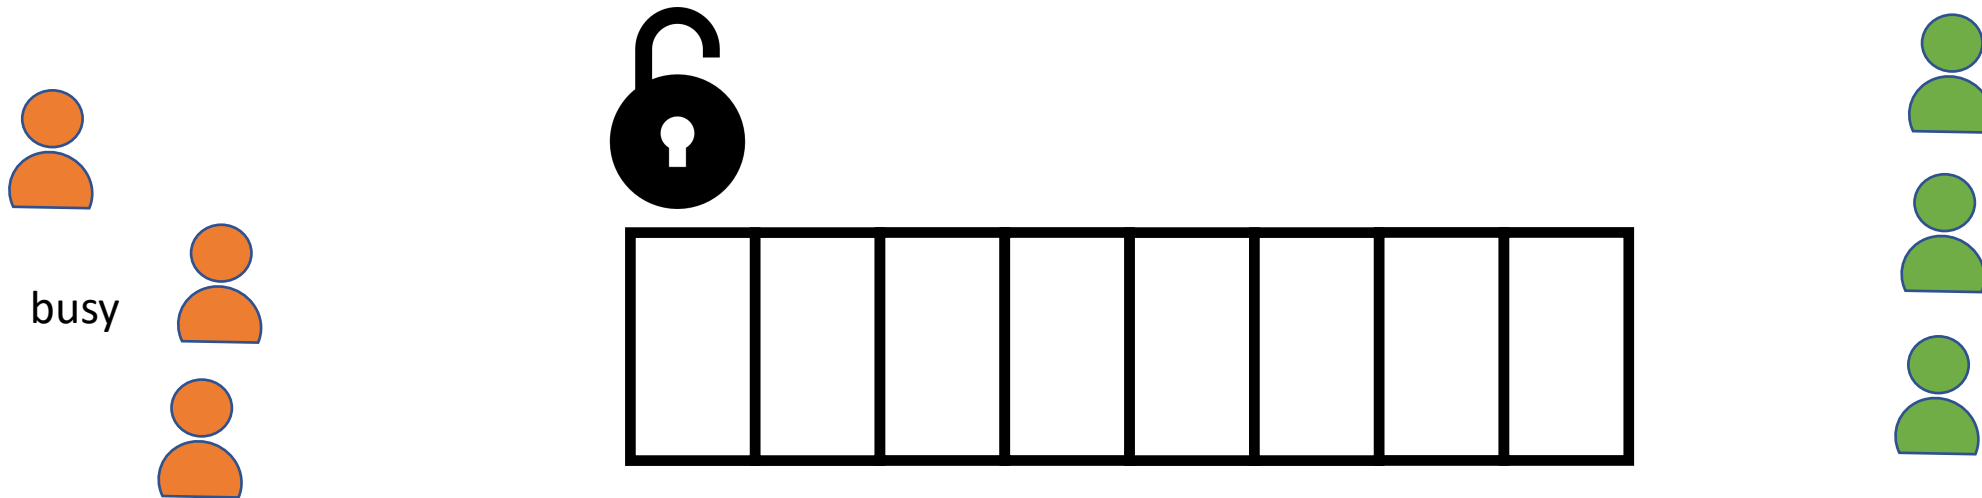
- Problem 1: Buffer is full and producer want to add entries



Packet needs to be dropped and problem requirement is violated
→ "No production when all N entries are full."

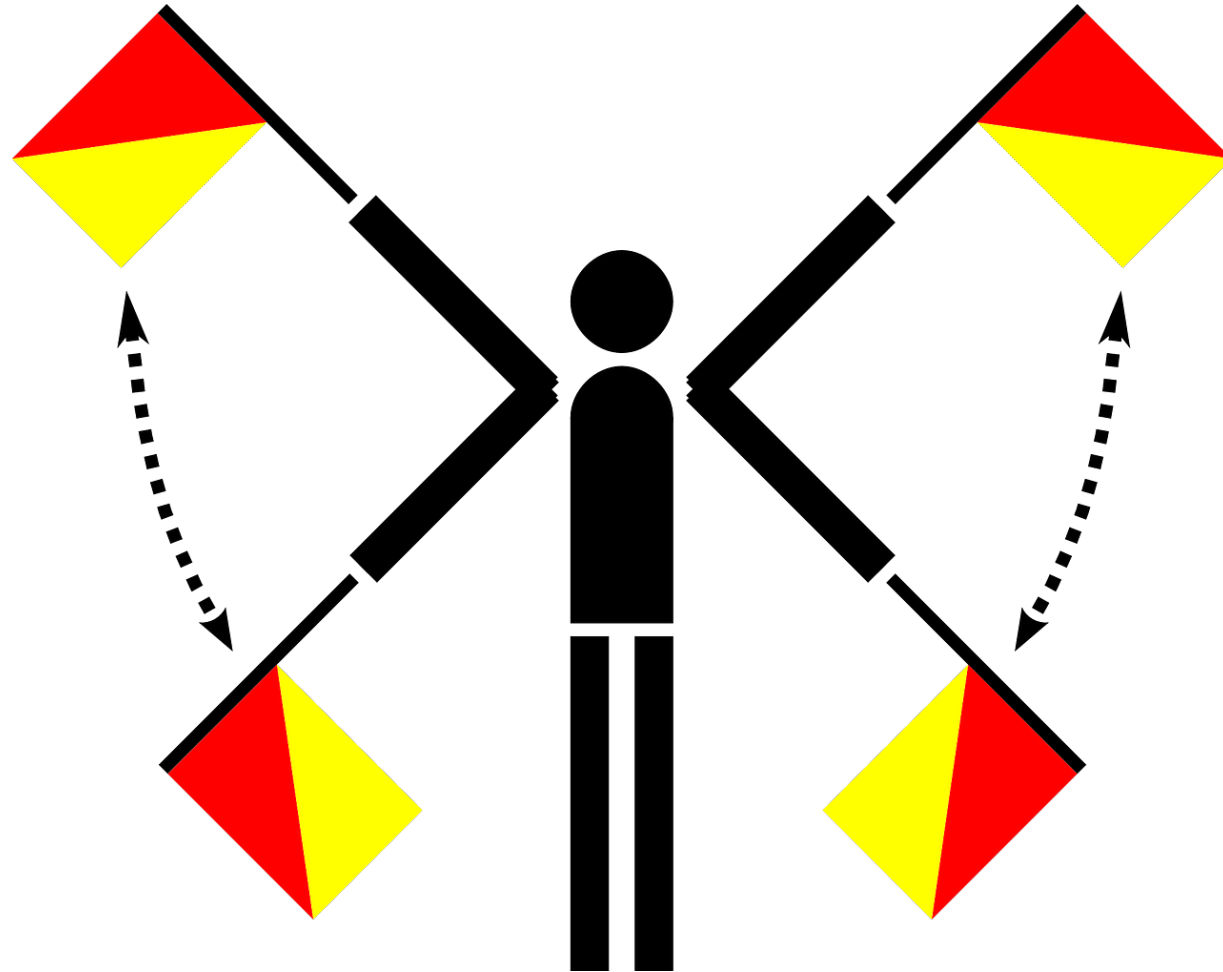
Problem: Threads can't know the state of the buffer before acquiring the lock

- Problem 2: Similarly, buffer is empty, and consumer want to consume

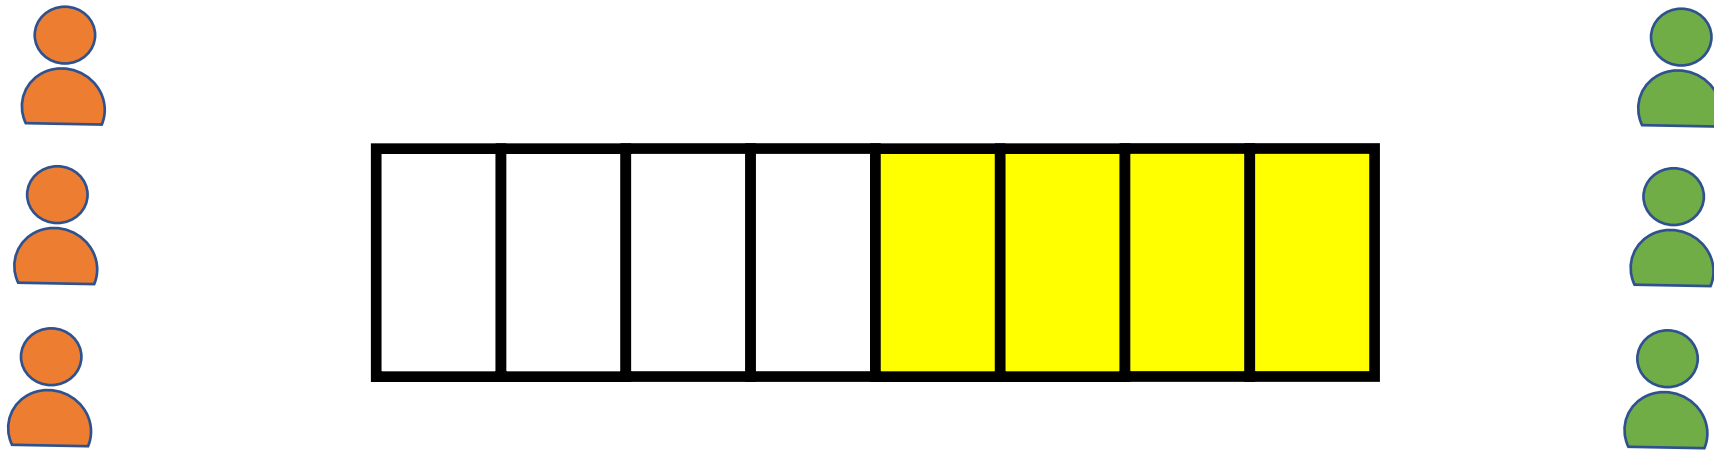


Problem requirement is *not* violated, but inefficient
Consumers waste CPU cycles locking the buffer just to see that it's empty

Solution: Semaphores

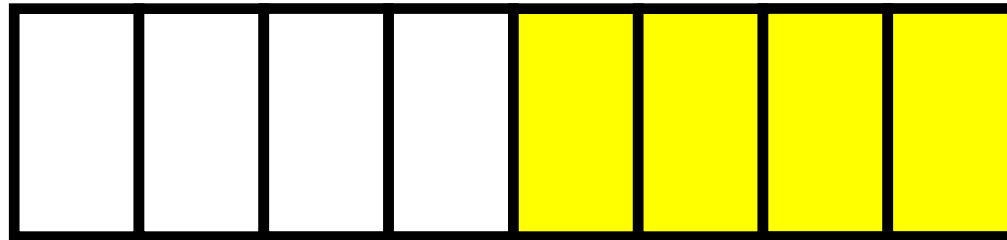


Solve Producer-Consumer with Semaphores?



Solve Producer-Consumer with Semaphores

Step 1: 1 producer, 1 consumer. Think!



Producer-Consumer: Semaphores (Step 1)

Circular Buffer, single producer thread, single consumer thread

- Shared buffer with **N** elements between producer and consumer

Requires 2 semaphores

- **emptyBuffer**: Initialize to ???
- **fullBuffer**: Initialize to ???

Producer-Consumer: Semaphores (Step 1)

Circular Buffer, single producer thread, single consumer thread

- Shared buffer with **N** elements between producer and consumer

Requires 2 semaphores

- **emptyBuffer**: Initialize to **N** \rightarrow N empty buffers; producer can run N times first
- **fullBuffer**: Initialize to **0** \rightarrow 0 full buffers; consumer can run 0 times first

Producer-Consumer: Semaphores (Step 1)

Circular Buffer, single producer thread, single consumer thread

- Shared buffer with **N** elements between producer and consumer

Requires 2 semaphores

- **emptyBuffer**: Initialize to **N** \rightarrow N empty buffers; producer can run N times first
- **fullBuffer**: Initialize to **0** \rightarrow 0 full buffers; consumer can run 0 times first

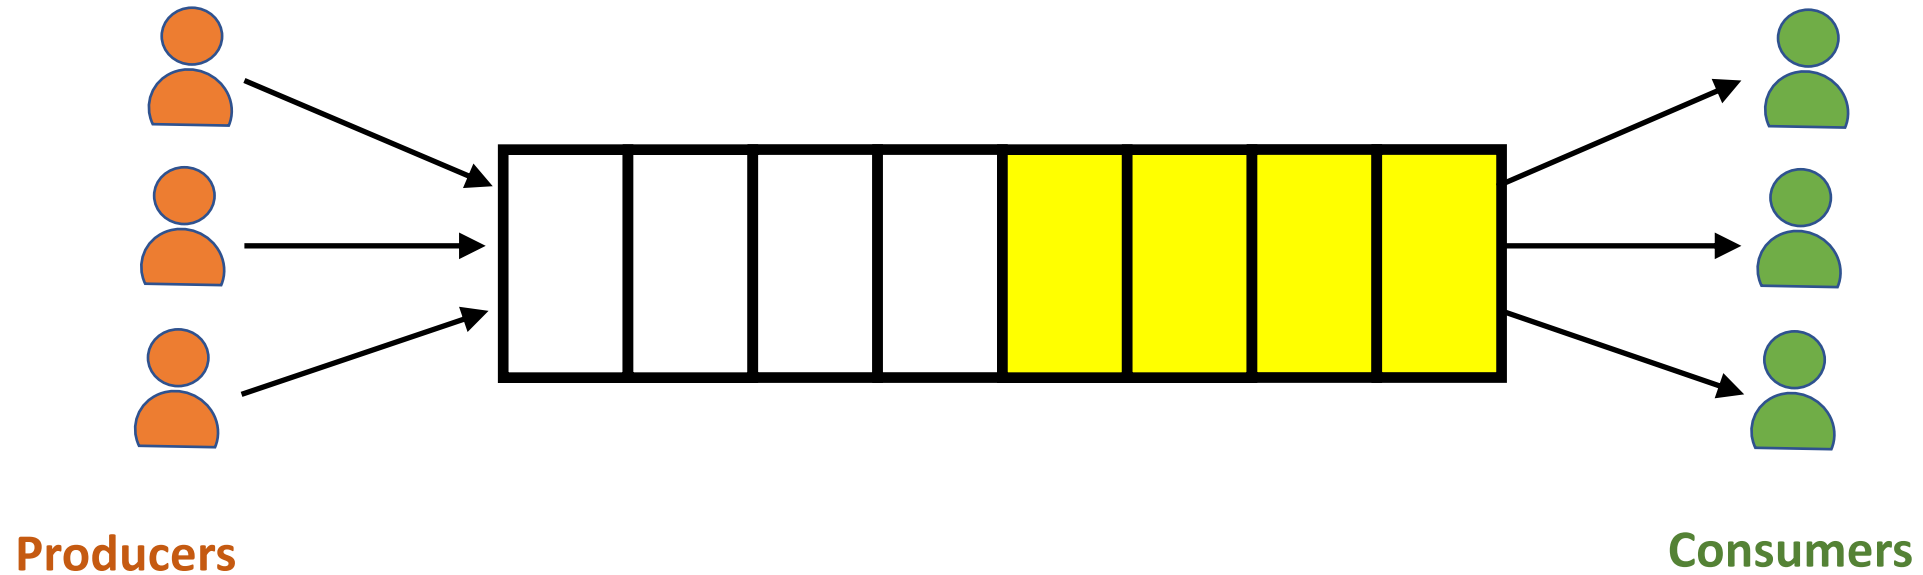
Producer

```
i = 0;
While (1) {
    wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    post(&fullBuffer);
}
```

Consumer

```
j = 0;
While (1) {
    wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    post(&emptyBuffer);
}
```

Multiple producers and multiple consumers



Producer-Consumer: Semaphore (Step 2)

Multiple producer threads, multiple consumer threads

- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will previous code (shown below) not work??? Think!**

Producer

```
i = 0;
While (1) {
    wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    post(&fullBuffer);
}
```

Consumer

```
j = 0;
While (1) {
    wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    post(&emptyBuffer);
}
```

Producer-Consumer: Semaphore (Step 2)

Multiple producer threads, multiple consumer threads

- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will previous code (shown below) not work???**

Producer

```
i = 0;
While (1) {
    wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    post(&fullBuffer);
}
```

Consumer

```
j = 0;
While (1) {
    wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    post(&emptyBuffer);
}
```

Are i and j private or shared? Need each producer to grab unique buffer

Producer-Consumer: Semaphore (Step 3)

Multiple producer threads, multiple consumer threads

- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will the code below still not work??? Think!**

Producer

```
While(1){  
    wait(&emptyBuffer);  
    myi = findempty(&buffer);  
    Fill(&buffer[myi]);  
    post(&fullBuffer);  
}
```

Consumer

```
While(1){  
    wait(&fullBuffer);  
    myj = findfull(&buffer);  
    Use(&buffer[myj]);  
    post(&emptyBuffer);  
}
```

Producer-Consumer: Semaphore (Step 3)

Multiple producer threads, multiple consumer threads

- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will the code below still not work???**

Producer

```
While(1){  
    wait(&emptyBuffer);  
    myi = findempty(&buffer);  
    Fill(&buffer[myi]);  
    post(&fullBuffer);  
}
```

Consumer

```
While(1){  
    wait(&fullBuffer);  
    myj = findfull(&buffer);  
    Use(&buffer[myj]);  
    post(&emptyBuffer);  
}
```

Are myi and myj private or shared? Where is mutual exclusion needed???

Producer-Consumer: Semaphore (Step 4)

Multiple producer threads, multiple consumer threads

- Consider possible locations for mutual exclusion (i.e., equivalent to semaphore initialized to 1)
- **Where is the problem with the code below???**

Producer

```
While(1){  
    lock(&mutex);  
    wait(&emptyBuffer);  
    myi = findempty(&buffer);  
  
    Fill(&buffer[myi]);  
    post(&fullBuffer);  
    unlock(&mutex);  
}
```

Consumer

```
While(1){  
    lock(&mutex);  
    wait(&fullBuffer);  
    myj = findfull(&buffer);  
    Use(&buffer[myj]);  
    post(&emptyBuffer);  
    unlock(&mutex);  
}
```

Producer-Consumer: Semaphore (Step 4)

Multiple producer threads, multiple consumer threads

- Consider possible locations for mutual exclusion (i.e., equivalent to semaphore initialized to 1)
- **Where is the problem with the code below???**

Producer

```
While(1){  
    lock(&mutex);  
    wait(&emptyBuffer);  
    myi = findempty(&buffer);  
  
    Fill(&buffer[myi]);  
    post(&fullBuffer);  
    unlock(&mutex);  
}
```

Consumer

```
While(1){  
    lock(&mutex);  
    wait(&fullBuffer);  
    myj = findfull(&buffer);  
    Use(&buffer[myj]);  
    post(&emptyBuffer);  
    unlock(&mutex);  
}
```

Problem: Deadlock at mutex (e.g., consumer runs first; won't release mutex)

Producer-Consumer Final Solution

Multiple producer threads, multiple consumer threads

- Consider possible locations for mutual exclusion (i.e., equivalent to semaphore initialized to 1)

Producer

```
While(1){  
    wait(&emptyBuffer);  
    lock(&mutex);  
    myi = findempty(&buffer);  
  
    Fill(&buffer[myi]);  
    unlock(&mutex);  
    post(&fullBuffer);  
}
```

Consumer

```
While(1){  
    wait(&fullBuffer);  
    lock(&mutex);  
    myj = findfull(&buffer);  
  
    Use(&buffer[myj]);  
    unlock(&mutex);  
    post(&emptyBuffer);  
}
```

Producer-Consumer Final Solution

Multiple producer threads, multiple consumer threads

- Consider possible locations for mutual exclusion (i.e., equivalent to semaphore initialized to 1)

Producer

```
While(1){  
    wait(&emptyBuffer);  
    lock(&mutex);  
    myi = findempty(&buffer);  
  
    Fill(&buffer[myi]);  
    unlock(&mutex);  
    post(&fullBuffer);  
}
```

Consumer

```
While(1){  
    wait(&fullBuffer);  
    lock(&mutex);  
    myj = findfull(&buffer);  
  
    Use(&buffer[myj]);  
    unlock(&mutex);  
    post(&emptyBuffer);  
}
```

Finally, works! 😊

But limits concurrency. Only 1 thread at a time can be using or filling different buffers

Let's practice:

Multithreaded Web Server

Let's practice:

Multithreaded Web Server

```
ListenerThread {  
    forever {  
        Receive( request )  
        pthread_create(...)  
    }  
}
```

```
WorkerThread( request ) {  
    read file from disk  
    Send( response )  
    pthread_exit()  
}
```

Note that clients are still
in separate processes

Shared Data?

- There is none!
- Process creation serves as synchronization

Multithreaded Web Server with Thread Pool

```
ListenerThread {  
    for( i=0; i<MAX_THREADS; i++ ) { Pthread_create(...) }  
    forever {  
        Receive( request )  
        hand request to thread[?]  
    }  
}  
  
WorkerThread[?] {  
    forever {  
        wait for available request  
        read file from disk  
        Send( reply )  
    }  
}
```

Shared Data?

- We need to create shared data
- Going to be some kind of queue
- Put lock/unlock around it

Multithreaded Web Server with Thread Pool (incorrect)

```
ListenerThread {  
    for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)  
    forever {  
        Receive( request )  
        pthread_mutex_lock( queuelock )  
        put request in queue  
        pthread_mutex_unlock( queuelock )  
    }  
}  
  
WorkerThread {  
    forever {  
        pthread_mutex_lock( queuelock )  
        take request out of queue  
        pthread_mutex_unlock( queuelock )  
        read file from disk  
        Send( reply )  
    }  
}
```

It will not work

- You need to tell worker(s) there is something for them to do (i.e., in the queue)
- Producer/consumer problem

Multithreaded Web Server with Thread Pool (incorrect)

```
ListenerThread {
    for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)
    forever {
        Receive( request )
        pthread_mutex_lock( queuelock )
        put request in queue
        pthread_cond_signal( notempty, queuelock )
        pthread_mutex_unlock( queuelock )
    }
}

WorkerThread {
    forever {
        pthread_mutex_lock( queuelock )
        pthread_cond_wait( notempty, queuelock )
        take request out of queue
        pthread_mutex_unlock( queuelock )
        read file from disk
        Send( reply )
    }
}
```

Incorrect

- All worker threads busy (none waiting)
- Listener does a signal
- No thread waiting: signal is no-op
- Worker thread finishes what it was doing
 - Will do a wait
 - Although request is waiting in queue

In General

- Signals have no memory
- Forgotten if no thread waiting
- So need an extra variable to remember them

Multithreaded Web Server with Thread Pool

```
ListenerThread {
    for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)
    forever {
        Receive( request )
        pthread_mutex_lock( queuelock )
        put request in queue
        avail++
        pthread_cond_signal( notempty, queuelock )
        pthread_mutex_unlock( queuelock )
    }
}
WorkerThread {
    forever {
        pthread_mutex_lock( queuelock )
        if( avail <= 0 ) pthread_cond_wait( notempty, queuelock )
        take request out of queue
        avail--
        pthread_mutex_unlock( queuelock )
        read file from disk
        Send( reply )
    }
}
```

Note

- Should now be clear why mutex must be held
- Avail is a shared data item
- Without mutex could have data race

Imagine Solution Without Locks

```
ListenerThread {
  for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)
  forever {
    Receive( request )
    pthread_mutex_lock( queuelock )
    put request in queue
    avail++
    pthread_cond_signal( notempty, queuelock )
    pthread_mutex_unlock( queuelock )
  }
}
WorkerThread {
  forever {
    pthread_mutex_lock( queuelock )
    if( avail <= 0 ) pthread_cond_wait( notempty, queuelock )
    take request out of queue
    avail--
    pthread_mutex_unlock( queuelock )
    read file from disk
    Send( reply )
  }
}
```

Imagine Solution Without Locks

```
ListenerThread {  
  for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)  
  forever {  
    Receive( request )  
    pthread_mutex_lock( queuelock )  
    put request in queue  
    avail++  
    pthread_cond_signal( notempty, queuelock )  
    pthread_mutex_unlock( queuelock )  
  }  
}  
WorkerThread {  
  forever {  
    pthread_mutex_lock( queuelock )  
    if( avail <= 0 ) pthread_cond_wait( notempty, queuelock )  
    take request out of queue  
    avail--  
    pthread_mutex_unlock( queuelock )  
    read file from disk  
    Send( reply )  
  }  
}
```

Worker checks avail and finds it to be 0
Worker interrupted and listener runs
Listener sets avail to 1 and signals
No thread is waiting, so signal is no-op
Listener interrupted and worker runs
Worker does a wait
Incorrect: worker waits with request in queue

Example incorrect execution: One Worker Thread

- Worker checks avail and finds it to be 0
- Worker interrupted and listener runs
- Listener sets avail to 1 and signals
- No thread is waiting, so signal is no-op
- Listener interrupted and worker runs
- Worker does a wait
- Incorrect: worker waits with request in queue

Back to Solution with Locks (correct)

```
ListenerThread {
    for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)
    forever {
        Receive( request )
        pthread_mutex_lock( queuelock )
        put request in queue
        avail++
        pthread_cond_signal( notempty, queuelock )
        pthread_mutex_unlock( queuelock )
    }
}
WorkerThread {
    forever {
        pthread_mutex_lock( queuelock )
        if( avail <= 0 ) pthread_cond_wait( notempty, queuelock )
        take request out of queue
        avail--
        pthread_mutex_unlock( queuelock )
        read file from disk
        Send( reply )
    }
}
```

Back to Solution with Locks

```
ListenerThread {
  for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)
  forever {
    Receive( request )
    pthread_mutex_lock( queuelock )
    put request in queue
    avail++
    pthread_cond_signal( notempty, queuelock )
    pthread_mutex_unlock( queuelock )
  }
}
WorkerThread {
  forever {
    pthread_mutex_lock( queuelock )
    if( avail <= 0 ) pthread_cond_wait( notempty, queuelock )
    take request out of queue
    avail--
    pthread_mutex_unlock( queuelock )
    read file from disk
    Send( reply )
  }
}
```

Can this execution happen?

Queue is empty, worker thread W1 waits

Listener thread L puts request in queue

Sets avail to 1

Signals

W1 is unblocked

Worker thread W2 runs and takes something out of queue

Sets avail to 0

Now W1 runs

It must check the value of avail again

Can this execution happen?

- Queue is empty, worker thread W1 waits
- Listener thread L puts request in queue
 - Sets avail to 1
 - Signals
 - W1 is unblocked
- Worker thread W2 runs and takes something out of queue
 - Sets avail to 0
- Now W1 runs
 - **It must check the value of avail again**

Answer: No

Remember pthreads Condition Variables

- Pthread_cond_wait(cond, mutex)
 - Wait for a signal on cond
 - Release mutex
- Pthread_cond_signal(cond, mutex)
- Pthread_cond_broadcast(cond, mutex)
- Must hold mutex when calling any of these!

Answer: No

Remember pthreads Condition Variables

- Pthread_cond_wait(cond, mutex)
- Pthread_cond_signal(cond, mutex)
 - Signal one thread waiting on cond
 - Signaled thread re-acquires mutex
 - immediately
 - If no thread waiting, no-op
- Pthread_cond_broadcast(cond, mutex)

Final Solution with Locks (correct)

```
ListenerThread {
    for( i=0; i<MAX_THREADS; i++ ) thread[i] = pthread_create(...)
    forever {
        Receive( request )
        pthread_mutex_lock( queuelock )
        put request in queue
        avail++
        pthread_cond_signal( notempty, queuelock )
        pthread_mutex_unlock( queuelock )
    }
}
WorkerThread {
    forever {
        pthread_mutex_lock( queuelock )
        if( avail <= 0 ) pthread_cond_wait( notempty, queuelock )
        take request out of queue
        avail--
        pthread_mutex_unlock( queuelock )
        read file from disk
        Send( reply )
    }
}
```

Summary

- Multithreading techniques
 - Division of work
 - Synchronization of shared data
 - Fine-grain locking
 - Privatization
 - Producer/consumer problem

Further Optional Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 25 – 31 (inclusive) <https://pages.cs.wisc.edu/~remzi/OSTEP/>

For a very helpful alternative explanation on the producer/consumer problem, with C code tutorial, check out this [link](#) from the CodeVault YouTube channel.

You are also encouraged to check the other CodeVault tutorials on multi-processing and multi-threading in C/Unix ([link](#)).

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney).