


# Week 6

## **Memory Management: Virtual Memory**

Oana Balmau

February 9, 2023

# Recap: OS Memory Management



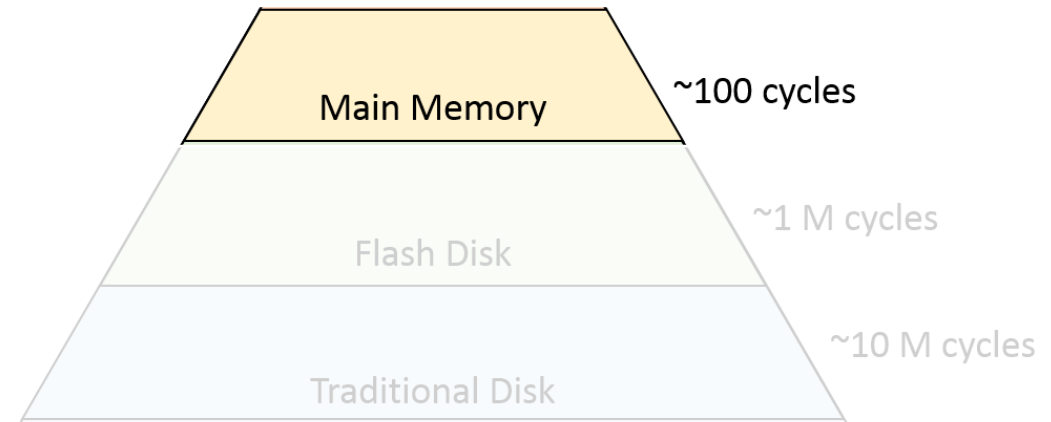
# Recap: Simplifying Assumption

Processor

**So for today:**

All of a program must be in  
main memory

Not concerned with disk



# Recap: Virtual vs. Physical address space

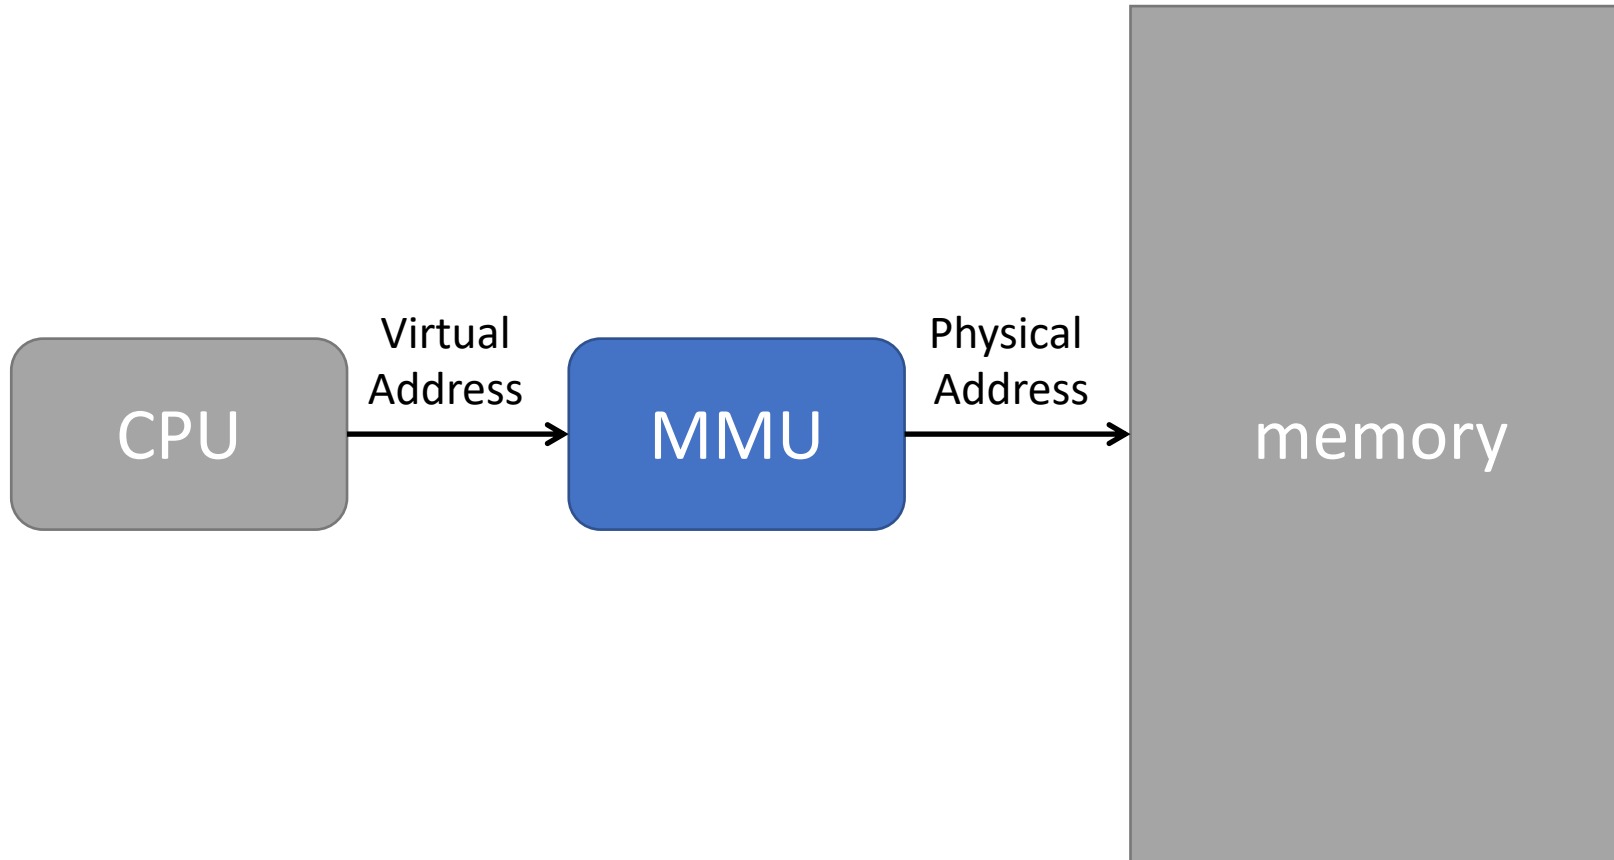
**Virtual/logical** address space = What the program(mer) thinks is its memory

**Physical** address space = Where the program actually is in physical memory



# Recap:

## MMU: Virtual to Physical



# Recap: Base and Bounds


## Virtual Address Space

- Linear address space : from 0 to MAX

## Physical Address Space

- Linear address space: from BASE to  $\text{BOUNDS} = \text{BASE} + \text{MAX}$

# Recap: Base and Bounds



# Recap: MMU for Base and Bounds

## MMU

**Relocation register:** holds the base value

**Limit register:** holds the bounds value

When a program starts running, the OS decides **where** in physical memory a process should be **loaded** (i.e., what the **base value** is).

Check for valid address:

$$0 \leq \text{virtual address} < \text{bound (in limit register)}$$

Address translation:

$$\text{physical address} = \text{virtual address} + \text{base (in relocation register)}$$




# Recap:

## Base and Bounds (External) Fragmentation

### Example:

Small holes become unusable  
Part of memory cannot be used  
Serious problem 😞



Cannot allocate a 20KB chunk, even if there are 20KB that are free in memory.

# Let's practice!

# Free Space Management (Part 1)

In this exercise we will look at memory allocation/free operations on the heap and **draw the heap and the free list at each step**. The simple memory allocator has 2 operations:

```
P = Alloc(n) // allocates n bytes to pointer P
```

```
Free(P) // frees memory that was allocated to pointer P
```

The heap of size is 20 bytes, starting at address 0. The free list is kept ordered by address (increasing). Finally, the allocator has a "best fit" free-list searching policy. The operations are:

```
1. P0 = Alloc(6);
```

```
2. P1 = Alloc(9);
```

```
3. Free(P1);
```

```
4. P2 = Alloc(6);
```

```
5. P3 = Alloc(3);
```

```
6. Free(P0);
```

```
7. P4 = Alloc(9) ;
```

```
8. Free(P3);
```

```
9. P5 = Alloc(7);
```

```
10. P6= Alloc(1);
```

# Free Space Management (Part 1)

Initialization

Free List

Addr:0; sz:20

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

Free List

Heap

```
1. P0 = Alloc(6);
```

Addr:0; sz:20

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
```

Free List

Addr:0; sz:20

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

Free List

Heap

```
1. P0 = Alloc(6);
```

Addr:6; sz:14

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);
```

Free List

Addr:6; sz:14

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	



# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);
```

Free List

Addr:15; sz:5

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P1
7	P1
8	P1
9	P1
10	P1
11	P1
12	P1
13	P1
14	P1
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);
```

Free List

Addr:15; sz:5


Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P1
7	P1
8	P1
9	P1
10	P1
11	P1
12	P1
13	P1
14	P1
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);
```

Free List




Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);
```

Free List




Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);
```

Free List




Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);
```

Free List



Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);
```

Free List

Addr:15; sz:5

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P3
13	P3
14	P3
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);
```

Free List

Addr:15; sz:5

Heap


0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P3
13	P3
14	P3
15	
16	
17	
18	
19	



# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P3
13	P3
14	P3
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9) ;
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P3
13	P3
14	P3
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);  
7. P4 = Alloc(9) ;
```

Free List



**Alloc failed!**


Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P3
13	P3
14	P3
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9) ;
8. Free(P3) ;
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P3
13	P3
14	P3
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);  
7. P4 = Alloc(9);  
8. Free(P3);
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9);
8. Free(P3);
9. P5 = Alloc(7);
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9);
8. Free(P3);
9. P5 = Alloc(7);
```

Free List



**Alloc failed!**


Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9);
8. Free(P3);
9. P5 = Alloc(7);
10. P6= Alloc(1);
```

Free List



Heap


0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	
13	
14	
15	
16	
17	
18	
19	



# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9);
8. Free(P3);
9. P5 = Alloc(7);
10. P6= Alloc(1);
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	P2
7	P2
8	P2
9	P2
10	P2
11	P2
12	P6
13	
14	
15	
16	
17	
18	
19	

# Different Virtual to Physical Mapping Schemes


- Base and bounds
- Segmentation
- (Simplified) Paging

# Base and Bounds: (Internal) Fragmentation



- **Big chunk of “free” space**
- “free” space **takes up** physical memory.
- Inefficient
- (Internal) memory fragmentation

# Base and Bounds: (Internal) Fragmentation



- **Big chunk of “free” space**
- “free” space **takes up** physical memory.
- Inefficient
- (Internal) memory fragmentation

# Segmentation

# Segmentation

## **Virtual Address Space**


- Two-dimensional
- Set of segments 0 .. n
- Each segment  $i$  is linear from 0 to  $MAX_i$

## **Physical Address Space**

- Set of segments, each linear


# What is a Segment?

- Anything you want it to be
- Typical examples:
  - **Code**
  - **Heap**
  - **Stack**




**Segments**

# Segmentation Example



Virtual Address Space



Physical Memory

Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K



# Segmentation: Virtual Address

Two-dimensional address:


- Segment number  $s$
- Offset  $d$  **within segment** (starting at 0)

It is like multiple base-and-bounds



# Segmentation Virtual Address example

Chop up the address space into segments based on the **top few bits** of virtual address.




How many segments?

What is the size of each segment?

# Segmentation Virtual Address example

Chop up the address space into segments based on the **top few bits** of virtual address.




How many segments?  
 **$2^2 = 4$  segments**


What is the size of each segment?  
 **$2^{12} = 4\text{KB}$**

# Segmentation Virtual Address example

Chop up the address space into segments based on the **top few bits** of virtual address.




Example: virtual address  $4200_{10}$  ( $01000001101000_2$ )



Segment	bits
Code	00
Heap	01
Stack	10
-	11


# Segmentation Virtual Address example

Chop up the address space into segments based on the **top few bits** of virtual address.



Why?

Example: virtual address  $4200_{10}$  ( $01000001101000_2$ )



Segment	bits
Code	00
Heap	01
Stack	10
-	11

# MMU for Segmentation

## Segment table

Indexed by segment number

Contains **(base, limit) pair**

- Base: physical address of segment in memory
- Limit: length of segment

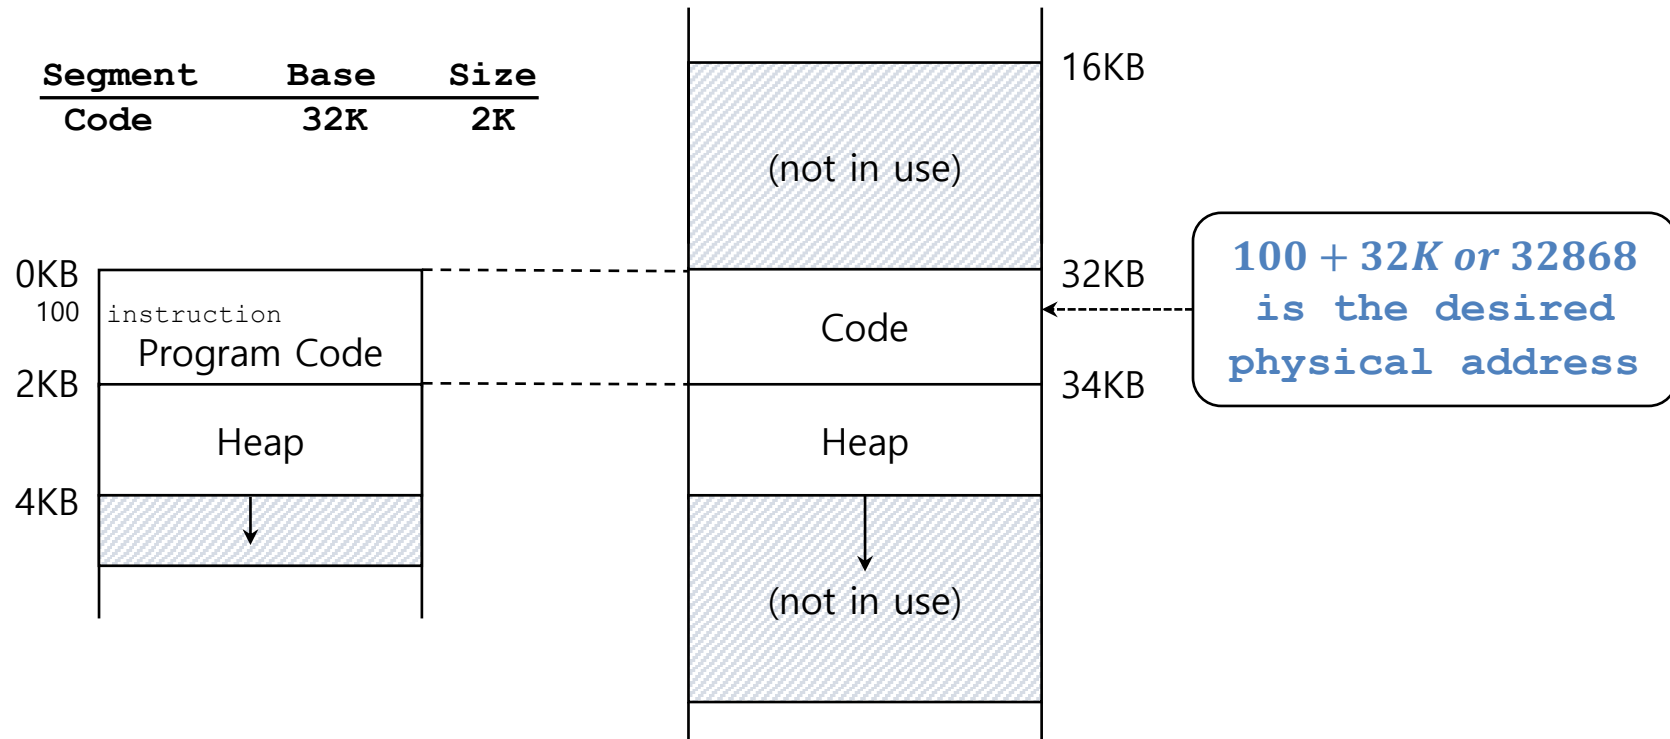
*Also need:*

**Pointer** to segment table in memory

**Length** of segment table

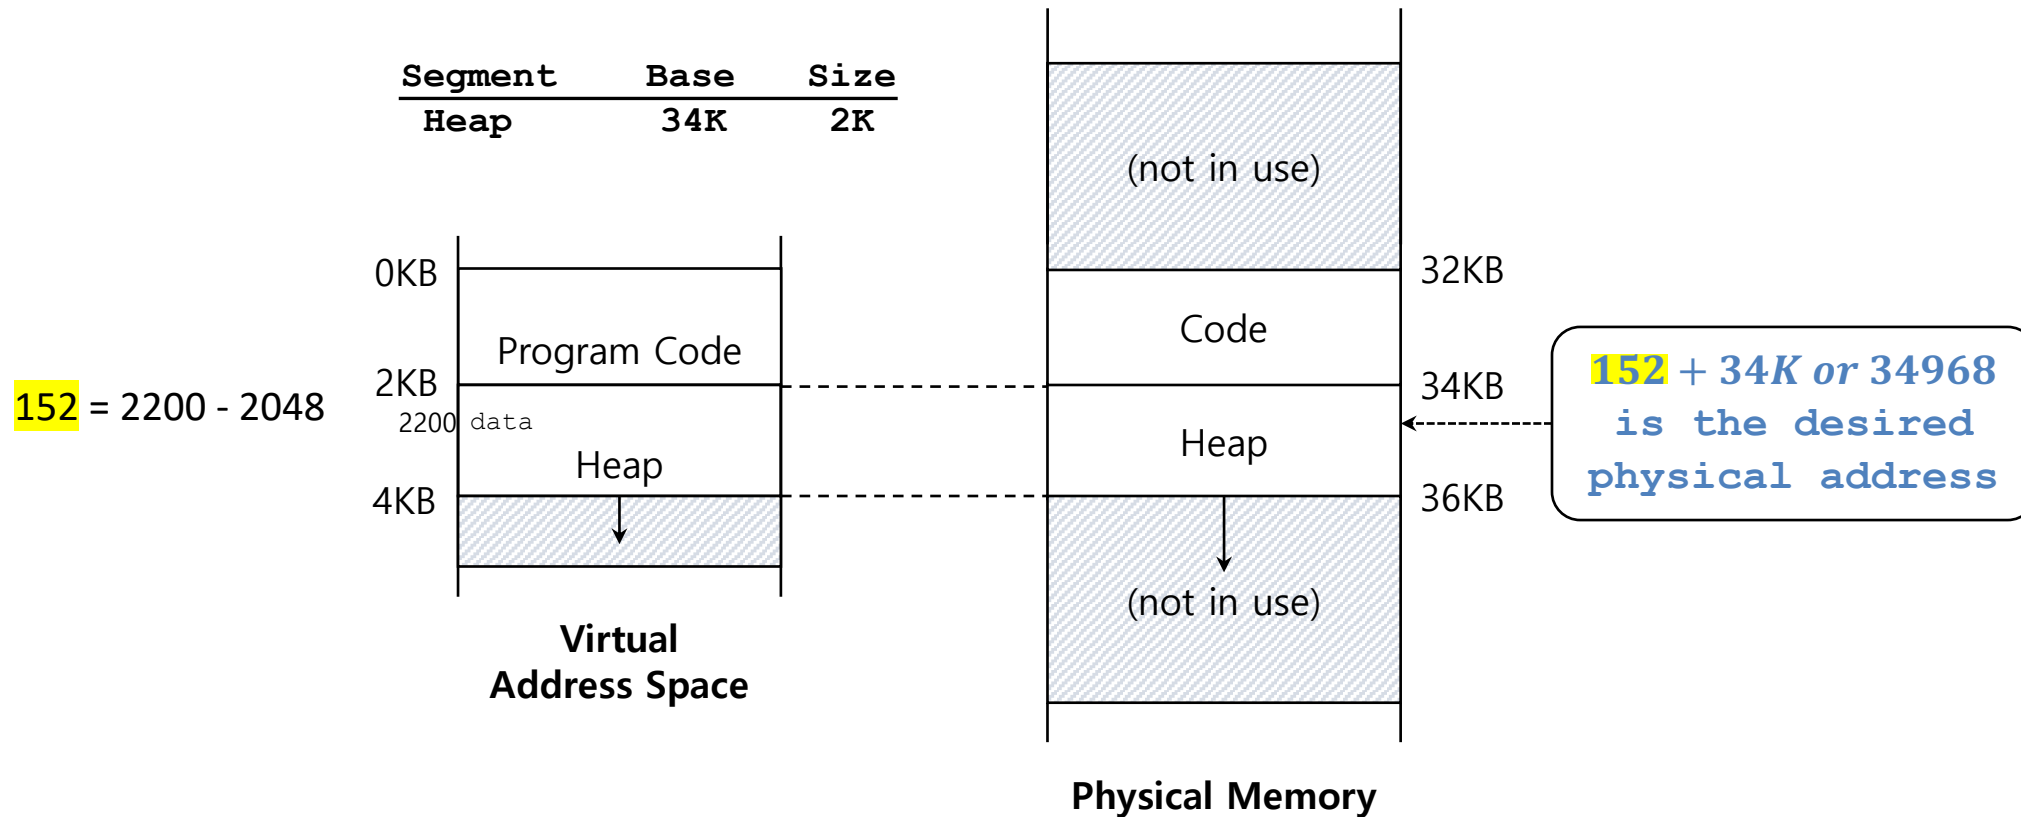
# Segmentation: Address Translation Example

*physical address = offset in segment + segment base*



# Segmentation: Address Translation (Cont.)

$$\text{physical address} = \text{offset in segment} + \text{segment base}$$





# Let's practice!

# Segmentation

Given a CPU with 3-bit instructions, and 16 bytes byte-addressable of physical memory, with the following 2 segments in main memory:

SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )

SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Compute the virtual to physical address translations (or Segmentation Faults), for the following virtual addresses:

$4_{10}$ ,  $5_{10}$ ,  $6_{10}$ ,  $0_{10}$ ,  $3_{10}$ .

\* $X_{10}$  means X in base 10 (decimal).

# Segmentation

CPU with 3-bit instructions, 16 bytes physical memory  
SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )  
SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

## Preliminaries:

1. How many bits does a virtual address have? What is the size of the virtual address space?
2. How many bits does a physical address have? What is the size of the physical address space?
3. What is the structure of the virtual address?
  - How many bits for the segment?
  - How many bits for the offset?
  - What is the maximum size of each segment?

# Segmentation

CPU with 3-bit instructions, 16 bytes physical memory  
SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )  
SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

## Preliminaries:

1. How many bits does a virtual address have? What is the size of the virtual address space?  
→ **3 bits, with a virtual address space of 8 addresses, because CPU has 3-bit instructions**
2. How many bits does a physical address have? What is the size of the physical address space?  
→ **4 bits, with a phys address space of 16 addresses, because  $16 = 2^4$**
3. What is the structure of the virtual address?
  - How many bits for the segment? → **1 bit**
  - How many bits for the offset? → **2 bits**
  - What is the maximum size of each segment? → **4 Bytes (each segment has max length 4 and each line in physical memory has 1 byte)**

# Segmentation

CPU with 3-bit instructions, 16 bytes physical memory  
SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )  
SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

**Virtual Address Space**

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

Physical Address Space

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

Physical Address Space

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space



# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

4

Virtual Address Space

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

4

VA 

1	0	0
---	---	---

Virt addr  $4_{10}$   
Is valid in SEG 1  
Phys addr  $3_{10}$

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

5

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

5

VA 

1	0	1
---	---	---

Virt addr  $5_{10}$   
Is valid in SEG 1  
Phys addr  $4_{10}$

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

6

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

**6**

VA 

1	1	0
---	---	---

Virt addr  $6_{10}$   
Is out of bounds

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

Physical Address Space

SEG 1

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Segmentation Fault!

SEG 0

Offset is not smaller than Bound register for SEG 1.

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

0

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space



# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

0

VA 

0	0	0
---	---	---

Virt addr  $0_{10}$   
Is valid in SEG 0  
Phys addr  $7_{10}$

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

3

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

SEG 1

SEG 0

- SEG 0 (base address:  $7_{10}^*$ , bound:  $3_{10}$ )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Physical Address Space

# Segmentation

Virtual Addresses (decimal): 4, 5, 6, 0, 3.

	s	offset	
0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Virtual Address Space

**3**

VA 

0	1	1
---	---	---

Virt addr  $3_{10}$   
Is out of bounds

0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

Physical Address Space

SEG 1

- SEG 0 (base address:  $7_{10}^*$ , **bound:  $3_{10}$** )
- SEG 1 (base address:  $3_{10}$ , bound:  $2_{10}$ )

Segmentation Fault!

SEG 0

Offset is not smaller than Bound register for SEG 0.

# Sharing Memory between Processes

Why would we want to do that?


For instance,

- Run twice the same program in different processes
- May want to share code
- Read twice the same file in different processes
- May want to share memory corresponding to file

**Sharing not possible with base and bounds, but is possible with segmentation**

# Segmentation Provides Easy Support for Sharing

- Create segment for shared data
- Add segment entry in segment table of both processes
- Points to shared segment in memory



# Segmentation Provides Easy Support for Sharing

Extra **hardware** support is need for form of **Protection bits**.

- **A few more bits** per segment to indicate **permissions** of **read**, **write** and **execute**.

## Example Segment Register Values(with Protection)

Segment	Base	Size	Protection
Code	32K	2K	Read-Execute
Heap	34K	2K	Read-Write
Stack	28K	2K	Read-Write

# Main memory allocation with Segmentation


Remember:

**Segmentation**  $\sim$  **multiple base-and-bounds**

No internal fragmentation inside each segment.

External fragmentation problem is similar.

- Pieces are typically smaller



# Main memory allocation with Segmentation

## Compaction:

Rearrange segments in physical memory to get rid of “holes”.

- **Stop** running process.

Inefficient! 😞

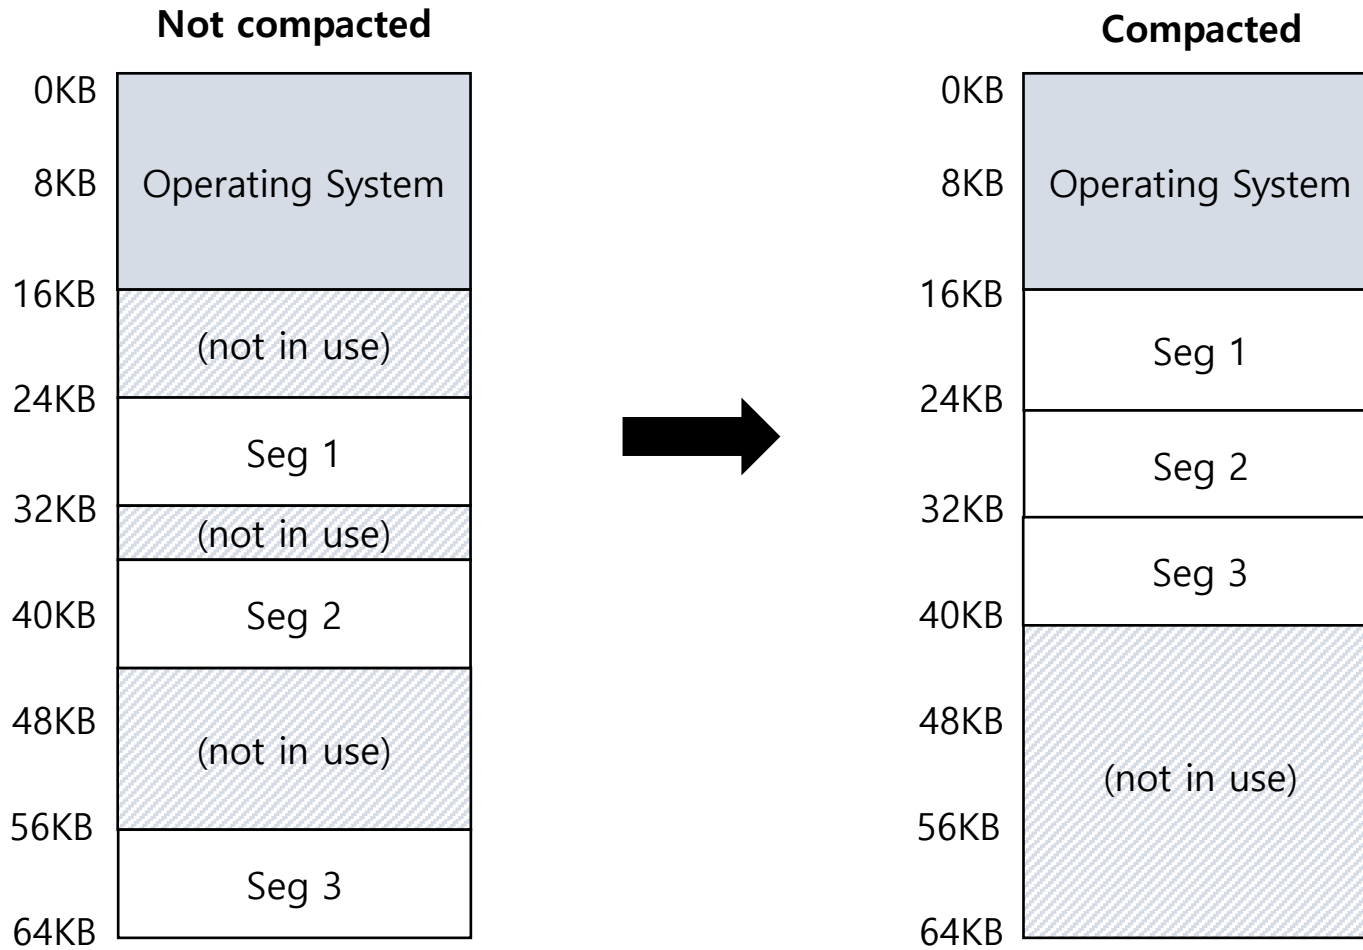
- **Copy** data to somewhere.

Expensive! 😞😞

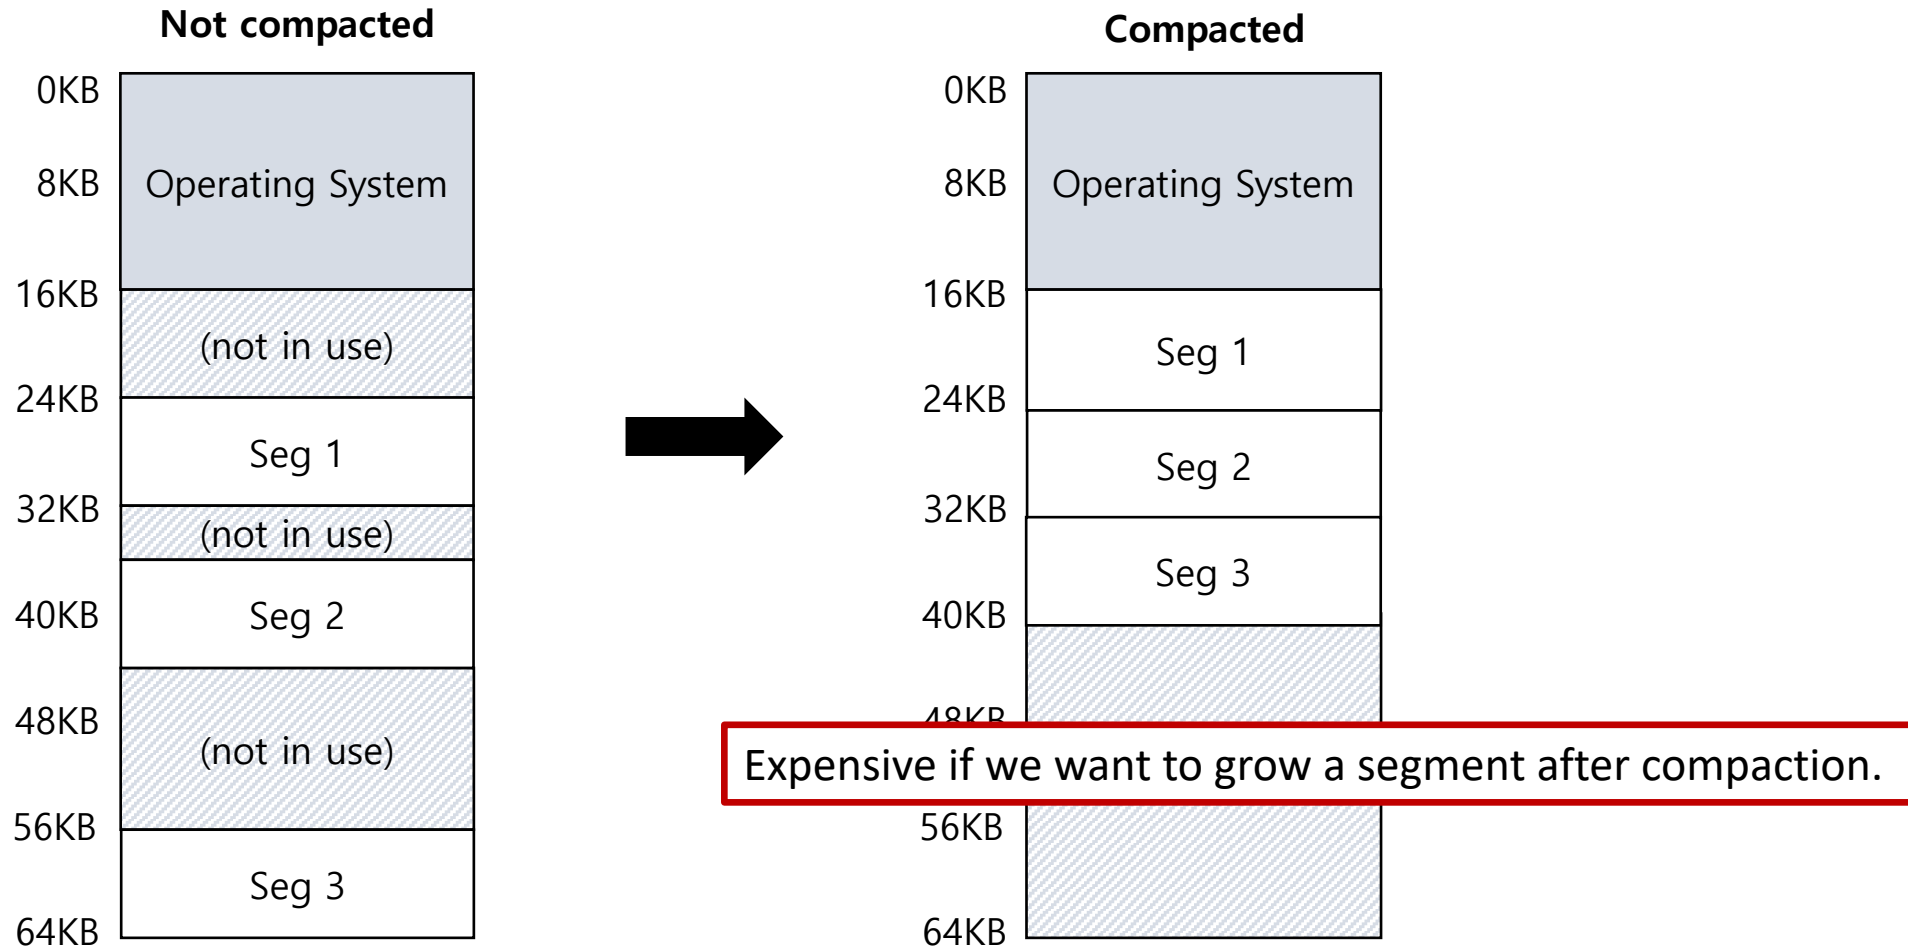
- **Change** segment register value.



# Memory Compaction Example



# Memory Compaction Example



# More practice!

# Free Space Management (Part 2)

Consider the same memory allocator as in Part 1, but this time **the allocator is 4-byte aligned**. This means that each allocated space rounds up to the nearest 4-byte free chunk in size. Draw the heap and the free list for each step.

Same as in part 1, the operations are:

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9);
8. Free(P3);
9. P5 = Alloc(7);
10. P6 = Alloc(1);
```

## Reminder:

- `P = Alloc(n)` // allocates `n` bytes to pointer `P`
- `Free(P)` // frees memory that was allocated to `P`
- The heap of size is 20 bytes, starting at address 0.
- The free list is kept ordered by address (increasing).
- "best fit" free-list searching policy.

# Free Space Management (Part 2)

Initialization

Free List

Addr:0; sz:20

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 2)

Free List

Heap

```
1. P0 = Alloc(6);
```

Addr:0; sz:20

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 2)

Free List

Heap

```
1. P0 = Alloc(6);
```

Addr:8; sz:12

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);
```

Free List

Addr:8; sz:12

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	



# Free Space Management (Part 1)

```
1. P0 = Alloc(6);
```

Free List

Addr:0; sz:20

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);
```

Free List

Addr:-1; sz:0

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	P1
9	P1
10	P1
11	P1
12	P1
13	P1
14	P1
15	P1
16	P1
17	P1
18	P1
19	P1

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);
```

Free List

Addr:-1; sz:0

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	P1
9	P1
10	P1
11	P1
12	P1
13	P1
14	P1
15	P1
16	P1
17	P1
18	P1
19	P1

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);
```

Free List

Addr:8; sz:12

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);
```

Free List

Addr:8; sz:12

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	
9	
10	
11	
12	
13	
14	
15	
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);
```

Free List

Addr:16; sz:4

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);
```

Free List

Addr:16; sz:4

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);
```

Free List

Addr:-1; sz:0

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P3
17	P3
18	P3
19	P3



# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);
```

Free List

Addr:-1; sz:0

Heap

0	P0
1	P0
2	P0
3	P0
4	P0
5	P0
6	P0
7	P0
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P3
17	P3
18	P3
19	P3

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);
```

Free List

Addr:0; sz:8

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P3
17	P3
18	P3
19	P3

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);  
7. P4 = Alloc(9) ;
```

Free List

Addr:0; sz:8

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P3
17	P3
18	P3
19	P3

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);  
7. P4 = Alloc(9) ;
```

Free List

Addr:0; sz:8

**Alloc failed!**

Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P3
17	P3
18	P3
19	P3

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);  
7. P4 = Alloc(9);  
8. Free(P3);
```

Free List

Addr:0; sz:8


Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P3
17	P3
18	P3
19	P3

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9) ;
8. Free(P3) ;
```

Free List




Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9);
8. Free(P3);
9. P5 = Alloc(7);
```

Free List



Heap

0	
1	
2	
3	
4	
5	
6	
7	
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9) ;
8. Free(P3);
9. P5 = Alloc(7);
```

Free List

Addr:16; sz:4

Heap

0	P5
1	P5
2	P5
3	P5
4	P5
5	P5
6	P5
7	P5
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	
17	
18	
19	



# Free Space Management (Part 2)

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9) ;
8. Free(P3);
9. P5 = Alloc(7);
10. P6= Alloc(1);
```

Free List

Addr:16; sz:4

Heap

0	P5
1	P5
2	P5
3	P5
4	P5
5	P5
6	P5
7	P5
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	
17	
18	
19	

# Free Space Management (Part 2)

```
1. P0 = Alloc(6);  
2. P1 = Alloc(9);  
3. Free(P1);  
4. P2 = Alloc(6);  
5. P3 = Alloc(3);  
6. Free(P0);  
7. P4 = Alloc(9) ;  
8. Free(P3);  
9. P5 = Alloc(7);  
10. P6= Alloc(1);
```

Free List

Addr:-1; sz:0

Heap

0	P5
1	P5
2	P5
3	P5
4	P5
5	P5
6	P5
7	P5
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P6
17	P6
18	P6
19	P6

# Free Space Management (Part 2)

## Free List

```
1. P0 = Alloc(6);
2. P1 = Alloc(9);
3. Free(P1);
4. P2 = Alloc(6);
5. P3 = Alloc(3);
6. Free(P0);
7. P4 = Alloc(9) ;
8. Free(P3);
9. P5 = Alloc(7);
10. P6= Alloc(1);
```

Addr:-1; sz:0

## Heap

0	P5
1	P5
2	P5
3	P5
4	P5
5	P5
6	P5
7	P5
8	P2
9	P2
10	P2
11	P2
12	P2
13	P2
14	P2
15	P2
16	P6
17	P6
18	P6
19	P6

**Advantages/disadvantages of aligned allocation?**

# Paging (simplified version)

# Paging (simplified version)

- **Page:** fixed-size portion of virtual memory
- **Frame:** fixed-size portion of physical memory
- **Page size = frame size**
- Typical size: 4k – 8k (always power of 2)

# Paging

## **Virtual Address Space**


- Linear from 0 up to a multiple of page size

## **Physical Address Space**

- Noncontiguous set of frames, one per page


# Paging: Example

Contiguous



**Virtual Address Space**

64B address space with four **16B** pages



**Physical Memory**


128B physical memory with eight **16B** frames

Not  
Contiguous

# Paging: Virtual Address

Two components in the virtual address

- VPN: virtual page number
- Offset: offset within the page; Page size =  $2^{\text{offset}}$






# Paging: Virtual Address Example


Virtual address 21 (binary 010101) in 64-byte address space

2 VPN bits  $\rightarrow$  Number of pages =  $2^2 = 4$  pages

4 offset bits  $\rightarrow$  Page size =  $2^4 = 16B$



# Paging: Virtual to Physical Address Example



# MMU for Paging

## Page table


- **Data structure** used to map the virtual address to physical address
- Indexed by page number
- Contains frame number of page in memory
- **Each process has a page table**

*Also need:*

**Pointer** to page table in memory

**Length** of page table

# Quiz: Fill in Page Table



Page Tables:

P1	P2	P3
3	0	8
1	4	5
7	2	9
10	6	11

# Page Tables can get large

32-bit address space with 4-KB pages, 20 bits for VPN


$$4MB = 2^{20} \text{ entries} * 4B \text{ per page table entry}$$

64-bit address space with 4-KB pages, 52 bits for VPN

$$2^{52} \text{ entries} * 4B \text{ per page table entry} = \text{Petabyte} - \text{scale} !$$

**True, but address space is often sparsely used**

# Problem?



Address space **sparsely used**

Access to unused portion will appear valid

**Would rather have an error**

# Solution: Valid/Invalid Bit

Page table has length  $2^p$

→ Page table does not cover the entire possible virtual address space, only the pages that the process has allocated.

Have valid bit in each page table entry

- Set to valid for used portions of address space
- Invalid for unused portions

→ **This is the common approach**

# Main Memory Allocation with Paging

Logical address space: fixed size pages

Physical address space: fixed size frames

New process

- Find frames for all of process's pages

**Easier problem 😊**

- **Fixed size**



# (Internal) Fragmentation in Paging

With paging

- Address space = multiple of page size

Part of last page may be unused

**With reasonable page size, not a big problem**

# Summary – Key Concepts

- Virtual and physical address spaces
- Mapping between virtual and physical address
- Different mapping methods:
  - Base and bounds
  - Segmentation
  - Paging
- Sharing, protection, memory allocation

# Further Reading

## **Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau**

Chapters 12–18

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

### **Credits:**

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Youjip Won (Hanyang University).