

# Week 7

## **Memory Management: Demand Paging**

Oana Balmau

February 14, 2023

# Announcements

Graded exercises released and  
Assignment 1 grades are released

Ungraded exercises will be released

Week 6 Memory Management	feb 6 C Review: C files	feb 7 Virtual Memory (1/2) Optional reading: <a href="#">OSTEP</a> Chapters 12 – 18	feb 8 Scheduling Assignment Released	feb 9 Virtual Memory (2/2) Scheduling Assignment Overview — with Jiaxuan	feb 10
Week 7 Memory Management	feb 13 OS Shell Assignment Due C Review: Working with pthreads I	feb 14 Demand Paging (1/3) Optional reading: <a href="#">OSTEP</a> Chapters 19 – 22	feb 15	feb 16 Demand Paging (2/3)	feb 17
Week 8 Memory Management	feb 20 C Review: Working with pthreads I	feb 21 Demand Paging (3/3) Optional reading: <a href="#">OSTEP</a> Chapters 19 – 22 Practice Exercises Sheet: Memory Management	feb 22	feb 23 Mid-semester Q&A – not recorded • Graded Exercises Sheet Released • Grades released for OS Shell Assignment	feb 24
Week 9	feb 27	feb 28	mar 1	mar 2	mar 3

# Key Concepts

- TLB
- Page Table for very large address spaces
- Demand paging
  - Page fault
- Page replacement policies

# Recap Week 6: Virtual vs. Physical Address Space

**Virtual/logical** address space = What the program(mer) thinks is its memory

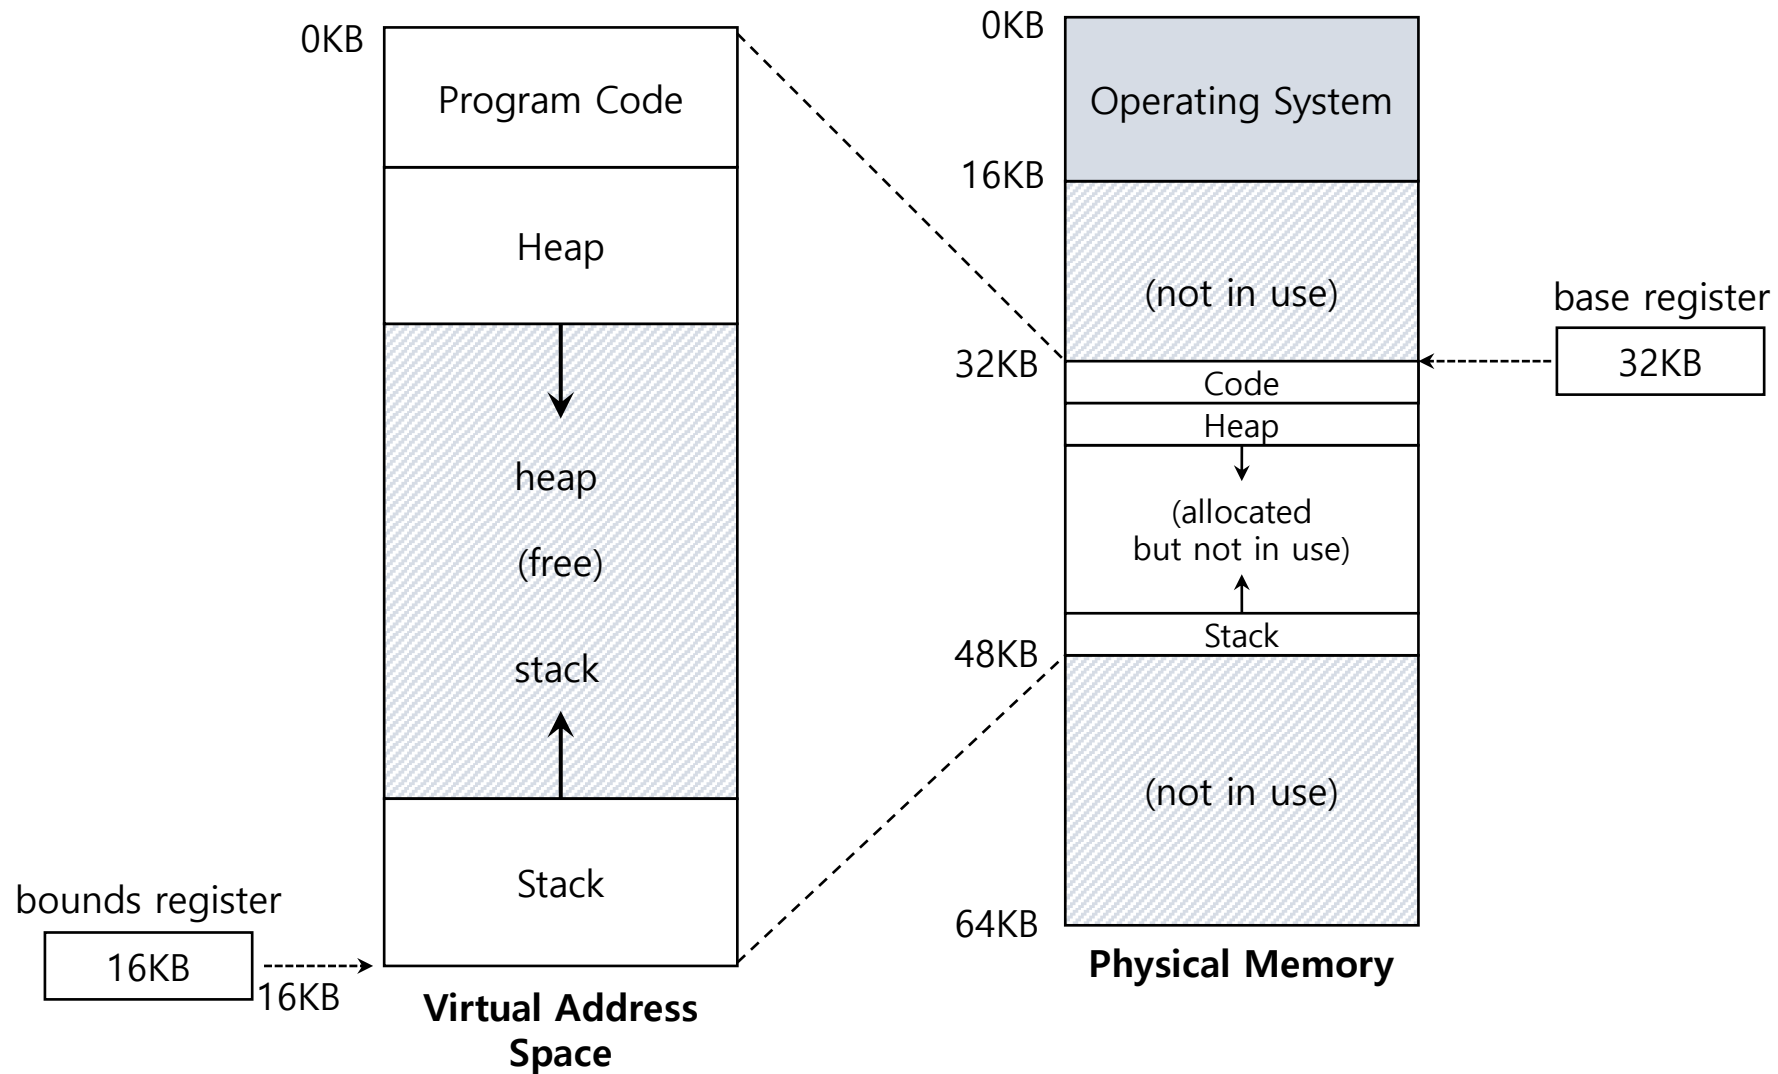
**Physical** address space = Where the program actually is in physical memory




# Recap Week 6: Memory Mapping Methods

- Base and bounds
- Segmentation


# Recap Week 6: Base and Bounds



# Recap Week 6: Segmentation



**Virtual Address Space**



**Physical Memory**

Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

# Sharing Memory between Processes

Why would we want to do that?

For instance,


- Run twice the same program in different processes
- May want to share code
- Read twice the same file in different processes
- May want to share memory corresponding to file

**Sharing not possible with base and bounds, but is possible with segmentation**



# Segmentation Provides Easy Support for Sharing

- Create segment for shared data
- Add segment entry in segment table of both processes
- Points to shared segment in memory



# Segmentation Provides Easy Support for Sharing

Extra **hardware** support is need for form of **Protection bits**.

- **A few more bits** per segment to indicate **permissions** of **read**, **write** and **execute**.

## Example Segment Register Values(with Protection)

Segment	Base	Size	Protection
Code	32K	2K	Read-Execute
Heap	34K	2K	Read-Write
Stack	28K	2K	Read-Write

# Main memory allocation with Segmentation

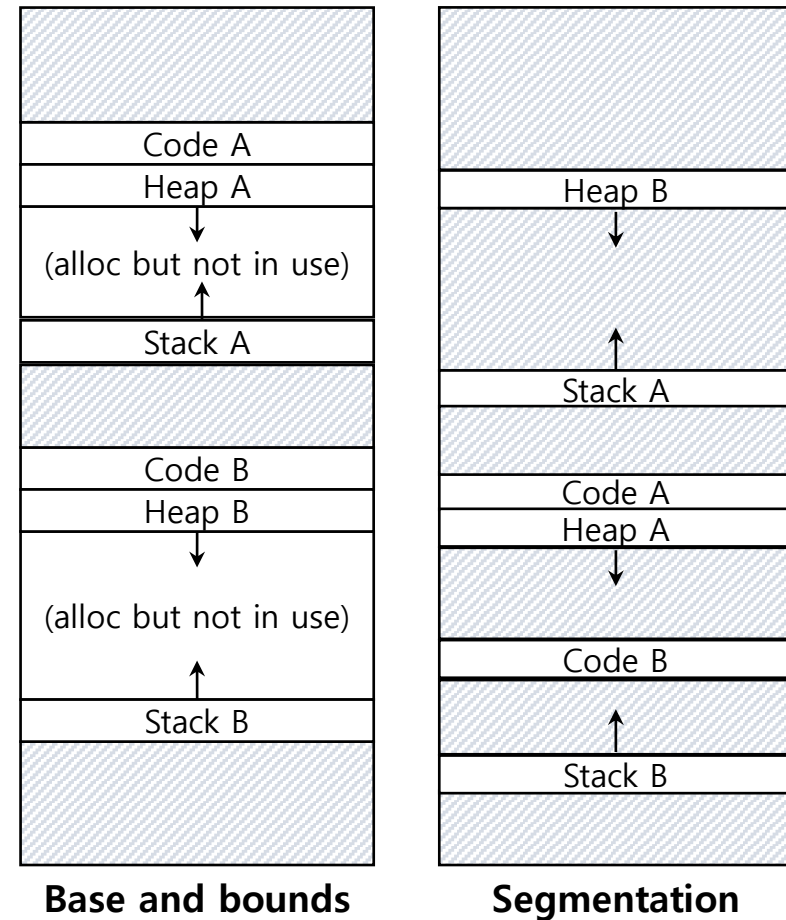
Remember:

**Segmentation**  $\sim$  **multiple base-and-bounds**

No internal fragmentation inside each segment.

External fragmentation problem is similar.

- Pieces are typically smaller



# Main memory allocation with Segmentation

## Compaction:


Rearrange segments in physical memory to get rid of “holes”.

- **Stop** running process.
- **Copy** data to somewhere.
- **Change** segment register value.

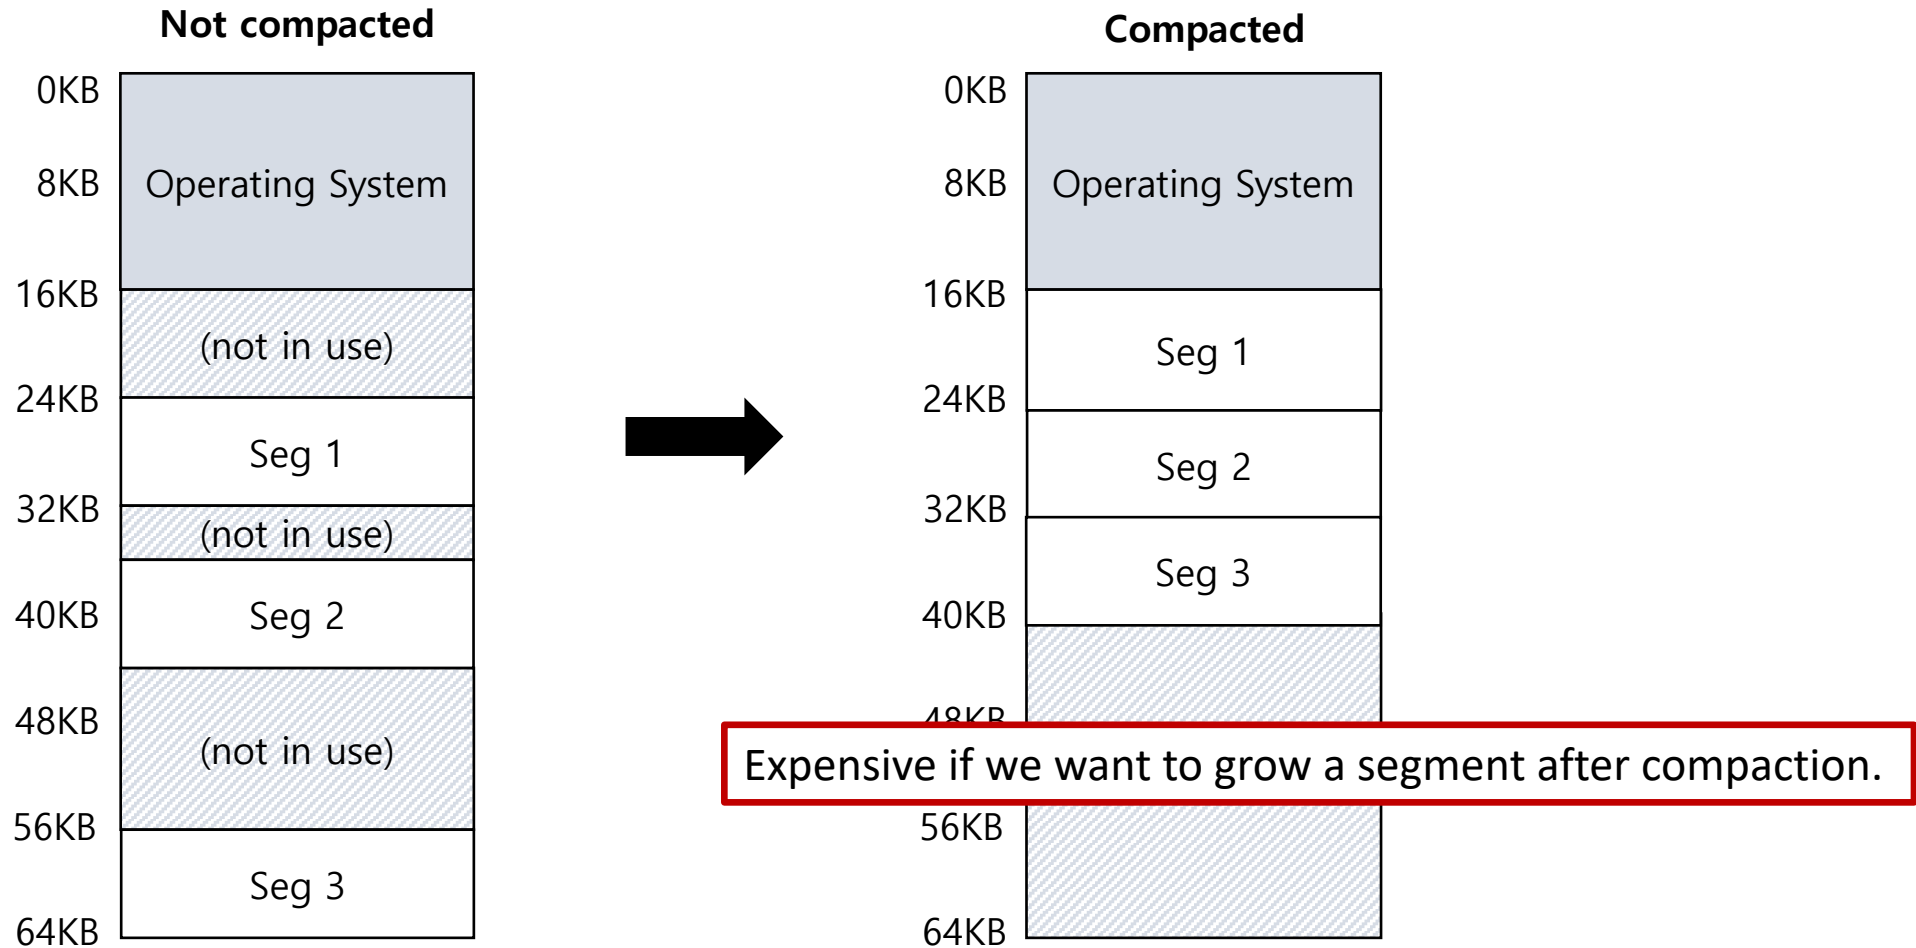
Inefficient! 😞

Expensive! 😞😞

# Memory Compaction Example



# Memory Compaction Example



# Paging (simplified version)

# Paging (simplified version)

- **Page:** fixed-size portion of virtual memory
- **Frame:** fixed-size portion of physical memory
- **Page size = frame size**
- Typical size: 4k – 8k (always power of 2)



# Paging

## **Virtual Address Space**


- Linear from 0 up to a multiple of page size

## **Physical Address Space**

- Noncontiguous set of frames, one per page


# Paging: Example

Contiguous



**Virtual Address Space**

64B address space with four **16B** pages



**Physical Memory**


128B physical memory with eight **16B** frames

Not  
Contiguous

# Paging: Virtual Address

Two components in the virtual address

- VPN: virtual page number
- Offset: offset within the page; Page size =  $2^{\text{offset}}$




# Paging: Virtual Address Example

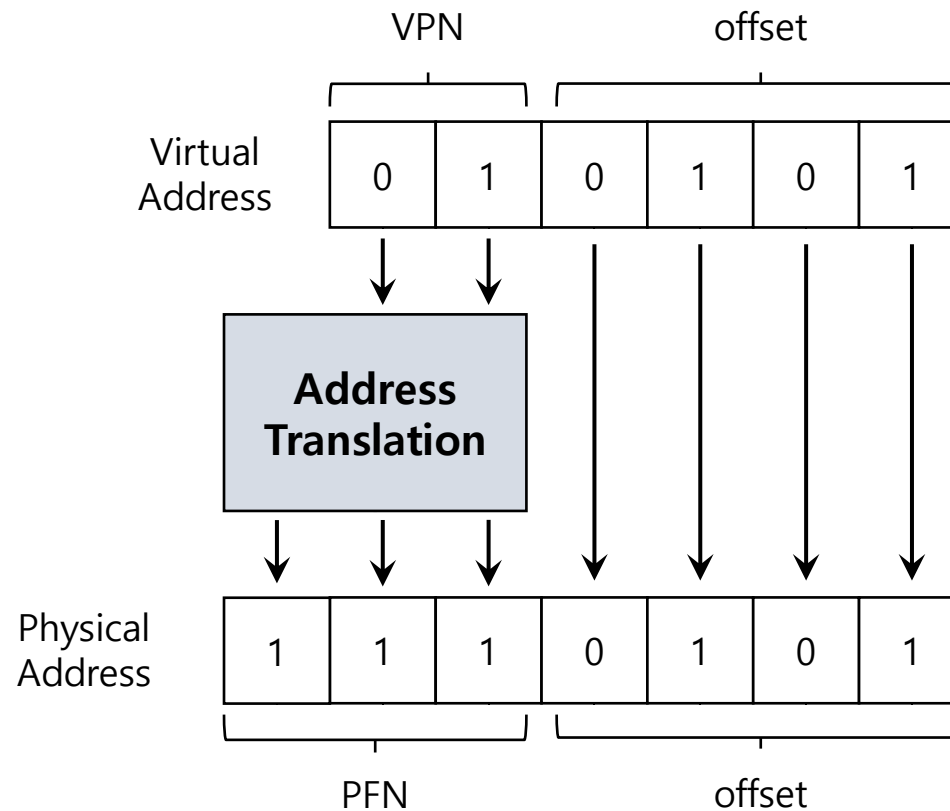
Virtual address 21 (binary 010101) in 64-byte address space

2 VPN bits  $\rightarrow$  Number of pages =  $2^2 = 4$  pages

4 offset bits  $\rightarrow$  Page size =  $2^4 = 16B$



# Paging: Virtual to Physical Address Example



# MMU for Paging

## Page table


- **Data structure** used to map the virtual address to physical address
- Indexed by page number
- Contains frame number of page in memory
- **Each process has a page table**

*Also need:*

**Pointer** to page table in memory

**Length** of page table

# Quiz: Fill in Page Table



Page Tables:

P1	P2	P3
3	0	8
1	4	5
7	2	9
10	6	11

# Page Tables can get large

32-bit address space with 4-KB pages, 20 bits for VPN

$$4MB = 2^{20} \text{ entries} * 4B \text{ per page table entry}$$


64-bit address space with 4-KB pages, 52 bits for VPN

$$2^{52} \text{ entries} * 4B \text{ per page table entry} = \text{Petabyte} - \text{scale} !$$

**True, but address space is often sparsely used**



# Problem?



Address space **sparsely used**

Access to unused portion will appear valid

**Would rather have an error**

# Solution: Valid/Invalid Bit

Page table has length  $2^p$

→ Page table does not cover the entire possible virtual address space, only the pages that the process has allocated.

Have valid bit in each page table entry

- Set to valid for used portions of address space
- Invalid for unused portions

→ **This is the common approach**

# Main Memory Allocation with Paging

Logical address space: fixed size pages

Physical address space: fixed size frames

New process

- Find frames for all of process's pages

**Easier problem 😊**

- **Fixed size**

# (Internal) Fragmentation in Paging

With paging

- Address space = multiple of page size


Part of last page may be unused

**With reasonable page size, not a big problem**

# In Reality


- Base-and-bounds only for niche
- Segmentation abandoned
  - High complexity for little gain
  - Effect approximated with paging + valid bits
- **Paging is now universal**

# Paging




**Page table**

frame	valid/ invalid
3	v
7	v
5	v
2	v




# Problem: Paging Address Translation Performance




Page table

frame	valid/ invalid
3	v
7	v
5	v
2	v




# Problem: Paging Address Translation Performance



Page table

frame	valid/ invalid
3	v
7	v
5	v
2	v


Page table is in memory





# Problem: Paging Address Translation Performance

1 virtual address →  
2 physical memory accesses




**Virtual Address Space**

**Page table**

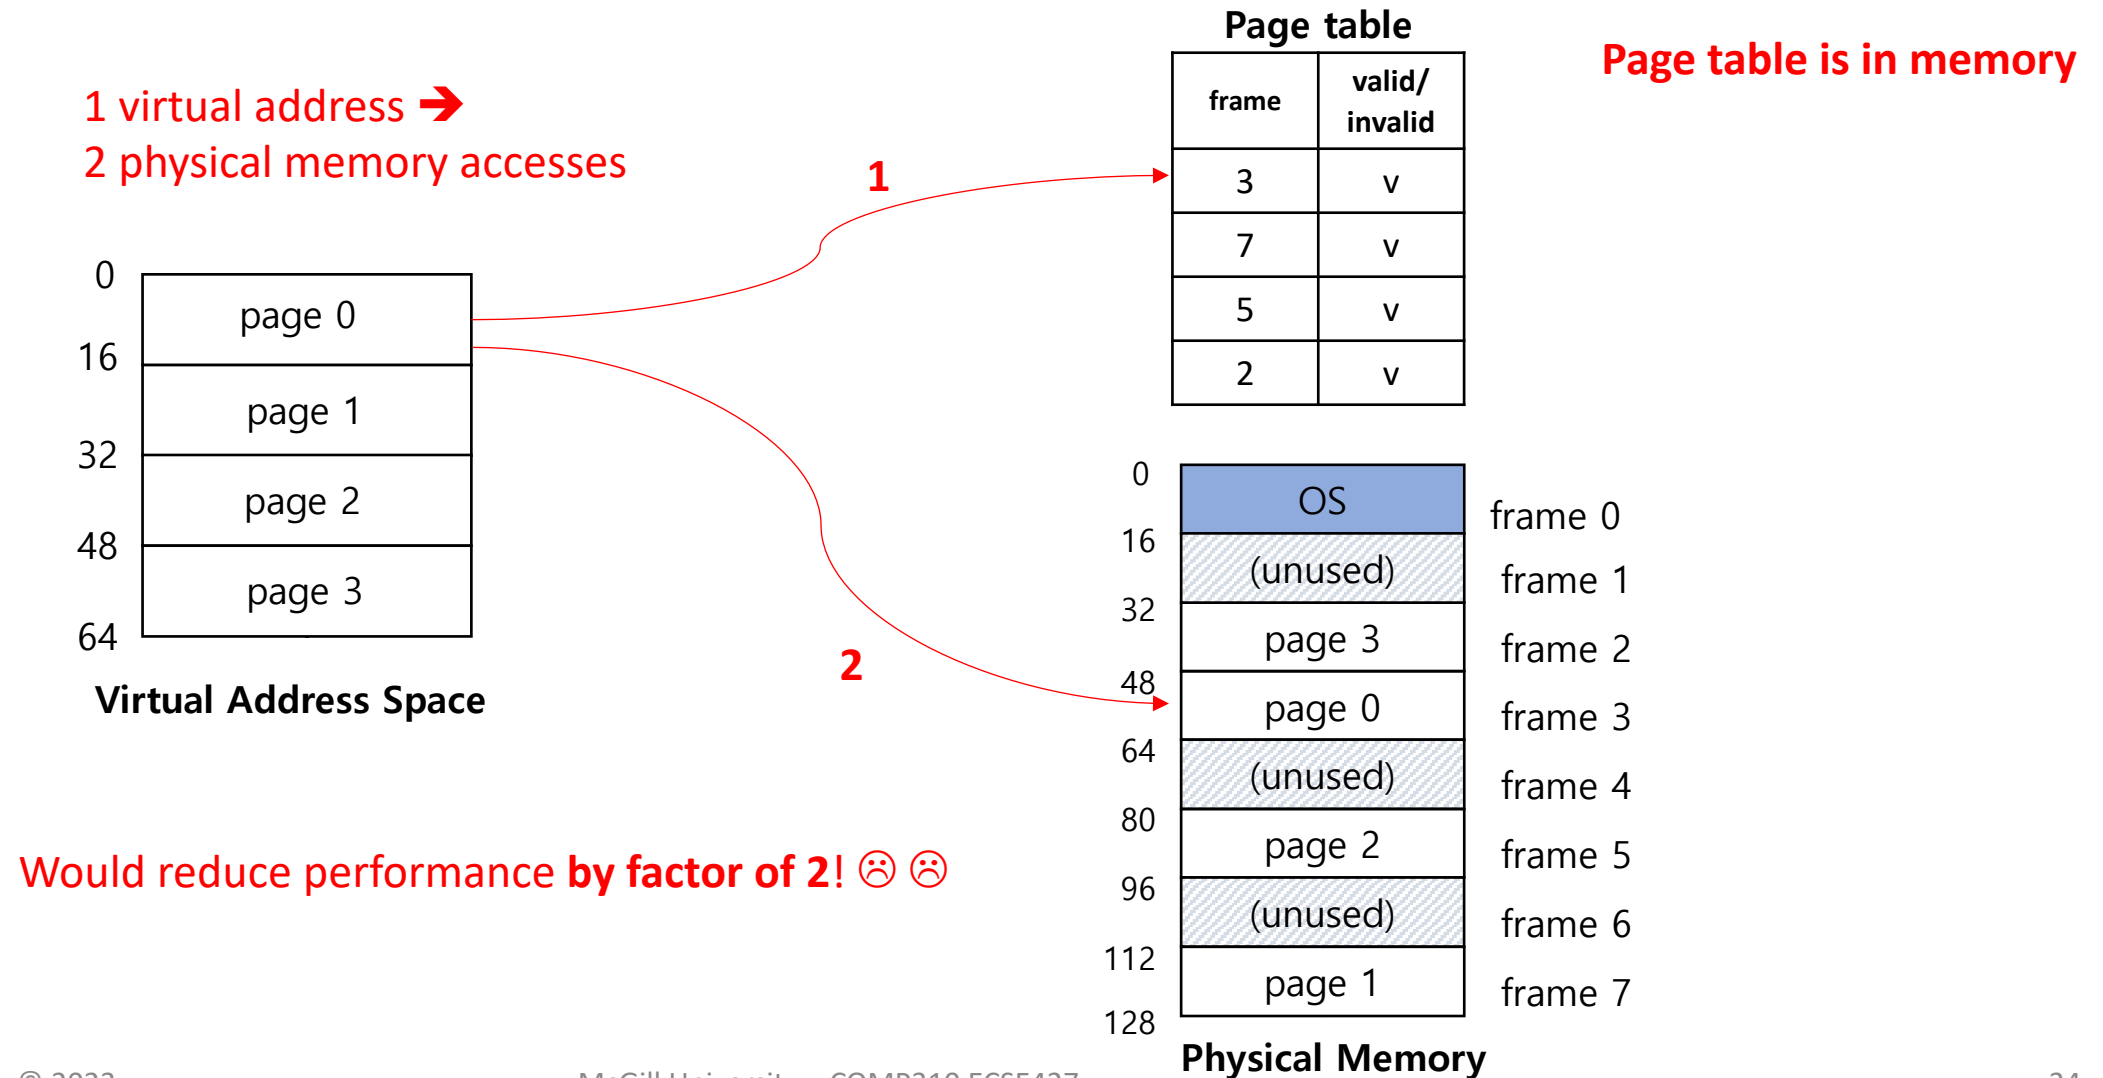
frame	valid/ invalid
3	v
7	v
5	v
2	v

**Page table is in memory**



**Physical Memory**

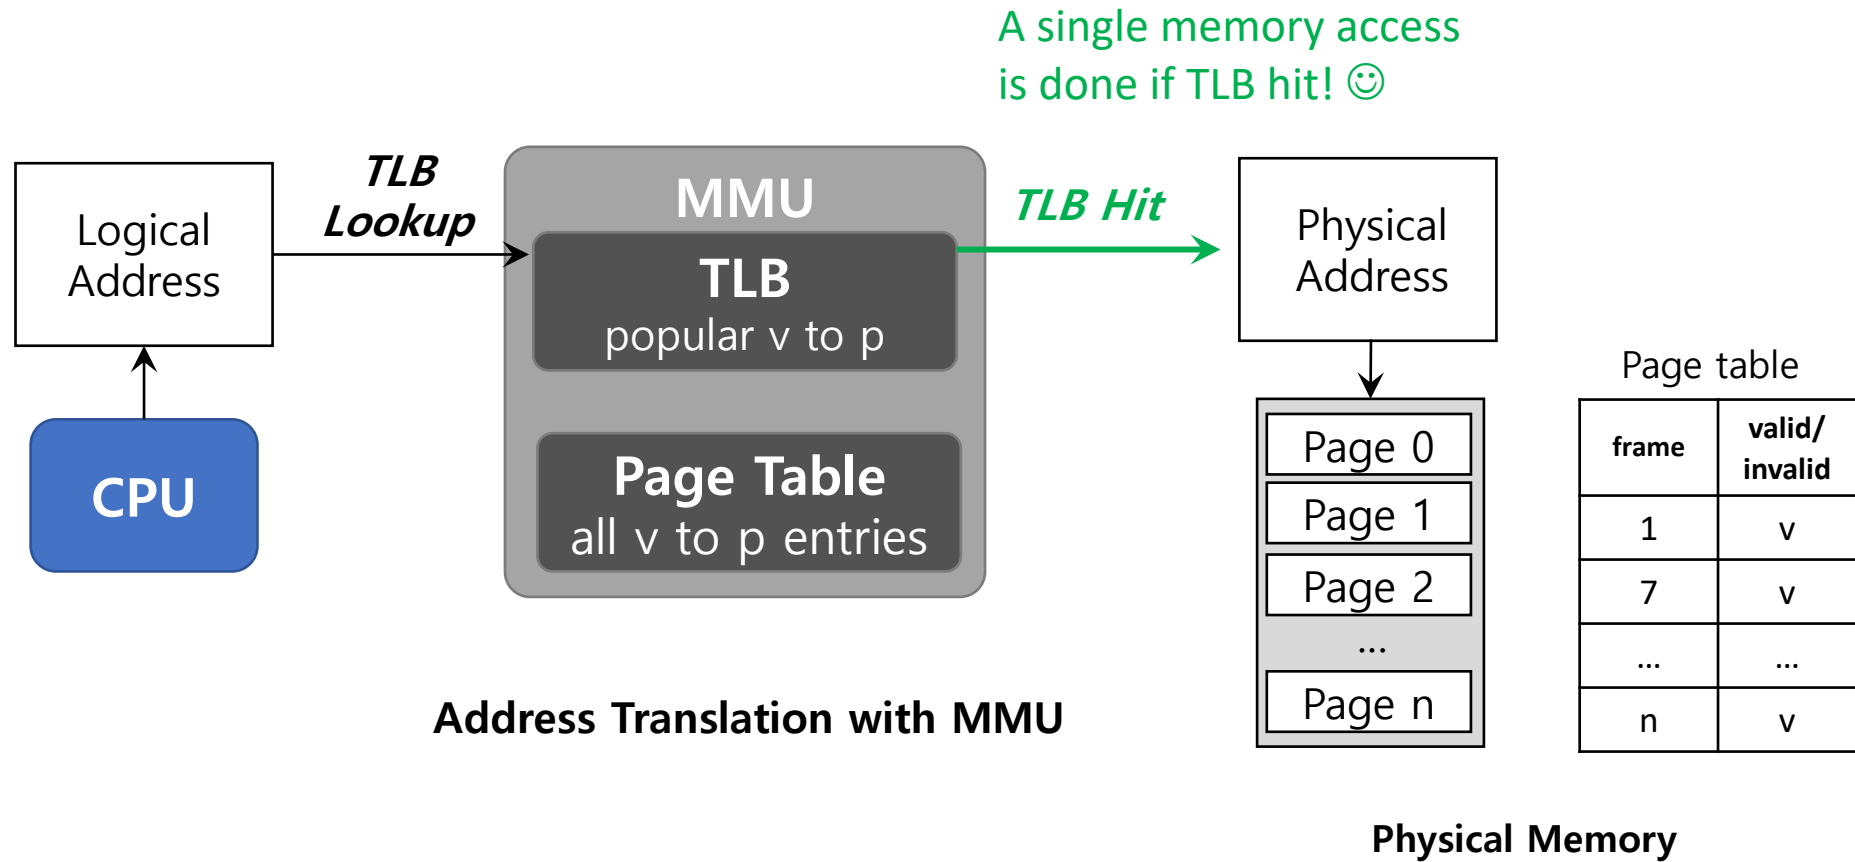
# Problem: Paging Address Translation Performance



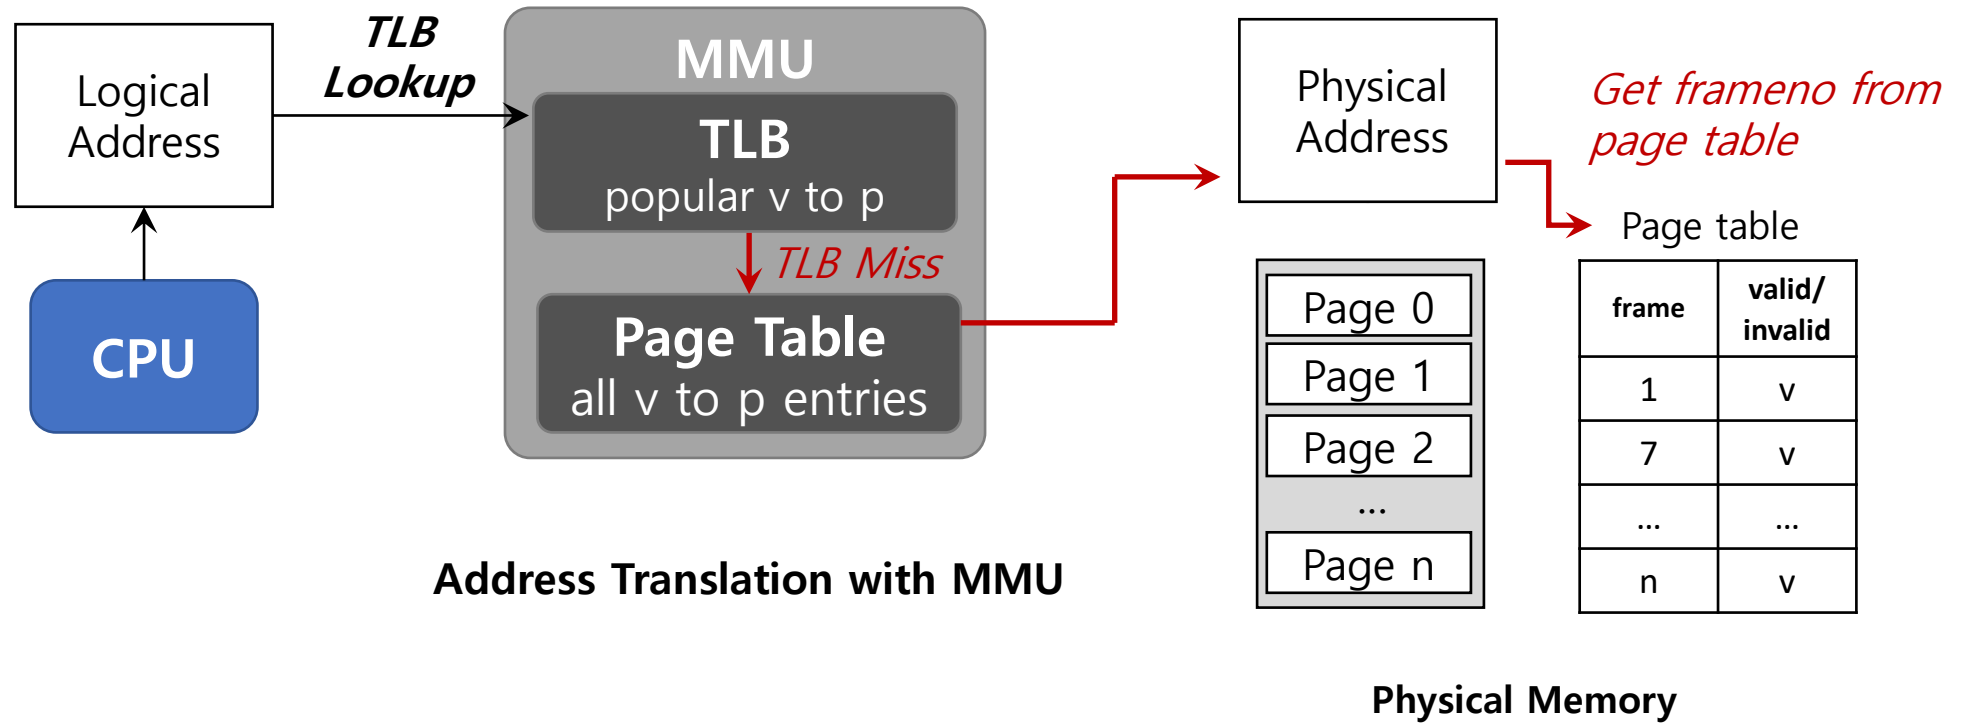
# Solution: Translation Lookaside Buffer (TLB)

- Small fast **hardware cache** of **popular (pageno, frameno) maps**.
- **Part of MMU**
- If mapping for pageno found in TLB
  - Use frameno from TLB
  - Abort mapping using page table
- If not
  - Perform mapping using page table
  - Insert (pageno, frameno) in TLB

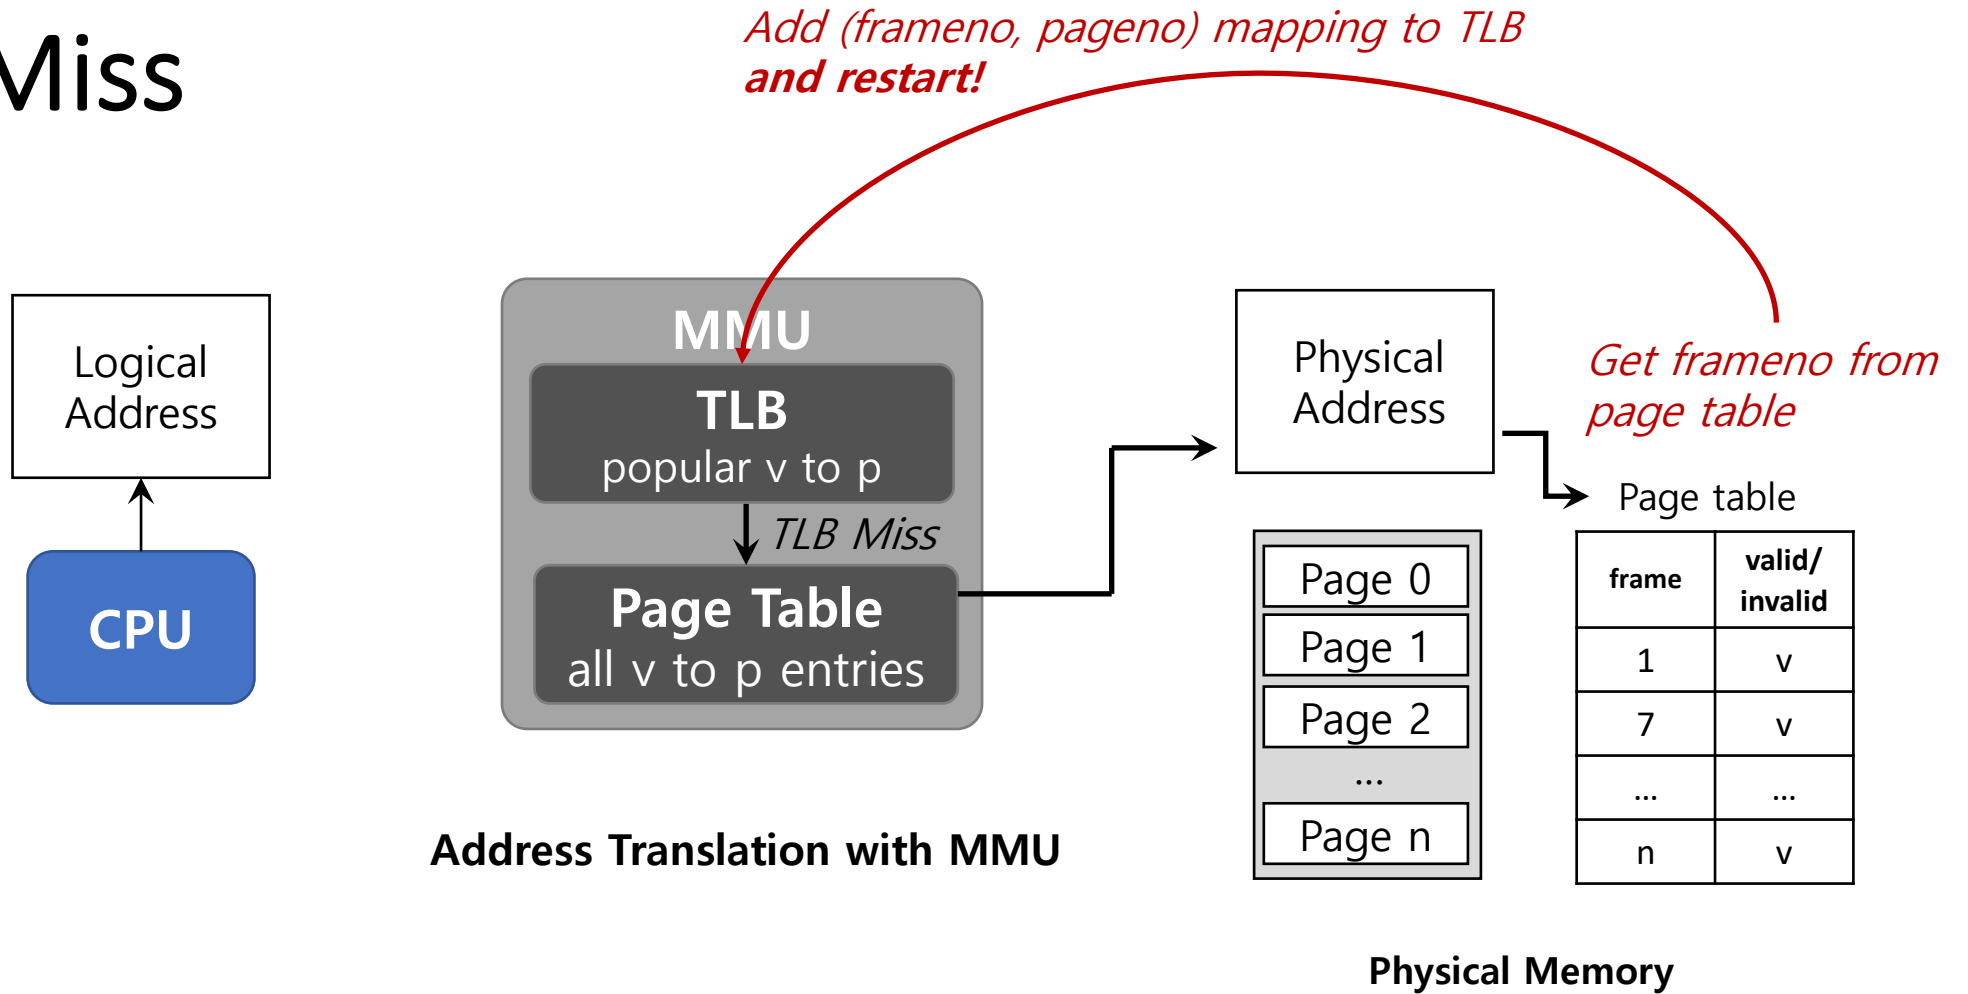
# TLB Hit



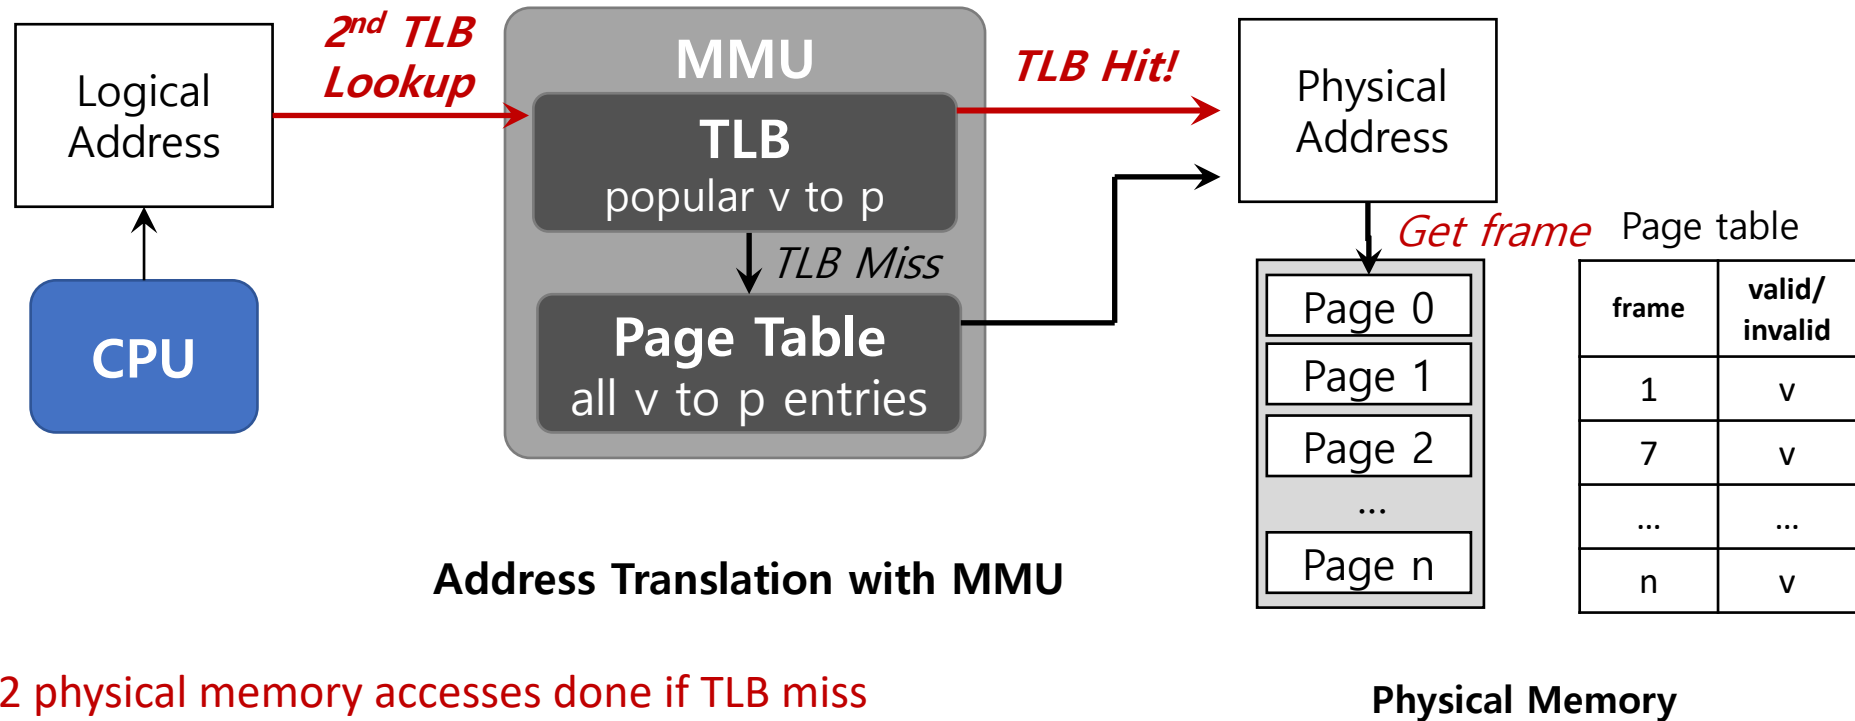
# TLB Miss



# TLB Miss



# TLB Miss




# How to make TLB fast?

Use **associative memory** (special hardware)

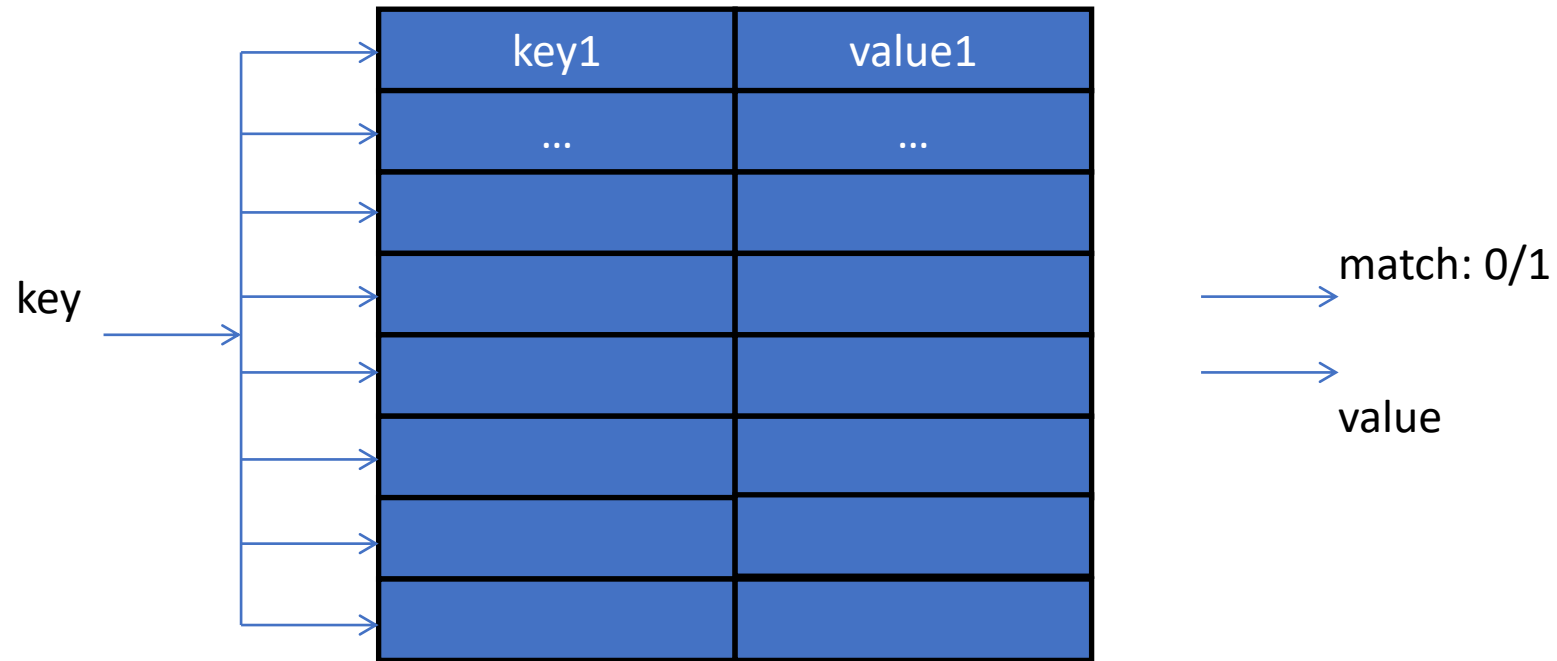
- Regular memory  
    **Lookup by address**
- Associative memory  
    **Lookup by contents**  
    Lookup in **parallel**



# Regular Memory



# Associative Memory



# TLB



# TLB Size


- Associative memory is very expensive
- Therefore, TLB small (64 – 1,024 entries)
- Want TLB hit-rate close to 100%
- If TLB full, **need to replace existing entry**. How?
  - Will discuss replacement policies later today.
  - Main idea: take advantage of locality.

# TLB Hit Rate and Locality

Want TLB hit rate close to 100%. To do this, try to take advantage of locality:


- **Temporal Locality**

An instruction or data item that has been recently accessed will likely be re-accessed soon in the future.



- **Spatial Locality**

If a program accesses memory at address  $x$ , it will likely soon access memory near  $x$ .



# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i];
3:      }
```

# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i];
3:      }
```

**TLB: miss**

# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i];
3:      }
```

TLB: miss, hit, hit



# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++) {
2:                  sum+=a[i];
3:      }
```

**TLB:** miss, hit, hit, miss

# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++) {
2:                  sum+=a[i];
3:      }
```

**TLB:** miss, hit, hit, miss, hit, hit, hit

# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++) {
2:                  sum+=a[i];
3:      }
```

TLB: miss, hit, hit, miss, hit, hit, hit  
miss

# Spatial Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++) {
2:                  sum+=a[i];
3:      }
```

TLB: miss, hit, hit, miss, hit, hit, hit  
miss, hit, hit

3 misses and 7 hits.  
Thus **TLB hit rate** is 70%.

**The TLB improves performance  
due to spatial locality**

# Temporal Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```

0:      int sum = 0 ;
1:      for( j=0; j<2; j++){
2:          for( i=0; i<10; i++)
3:              sum+=a[i];
4:      }

```

**j=0**

TLB: miss, hit, hit, miss, hit, hit, hit  
miss, hit, hit

3 misses and 7 hits.  
Thus **TLB hit rate** is 70%.

# Temporal Locality Example: Accessing an Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```

0:      int sum = 0 ;
1:      for( j=0; j<2; j++){
2:          for( i=0; i<10; i++)
3:              sum+=a[i];
4:      }
    
```

**j=0** TLB: miss, hit, hit, miss, hit, hit, hit 3 misses and 7 hits.  
 miss, hit, hit Thus **TLB hit rate** is 70%.

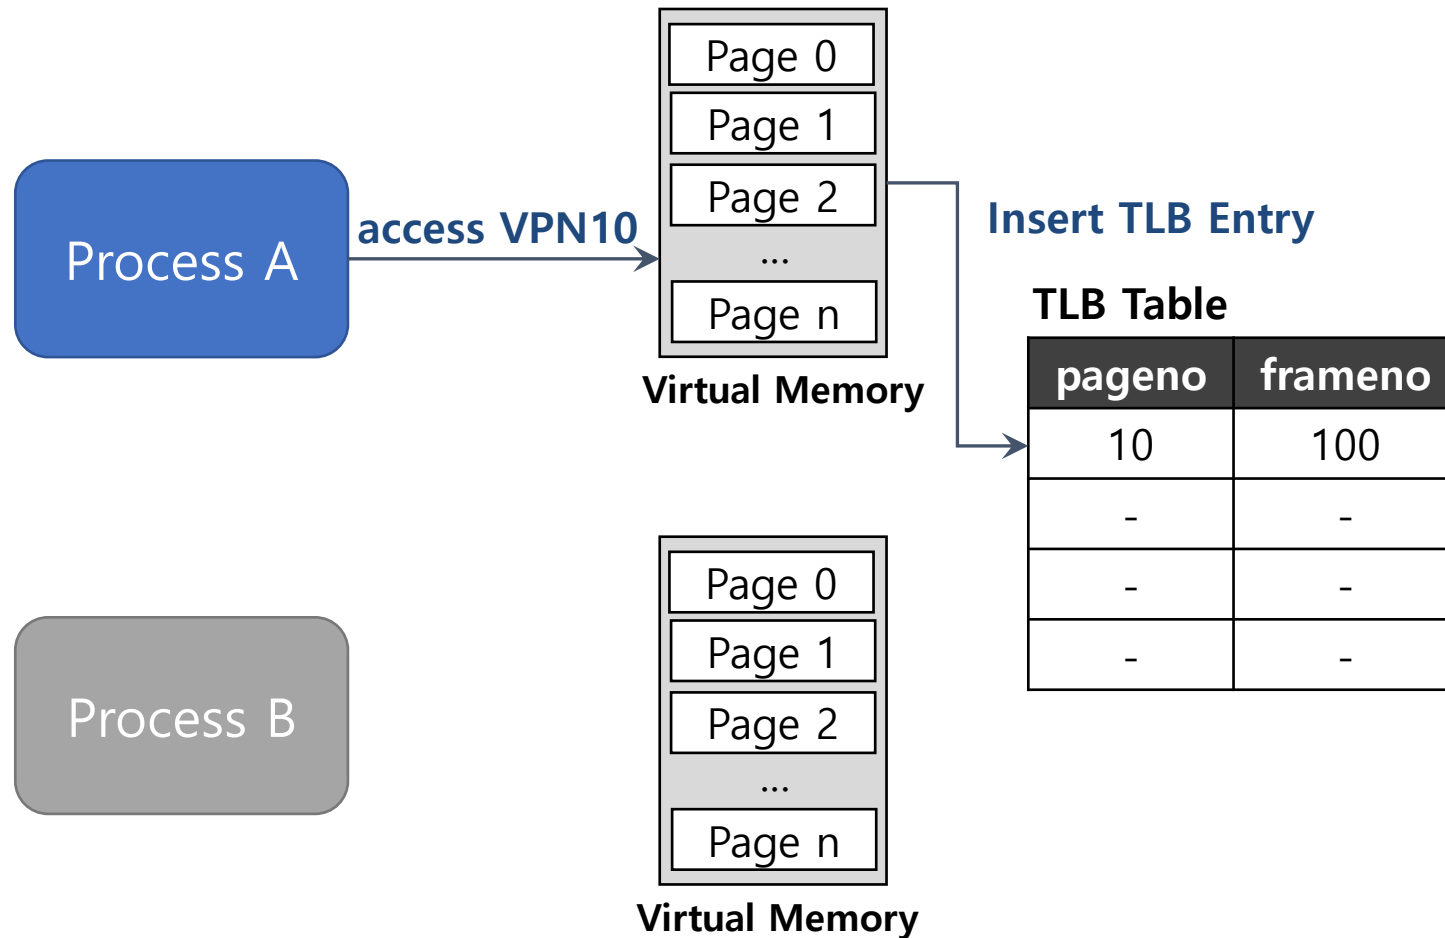
**j=1** TLB: hit, hit, hit, hit, hit, hit, hit 0 misses and 10 hits.  
 hit, hit, hit Thus **TLB hit rate** is 100%.

The TLB improves performance  
 due to **temporal locality**

# Revisiting Process Switching

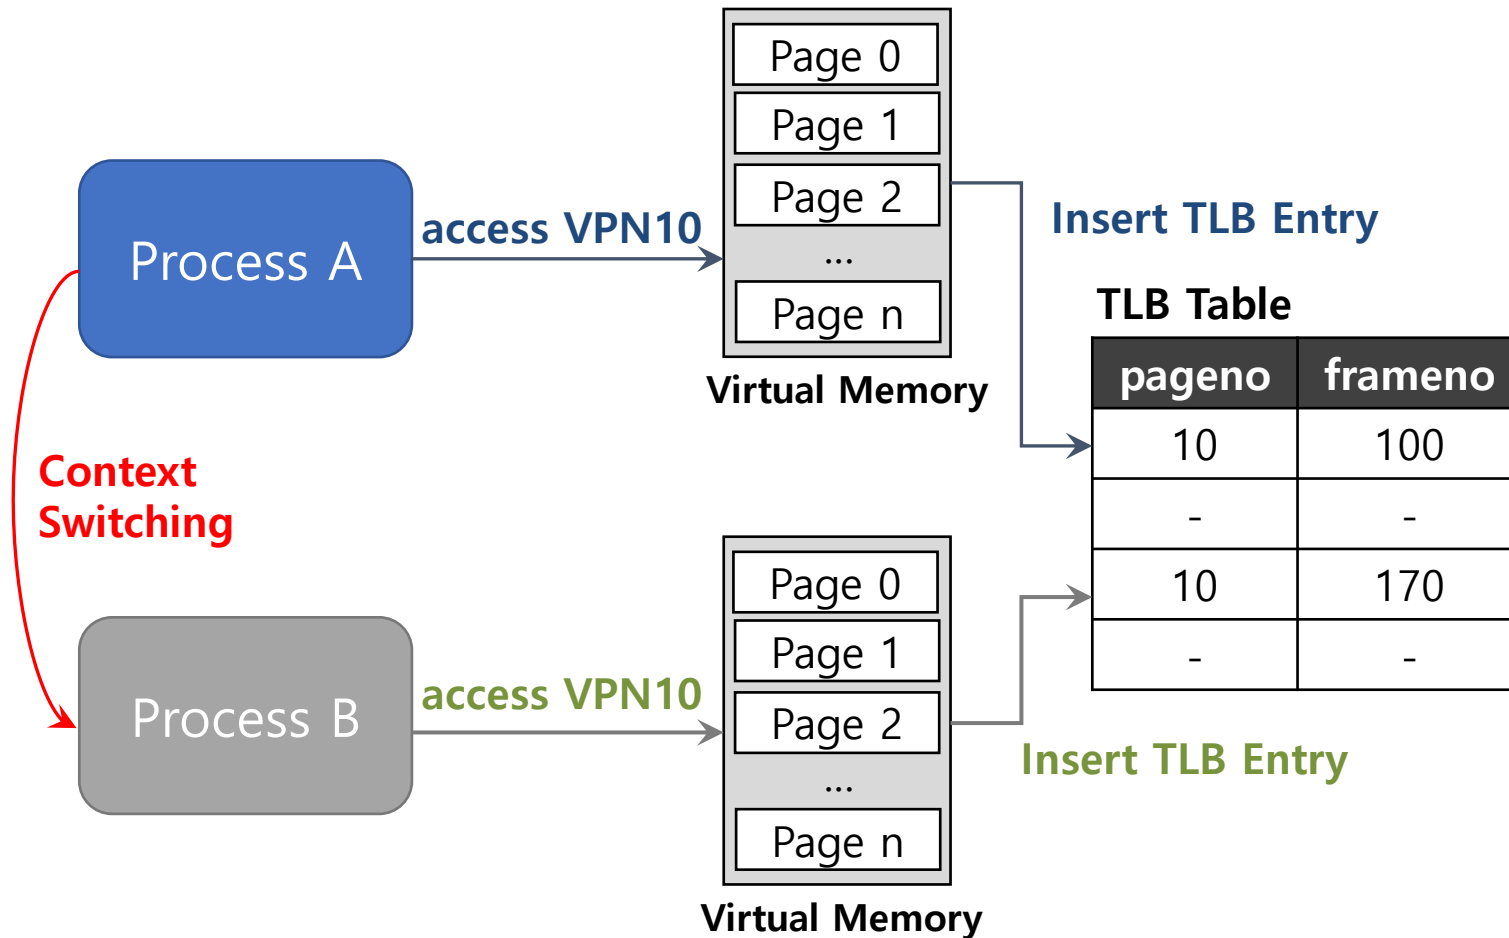
- Suppose
  - Process P1 is running
  - Entry (**pageno**, frameno) in TLB
  - Switch from P1 to P2
  - P2 issues virtual address in page **pageno**
- P2 accesses P1's memory!
- Or unable to distinguish between P1 and P2 mapping

# TLB Issue: Context Switching

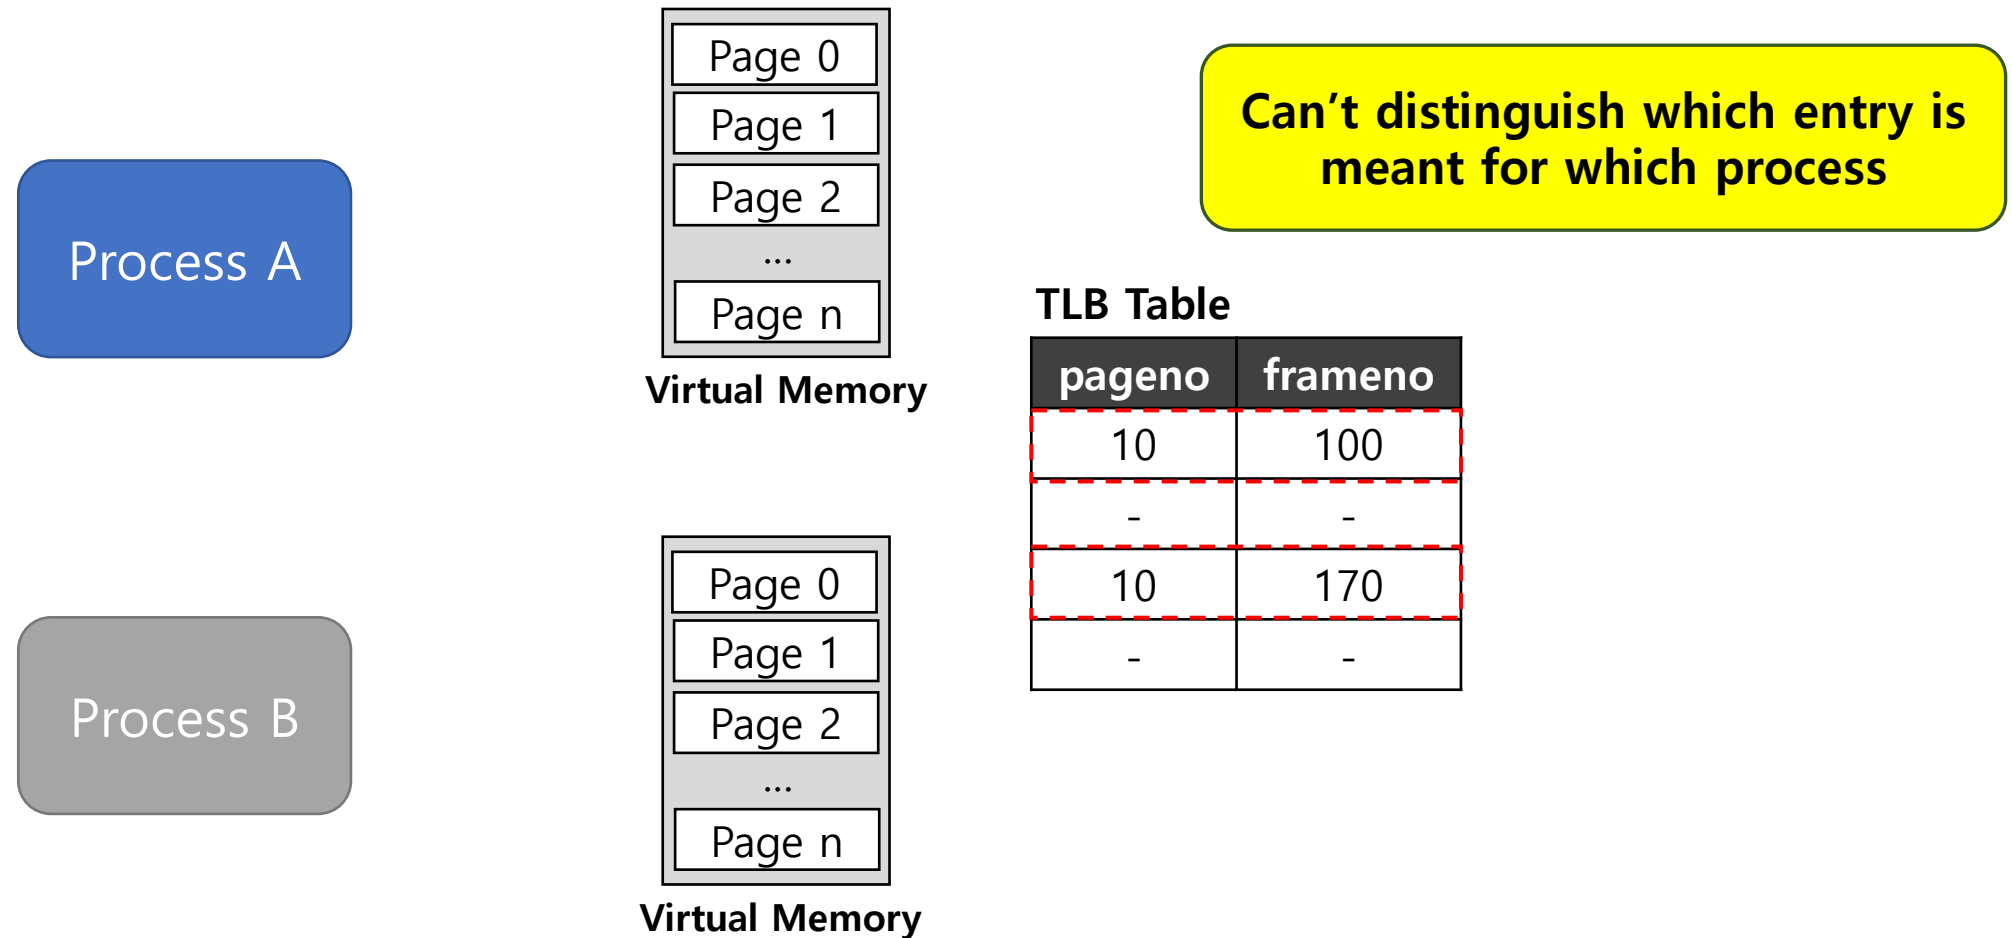




# TLB Issue: Context Switching



# TLB Issue: Context Switching

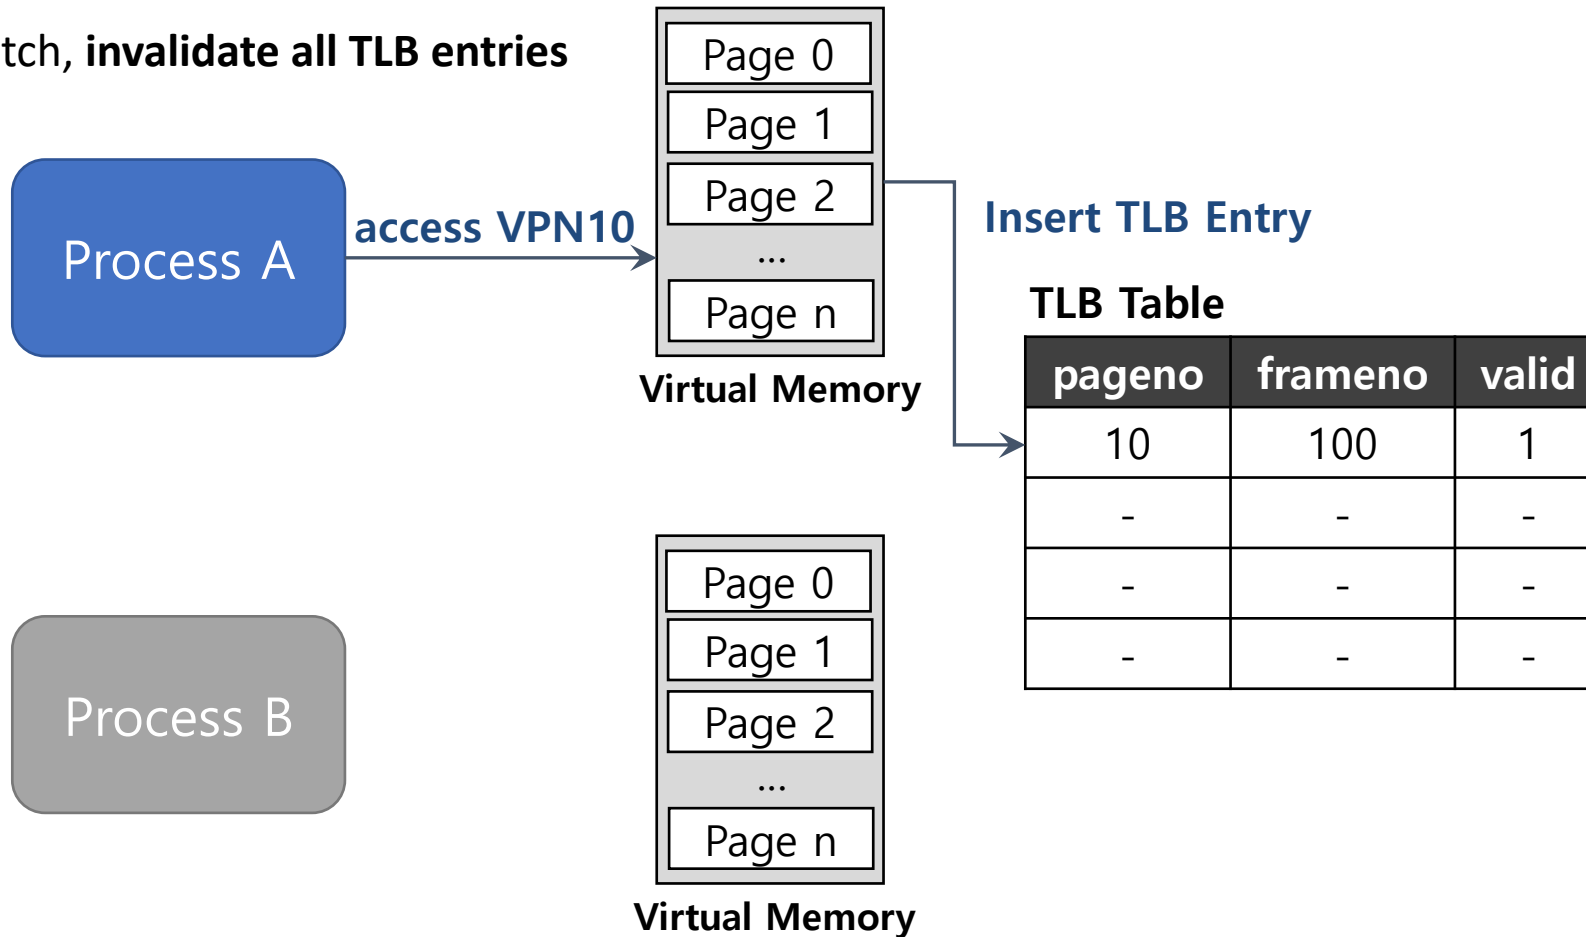


# TLB Issue: Context Switching – Solution 1

On process switch, **invalidate all TLB entries**




# TLB Issue: Context Switching – Solution 1

On process switch, **invalidate all TLB entries**



# TLB Issue: Context Switching – Solution 1

On process switch, **invalidate** all TLB entries

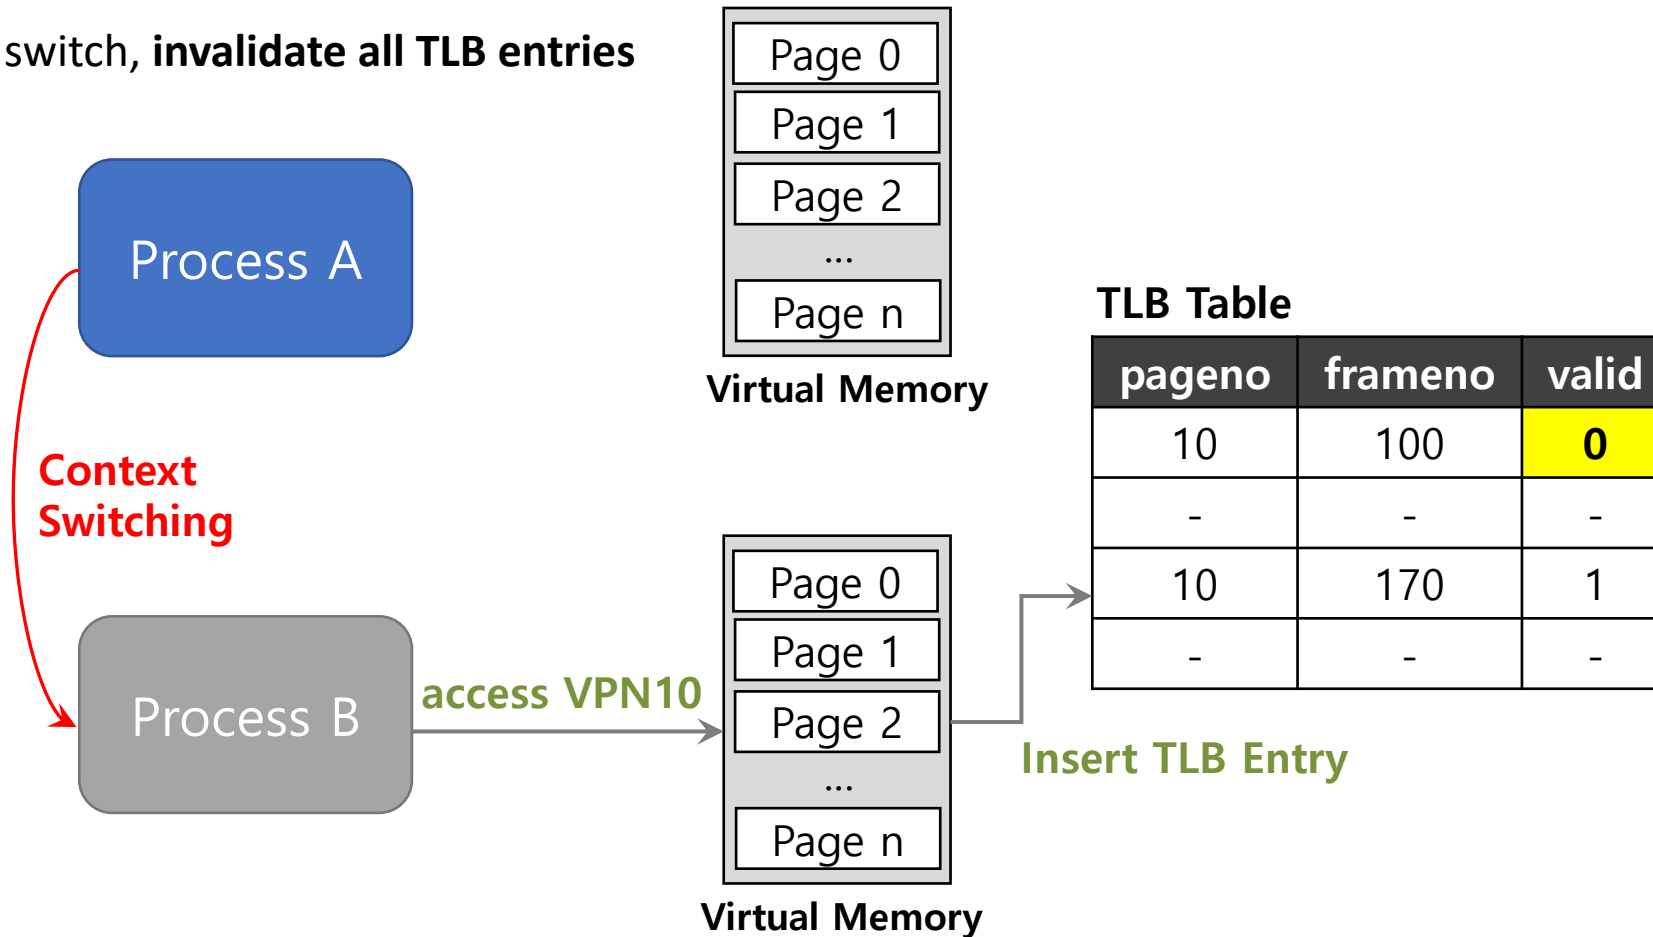


TLB Table

pageno	frameno	valid
10	100	0
-	-	-
-	-	-
-	-	-



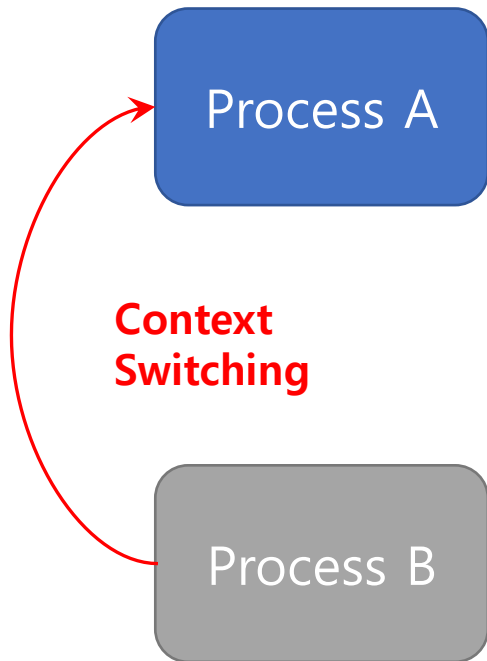
# TLB Issue: Context Switching – Solution 1

On process switch, **invalidate** all TLB entries



# TLB Issue: Context Switching – Solution 1

On process switch, **invalidate all TLB entries**



+ Simply requires invalid bit in each TLB entry

☹ Makes process switch expensive

☹ ☹ New process initially incurs 100% TLB misses

TLB Table

pageno	frameno	valid
10	100	0
-	-	-
10	170	0
-	-	-

# TLB Issue: Context Switching – Solution 2


- Add **process identifier** to TLB entries
  - **Match = match on pid AND match on pageno**
  - ☹️ **Makes TLB more complicated and expensive**
- Process switch
  - 😊 **Nothing to do**
  - 😊😊 **Cheaper**
- All modern machines have this feature



# TLB Issue: Context Switching – Solution 2


TLB match on pid AND match on pageno

Process A



Virtual Memory

Process B



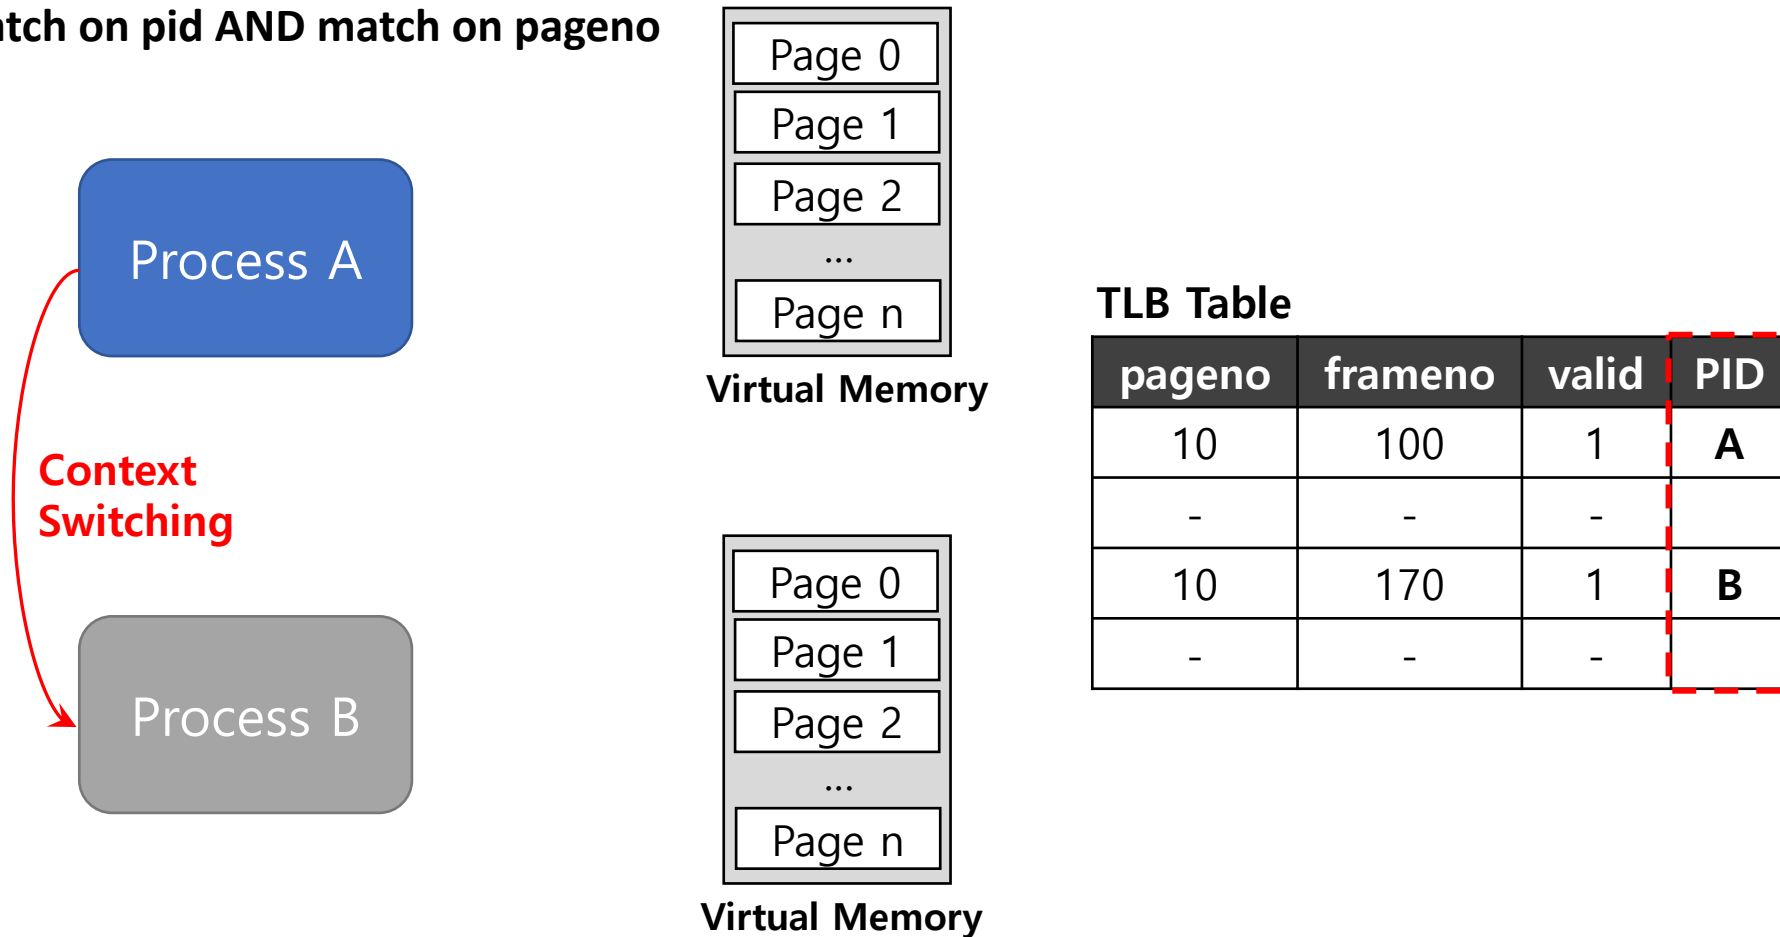
Virtual Memory

TLB Table

pageno	frameno	valid	PID
10	100	1	A
-	-	-	
10	170	1	B
-	-	-	

# TLB Issue: Context Switching – Solution 2

TLB match on pid AND match on pageno



# Further Reading

## **Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau**

Chapters 19–22

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

### **Credits:**

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Youjip Won (Hanyang University).