

# Week 8

## **Memory Management: Demand Paging**

Oana Balmau

February 21, 2023

# Announcements

Week 8 <b>Memory Management</b>	feb 20 C Review: Working with pthreads II	feb 21 <b>Demand Paging (3/3)</b> Optional reading: <a href="#">OSTEP</a> Chapters 19 – 22 Practice Exercises Sheet: Memory Management	feb 22	feb 23 Mid-semester Q&A – not recorded • Graded Exercises Sheet Released • Grades released for OS Shell Assignment	feb 24
Week 9 <b>Reading week</b>	feb 27 No class	feb 28 No class	mar 1 No class	mar 2 No class	mar 3 No class
Week 10 <b>File Systems</b>	mar 6 No lab. Work on Assignment 2 <b>Scheduling Assignment Due</b>	mar 7 <b>Intro to File Systems (1/2)</b> <b>Recorded lecture. Do not come to class.</b> Optional reading: <a href="#">OSTEP</a> Chapters 36, 37, 39	mar 8 <b>Memory Management Assignment Released</b>	mar 9 <b>Intro to File Systems (2/2)</b> Memory Management Assignment Overview – with Jiaxuan	mar 10
Week 11 <b>File Systems</b>	mar 13 <b>Graded Exercises Due</b> C Review: Complex structs	mar 14 <b>Basic File System Implementation (1/2)</b> Optional reading: <a href="#">OSTEP</a> Chapters 40, 41, 45	mar 15	mar 16 <b>Basic File System Implementation (2/2)</b> • Grades released for Scheduling Assignment	mar 17
Week 12 <b>File Systems</b>	mar 20 C Review: Pointers & Memory Allocation II	mar 21 <b>Advanced File System Implementation (1/2)</b>	mar 22	mar 23 <b>Advanced File System Implementation (2/2)</b>	mar 24
Week 13 <b>File Systems</b>	mar 27 C Review: Advanced debugging	mar 28 <b>Handling Crashes &amp; Performance (1/2)</b> Optional reading: <a href="#">OSTEP</a> Chapters 38, 43	mar 29	mar 30 <b>Handling Crashes &amp; Performance (2/2)</b> • Grades released for Exercises Sheet • Practice Exercises Sheet: File Systems	mar 31
Week 14 <b>Advanced Topics</b>	apr 3 No lab. Work on Assignment 3 <b>Memory Management Assignment Due</b>	apr 4 <b>Advanced topics: Virtualization</b>	apr 5	apr 6 <b>Advanced topics: Operating Systems Research</b> (Invited Speaker: TBD)	apr 7
Week 15 <b>Wrap-up</b>	apr 10 No Lab. Prepare for end-of-semester.	apr 11 <b>End-of-semester Q&amp;A</b> – not recorded	apr 12	apr 13 <b>End-of-semester Q&amp;A</b> – not recorded. <b>Last class!</b> • Grades released for Memory Management Assignment	apr 14

# Announcements

Week 8 <b>Memory Management</b>	feb 20 C Review: Working with pthreads II	feb 21 <b>Demand Paging (3/3)</b> Optional reading: <a href="#">OSTEP</a> Chapters 19 – 22 Practice Exercises Sheet: Memory Management	feb 22	feb 23 Mid-semester Q&A – not recorded • Graded Exercises Sheet Released • Grades released for OS Shell Assignment	feb 24
Week 9 <b>Reading week</b>	feb 27 No class	feb 28 No class	mar 1 No class	mar 2 No class	mar 3 No class
Week 10 <b>File Systems</b>	mar 6 No lab. Work on Assignment 2 <b>Scheduling Assignment Due</b>	mar 7 <b>Intro to File Systems (1/2)</b> Recorded lecture. Do not come to class. Optional reading: <a href="#">OSTEP</a> Chapters 36, 37, 39	mar 8 <b>Memory Management Assignment Released</b>	mar 9 <b>Intro to File Systems (2/2)</b> Memory Management Assignment Overview – with Jiaxuan	mar 10
Week 11 <b>File Systems</b>	mar 13 <b>Graded Exercises Due</b> C Review: Complex structs	mar 14 <b>Basic File System Implementation (1/2)</b> Optional reading: <a href="#">OSTEP</a> Chapters 40, 41, 45	mar 15	mar 16 <b>Basic File System Implementation (2/2)</b> • Grades released for Scheduling Assignment	mar 17
Week 12 <b>File Systems</b>	mar 20 C Review: Pointers & Memory Allocation II	mar 21 <b>Advanced File System Implementation (1/2)</b>	mar 22	mar 23 <b>Advanced File System Implementation (2/2)</b>	mar 24
Week 13 <b>File Systems</b>	mar 27 C Review: Advanced debugging	mar 28 <b>Handling Crashes &amp; Performance (1/2)</b> Optional reading: <a href="#">OSTEP</a> Chapters 38, 43	mar 29	mar 30 <b>Handling Crashes &amp; Performance (2/2)</b> • Grades released for Exercises Sheet • Practice Exercises Sheet: File Systems	mar 31
Week 14 <b>Advanced Topics</b>	apr 3 No lab. Work on Assignment 3 <b>Memory Management Assignment Due</b>	apr 4 <b>Advanced topics: Virtualization</b>	apr 5	apr 6 <b>Advanced topics: Operating Systems Research</b> (Invited Speaker: TBD)	apr 7
Week 15 <b>Wrap-up</b>	apr 10 No Lab. Prepare for end-of-semester.	apr 11 <b>End-of-semester Q&amp;A</b> – not recorded	apr 12	apr 13 <b>End-of-semester Q&amp;A</b> – not recorded. <b>Last class!</b> • Grades released for Memory Management Assignment	apr 14



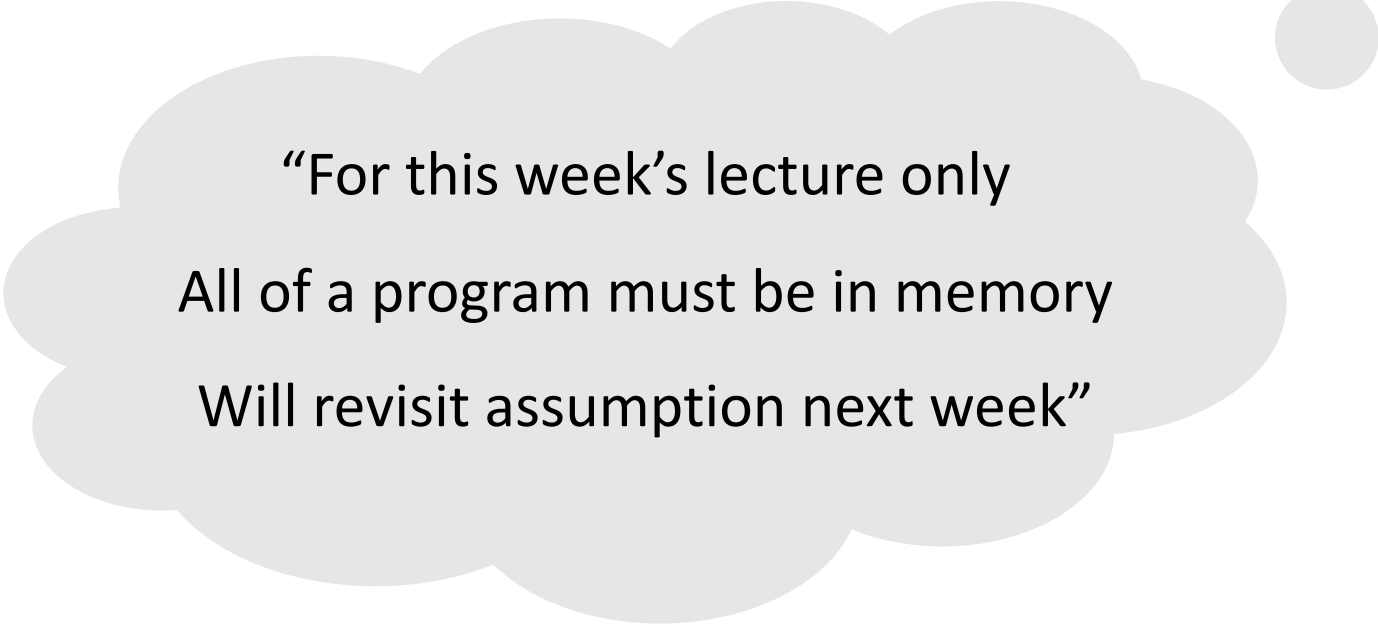
# Announcements

Week 8 <b>Memory Management</b>	feb 20 C Review: Working with pthreads II	feb 21 <b>Demand Paging (3/3)</b> Optional reading: <a href="#">OSTEP</a> Chapters 19 – 22 Practice Exercises Sheet: Memory Management	feb 22	feb 23 Mid-semester Q&A – not recorded • Graded Exercises Sheet Released • Grades released for OS Shell Assignment	feb 24
Week 9 <b>Reading week</b>	feb 27 No class	feb 28 No class	mar 1 No class	mar 2 No class	mar 3 No class
Week 10 <b>File Systems</b>	mar 6 No lab. Work on Assignment 2 <b>Scheduling Assignment Due</b>	mar 7 <b>Intro to File Systems (1/2)</b> <b>Recorded lecture. Do not come to class.</b> Optional reading: <a href="#">OSTEP</a> Chapters 36, 37, 39	mar 8 <b>Memory Management Assignment Released</b>	mar 9 <b>Intro to File Systems (2/2)</b> Memory Management Assignment Overview – with Jiaxuan	mar 10
Week 11 <b>File Systems</b>	mar 13 <b>Graded Exercises Due</b> C Review: Complex structs	mar 14 <b>Basic File System Implementation (1/2)</b> Optional reading: <a href="#">OSTEP</a> Chapters 40, 41, 45	mar 15	mar 16 <b>Basic File System Implementation (2/2)</b> • Grades released for Scheduling Assignment	mar 17
Week 12 <b>File Systems</b>	mar 20 C Review: Pointers & Memory Allocation II	mar 21 <b>Advanced File System Implementation (1/2)</b>	mar 22	mar 23 <b>Advanced File System Implementation (2/2)</b>	mar 24
Week 13 <b>File Systems</b>	mar 27 C Review: Advanced debugging	mar 28 <b>Handling Crashes &amp; Performance (1/2)</b> Optional reading: <a href="#">OSTEP</a> Chapters 38, 43	mar 29	mar 30 <b>Handling Crashes &amp; Performance (2/2)</b> • Grades released for Exercises Sheet • Practice Exercises Sheet: File Systems	mar 31
Week 14 <b>Advanced Topics</b>	apr 3 No lab. Work on Assignment 3 <b>Memory Management Assignment Due</b>	apr 4 <b>Advanced topics: Virtualization</b>	apr 5	apr 6 <b>Advanced topics: Operating Systems Research</b> (Invited Speaker: TBD)	apr 7
Week 15 <b>Wrap-up</b>	apr 10 No Lab. Prepare for end-of-semester.	apr 11 <b>End-of-semester Q&amp;A</b> – not recorded	apr 12	apr 13 <b>End-of-semester Q&amp;A</b> – not recorded. <b>Last class!</b> • Grades released for Memory Management Assignment	apr 14

# Demand Paging

# *Remember from Last Week*

## Simplifying Assumption



“For this week’s lecture only  
All of a program must be in memory  
Will revisit assumption next week”


***We are now going to drop this assumption***

# What if “out of memory”?

Need to get rid of one or more processes

Store them temporarily on disk


This is called **Swapping**



# Process Switch to a Swapped Process?

Latency can be **very high**

Need to read image from disk





# Process Switch to a Swapped Process?

A better solution:

- **Demand paging**
- **Not all of a process needs to be in memory**

# Main Reason for Demand Paging

## **Virtual address spaces >> physical address space**

- No machine has  $2^{64}$  bytes (16 exabytes) of memory (yet).

## Why such large virtual address space?

- Convenient for programmer
- Don't have to worry about running out

# Additional Benefits of Demand Paging

## + Shorter process startup latency

- Can start process without all of it in memory
- Even 1 page suffices

## + Better use of main memory

- Program often does not use certain parts
  - E.g., error handling routines
- Program often goes through different parts
  - E.g., initialization, computation, termination

# If the program is not in memory, then where is it?

Part of it is in memory

(Typically) all of it is on disk

- In a special partition called the **Backing Store**

# If the program is not in memory, then where is it?

Part of it is in memory

(Typically) all of it is on disk

- In a special partition called the **Backing Store**

Note the difference with swapping:

**Swapping** = *all* of program is **in memory** OR *all* of program is **on disk**

**Demand paging** = **part** of program is **in memory**

# If the program is not in memory, then where is it?

Part of it is in memory

(Typically) all of it is on disk

- In a special partition called the **Backing Store**

## Remember:

- CPU can only directly access memory
- CPU can only access data on disk through OS

Note the difference with swapping:

**Swapping** = *all* of program is **in memory** OR *all* of program is **on disk**

**Demand paging** = **part** of program is **in memory**



# Demand Paging Mechanism High Level

- What if program needs to access part only on disk?

# Demand Paging Mechanism High Level

- What if program needs to access part only on disk?
- Program is suspended
- OS runs, gets page from disk
- Program is restarted

# Demand Paging Mechanism High Level

- What if program needs to access part only on disk?

This is called a **page fault**

- Program is suspended
- OS runs, gets page from disk
- Program is restarted

This is called **page fault handling**

# Demand Paging Issues

1. How to discover a page fault?
2. How to suspend process?
3. How to get a page from disk?
  - 3'. How to find a free frame in memory?
4. How to restart process?

# 1. Discover Page Fault

## Idea: Use the valid bit in page table

- Without demand paging:
  - Valid bit = 0: page is invalid
  - Valid bit = 1: page is valid
- **With demand paging**
  - Valid bit = 0: page is invalid **OR page is on disk**
  - Valid bit = 1: page is valid **AND page is in memory**
  - OS needs additional table: invalid / on-disk?

## 2. Suspending the Faulting Process

### Idea: Trap into the OS

- Invalid bit access **generates trap**
- Save process information in PCB when trapping into the OS



# 3. Getting the Page from Disk

## Idea: OS handles fetch from Disk

- Assume (for now) there is at least one free frame in memory
- Allocate a free frame to process
- Find page on disk
  - OS needs to remember the swap space in page-sized units.
  - Need an extra table for that in OS.
- Get disk to transfer page from disk to frame

# 3. Getting the Page from Disk

## Idea: OS handles fetch from Disk

- Assume (for now) there is at least one free frame in memory
- Allocate a free frame to process
- Find page on disk
  - OS needs to remember the swap space in page-sized units.
  - Need an extra table for that in OS.
- Get disk to transfer page from disk to frame



### 3. While the Disk is Busy



- Shouldn't waste CPU cycles
- Invoke scheduler to run another process
- When **disk interrupt** arrives
  - Suspend running process
  - Get **back to page fault handling**

### 3. While the Disk is Busy



Why OK to go through OS (and not in hardware)?

- Disk is so slow that it masks the latency of handling this in OS.

### 3. Completing Page Fault Handling

- `Pagetable[pageno].framenno = new framenno`
- **`Pagetable[pageno].valid = 1`**
- Set process state to ready
- Invoke scheduler

## 4. When Process Runs Again

**Idea: Restarts the previously faulting instruction**

Process now finds

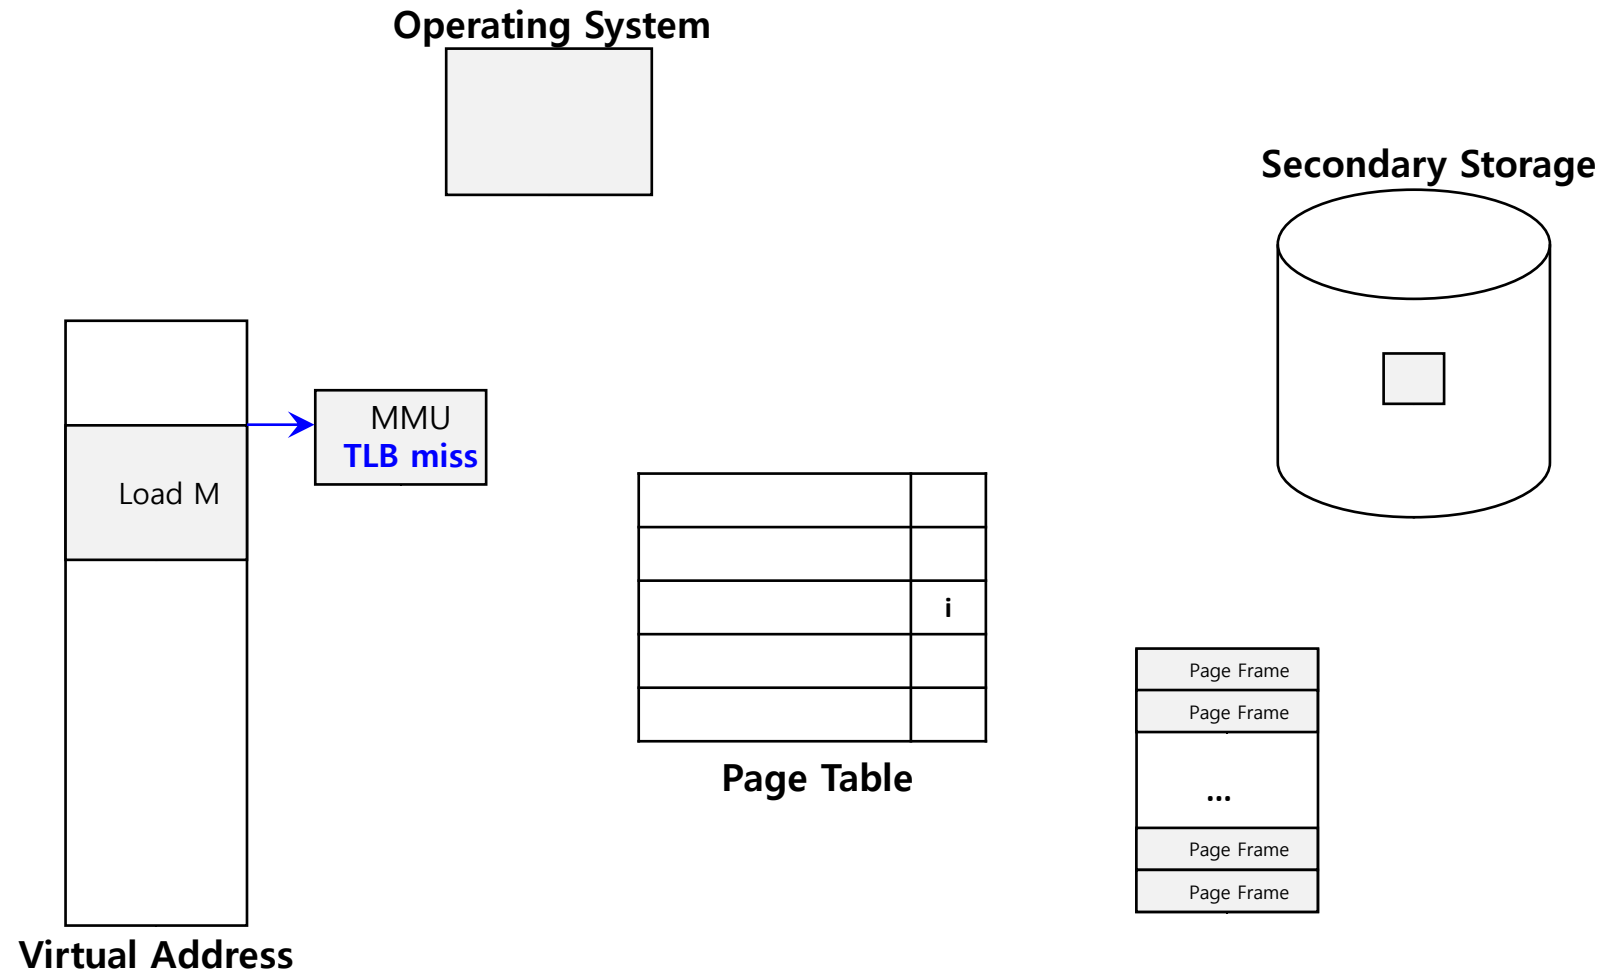
- Valid bit to be set to 1
- Page in corresponding frame in memory
- Note: different from context switch, which continues with the next instruction



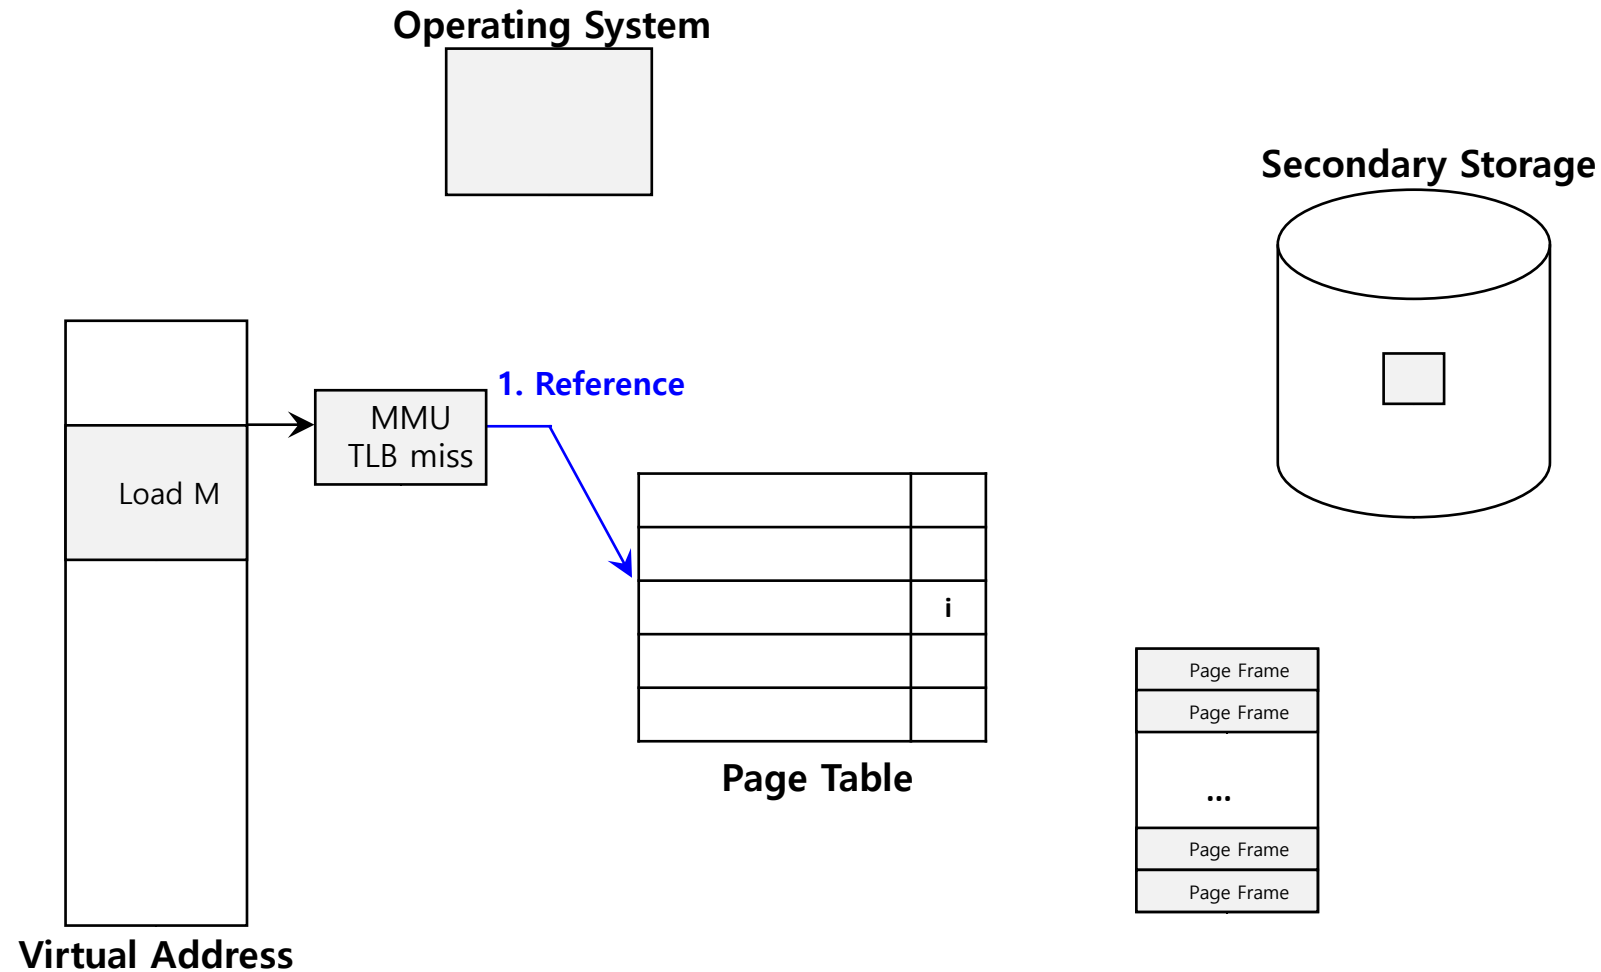
# Demand Paging Summary

1. How to discover a page fault?      ← **Use valid bit in page table**
2. How to suspend process?      ← **Invalid bit traps into OS**
3. How to get a page from disk?      ← **OS handles fetch from disk**
  - 3'. How to find a free frame in memory?
4. How to restart process?      ← **Restart faulting instruction**

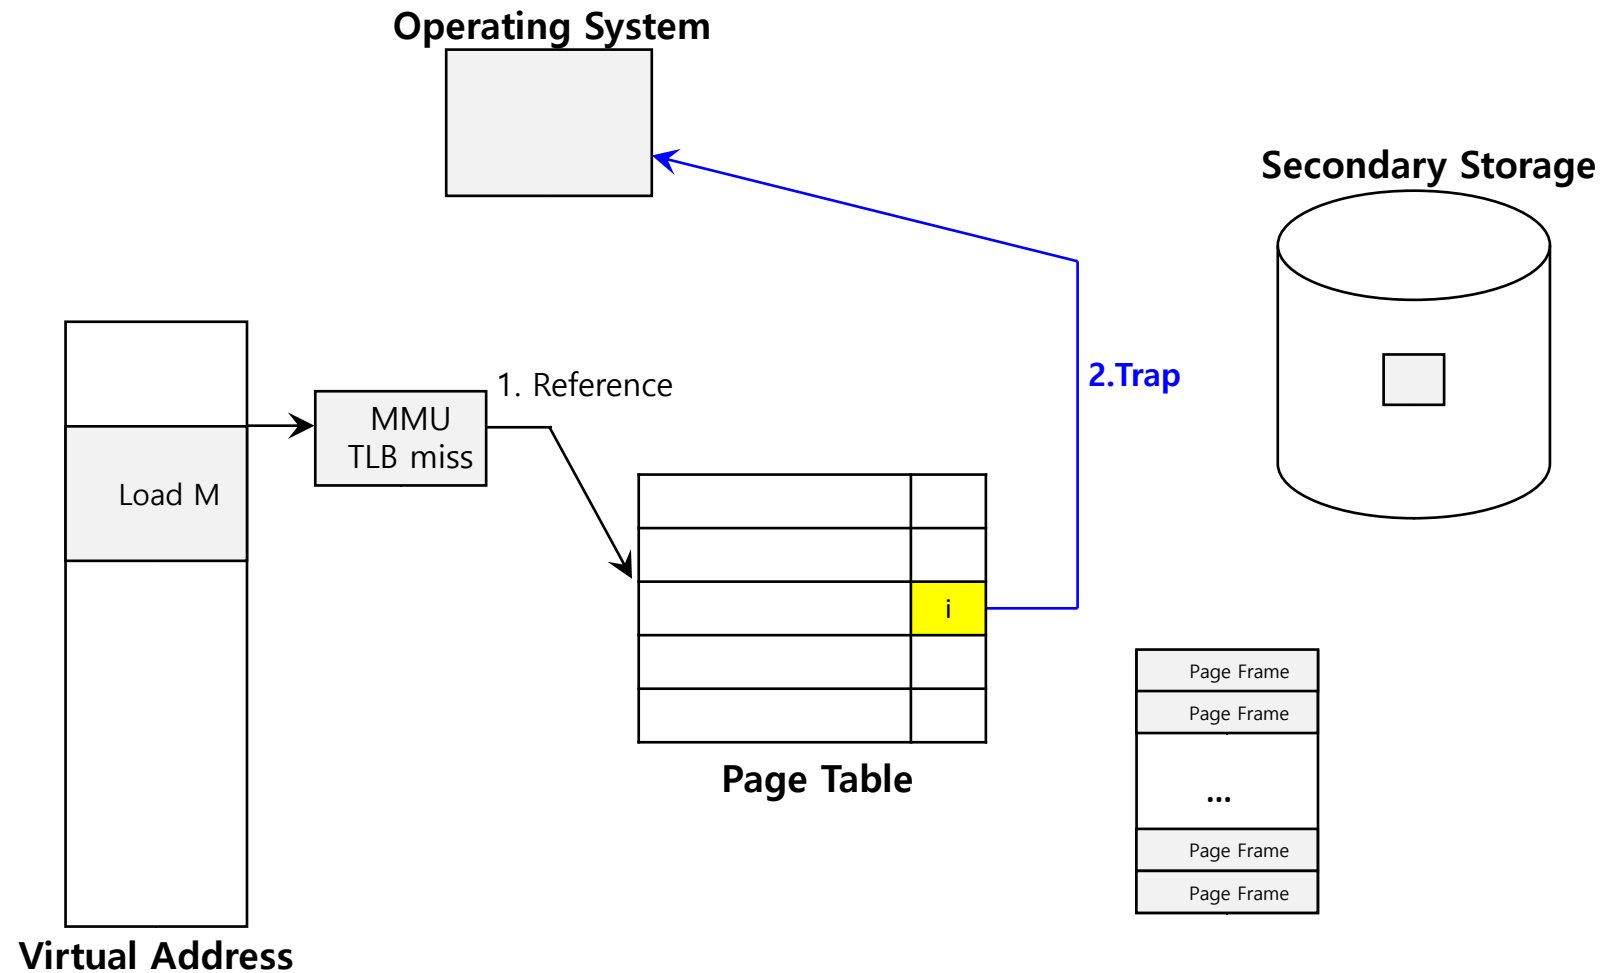
# Demand Paging Summary



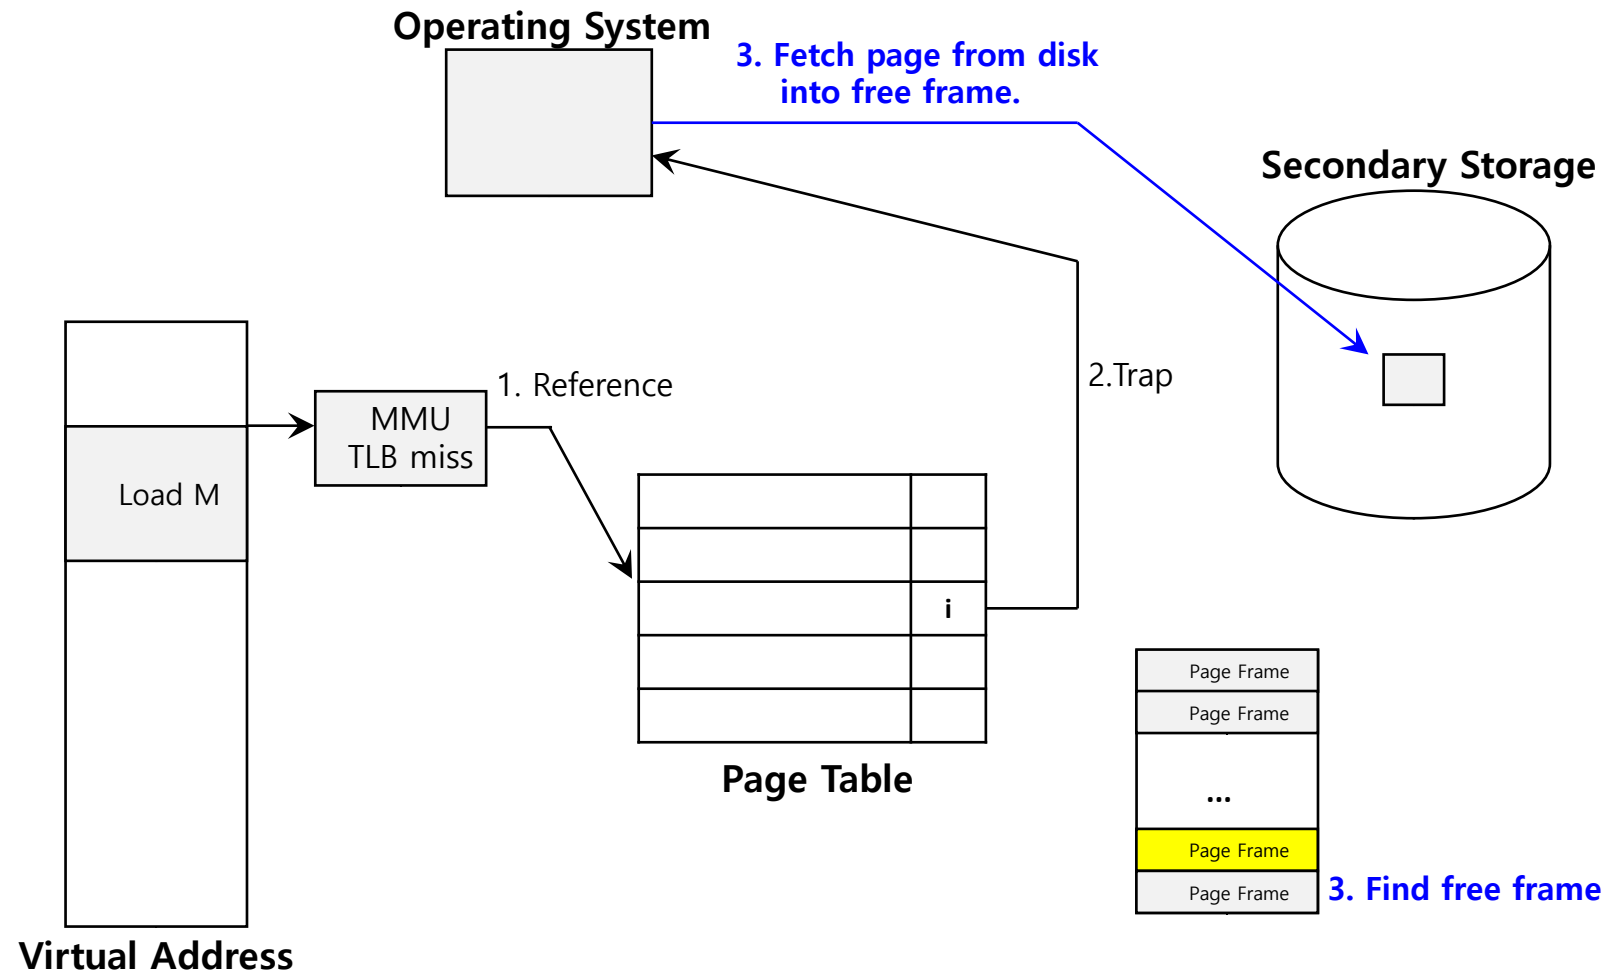
# Demand Paging Summary



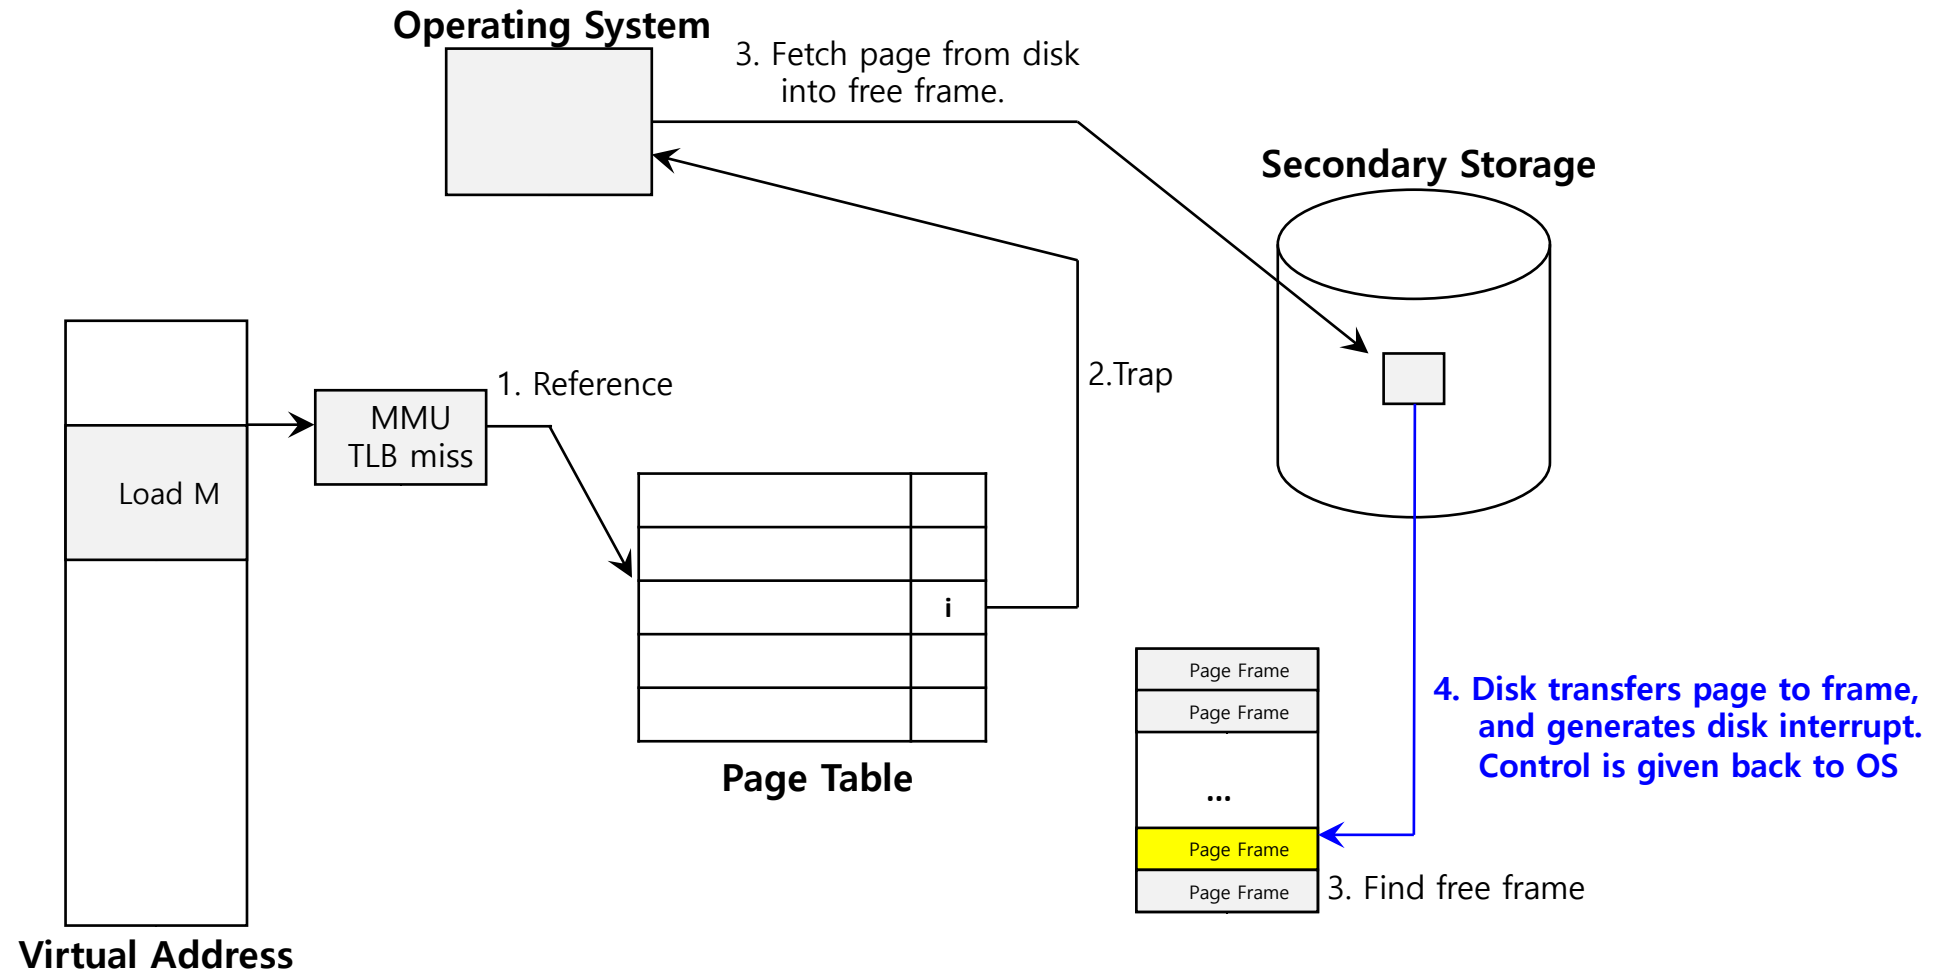
# Demand Paging Summary



# Demand Paging Summary

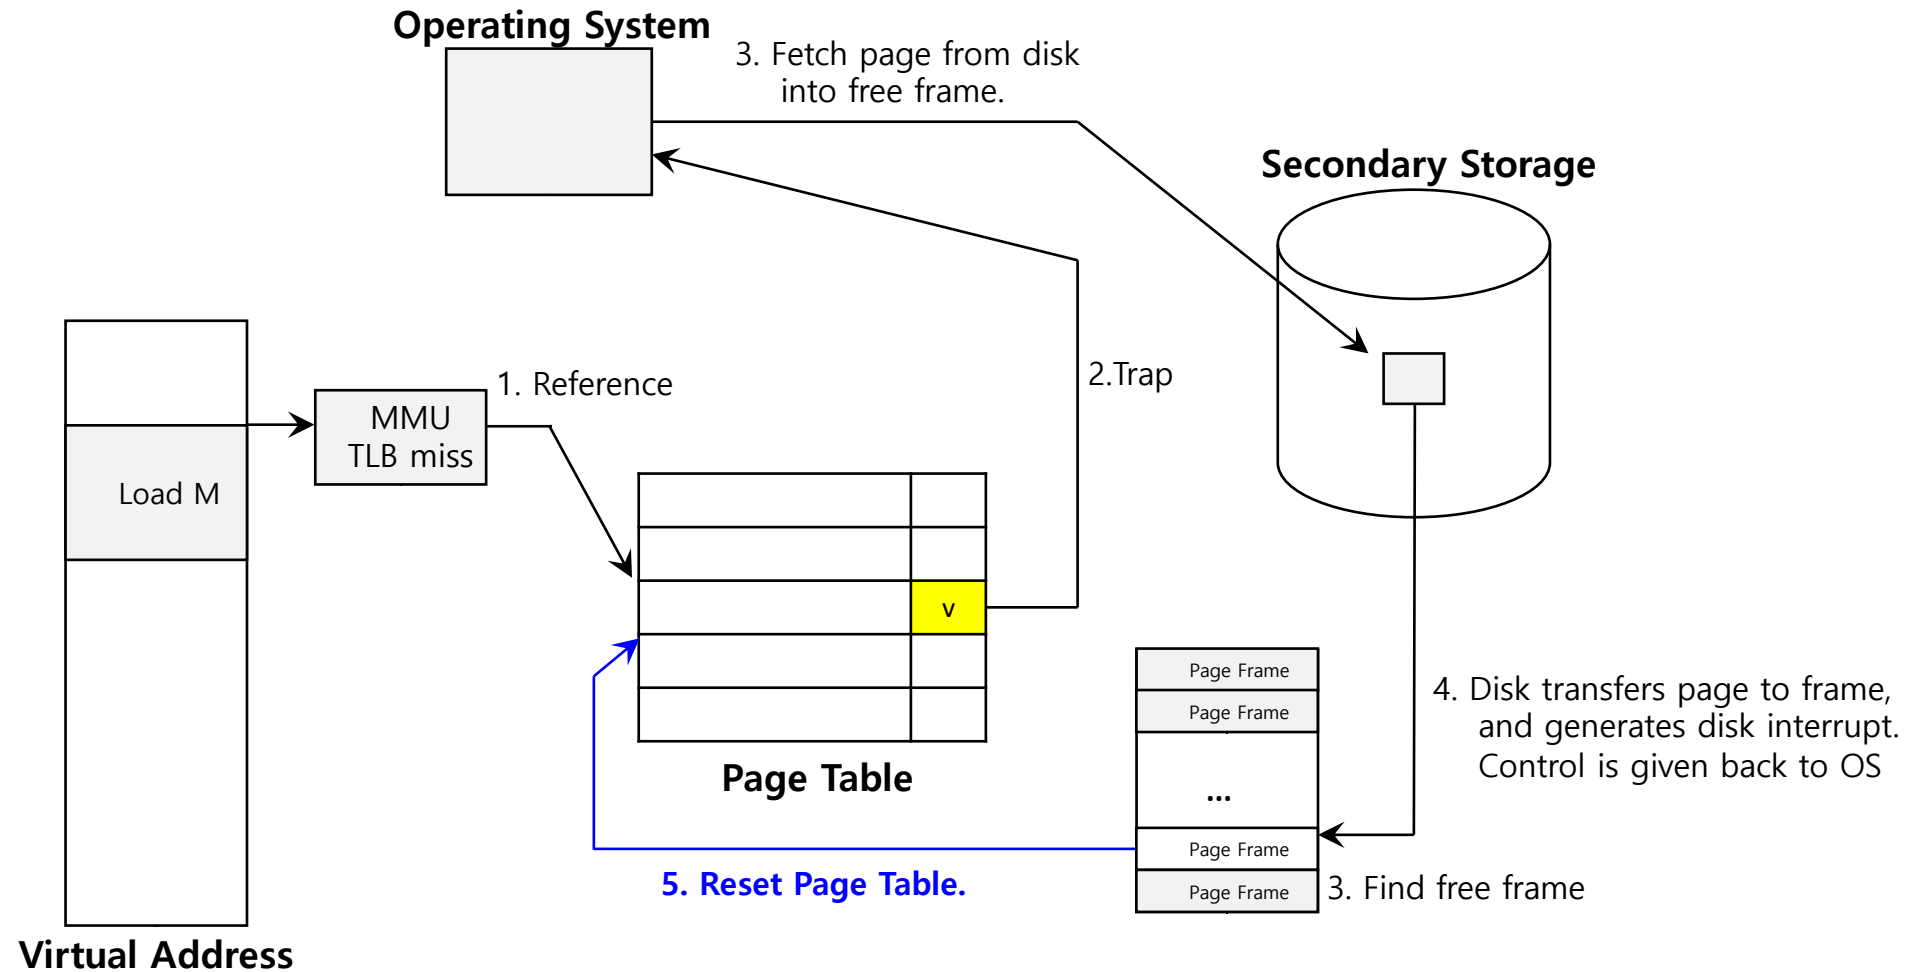


# Demand Paging Summary

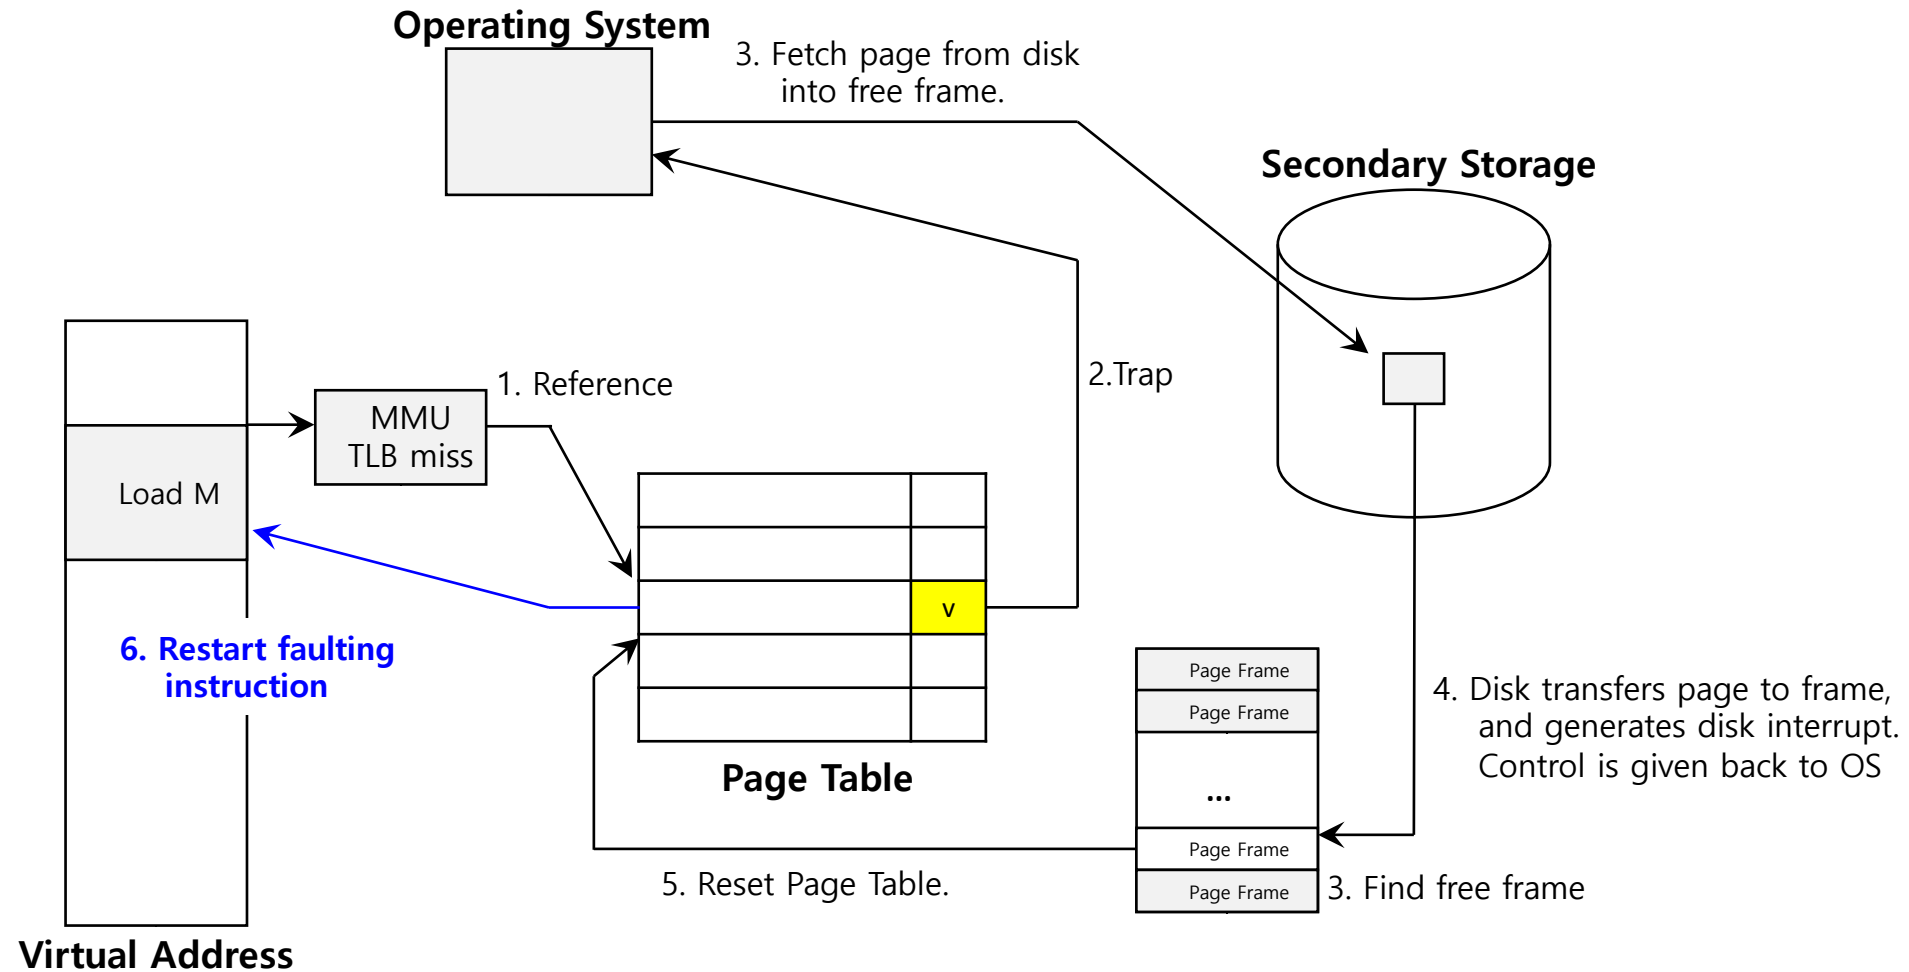




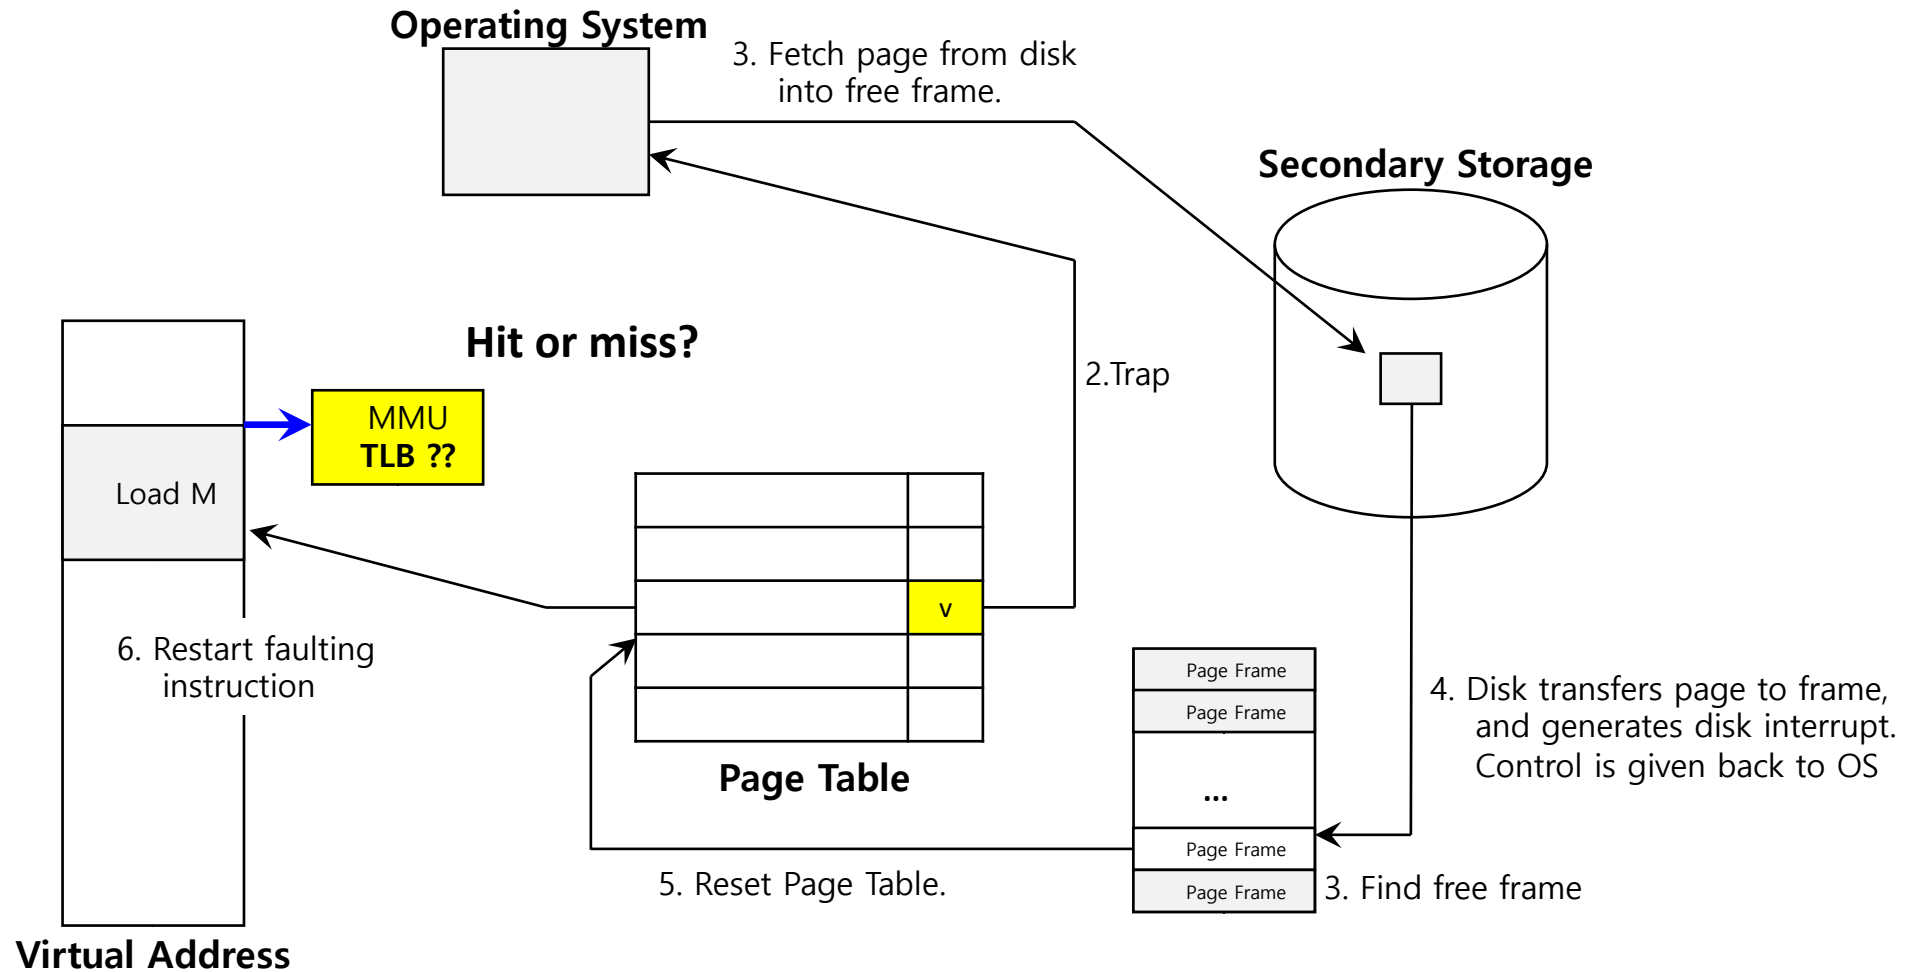
# Demand Paging Summary



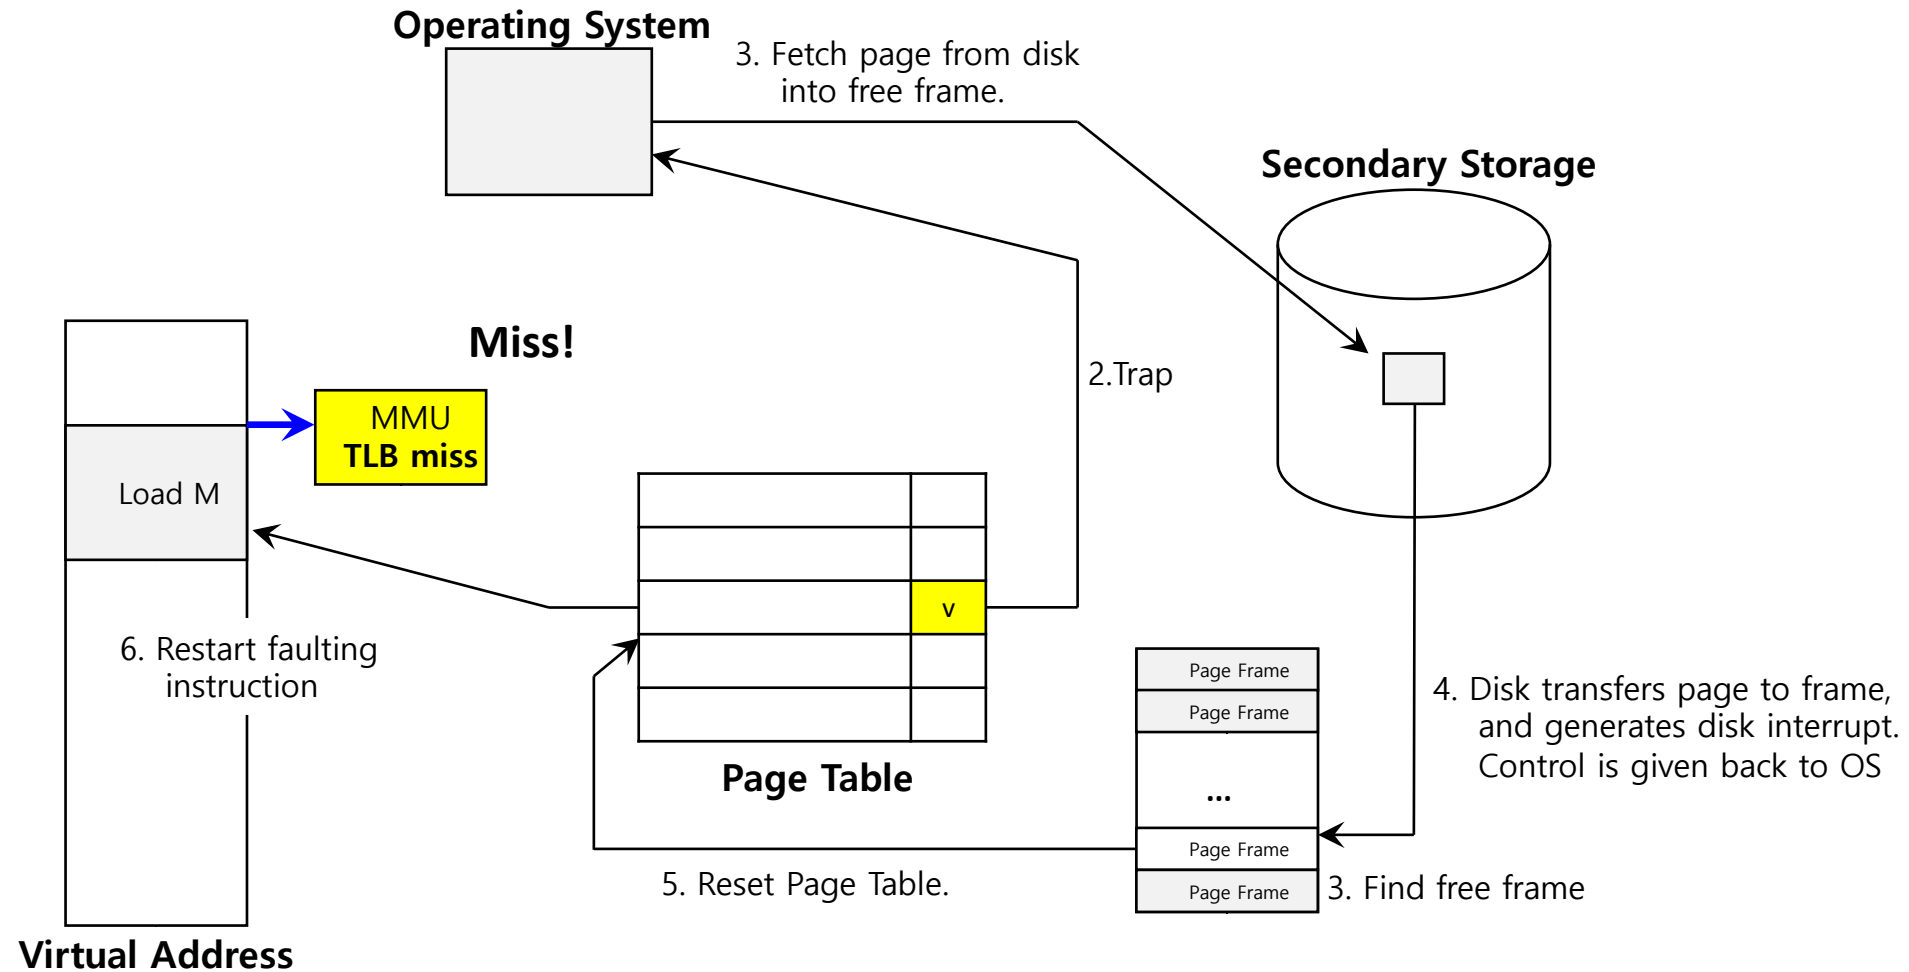
# Demand Paging Summary



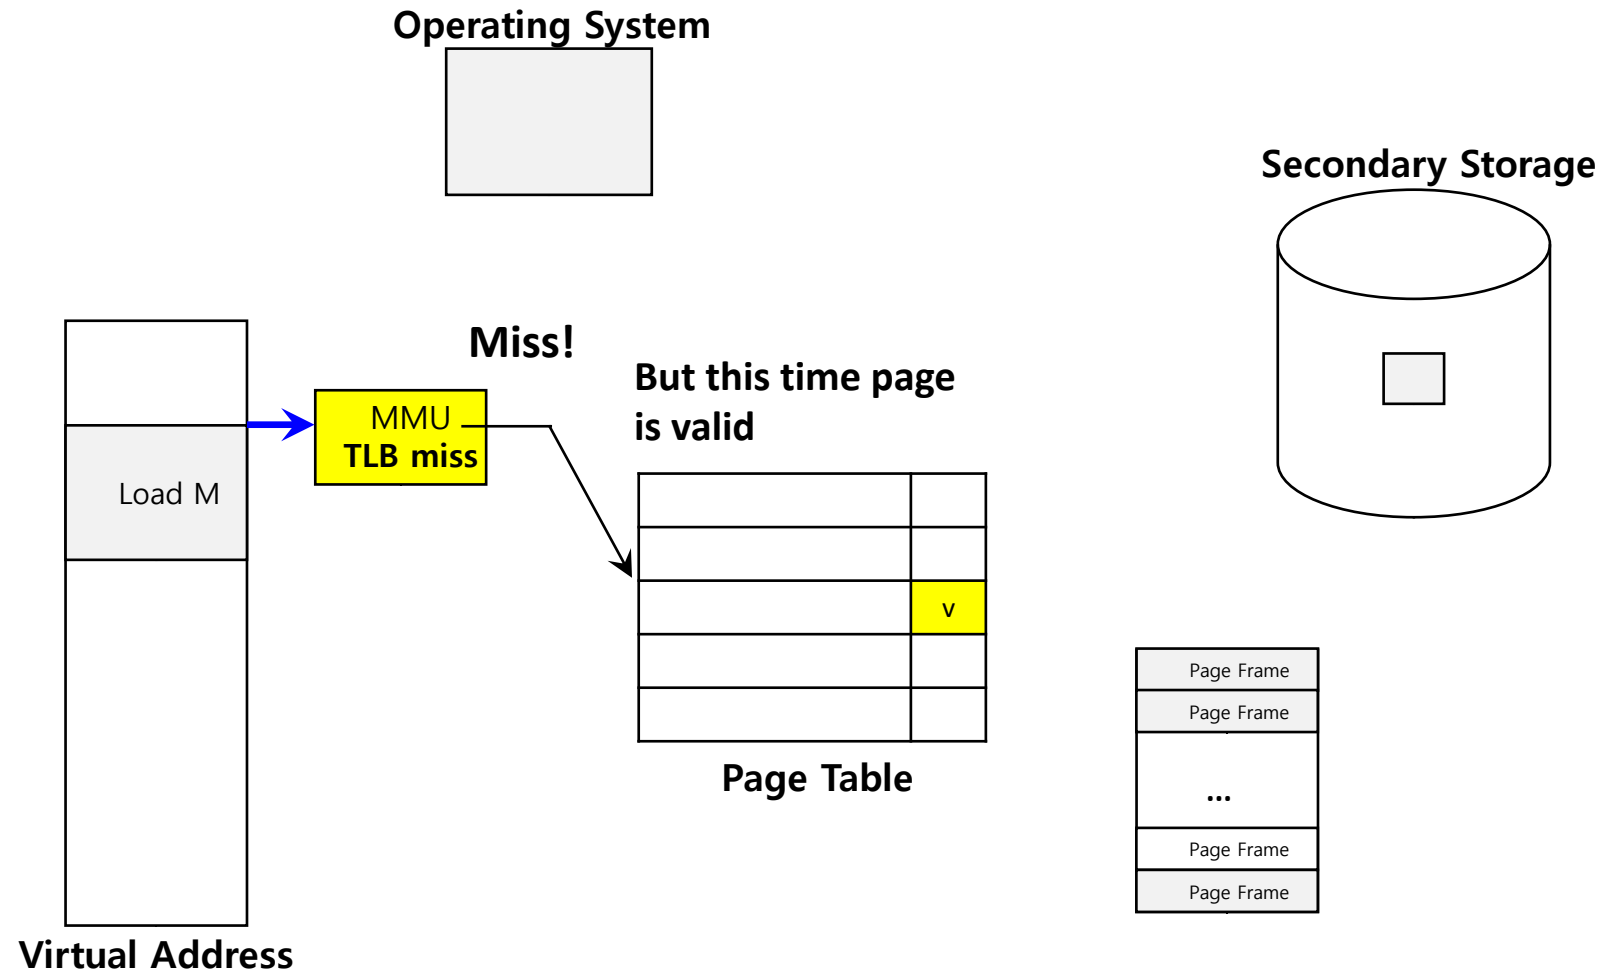
# Demand Paging Summary



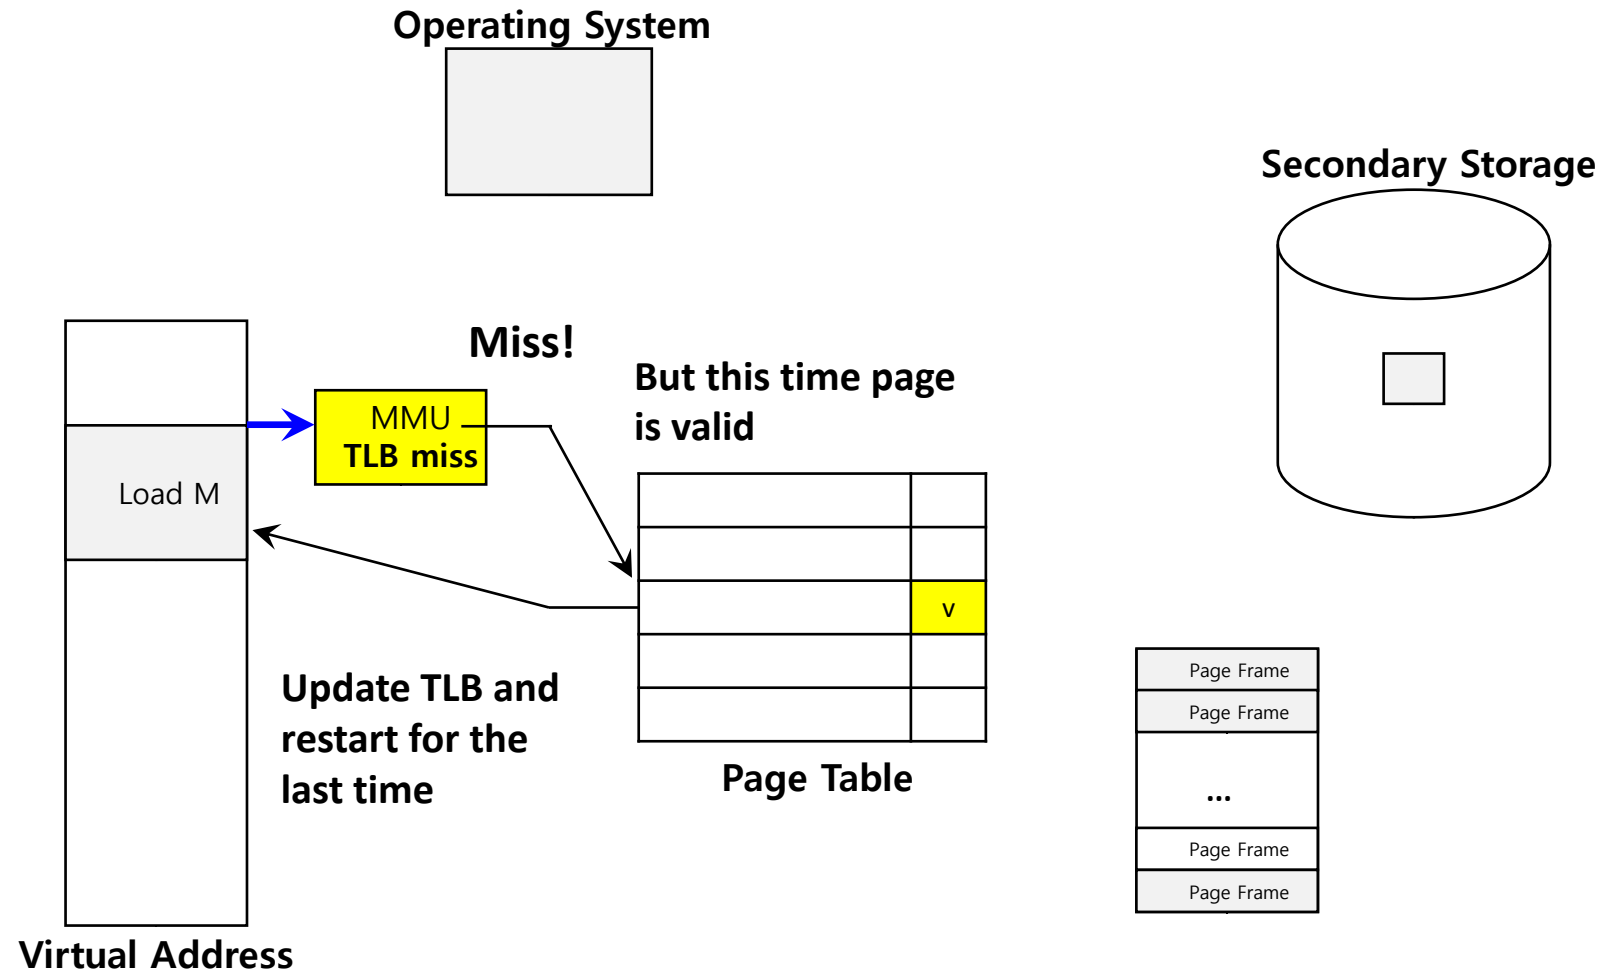
# Demand Paging Summary



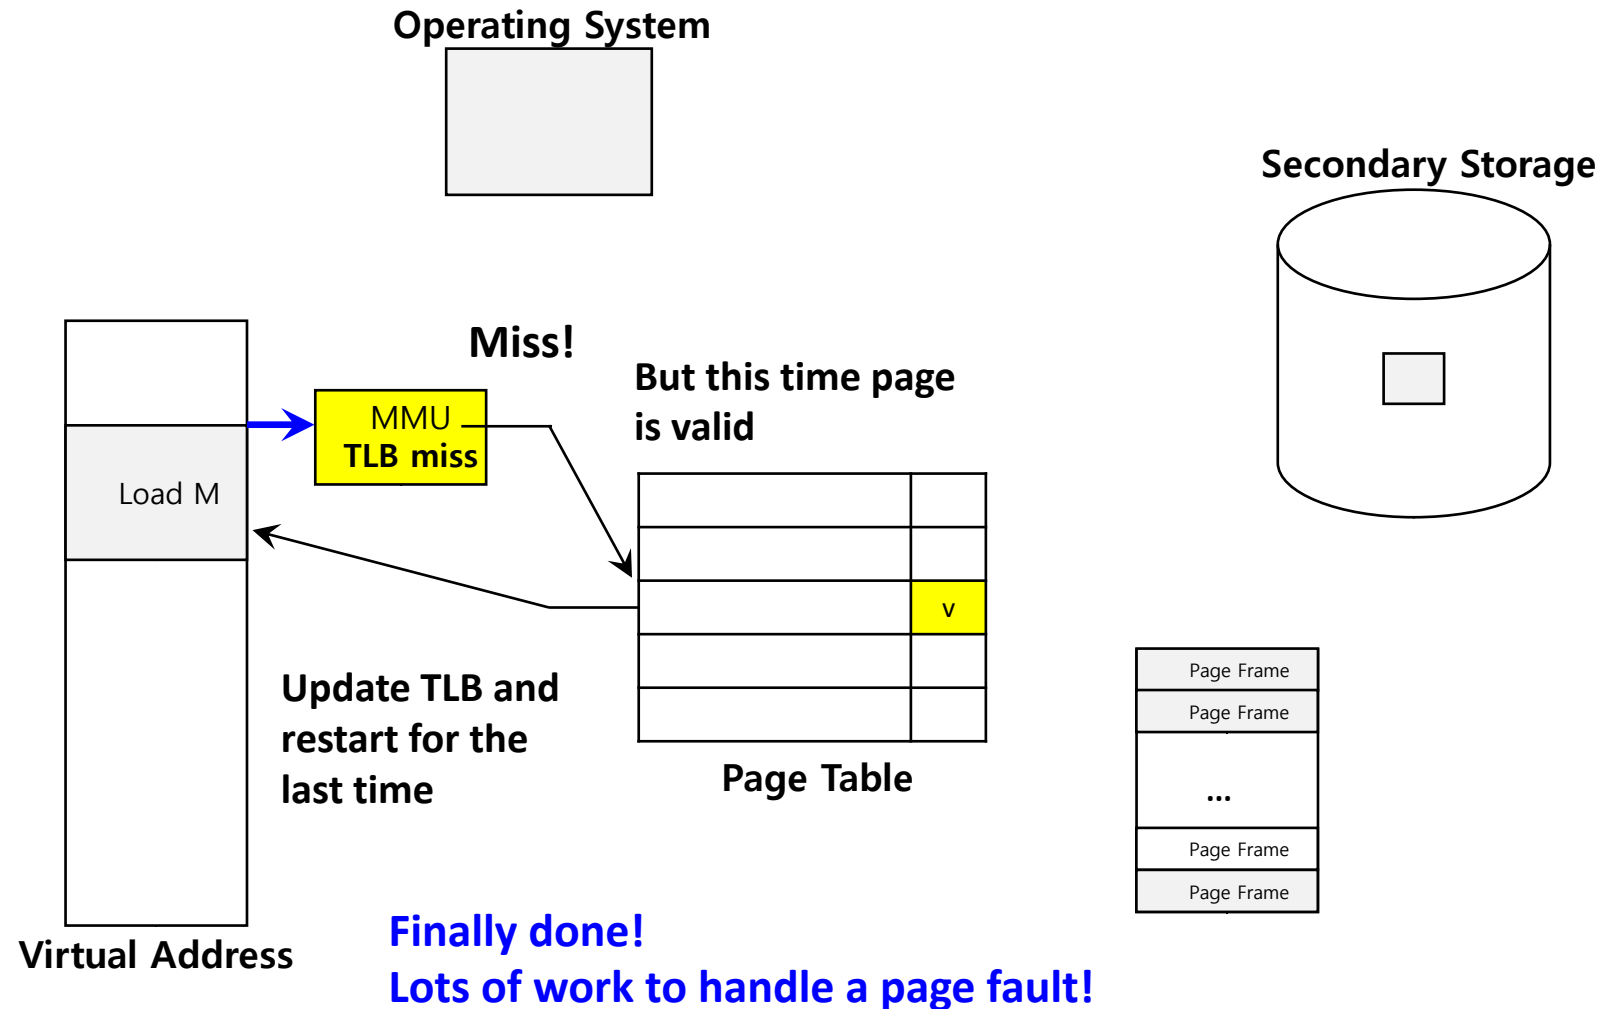
# Demand Paging Summary



# Demand Paging Summary



# Demand Paging Summary



# *Remember Assumption:* Getting the Page from Disk

“Assume (for now) there is at least one free frame in memory”

If no free frame available:

- Pick a frame to be replaced
- Invalidate its page table entry (and TLB entry)
- You may have to write that frame to disk



# *Remember Assumption:* Getting the Page from Disk

“Assume (for now) there is at least one free frame in memory”

If no free frame available:

- Pick a frame to be replaced
- Invalidate its page table entry (and TLB entry)
- You may have to write that frame to disk
  - Page table has a modified bit
  - If set, write out page to disk
  - If not, proceed with page fault handling

# *Remember Assumption:* Getting the Page from Disk

“Assume (for now) there is at least one free frame in memory”

If no free frame available:

- Pick a frame to be replaced
  - **How?**
- Invalidate its page table entry (and TLB entry)
- You may have to write that frame to disk
  - Page table has a modified bit
  - If set, write out page to disk
  - If not, proceed with page fault handling

# How to pick with page/frame to replace?

Different page replacement policies:

- Random
- FIFO (First In, First Out)
- OPT
- LRU

# Page Faults and Performance

- Normal memory access
  - ~ nanoseconds
- Faulting memory access
  - Disk i/o ~ 10 milliseconds
- Too many page faults -> **program very slow**
- **Hence, importance of good page replacement policy**

# Page Replacement Policies

- Random
- FIFO (First In, First Out)
- OPT
- LRU

# Page Replacement Policies

- Random
- FIFO (First In, First Out)
- OPT
- LRU

Plus, in general, prefer replacing clean over dirty

- 1 disk i/o instead of 2

# Random

## Random page is replaced

+ Easy to implement

☹ Does not take advantage of spatial/temporal locality.

# FIFO

## Oldest page is replaced

- Age = Time since brought into memory

## + Easy to implement

- Keep a queue of pages
- Bring in a page: stick at the end of the queue
- Need replacement: pick head of queue

## + Fair

- All pages receive equal residency

☹ Does not take into account “hot” pages that may always be needed



# Quiz: FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7
	0	0
		1

page frames

?? page faults (not counting initial paging in)

# Quiz: FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																
	0	0	0	2	2	4	4	4	0				0	0			7	7	7
				3	3	3	2	2	2				1	1			1	0	0
		1	1	1	0	0	0	3	3				3	2			2	2	1

page frames

**12** page faults (not counting initial paging in)

# OPT: An Optimal Algorithm

Replace the page that will be referenced the furthest in the future

+ Provably optimal

☹ Can't implement (Can't predict the future)

A basis of comparison for other algorithms



# Quiz: Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7
	0	0
		1

page frames

?? page faults (not counting initial paging in)

# Quiz: Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2	2		2		2							7
	0	0	0		0	4		0		0							0
		1	1		3	3		3		1							1

page frames

**6** page faults (not counting initial paging in)

# LRU: Least Recently Used

Cannot look into the future, but can try to predict future using past  
Replace least recently accessed page

+ With locality, LRU approximates OPT

☹ Harder to implement, must track which pages have been accessed

☹ Does not handle all workloads well

- Example: Large array scans that repeat. Popular in DBMS.

# LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7
	0	0
		1

page frames

?? page faults (not counting initial paging in)

# LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

**9** page faults (not counting initial paging in)



# LRU Implementation

- Too expensive to implement exactly
  - Need to timestamp every memory reference
- But can be (well) approximated

# LRU Approximation with Hardware Support

## Use a **reference bit**

- Bit in page table
- Hardware sets bit when page is referenced

## Periodically

- Read out and store all reference bits
- Reset all reference bits to zero

## Keep all reference bits for some time

- The more bits kept, the better approximation

## Replacement

- **Page with smallest value of reference bit history**

# Let's practice!


# Page Replacement Policies

Consider a **cache of size 5**.

1. Generate worst-case address reference streams for FIFO and LRU. Worst-case reference streams cause the most misses possible.
2. Compare with OPT in all scenarios from point 1.
3. For the worst case reference streams, how much bigger of a cache is needed to improve performance dramatically and approach OPT?

# Page Replacement Policies

Worst case FIFO:



# Page Replacement Policies

Worst case FIFO:

1      2      3      4      5      6      1      2      3      4      5      6

1	1	1	1	1
	2	2	2	2
		3	3	3
			4	4
				5

# Page Replacement Policies


Worst case FIFO:

1 2 3 4 5 6 1 2 3 4 5 6

1	1	1	1	1	6	6	6	6	6	5	5
	2	2	2	2	2	1	1	1	1	1	6
		3	3	3	3	3	2	2	2	2	2
			4	4	4	4	4	3	3	3	3
				5	5	5	5	5	4	4	4

# Page Replacement Policies

Worst case LRU:





# Page Replacement Policies

Worst case LRU:

1      2      3      4      5      6      1      2      3      4      5      6

1	1	1	1	1
	2	2	2	2
		3	3	3
			4	4
				5

# Page Replacement Policies

Worst case LRU:

1 2 3 4 5 6 1 2 3 4 5 6

1	1	1	1	1	6	6	6	6	6	5	5
	2	2	2	2	2	1	1	1	1	1	6
		3	3	3	3	3	2	2	2	2	2
			4	4	4	4	4	3	3	3	3
				5	5	5	5	5	4	4	4

# Page Replacement Policies

OPT:

**1** **2** **3** **4** **5** 6 1 2 3 4 5 6

<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>
	<b>2</b>	<b>2</b>	<b>2</b>	<b>2</b>
		<b>3</b>	<b>3</b>	<b>3</b>
			<b>4</b>	<b>4</b>
				<b>5</b>

# Page Replacement Policies


OPT:

1 2 3 4 5 6 1 2 3 4 5 6

1	1	1	1	1	1							5
	2	2	2	2	2							2
		3	3	3	3							3
			4	4	4							4
				5	6							6

# Page Replacement Policies


OPT:



**For the worst case reference streams, how much bigger of a cache is needed to improve performance dramatically and approach OPT?**

# Page Replacement Policies

OPT:



**For the worst case reference streams, how much bigger of a cache is needed to improve performance dramatically and approach OPT?**

→ Need a cache of 6

# Page Replacement Policies

FIFO/LRU, cache of 6:

1 2 3 4 5 6 1 2 3 4 5 6

1	1	1	1	1	1
	2	2	2	2	2
		3	3	3	3
			4	4	4
				5	5
					6

# Summary – Key Concepts

- TLB
- Page Table for very large address spaces
- Demand paging
  - Page fault
- Page replacement policies:
  - Random
  - FIFO (First In, First Out)
  - OPT
  - LRU (Least Recently Used)



# Further Reading

## **Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau**

Chapters 19–22

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

### **Credits:**

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Youjip Won (Hanyang University).