

Week 10

Persistent Storage: Intro to File Systems

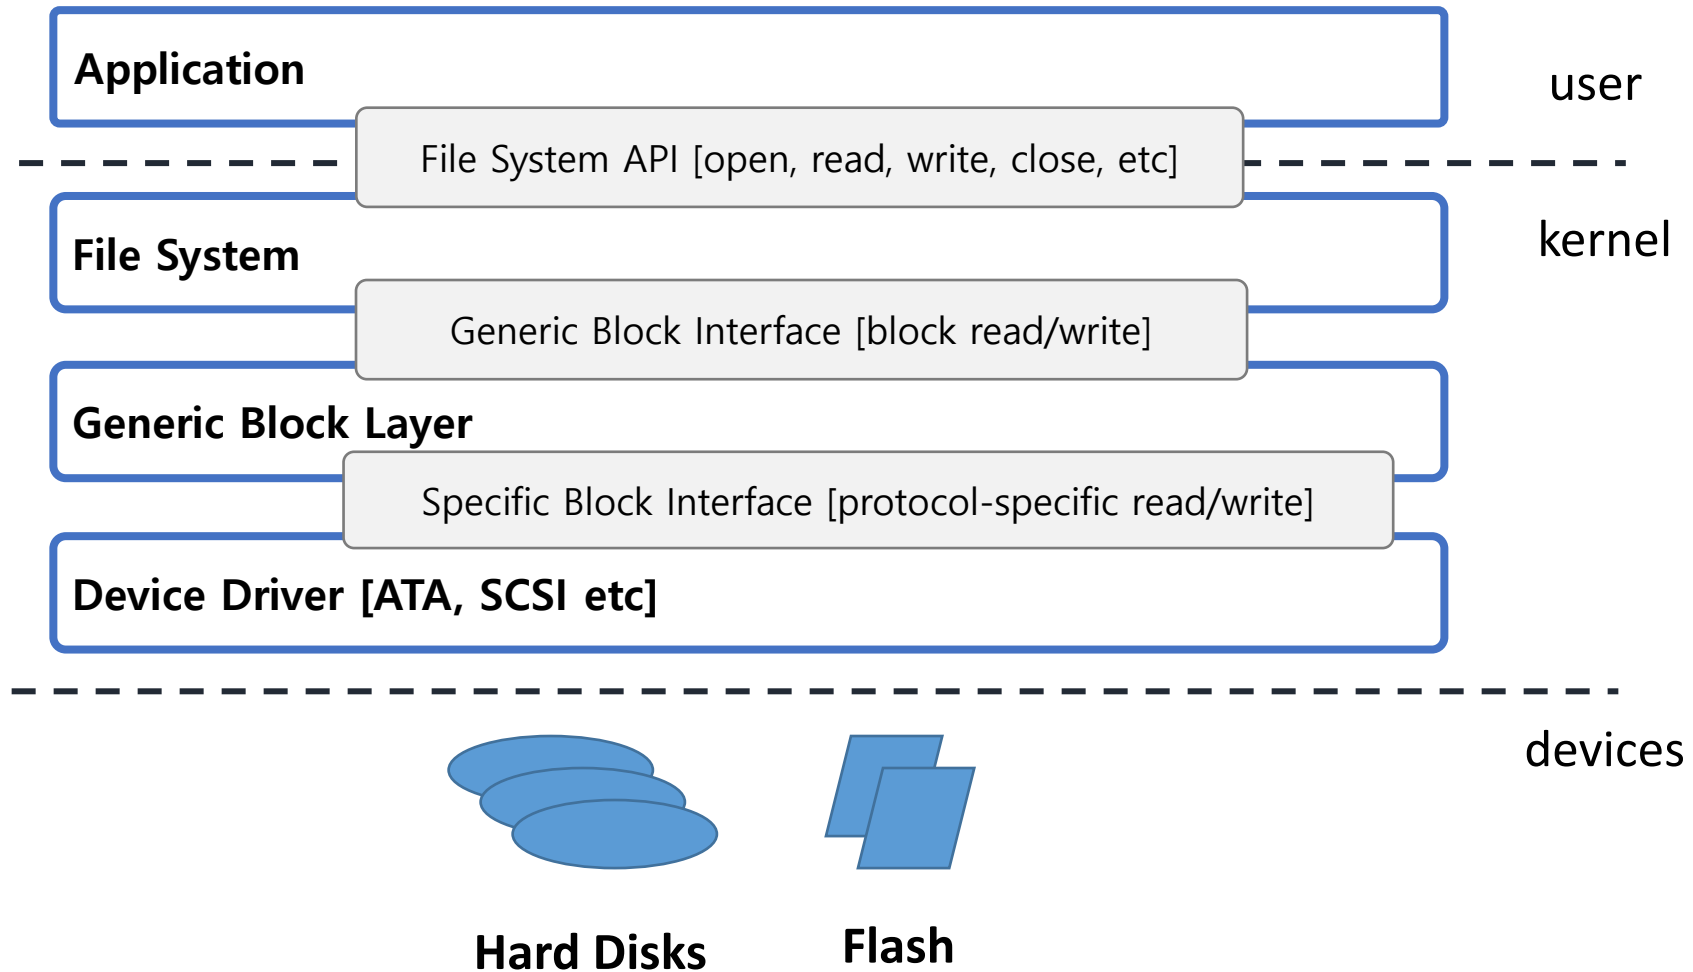
Oana Balmau

March 9, 2023


Class Admin

Week 10 File Systems	mar 6 No lab. Work on Assignment 2 Scheduling Assignment Due	mar 7 Intro to File Systems (1/2) Recorded lecture. Do not come to class. Optional reading: OSTEP Chapters 36, 37, 39	mar 8 Memory Management Assignment Released	mar 9 Intro to File Systems (2/2) Memory Management Assignment Overview — with Jiaxuan	mar 10
Week 11 File Systems	mar 13 Scheduling Assignment Due Graded Exercises Due C Review: Complex structs	mar 14 Basic File System Implementation (1/2) Optional reading: OSTEP Chapters 40, 41, 45	mar 15	mar 16 Basic File System Implementation (2/2) * Grades released for Scheduling Assignment	mar 17
Week 12 File Systems	mar 20 Graded Exercises Due C Review: Pointers & Memory Allocation II	mar 21 Advanced File System Implementation (1/2)	mar 22	mar 23 Advanced File System Implementation (2/2) * Grades released for Scheduling Assignment	mar 24
Week 13 File Systems	mar 27 C Review: Advanced debugging	mar 28 Handling Crashes & Performance (1/2) Optional reading: OSTEP Chapters 38, 43	mar 29	mar 30 Handling Crashes & Performance (2/2) * Grades released for Exercises Sheet * Practice Exercises Sheet: File Systems	mar 31
Week 14 Advanced Topics	apr 3 No lab. Work on Assignment 3 Memory Management Assignment Due	apr 4 Advanced topics: Virtualization	apr 5	apr 6 Advanced topics: Operating Systems Research (Invited Speaker: TBD) Grades released for Exercises Sheet	apr 7
Week 15 Wrap-up	apr 10 No Lab. Prepare for end-of-semester. Memory Management Assignment Due	apr 11 End-of-semester Q&A— not recorded	apr 12	apr 13 End-of-semester Q&A — not recorded. Last class!	apr 14 Grades released for Memory Management Assignment

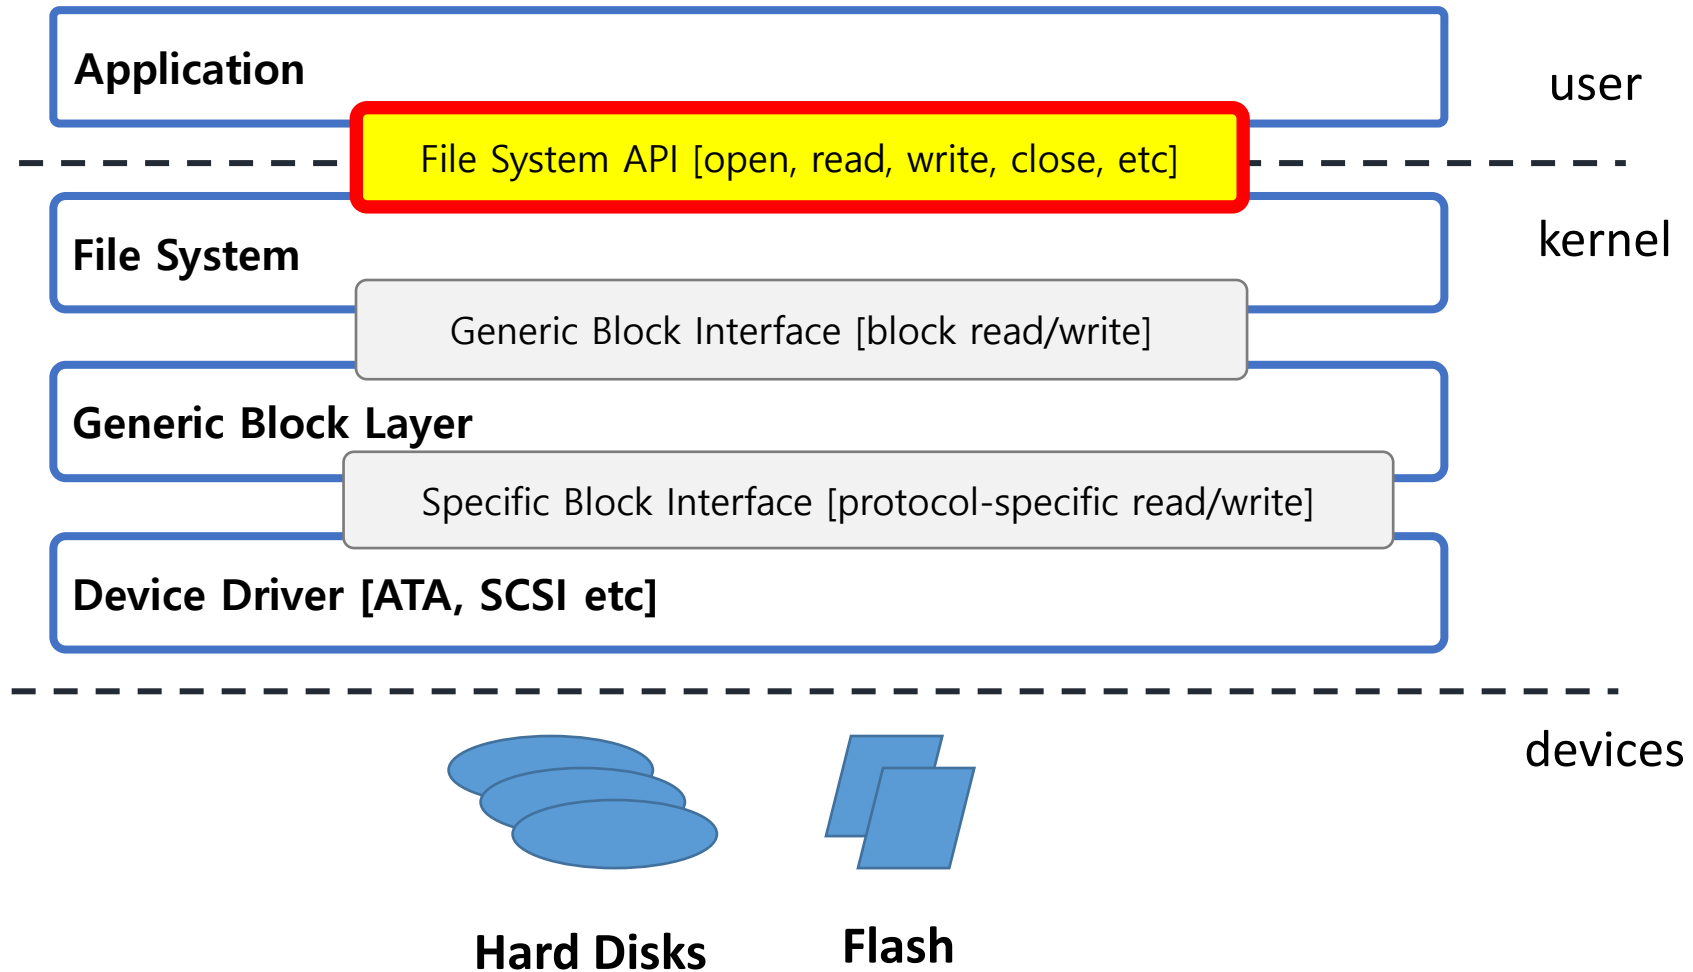
Recap: Overall Picture




On Tuesday we talked about this part



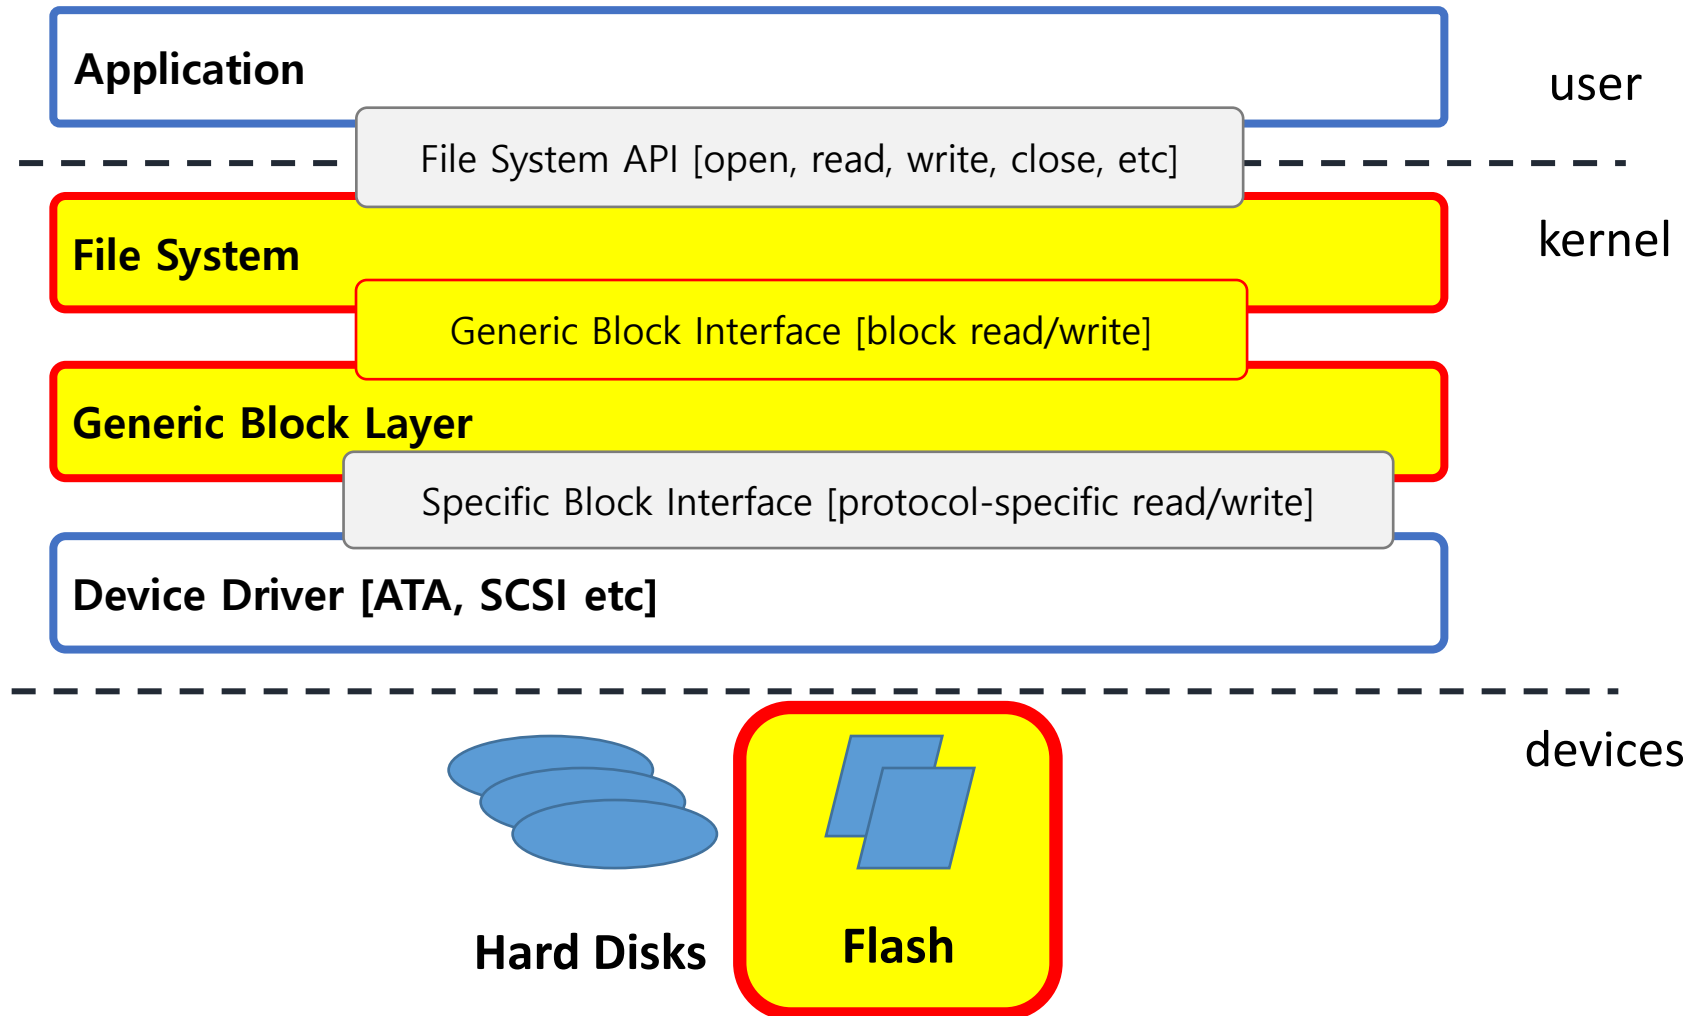
Then we talked about this part




And this part. Today we continue with this part.



Next Weeks' Lectures



Recap: Disk Terminology



Disk Access Time

Component	Time
Head Selection	nanoseconds
Seek Time	3-12 milliseconds
Rotational Latency	2-7 milliseconds
Transfer Time	microseconds
Controller Overhead	< 1 millisecond

Disk Access Time

Seek time dominates



Component	Time
Head Selection	nanoseconds
Seek Time	3-12 milliseconds
Rotational Latency	2-7 milliseconds
Transfer Time	microseconds
Controller Overhead	< 1 millisecond


Note: Disk access time >> memory access time (nanoseconds)

Let's Practice! Disk Access Time

Consider a disk with a sector size of 512 bytes, 1,000 tracks per surface, 100 sectors per track, 5 double-sided platters and a block size of 2,048 bytes. Suppose that the average seek time is 10ms, the average rotational delay is 5 ms, and the transfer rate is 200 MB per second. Suppose that a file containing 1,000,000 records of 100 bytes each is to be stored on such a disk and that no record is allowed to span two blocks.

- A. How many blocks are required to store the entire file?
- B. If the file is arranged sequentially on disk, how many cylinders are needed?
- C. What is the time required to read the file sequentially?

Let's Practice! Disk Access Time



sector size of 512 bytes, 1,000 tracks per surface, 100 sectors per track, 5 double-sided platters and a block size of 2,048 bytes.

Let's Practice! Disk Access Time

Consider a disk with a sector size of 512 bytes, 1,000 tracks per surface, 100 sectors per track, 5 double-sided platters and a block size of 2,048 bytes. Suppose that the average seek time is 10ms, the average rotational delay is 5 ms, and the transfer rate is 200 MB per second. Suppose that a file containing 1,000,000 records of 100 bytes each is to be stored on such a disk and that no record is allowed to span two blocks.

A. How many blocks are required to store the entire file?

Let's Practice! Disk Access Time

Consider a disk with a sector size of 512 bytes, 1,000 tracks per surface, 100 sectors per track, 5 double-sided platters and a block size of 2,048 bytes. Suppose that the average seek time is 10ms, the average rotational delay is 5 ms, and the transfer rate is 200 MB per second. Suppose that a file containing 1,000,000 records of 100 bytes each is to be stored on such a disk and that **no record is allowed to span two blocks**.

A. How many blocks are required to store the entire file?

→ How many records fit in a block?


20 records (2048B can fully fit 20 records of 100B each)

→ 1,000,000 records are stored in $1,000,000 / 20 =$ **50,000 blocks**

Let's Practice! Disk Access Time


Consider a disk with a sector size of 512 bytes, 1,000 tracks per surface, 100 sectors per track, 5 double-sided platters and a block size of 2,048 bytes. Suppose that the average seek time is 10ms, the average rotational delay is 5 ms, and the transfer rate is 200 MB per second. Suppose that a file containing 100,000 records of 100 bytes each is to be stored on such a disk and that no record is allowed to span two blocks.

- A. How many blocks are required to store the entire file? → **50,000 blocks**
- B. If the file is arranged sequentially on disk, how many cylinders are needed?



2.B. If the file is arranged **sequentially on disk**, how many cylinders are needed?


Need to place 50,000 blocks (computed in 2.A).



2.B. If the file is arranged **sequentially on disk**, how many cylinders are needed?

Need to place 50,000 blocks (computed in 2.A).

→ How do we arrange blocks for sequential access?




2.B. If the file is arranged **sequentially on disk**, how many cylinders are needed?

Need to place 50,000 blocks (computed in 2.A).

→ How do we arrange blocks for sequential access?

Next block concept:

- blocks on same track, followed by
- blocks on same cylinder, followed by
- blocks on adjacent cylinder



2.B. If the file is arranged **sequentially on disk**, how many cylinders are needed?

Need to place 50,000 blocks (computed in 2.A).

→ How do we arrange blocks for sequential access?


Next block concept:

- blocks on same track, followed by
- blocks on same cylinder, followed by
- blocks on adjacent cylinder

→ How many blocks can we place in a track?

→ How many blocks can we place in a cylinder?

→ How many cylinders do we need?



2.B. If the file is arranged **sequentially on disk**, how many cylinders are needed?

Need to place 50,000 blocks (computed in 2.A).

→ How do we arrange blocks for sequential access?

Next block concept:

- blocks on same track, followed by
- blocks on same cylinder, followed by
- blocks on adjacent cylinder

→ How many blocks can we place in a track?

Sector size: 512B, 100 sectors per track

→ **25 blocks per track**

→ How many blocks can we place in a cylinder?


$25 \times 2 \times 5 =$ **250 blocks per cylinder**


→ How many cylinders do we need?

$50,000 / 250 =$ **200 cylinders**

2.C. What is the time required to read the file sequentially?

Need to read 50,000 blocks; each block has size 2048B.






2.C. What is the time required to read the file sequentially?

Need to read 50,000 blocks; each block has size 2048B.

Seek time = 10ms

Avg rotational delay = 5ms

Transfer rate = 200 MB / second.



2.C. What is the time required to read the file sequentially?

Need to read 50,000 blocks; each block has size 2048B.

Seek time = 10ms


Avg rotational delay = 5ms

Transfer rate = 200 MB / second.

→ What are the components of disk access?

→ How long do each of them take?

→ What is the total read time?



2.C. What is the time required to read the file sequentially?

Need to read 50,000 blocks; each block has size 2048B.

Seek time = 10ms

Avg rotational delay = 5ms

Transfer rate = 200 MB / second.

→ What are the components of disk access?

Head selection (negligible)

Seek (one time cost because sequential access)

Rotational Delay (one time cost because sequential access)

Transfer time

Controller overhead (negligible)

→ What is the transfer time?

$$(5 * 5^4 * 2^4 * 2^{11} \text{ B}) / (200 * 2^{20} \text{ B/s}) = 0.488 \text{ s} \\ = 488 \text{ ms}$$

→ What is the total read time?

$$10 + 5 + 488 = \mathbf{503 \text{ ms}}$$

Optimizing disk access

Remember:

- Disk access time \gg memory access time
- If we go to disk, seek time dominates

Optimize Disk Access

Rule 1:

Do not access disk, use a cache

File System Cache (Buffer Cache)

What?

- Keep recently accessed blocks in memory

Why?

- Reduce latency
- Reduce disk load

How?

- Reserve kernel memory for cache
- Cache entries: file blocks (of block size)

Read with a Cache

If in cache

- Return data from cache

If not

- Find free cache slot
- Initiate disk read
- When disk read completes, return data

Write with a Cache

Always write in cache

How does it get to disk?

- **Write-through**
- **Write-behind**

Write with a Cache

Always write in cache

How does it get to disk?

- **Write-through**
- **Write-behind**

Write-through

1. Write to cache
2. **Write to disk**
3. Return to user

Write with a Cache

Always write in cache

How does it get to disk?

- **Write-through**
- **Write-behind**

Write-through

1. Write to cache
2. **Write to disk**
3. Return to user

Write-behind

1. Write to cache
2. Return to user
3. **Later: write to disk**

Write-Through vs. Write-Behind

Response time:

- **Write-behind** is (much) better

Disk load:

- **Write-behind** is (much) better
- Much data overwritten before it gets to disk

Crash:

- **Write-through** is much better
- No “window of vulnerability”

Write-Through vs. Write-Behind

Response time:

- **Write-behind** is (much) better

Disk load:

- **Write-behind** is (much) better
- Much data overwritten before it gets to disk

Crash:

- **Write-through** is much better
- No “window of vulnerability”

In practice:

- **Write-behind**
- Periodic cache flush
- User primitive to flush data

Optimize Disk Access

Rule 2:

Do not wait for disk, read ahead

- Also called prefetching
- Only for sequential access

Read-Ahead

What?

- User request for block i of a file
- Also read block $i+1$ from disk

Why?

- No disk I/O on (expected) user access to block $i+1$

How?

- Put block $i+1$ in the buffer cache

→ **Remember:** Pre-paging uses read-ahead if neighboring virtual pages are also neighbors in physical memory.

Read-Ahead

- **Works for sequential access**
- Most access is sequential
- In Linux it is the default

Caveat about Read-Ahead

- Does not reduce number of disk I/Os
- In fact, could increase them (if not sequential)
- In practice, very often a win
- Linux always reads one block ahead

Optimize Disk Access

Rule 3:

Minimize seeks

2 Approaches

- Clever disk allocation
- Clever scheduling

Clever Disk Allocation

Idea:

- Locate related data (same file) on same cylinder
- Allocate “related” blocks “together”

“together”

- On the same cylinder
- On a nearby cylinder

“related”

- Consecutive blocks in the same file
- Sequential access

Disk Scheduling

Idea: Reorder requests to seek as little as possible

Different disk scheduling policies:

- FCFS – First-Come-First-Served
- SSTF – Shortest-Seek-Time-First
- SCAN
- C-SCAN
- LOOK
- C-LOOK

Disk Scheduling Illustration

Initial position of the **head** = **cylinder 53**

Queue of requests:

98, 193, 37, 122, 14, 124, 65, 67

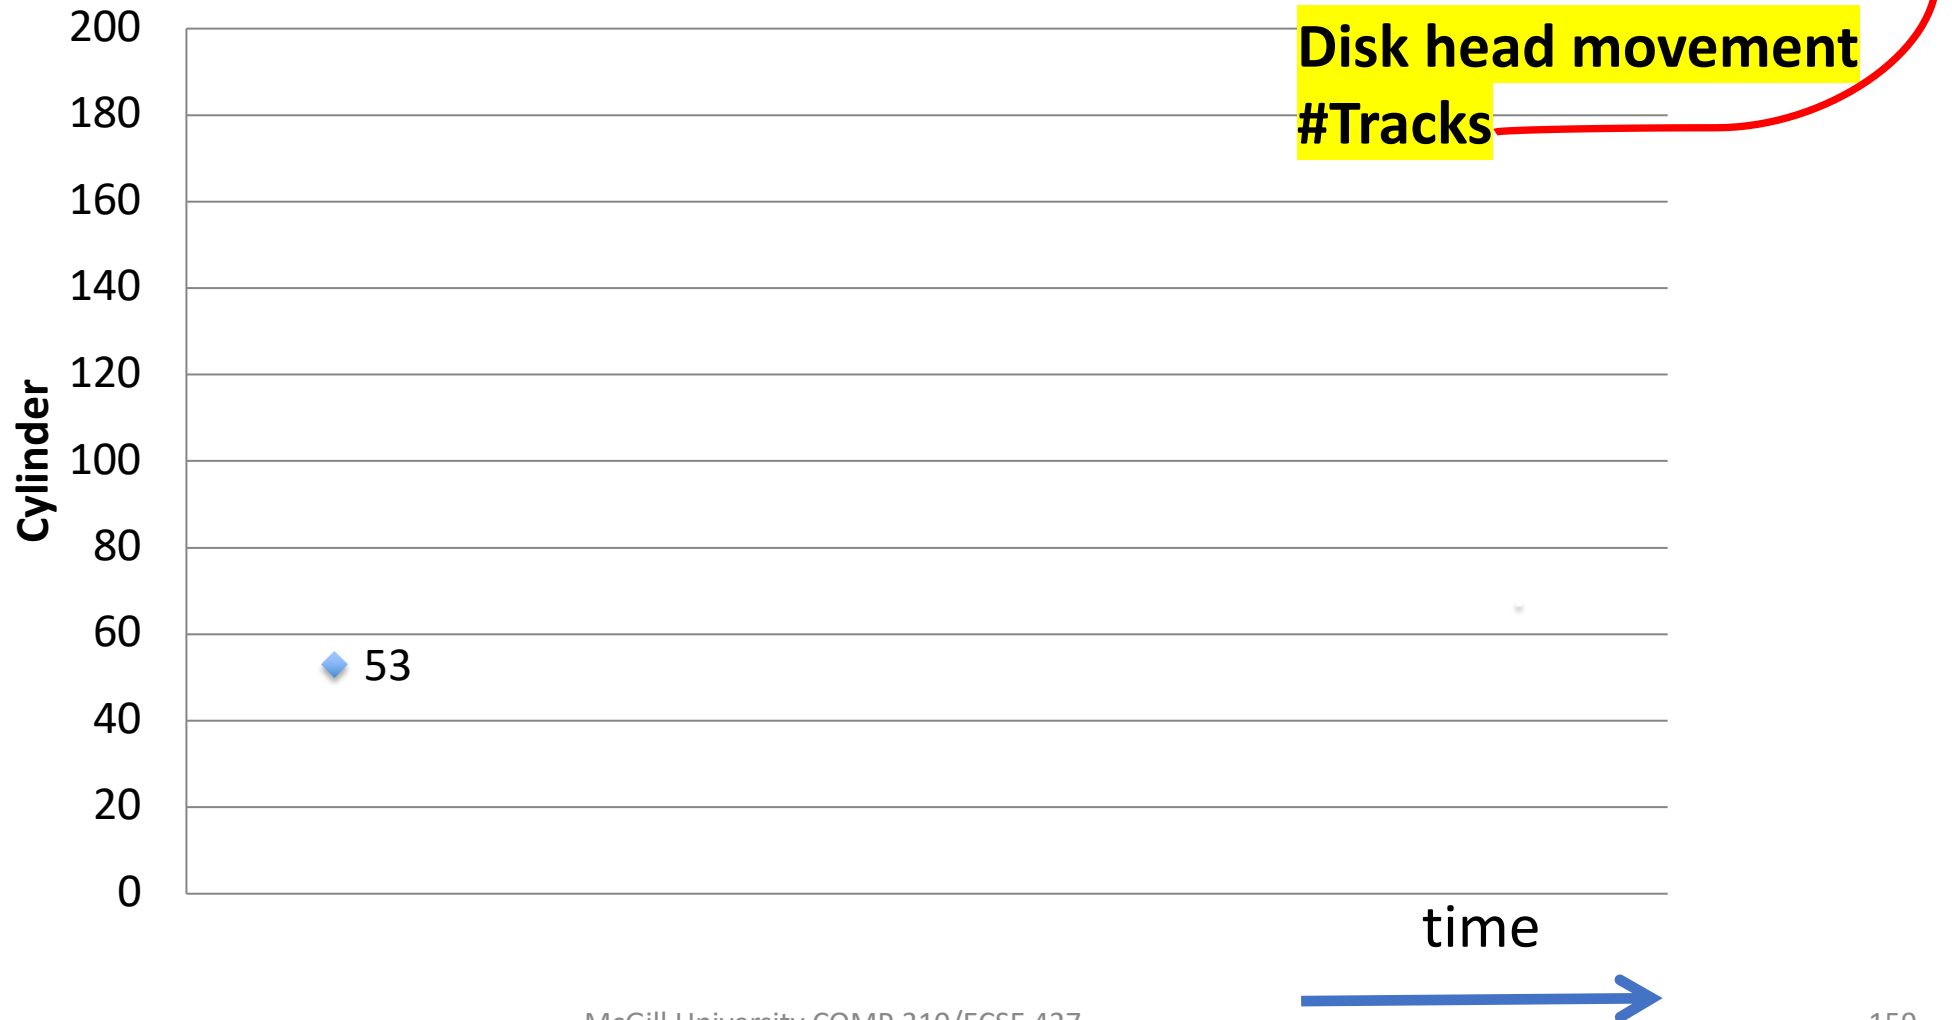
First Come, First Served (FCFS)

Serve **next request** in the queue

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

FCFS




45

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

FCFS




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

130

FCFS




276

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

FCFS




361

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

FCFS




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

469

FCFS




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

579

FCFS




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

638

FCFS




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

640

FCFS




Shortest Seek Time First (SSTF)

Pick “**nearest**” **request** in queue


- “nearest” = closest to current head position

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67



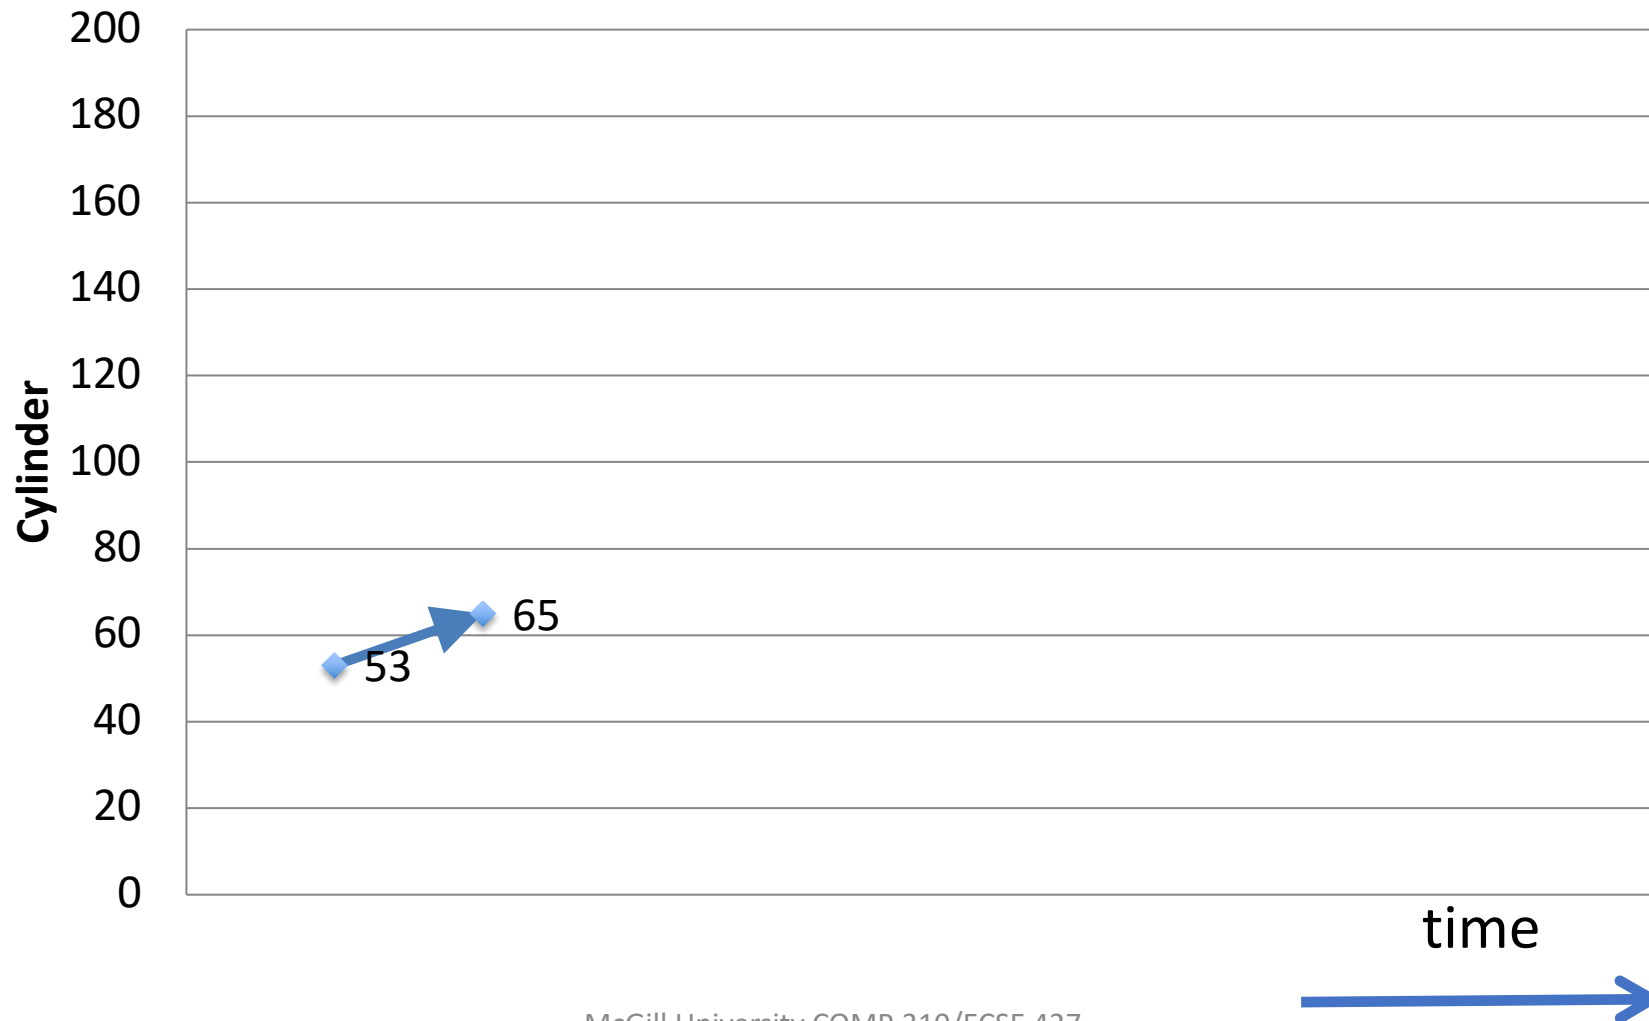
SSTF



Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SSTF




14

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67


SSTF



Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SSTF




67

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SSTF




151

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SSTF




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

175

SSTF




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

177

SSTF




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

236

SSTF




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

236

< FCFS
(640)

SSTF



SSTF

+ Very good seek times

☹ Subject to starvation


- Request on inside or outside can get starved

SCAN


- Continue moving head in one direction
 - From 0 to MAX_CYL
 - Then, from MAX_CYL to 0
- Pick up requests as you move head

Head = 53, moving down

Queue = 98, 183, 37, 122, 14, 124, 65, 67




SCAN



Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SCAN




39

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SCAN




53

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SCAN




118

Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

SCAN




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

120

SCAN




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

151

SCAN




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

175

SCAN




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

177

SCAN




Head = 53

Queue = 98, 183, 37, 122, 14, 124, 65, 67

236

SCAN



Here we do not go all the way to MAX_CYL, because this is the last request in the queue.

C-SCAN


- Similar to SCAN
- Always move head
 - From MAX_CYL to 0; pick up requests as head moves
 - From 0 to MAX_CYL; no requests served
- C-SCAN can also be implemented in reverse
 - From MAX_CYL to 0; no requests served
 - From 0 to MAX_CYL; pick up requests as head moves

388

Head = 53, moving down

Queue = 98, 183, 37, 122, 14, 124, 65, 67

C-SCAN



C-SCAN

☹ Number of cylinders slightly higher

+ More uniform wait time

C-LOOK


- Similar to C-SCAN
- Always move head
 - From MAX_CYL_IN_QUEUE to MIN_CYL; serve requests as head moves
 - From MIN_CYL to MAX_CYL_IN_QUEUE; no requests served
- C-LOOK can also be implemented in reverse

326

Head = 53, moving down

Queue = 98, 183, 37, 122, 14, 124, 65, 67

C-LOOK



In Practice

- Some variation of C-LOOK

Optimize Disk Access

Rule 4:

Avoid rotational latency

- Clever disk allocation
- Locate consecutive blocks of file on consecutive sectors in a cylinder

When does what work well?

- Low load: clever allocation
- High load: disk scheduling

Why? – Under High Load

- Many scheduling opportunities
 - Many requests in the queue
- Allocation gets defeated
 - By interleaved requests for different files

Why? – Under Low Load

- Not much scheduling opportunity
 - Not many requests in the queue
- Sequential user access -> sequential disk access
- Cache tends to reduce load

Summary – Disk Management

- Disk characteristics
 - Access disk >> access memory
 - Seek > Rotational Latency > Transfer
- Optimizations
 - Cache
 - Read-ahead
 - Disk allocation
 - Disk scheduling

Summary – Key Concepts

- I/O devices
 - OS role for integrating I/O devices in systems
 - Polling, Interrupts
- Notion of “permanent” storage
- File system interface
- Disk Management for HDDs

Further Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 36, 37, 39.

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Youjip Won (Hanyang University).