

Week 3

Synchronization Primitives

Oana Balmau
January 17, 2023

Schedule Highlights

Week 3 Process Management	jan 16 C Review: C Basics	jan 17 Synchronization Primitives (1/2) Optional reading: OSTEP Chapters 25 – 32 add/drop deadline	jan 18 OS Shell Assignment Released	jan 19 Synchronization Primitives (2/2) OS Shell Assignment Overview – with Jiaxuan	jan 20
Week 4 Process Management	jan 23 C Tools: GDB basics	jan 24 Multi-process Structuring (1/2) Team registration deadline	jan 25	jan 26 Multi-process Structuring (2/2)	jan 27
Week 5 Process Management	jan 30 C Review: Pointers & Memory Allocation I	jan 31 Multithreading (1/2) Practice Exercises Sheet: Process Management	feb 1	feb 2 Multithreading (2/2)	feb 3

One more week to create teams.

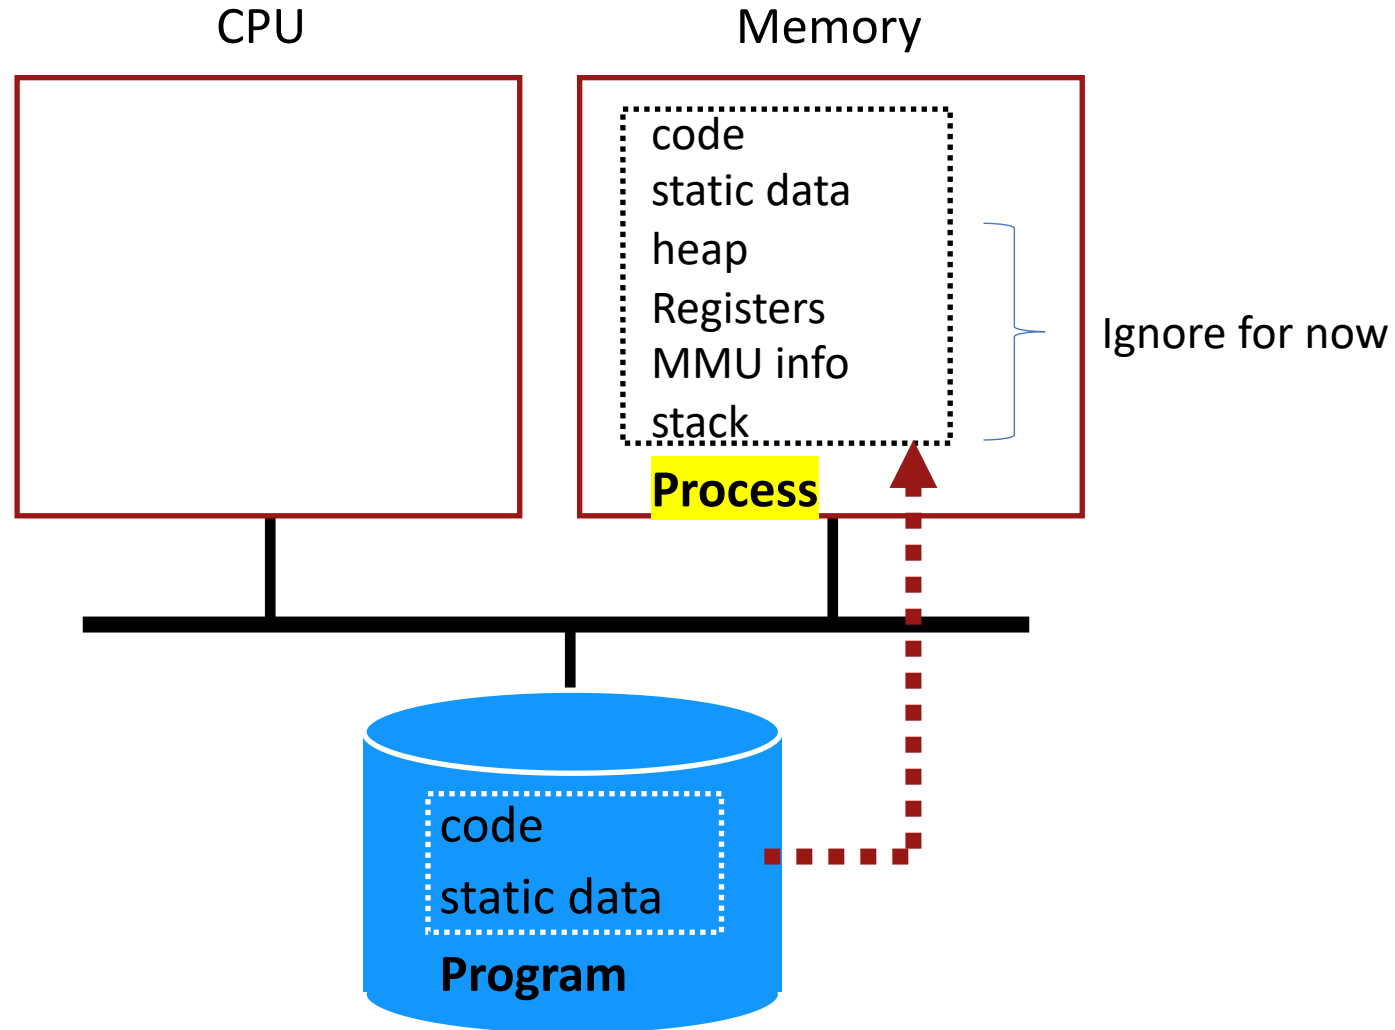
Jiaxuan will explain the assignment during part of this class.

Recap from Week 2

- Process
- Linux process tree
- Process switch
- Process scheduler

Recap from Week 2

Process = Program in execution




Recap from Week 2

Linux Process Primitives

- `pid = fork()`
- `exec(filename)`
- `exit()`
- `wait()`


Recap from Week 2

Process = Environment + Code




Recap from Week 2

After a fork()



Recap from Week 2

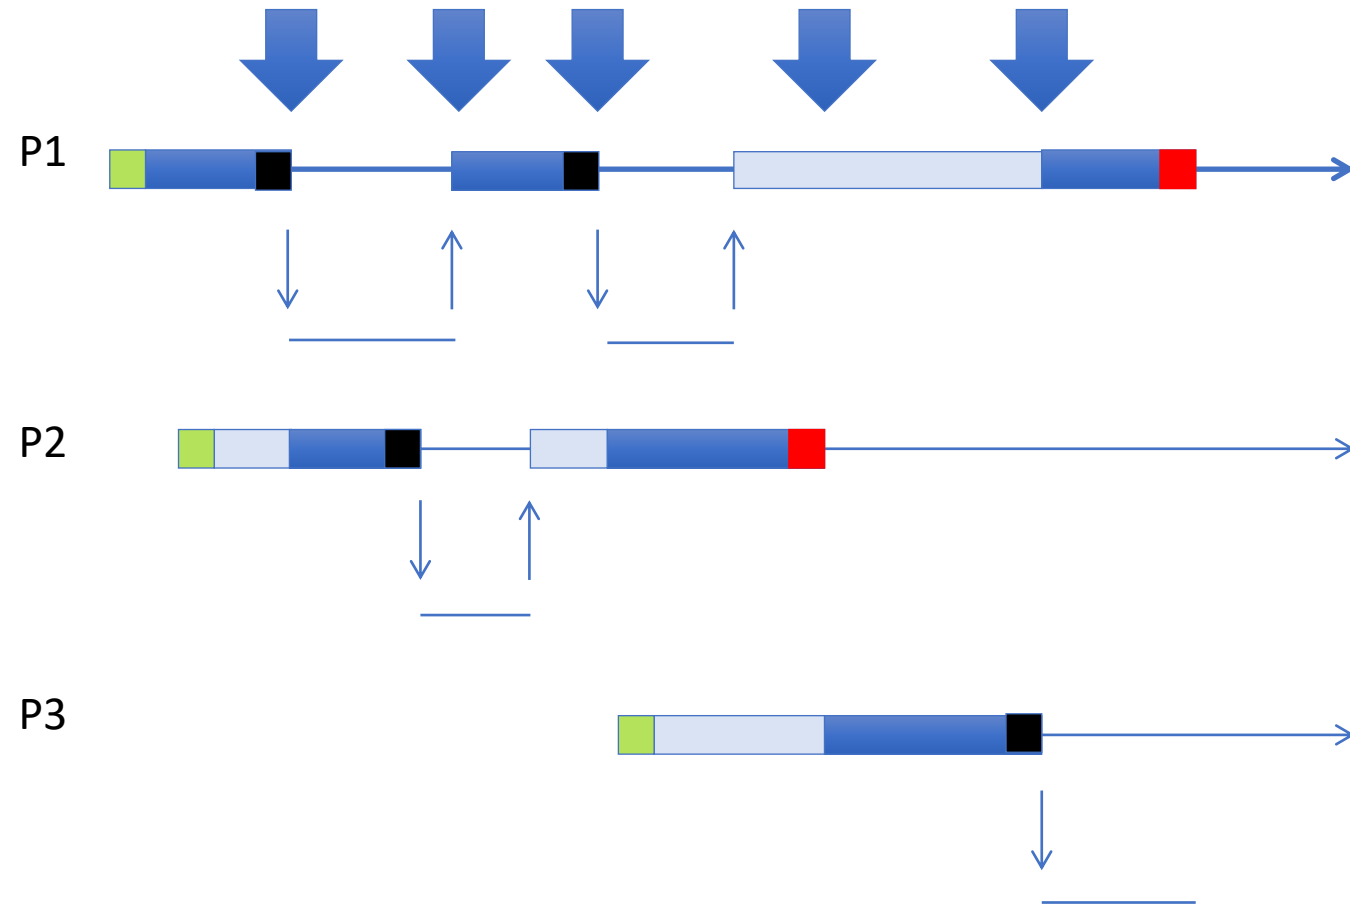
After an exec() in the Child



Recap from Week 2

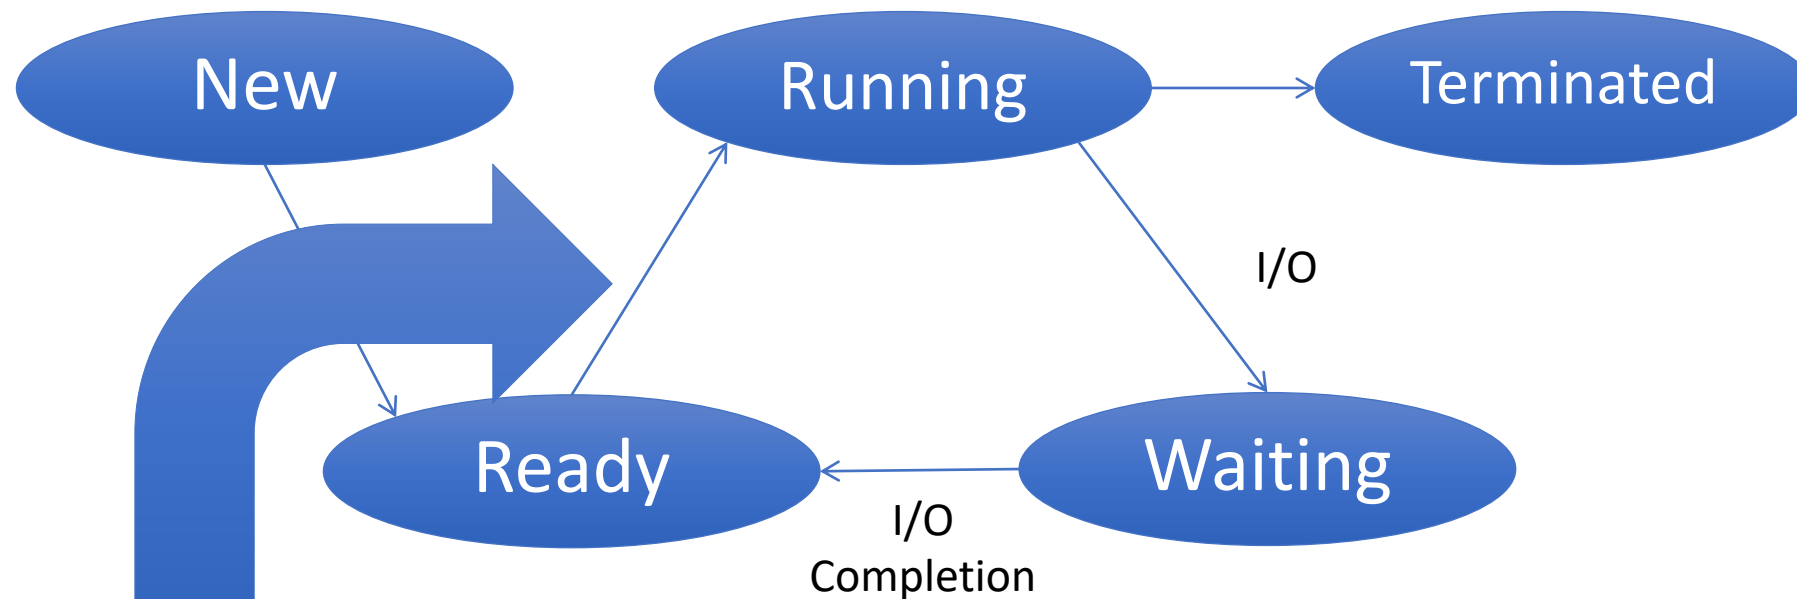
Process Switch

- Change of process using the CPU
- Save and restore registers and other info



Recap from Week 2

Process Scheduling



Many processes may be ready.
Process scheduler picks one.

Questions from last week?

Before we begin with today's topic

- Concurrency is a large sub-field of computer science
- In this course, we get a small taste of it
- If you enjoy this lecture, consider [COMP-409 Concurrent programming](#)
 - Highly recommend for a strong systems background

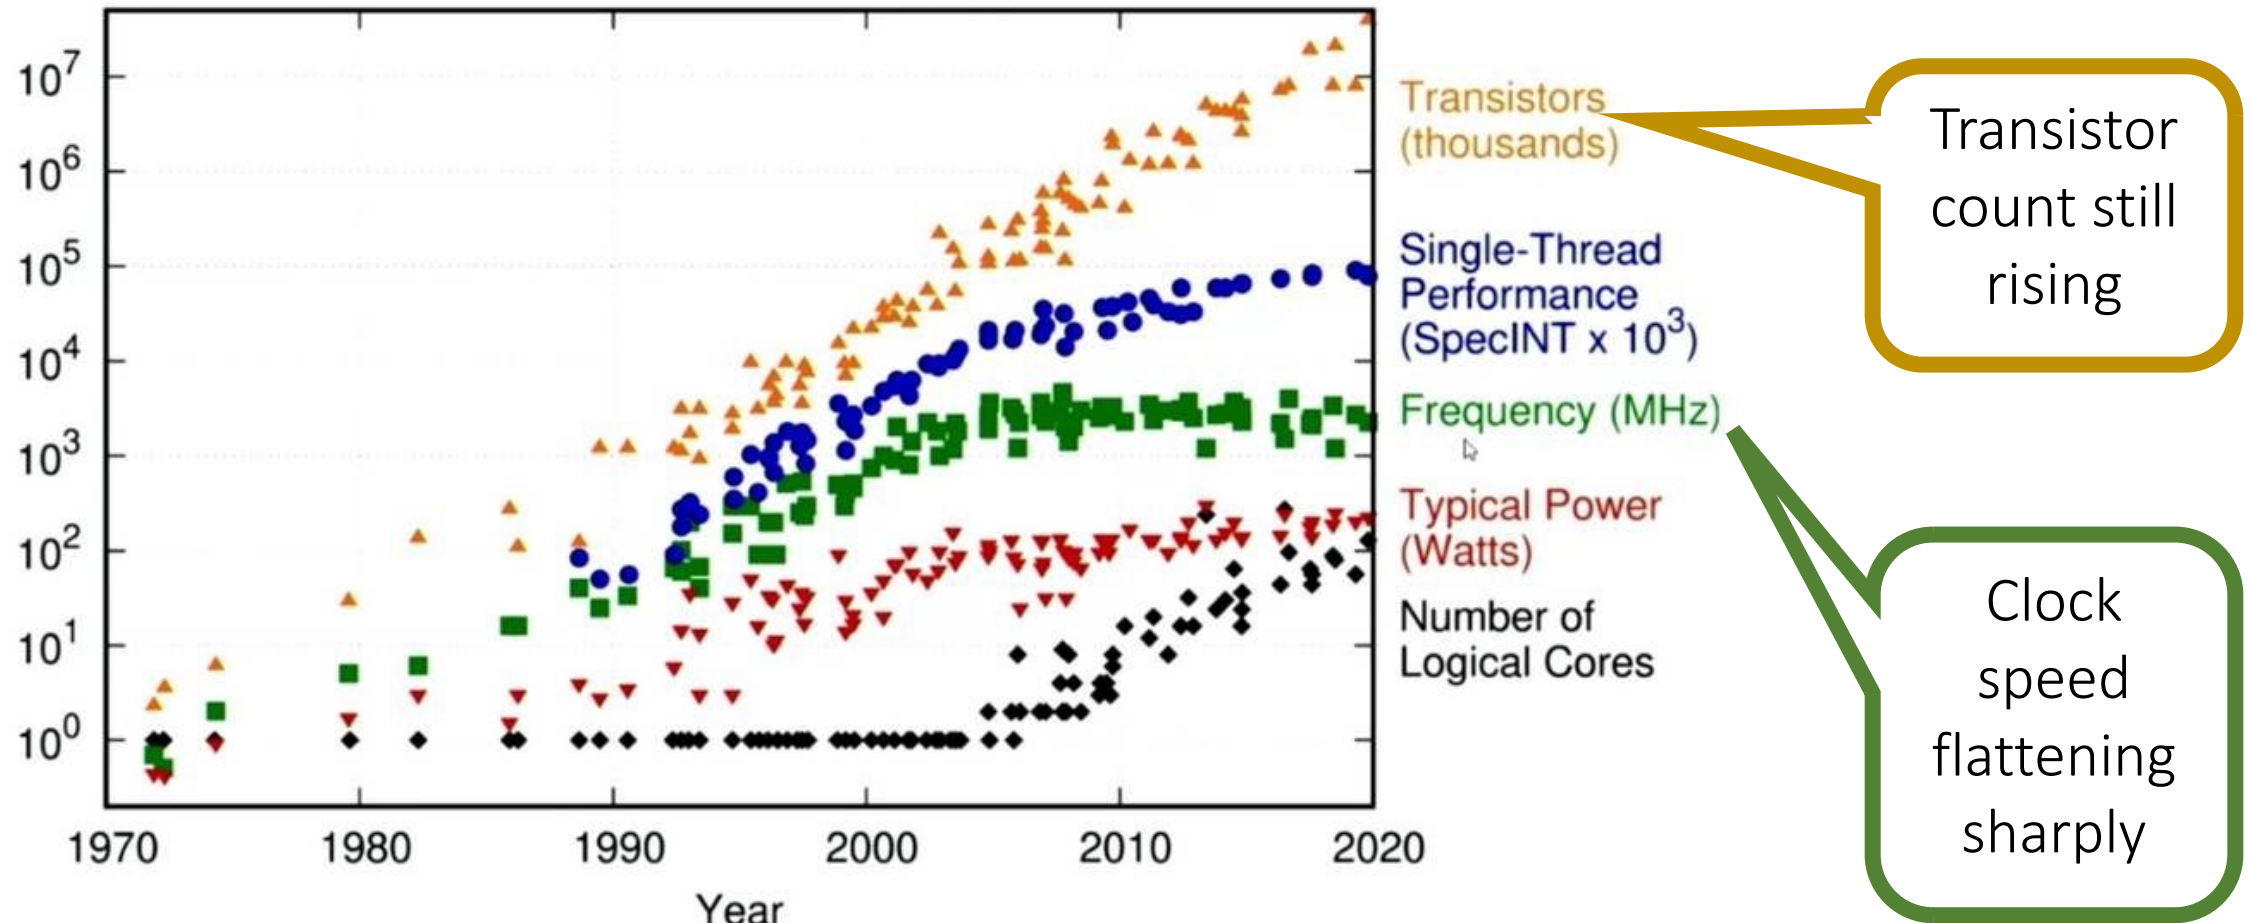
Real world concurrency

- Millions of drivers on highway at once.
- Student does homework while watching Netflix.
- Faculty has lunch while grading papers and watching Netflix.

Key Concepts for Today


- Process vs Thread
- Mutual Exclusion
- Locking
- Deadlocks
- Conditional variables

Motivation for Concurrency – Moore's Law




Motivation for Concurrency

- CPU trend: Same speed, but multiple cores
- Goal: write applications that fully utilize many cores




Concurrency Abstraction vs Reality


Programmer Abstraction



Physical Reality



Ready queue




Concurrency – Option 1

- Build apps from many communicating **processes**
- Communicate through **message passing**
 - No shared memory
- Pros
 - If one process crashes, other processes unaffected
- Cons
 - High communication overheads
 - Expensive context switching

Concurrency – Option 1

- Build apps from many communicating **processes**
- Communicate through **message passing**
 - No shared memory
- Pros
 - If one process crashes, other processes unaffected
- Cons
 - High communication overheads
 - Expensive context switching



Will see
next week

Concurrency – Option 2

- New abstraction: **thread**
- Multiple threads in a process
- Threads are like processes except
 - Multiple threads in the same process share an address space
 - **Communicate through shared address space**
 - If one thread crashes,
 - the entire process, including other threads, crashes

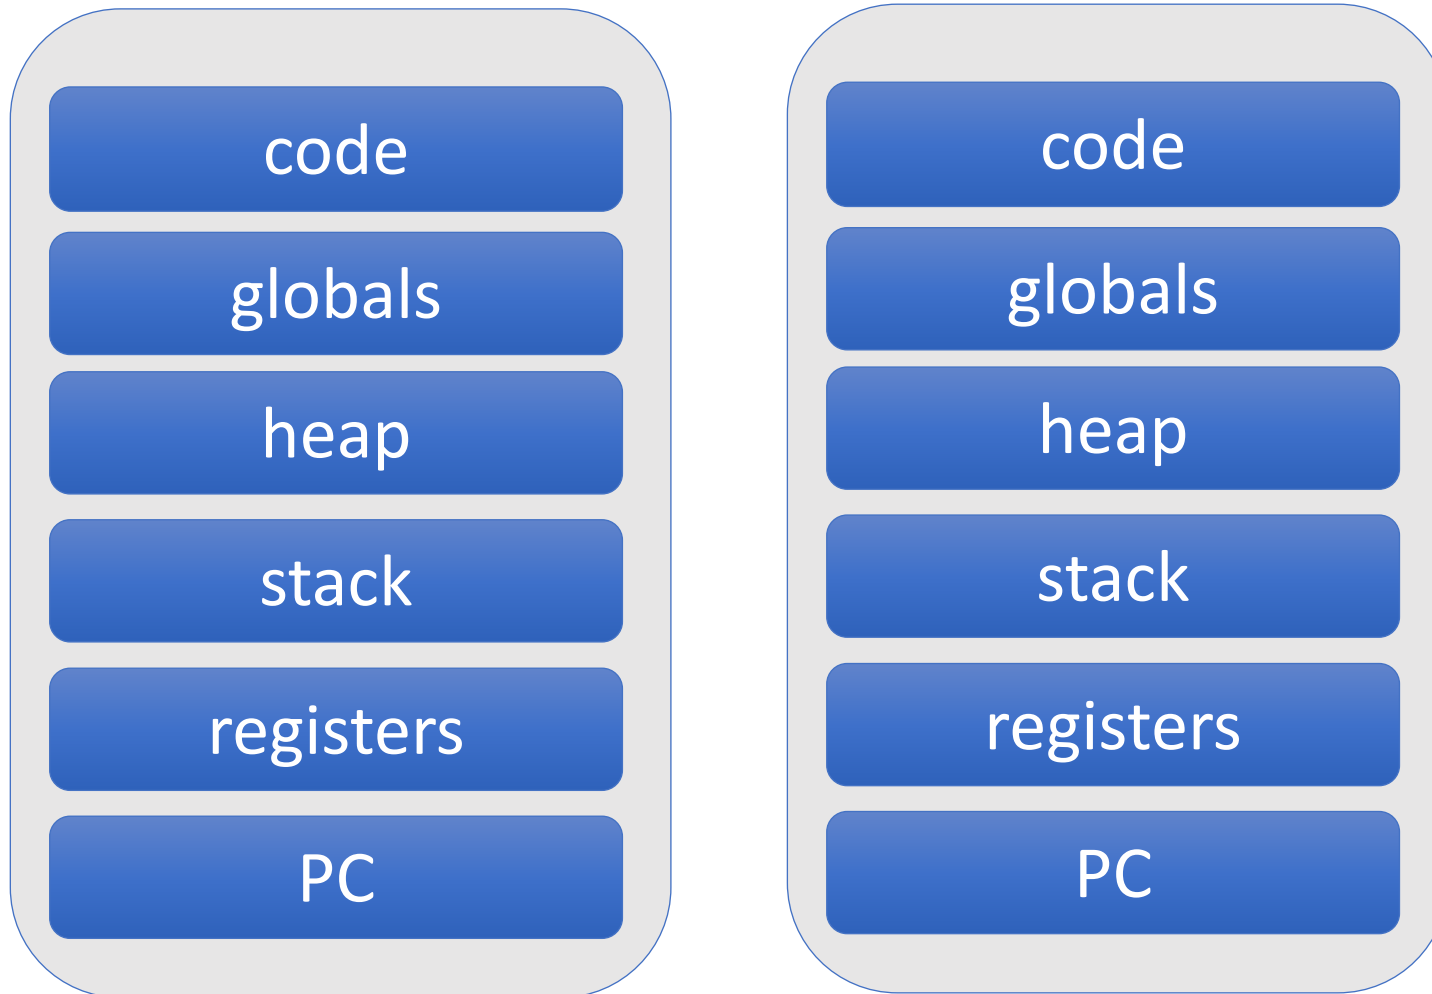
Concurrency – Option 2

- New abstraction: **thread**
- Multiple threads in a process
- Threads are like processes except
 - Multiple threads in the same process share an address space
 - **Communicate through shared address space**
 - If one thread crashes,
 - the entire process, including other threads, crashes

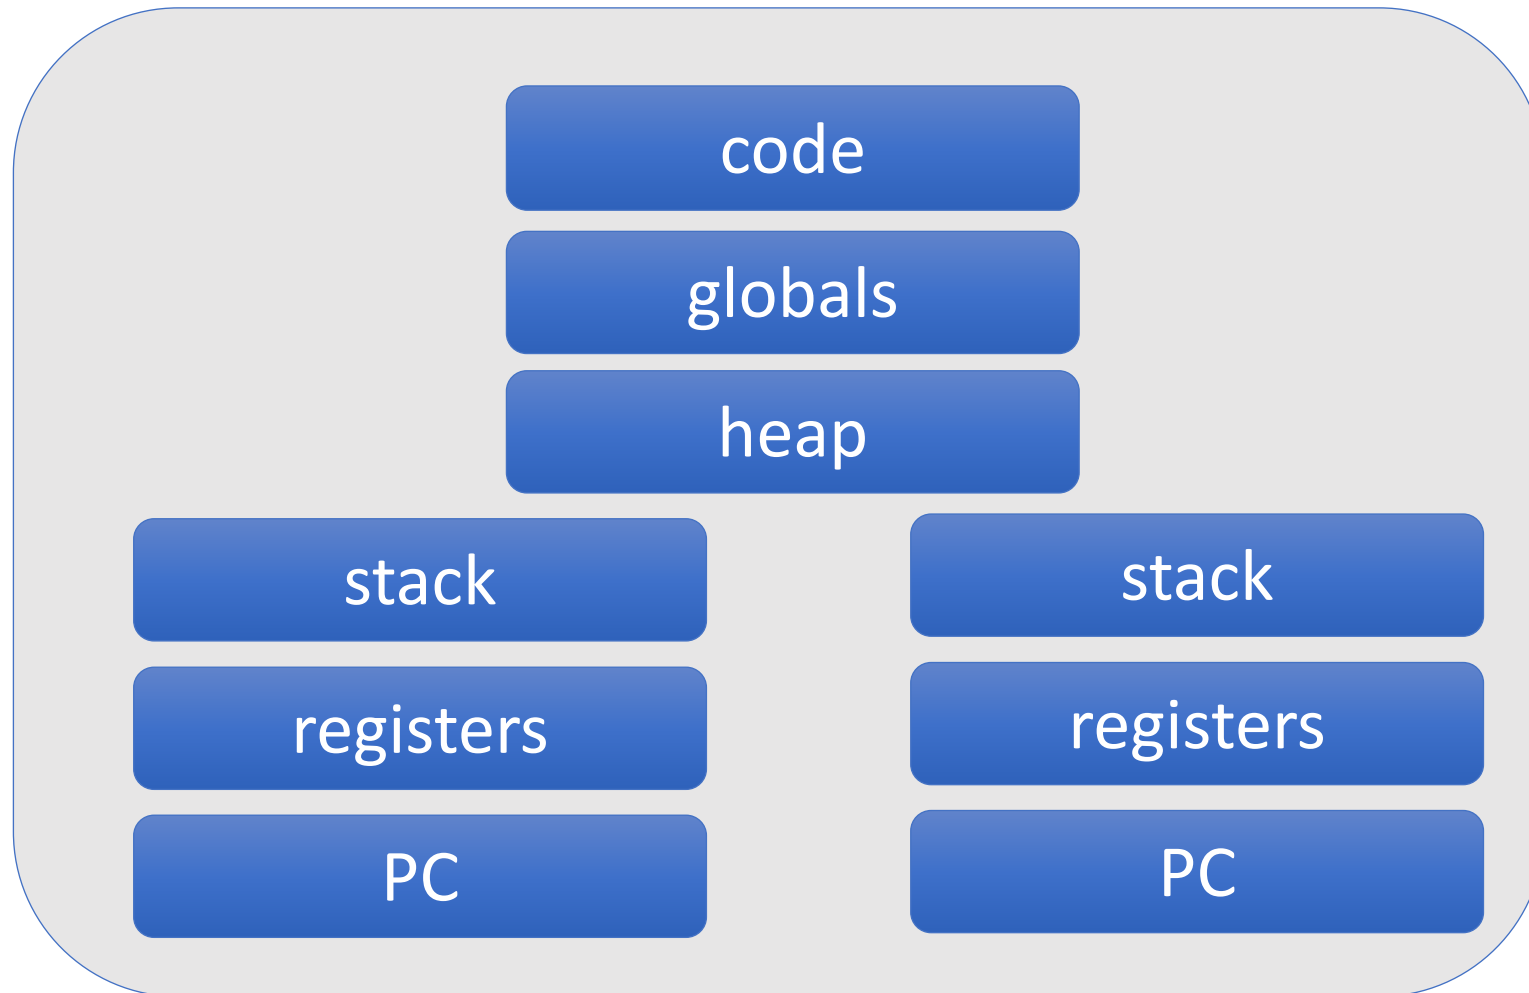
Will see synchronization principles **today**

Will see practical examples in two weeks

Two Processes




Two Threads in a Process



In General

- Processes provide separation
 - In particular, memory separation (no shared data)
 - Suitable for coarse-grain interaction
- Threads do not provide separation
 - In particular, threads share memory (shared data)
 - Suitable for tighter integration

Where Things Reside



Shared Data

- Advantage:
 - Many threads can read/write it
- Disadvantage:
 - Many threads can read/write it
 - Can lead to *data races*

Data Race

- Unexpected/unwanted access to shared data
- Result of *interleaving* of thread executions
- **Program must be correct for all interleavings**

A Common Mistake/Misunderstanding: A Single Line of Code is not Atomic

- $a = a + 1$
- Is in reality
 - Load a from memory into register
 - Increment register
 - Store register value in memory
- Instruction sequence may be interleaved
- (Some machines have atomic increments)

Thread Schedule #1

```
balance = balance + 1; // balance in shared memory at 0x9cd4
```

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: ?`

`%rip = 0x195`

process
control
blocks:

Thread 1

`%eax: ?`
`%rip: 0x195`

Thread 2

`%eax: ?`
`%rip: 0x195`



- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: 100`

`%rip = 0x19a`

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: 101`

`%rip = 0x19d`

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x1a2`

process
control
blocks:

Thread 1

`%eax: ?`
`%rip: 0x195`

Thread 2

`%eax: ?`
`%rip: 0x195`

- `0x195 mov 0x9cd4, %eax`
- `0x19a add $0x1, %eax`
- `0x19d mov %eax, 0x9cd4`

T1 

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x1a2`

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

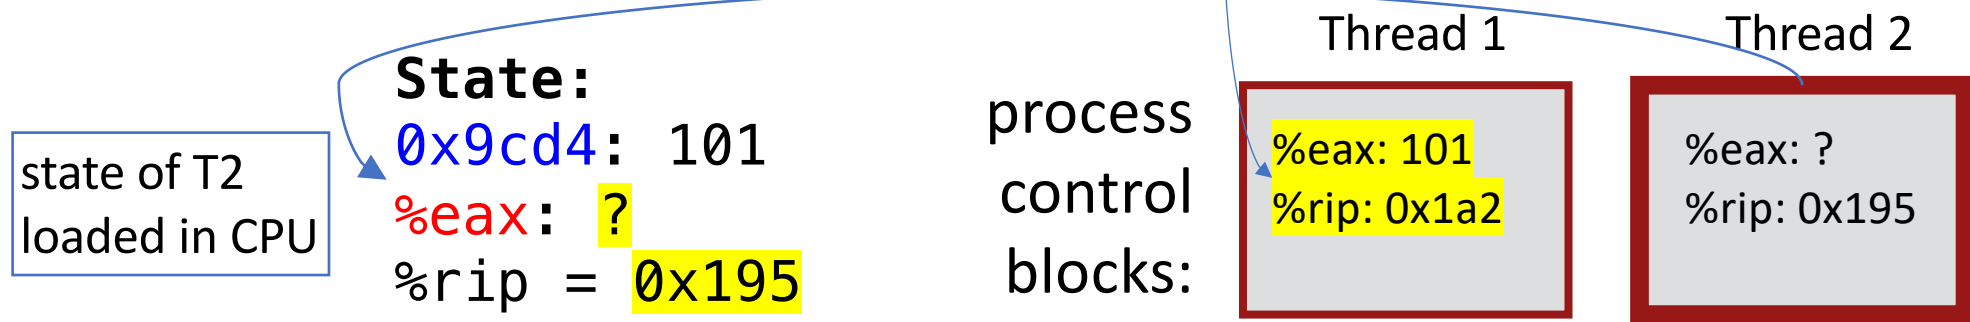
- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`



Thread context switch!

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4



T2 →

- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

note that code region is common to T1 & T2

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: ?`

`%rip = 0x195`


process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x1a2`

Thread 2

`%eax: ?`
`%rip: 0x195`



T2

- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x19a`


process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x1a2`

Thread 2

`%eax: ?`
`%rip: 0x195`



- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 102`

`%rip = 0x19d`


process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x1a2`

Thread 2

`%eax: ?`
`%rip: 0x195`



- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #1

`balance = balance + 1; // balance in shared memory at 0x9cd4`

Desired
result 😊

State:

`0x9cd4: 102`

`%eax: 102`

`%rip = 0x1a2`

process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x1a2`

Thread 2

`%eax: ?`
`%rip: 0x195`

- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

T2 →

Another schedule

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: ?`

`%rip = 0x195`

process
control
blocks:

Thread 1

`%eax: ?`
`%rip: 0x195`

Thread 2

`%eax: ?`
`%rip: 0x195`



- `0x195 mov 0x9cd4, %eax`
- `0x19a add $0x1, %eax`
- `0x19d mov %eax, 0x9cd4`

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: 100`

`%rip = 0x19a`


process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: 101`

`%rip = 0x19d`

process
control
blocks:

Thread 1

`%eax: ?`
`%rip: 0x195`

Thread 2

`%eax: ?`
`%rip: 0x195`



- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

Thread context switch!

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: ?`

`%rip = 0x195`


process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: ?`
`%rip: 0x195`



T2

- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: 100`

`%rip = 0x19a`


process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: ?`
`%rip: 0x195`



- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 100`

`%eax: 101`

`%rip = 0x19d`


process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: ?`
`%rip: 0x195`



- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x1a2`

process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: ?`
`%rip: 0x195`

- `0x195 mov 0x9cd4, %eax`
- `0x19a add $0x1, %eax`
- `0x19d mov %eax, 0x9cd4`

T2



Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x1a2`

process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: ?`
`%rip: 0x195`

- `0x195 mov 0x9cd4, %eax`
- `0x19a add $0x1, %eax`
- `0x19d mov %eax, 0x9cd4`

T2



Thread context switch!

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x19d`

process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: 101`
`%rip: 0x1a2`



- `0x195` `mov 0x9cd4, %eax`
- `0x19a` `add $0x1, %eax`
- `0x19d` `mov %eax, 0x9cd4`

Thread Schedule #2

`balance = balance + 1; // balance in shared memory at 0x9cd4`

WRONG Result! ☹️
**Final value of
balance is 101**

State:

`0x9cd4: 101`

`%eax: 101`

`%rip = 0x1a2`

process
control
blocks:

Thread 1

`%eax: 101`
`%rip: 0x19d`

Thread 2

`%eax: 101`
`%rip: 0x1a2`

- 0x195 `mov 0x9cd4, %eax`
- 0x19a `add $0x1, %eax`
- 0x19d `mov %eax, 0x9cd4`

T1



Non-Determinism

Concurrency leads to non-deterministic results

- Different results even with same inputs
- Race conditions

Whether bug manifests or not depends on CPU schedule!

How to program?

- Imagine scheduler is malicious
- Assume scheduler will pick bad ordering at some point...

Basic Approach to Multithreading

1. Divide “work” among multiple threads &
2. Share data
 - Which data is shared?
 - **Global variables and heap**
 - Not local variables, not read-only variables
 - Where is shared data accessed?
 - Put shared data access in **critical section**

Critical Section

- Want 3 instructions to execute as an uninterruptable group
- That is, we want them to be **atomic**

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

critical section

Need **mutual exclusion** for critical sections

- If thread A is in critical section C, thread B can't enter C
- Ok if other processes do unrelated work

Mutual Exclusion

- Prevents simultaneous access to a shared resource.
 - In this case, shared resource = shared memory region
- How can we achieve mutual exclusion?
 - Today we will first see library support (pthreads)
 - Then, we will see implementation of synchronization primitives.

Why this (mostly) works

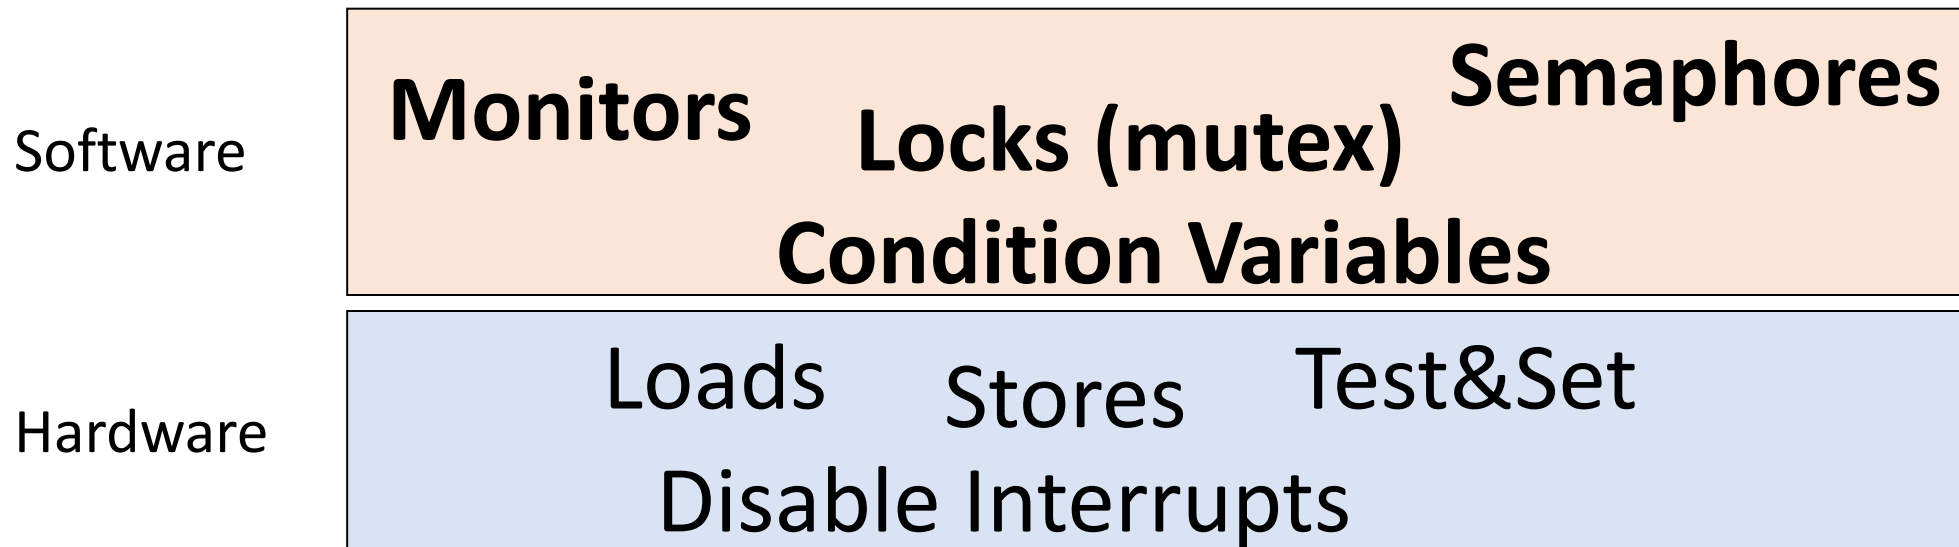
- Critical section:
 - No other thread can change data
- So you are (mostly) ok

Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



POSIX Thread Libraries (pthreads)

- Thread API for C/C++
- User-level library:
 - **#include <pthread.h>**
 - Compile and link with *-pthread*.
- Support for thread creation, termination, synchronization.
- See more details here: <https://man7.org/linux/man-pages/man3/>

Pthreads: Thread Creation and Destruction

- `pthread_create()`
- `pthread_exit()`
- `pthread_join()`

pthread_create()

```
int pthread_create(pthread_t * thread, pthread_attr_t * attr, void  
*(*start_routine)(void *), void * arg);
```

- Create thread, in *thread*.
- Run *start_routine* with arguments *arg*.
- *attr* points to a *pthread_attr_t* structure. If *attr* is NULL, then the thread is created with default attributes (ok in most cases).
- On success, return 0; on error, return an error number.

pthread_exit()

```
void pthread_exit(void *retval);
```


- Terminate calling thread.
- Returns a value via *retval*.

pthread_join()

```
int pthread_join(pthread_t thread, void **retval);
```

- Join with a terminated thread.
- Waits for the thread specified by *thread* to exit.
- If *retval* is not NULL, then **pthread_join()** copies the exit status of the target thread into the location pointed to by *retval*.
- On success, return 0; on error, return an error number.

Fork-Join Pattern for threads



Main thread creates (forks) collection of sub-threads passing them args to work on...
... and then joins with them, collecting the results

Note: In this example, start_routine is the same for all threads; only args are different.

Fork-join example

```
void *mythread(void *arg) {
    printf("%s\n", (char *) arg);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p1, p2;
    printf("main: begin\n");
    pthread_create(&p1, NULL, mythread, "A");
    pthread_create(&p2, NULL, mythread, "B");
    // join waits for the threads to finish
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("main: end\n");
}
```

Counting example – What is the final answer?

```
int count;
void *mythread(void *arg) {
    int j;
    for (j = 0; j < 1000000; j++){
        count +=1;
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p1, p2;
    count = 0;
    pthread_create(&p1, NULL, mythread, NULL);
    pthread_create(&p2, NULL, mythread, NULL);
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("%d \n", count);
}
```


Pthreads: Locks

- `Pthread_mutex_lock(mutex)`
- `Pthread_mutex_unlock(mutex)`

Pthread_mutex_lock(mutex)

- If lock is held by another thread, block
- If lock is not held by another thread
 - Acquire lock
 - Proceed

Pthread_mutex_unlock(mutex)

- Release lock

Counting example revisited – What is the final answer?

```
pthread_mutex_t count_mutex;  
int count;  
  
void *mythread(void *arg) {  
    int j;  
    for (j = 0; j < 1000000; j++){  
        pthread_mutex_lock(&count_mutex);  
        count +=1;  
        pthread_mutex_unlock(&count_mutex);  
    }  
    return NULL;  
}  
  
int main(int argc, char *argv[]) {  
    pthread_t p1, p2;  
    pthread_mutex_init(&count_mutex, NULL);  
  
    count = 0;  
    pthread_create(&p1, NULL, mythread, NULL);  
    pthread_create(&p2, NULL, mythread, NULL);  
    pthread_join(p1, NULL);  
    pthread_join(p2, NULL);  
    printf("%d \n", count);  
}
```

Deadlocks

- Threads are stuck waiting for blocked resources and no amount of retry (backoff) will help.
- Classic example:

Thread A

```
1 lock(object1)
2 lock(object2)
3 //do stuff
4 unlock(object2)
5 unlock(object1)
...
```

Thread B

```
1 lock(object2)
2 lock(object1)
3 //do stuff
4 unlock(object1)
5 unlock(object2)
...
```



Deadlocks


- Threads are stuck waiting for blocked resources and no amount of retry (backoff) will help.
- Classic example:

Thread A

```
1 lock(object1)
2 lock(object2)
3 //do stuff
4 unlock(object2)
5 unlock(object1)
...
```

Thread B

```
1 lock(object2)
2 lock(object1)
3 //do stuff
4 unlock(object1)
5 unlock(object2)
...
```



Deadlocks


- Threads are stuck waiting for blocked resources and no amount of retry (backoff) will help.
- Classic example:

Thread A

```
1 lock(object1)
2 lock(object2)
3 //do stuff
4 unlock(object2)
5 unlock(object1)
...
```

Thread B

```
1 lock(object2)
2 lock(object1)
3 //do stuff
4 unlock(object1)
5 unlock(object2)
...
```



Deadlock example

```
pthread_mutex_t lock1;
pthread_mutex_t lock2;

void *a_func(void *arg) {
    long j;
    for (j = 0; j < 1000000000; j++) {
        pthread_mutex_lock(&lock1);
        pthread_mutex_lock(&lock2);
        printf("A");
        pthread_mutex_unlock(&lock2);
        pthread_mutex_unlock(&lock1);
    }
    return NULL;
}
```

```
void *b_func(void *arg) {
    long j;
    for (j = 0; j < 1000000000; j++) {
        pthread_mutex_lock(&lock2);
        pthread_mutex_lock(&lock1);
        printf("B");
        pthread_mutex_unlock(&lock1);
        pthread_mutex_unlock(&lock2);
    }
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t a, b;
    pthread_mutex_init(&lock1, NULL);
    pthread_mutex_init(&lock2, NULL);
    pthread_create(&a, NULL, a_func, NULL);
    pthread_create(&b, NULL, b_func, NULL);
    pthread_join(a, NULL);
    pthread_join(b, NULL);
    printf("End!\n");
}
```


Condition Variables

- Used when thread A needs to wait for an event done by thread B
- A waits until a certain condition is true
 - First test condition,
 - If condition not true, call `pthread_cond_wait()`
 - A blocks until condition is true.
- At some point B makes the condition true
 - Then B calls `pthread_cond_signal()`, which unblocks A.

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

Thread B

```
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv);
```

Find the data race.

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal (&cv);
```

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f ( a , b ) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal (&cv);
```

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal(&cv);
```

:(Broadcast missed by A

Condition variables use (incorrect)

Thread A

```
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

A waits forever...

Thread B

```
//change a and b;  
x = f ( a , b ) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv);
```

Condition variables use (still incorrect)

Thread A

```
pthread_mutex_lock(&mutex );  
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv, &mutex);  
pthread_mutex_unlock(&mutex) ;
```

Thread B

```
pthread_mutex_lock(&mutex );  
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv , &mutex);  
pthread_mutex_unlock(&mutex);
```

Every time you use a condition variable you must also use a mutex to prevent the race condition.

One more issue...

Sometimes, the wait function might return even though the condition variable has not actually been signaled.

Thread A

```
pthread_mutex_lock(&mutex );  
x = f ( a , b ) ;  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv, &mutex);  
pthread_mutex_unlock(&mutex) ;
```

Thread B

```
pthread_mutex_lock(&mutex );  
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv , &mutex);  
pthread_mutex_unlock(&mutex);
```

Example:

If process P running A and B receives an OS signal

- Any thread in P can be chosen to process the signal.
- → A might be chosen to process the signal handling function
- → wait returns with an error code → A runs even if condition is not true...

How can we fix this?

Condition variables use (correct)

- Retest the condition after `pthread_cond_wait()` returns.
 - This is most easily done using a loop.

Thread A

```
pthread_mutex lock(&mutex ) ;
while (1){
    x = f ( a , b ) ;
    if ( x < 0 || x > 9)
        pthread_cond_wait (&cv, &mutex);
    else break;
}
pthread_mutex_unlock(&mutex) ;
```

Thread B

```
pthread_mutex lock(&mutex ) ;
//change a and b;
x = f (a , b) ;
if ( x >= 0 && x <= 9)
    pthread_cond_signal (&cv, &mutex);
pthread_mutex_unlock(&mutex);
```

Conditional Variables Interface

- `pthread_cond_init(pthread_cond_t *cv, pthread_condattr_t *cattr)`
 - Initialize the conditional variable, cattr can be NULL
- `pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex)`
 - Block thread until condition is true, and atomically unblock mutex.
- `pthread_cond_signal(pthread_cond_t *cv)`
 - Unblock one thread at random that is blocked by the condition variable
- `pthread_cond_broadcast(pthread_cond_t *cv)`
 - Unblock all threads that are blocked on the condition variable pointed to by cv.

Condition Variable Example

```
pthread_cond_t is_zero;
pthread_mutex_t mutex;
int shared_data = 100;

void *thread_func(void *arg){
    while(shared_data > 0) {
        pthread_mutex_lock(&mutex);
        shared_data--;
        printf("%d ", shared_data);
        pthread_mutex_unlock(&mutex);
    }

    printf("Signaling main\n");
    pthread_cond_signal(&is_zero);
    return NULL;
}
```

```
int main (void){

    pthread_t tid;
    void * exit_status;
    int i;

    pthread_cond_init(&is_zero, NULL);
    pthread_mutex_init(&mutex, NULL);

    pthread_create(&tid, NULL, thread_func, NULL);

    pthread_mutex_lock(&mutex);
    printf("Start waiting in main\n");
    while(shared_data!=0)
        pthread_cond_wait(&is_zero, &mutex);
    pthread_mutex_unlock(&mutex);

    printf("Done waiting in main!\n");

    pthread_join(tid, &exit_status);
    return 0;
}
```

Further Optional Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 25 – 31 (inclusive)

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

Reading on concurrency: Herlihy & Shavit: The Art of Multiprocessor Programming, 2nd edition.

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Maurice Herlihy (Brown University), Prof. Natacha Crooks (UC Berkeley).

