



# **Deadlocks**

# What is a *deadlock*?

- **Deadlock** -- *permanent* blocking of a set of processes that compete for system resources
- **No efficient solution** to the deadlock problem in the general case.
- In deadlock
  - ◆ a set of **processes are in a wait** state, because each process is waiting for a resource that is held by some other waiting process
- **All deadlocks involve conflicting resource needs by two or more processes**

# Classification of resources—I

- Resources grouped into two classes:
  - ◆ **Reusable**: something that can be safely used by one process at a time and is ***not depleted*** by that use. Processes obtain resources that they later release for reuse by others.  
E.g., CPU, memory, specific I/O devices, or files.
  - ◆ **Consumable**: these can be ***created and destroyed***. When a resource is acquired by a process, the resource ceases to exist.  
E.g., interrupts, signals, or messages.

# Classification of resources—II

- Another classification of resources:
  - ◆ **Preemptable**: these can be taken away from the process owning it with no ill effects (needs save/restore).  
E.g., memory or CPU.
  - ◆ **Non-preemptable**: cannot be taken away from its current owner without causing the computation to fail.  
E.g., printer or floppy disk.
- Deadlocks mostly occur when sharing **reusable** and **non-preemptable** resources.

# Conditions for deadlock

- Four conditions that must hold for a deadlock to be possible:
  - ◆ **Mutual exclusion:** processes require exclusive control of its resources (not sharing)
  - ◆ **Hold and wait:** process may wait for a resource while holding others
  - ◆ **No preemption:** process will not give up a resource until it is finished with it
  - ◆ **Circular wait:** each process in the chain holds a resource requested by another

# Conditions for deadlock

- ***Also, a process cannot be reset to an earlier state where resources not held***
- Conditions 1—3 are necessary, but not sufficient. All 4 are needed

# Solving deadlocks

- If a necessary condition is prevented a deadlock cannot occur. For example:
  - ◆ Systems with only shared resources cannot deadlock.
    - Negates *mutual exclusion*.
  - ◆ Systems that abort processes which request a resource that is in use.
    - Negates *hold and wait*.

# Solving deadlocks (contd.)

- If a necessary condition is prevented a deadlock cannot occur. For example:
  - ◆ Systems that detect or avoid deadlocks.
    - Prevents *cycle*.
  - ◆ Transaction processing systems provide checkpoints so that processes may back out of a transaction.
    - Negates *irreversible process*.



# Resource allocation graphs

Set of Processes



$$P = \{P_1, P_2, \dots, P_n\}$$

Set of Resources


$$R = \{R_1, R_2, \dots, R_m\}$$

*Some resources come in multiple units.*

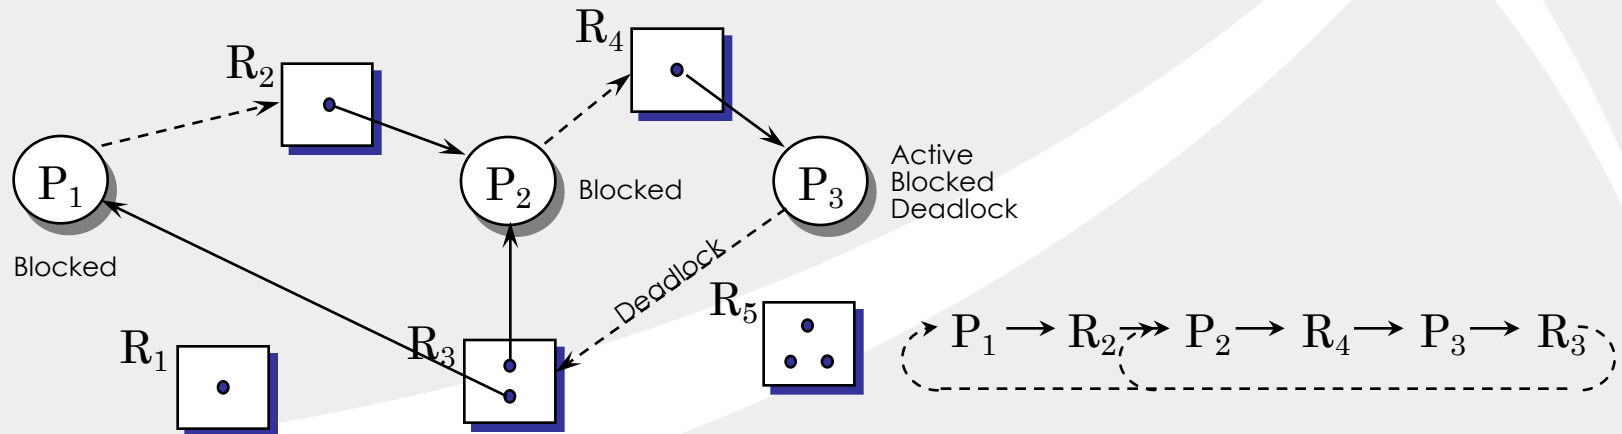
$R_j$  has 2 units




Process  $P_i$  waits for (has requested)  $R_j$



Resource  $R_j$  has been allocated to  $P_i$



# Cycle is necessary, but ...



Multiple resource unit case:

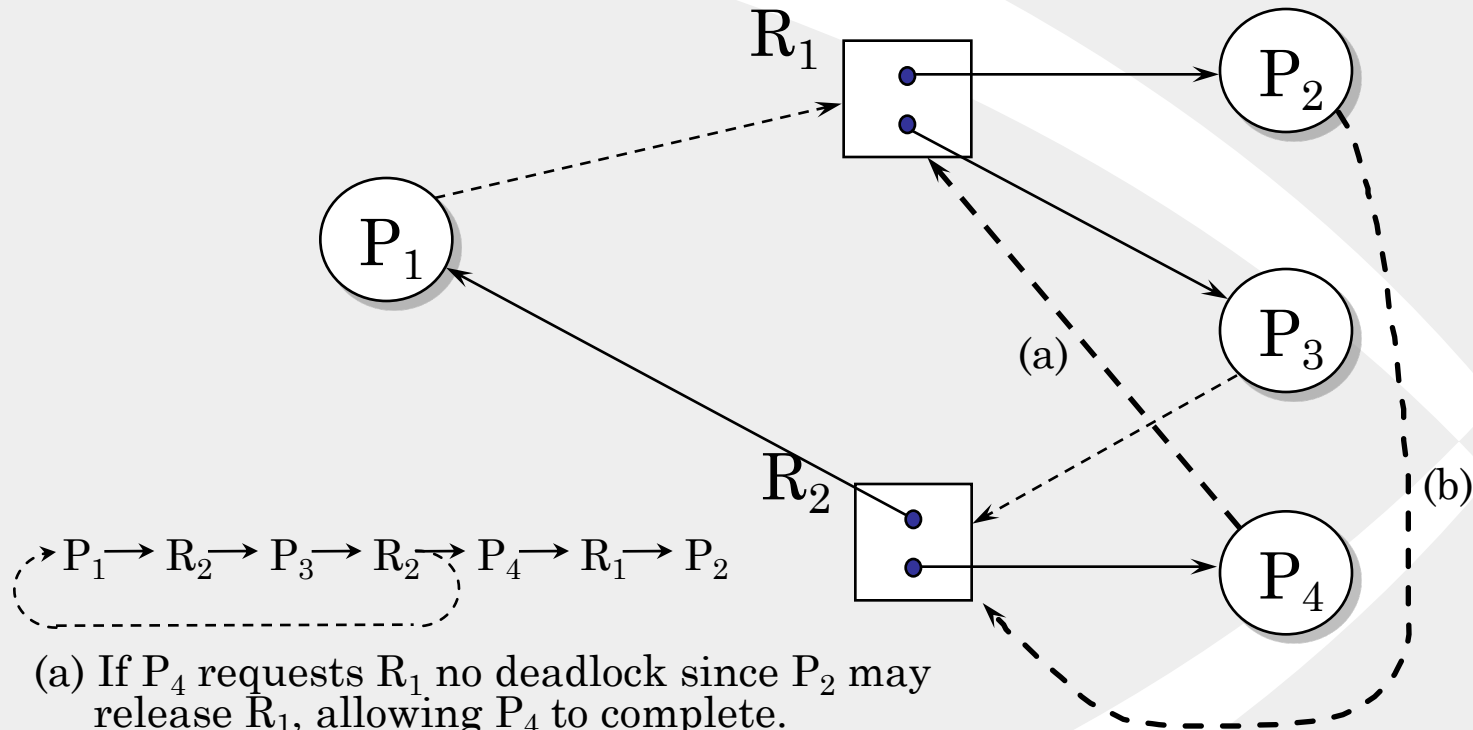
*No Deadlock—yet!*

Because, either  $P_2$  or  $P_4$  could relinquish a resource allowing  $P_1$  or  $P_3$  (which are currently blocked) to continue.  $P_2$  is still executing, even if  $P_4$  requests  $R_1$ .

## ... a knot is required

- Cycle is a *necessary condition* for a deadlock. With multiple unit resources—*not sufficient*.
- A *knot* must exist—a cycle with no non-cycle outgoing path from any involved node.
- At the moment assume that:
  - ◆ a process *halts* as soon as it waits for one resource, and
  - ◆ processes can wait for only *one* resource at a time

# Further requests



# Strategies for deadlocks

- In general, four strategies are used for dealing with deadlocks:
  - ◆ **Ignorance:** pretend there is no problem at all.
  - ◆ **Prevention:** design a system in such a way deadlock is excluded a priori.
  - ◆ **Avoidance:** make a decision dynamically checking whether a request will, if granted, potentially lead to a deadlock or not.
  - ◆ **Detection:** let the deadlock occur and detect when it happens, and take some action to recover after the fact.

# Dealing with Deadlocks

- Different people react to this strategy in different ways:
  - ◆ **Mathematicians:** find deadlock totally unacceptable, and say that it must be prevented at all costs.
  - ◆ **Engineers:** ask how serious it is, and do not want to pay a penalty in performance and convenience.
- UNIX approach
  - ◆ ignore the problem
  - ◆ Prevention price is high – inconvenient restrictions

# Deadlock prevention

- Deadlock prevention is to design a system to exclude the possibility of a deadlock
  - ◆ *indirect methods* -- prevent the occurrence of one of the necessary conditions listed earlier.
  - ◆ *direct methods* -- prevent the occurrence of a circular wait condition.
- Deadlock prevention strategies are very **conservative** -- solve the problem of deadlock by limiting access to resources and by imposing restrictions on processes

# More on deadlock prevention

## ◆ Mutual exclusion

- In general, this condition cannot be disallowed.

## ◆ Hold-and-wait

- The hold and-wait condition can be prevented by requiring that a process request all its required resources at one time, and blocking the process until all requests can be granted simultaneously.

## ◆ No preemption

- One solution is that if a process holding certain resources is denied a further request, that process must release its unused resources and request them again, together with the additional resource.

## ◆ Circular Wait

- The circular wait condition can be prevented by defining a linear ordering of resource types. If a process has been allocated resources of type R, then it may subsequently request only those resources of types following R in the ordering.



# Deadlock avoidance

- Allows the necessary conditions
- Makes judicious resource allocation choices to ensure a deadlock-free system
  - ◆ System evaluates current request and denies it if granting it will lead to potential deadlock
  - ◆ Requires knowledge of future requests
- Deadlock avoidance approaches:
  - ◆ Resource trajectories
  - ◆ Safe/unsafe states
  - ◆ Dijkstra's Banker's algorithm

# Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare the ***maximum number*** of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes


# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$
- That is:
  - ◆ If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - ◆ When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - ◆ When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.

# Safe, Unsafe, Deadlock State




# Avoidance Algorithms

- Single instance of a resource type
  - ◆ Use a resource-allocation graph
- Multiple instances of a resource type
  - ◆ Use the banker's algorithm

# Resource-Allocation Graph Scheme


- **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_i$  may request/wait resource  $R_j$ ; represented by a dashed line
- Claim edge converts to request/wait edge when a process requests a resource
- Request edge converted to an assignment/allocation edge when the resource is allocated to the process
- When a resource is released by a process, assignment/allocation edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system

# Resource-Allocation Graph





# Unsafe State In Resource-Allocation Graph



# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# Banker's Algorithm

- Multiple instances
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

# Banker's algorithm—*definitions*

Assume     $N$  Processes  $\{P_i\}$   
           $M$  Resources  $\{R_j\}$

Availability vector **Avail<sub>j</sub>**, units of each resource (initialized to maximum, changes dynamically).

Let [**Max<sub>ij</sub>**] be an  $N \times M$  matrix.

**Max<sub>ij</sub>** =  $L$  means Process  $P_i$  will request at most  $L$  units of  $R_j$ .

[**Hold<sub>ij</sub>**] Units of  $R_j$  currently held by  $P_i$

[**Need<sub>ij</sub>**] Remaining need by  $P_i$  for units of  $R_j$

**Need<sub>ij</sub>** = **Max<sub>ij</sub>** - **Hold<sub>ij</sub>**, for all  $i$  &  $j$

# Banker's algorithm—*resource request*

- At any instance,  $P_i$  posts its request for resources in vector  $REQ_j$ .
- **Step 1:** verify that a process matches its needs.  
■ *if*  $REQ_j > Need_{ij}$  *abort—error, impossible*
- **Step 2:** check if the requested amount is available.  
■ *if*  $REQ_j > Avail_j$  *goto Step 1— $P_i$  must wait*
- **Step 3:** provisional allocation.  
■  $Avail_j = Avail_j - REQ_j$   
■  $Hold_{ij} = Hold_{ij} + REQ_j$   
■  $Need_{ij} = Need_{ij} - REQ_j$   
■ *if*  $isSafe()$  *then grant resources—system is safe*  
■ *else cancel allocation; goto Step 1— $P_i$  must wait*

# Banker's algorithm—*isSafe*

Find out whether the system is in a safe state.

**Work** and **Finish** are two temporary vectors.

Step 1: initialize.

$\text{Work}_j = \text{Avail}_j$  for all  $j$ ;  $\text{Finish}_i = \text{false}$  for all  $i$ .

Step 2: find a process  $P_i$  such that

$\text{Finish}_i = \text{false}$  and  $\text{Need}_{ij} \leq \text{Work}_j$   
if no such process, *goto Step 4*.

Step 3:  $\text{Work}_j = \text{Work}_j + \text{Hold}_{ij}$

$\text{Finish}_i = \text{true}$

*goto Step 2*.

Step 4: *if*  $\text{Finish}_i = \text{true}$  for all  $i$

*then return true—yes, the system is safe*

*else return false—no, the system is NOT safe*

# Banker's algorithm—*what is safe?*

## ■ Safe with respect to some resource allocation.

### ◆ very safe

$NEED_i \leq AVAIL$  for all Processes  $P_i$ .

*Processes can run to completion in any order.*

### ◆ safe (but take care)

$NEED_i > AVAIL$  for some  $P_i$

$NEED_i \leq AVAIL$  for at least one  $P_i$  such that

*There is at least one correct order in which the processes may complete their use of resources.*

### ◆ unsafe (deadlock inevitable)

$NEED_i > AVAIL$  for some  $P_i$

$NEED_i \leq AVAIL$  for at least one  $P_i$

*But some processes cannot complete successfully.*

### ◆ deadlock

$NEED_i > AVAIL$  for all  $P_i$

*Processes are already blocked or will become so as they request a resource.*

# Example—safe allocation

	Max	Hold	Need	Finish	Avail	Work
$P_1$	5	<del>2</del> <sub>3</sub>	<del>3</del> <sub>2</sub>	F	2	2
$P_2$	4	1	3	F		
$P_3$	2	1	1	F		

For simplicity, assume that all the resources are identical.

Assume  $P_1$  acquires one unit. *Very safe?* **No!  $Need_2 > 2$**

*Safe?* Let us see with the safe/unsafe algorithm...

$i = 1$ ; does  $P_1$  agree with **Step 2**? No.

$i = 2$ ; does  $P_2$  agree with **Step 2**? No.

$i = 3$ ; does  $P_3$  agree with **Step 2**? Yes. **Work** = **Work**+**Hold**<sub>3</sub>; **Finish**<sub>3</sub> = **T**

$i = 1$ ; does  $P_1$  agree with **Step 2**? Yes. **Work** = **Work**+**Hold**<sub>1</sub>; **Finish**<sub>1</sub> = **T**

$i = 2$ ; does  $P_2$  agree with **Step 2**? Yes. **Work** = **Work**+**Hold**<sub>2</sub>; **Finish**<sub>2</sub> = **T**

No more (unfinished)  $P_i$ , therefore *safe*.




# Example—safe allocation

*cont.*

	Max	Hold	Need	Finish	Avail	Work
P <sub>1</sub>	5	<del>2</del> <sup>3</sup>	<del>3</del> <sup>2</sup>	F	2	<del>2</del> <sub>1</sub>
P <sub>2</sub>	4	1	3	F		
P <sub>3</sub>	2	1	1	F		

Assume P<sub>1</sub> acquires one unit.




# Example—safe allocation

*cont.*

	Max	Hold	Need	Finish	Avail	Work
P <sub>1</sub>	5	3	2	F	1	<del>1</del> 0
P <sub>2</sub>	4	1	3	F		
P <sub>3</sub>	2	<del>1</del> 2	<del>1</del> 0	F		

P<sub>3</sub> can acquire the last unit and finish.




# Example—safe allocation

*cont.*

	Max	Hold	Need	Finish	Avail	Work
$P_1$	5	5	0	F	2	<del>2</del> <sub>0</sub>
$P_2$	4	1	3	F		
$P_3$	2	0	0	T		

Then,  $P_1$  can acquire two more units and finish.




# Example—safe allocation

*cont.*

	Max	Hold	Need	Finish	Avail	Work
$P_1$	5	0	0	T	5	$\cancel{P}_2$
$P_2$	4	$\cancel{1}^4$	$\cancel{3}^0$	F		
$P_3$	2	0	0	T		

Finally,  $P_2$  can acquire three more units and finish.



# Example—unsafe allocation

	Max	Hold	Need	Finish	Avail	Work
P <sub>1</sub>	5	2	3	F	2	2 <sub>1</sub>
P <sub>2</sub>	5	1	4	F		
P <sub>3</sub>	2	1	1	F		

Assume P<sub>2</sub> acquires one unit.

As before, P<sub>3</sub> can finish and release its resources.

*BUT...*

i = 1; does P<sub>1</sub> agree with **Step 2**? No.

i = 2; does P<sub>2</sub> agree with **Step 2**? No.

i = 3; does P<sub>3</sub> agree with **Step 2**? Yes. **Work** = **Work**+**Hold**<sub>2</sub>; **Finish**<sub>2</sub> = **T**

Any more unfinished P<sub>i</sub>? Yes.


P<sub>1</sub> and P<sub>2</sub> cannot finish. Therefore *unsafe*.

# Example—unsafe allocation

*cont.*

	Max	Hold	Need	Finish	Avail	Work
P <sub>1</sub>	5	2	3	F	2	2
P <sub>2</sub>	5	2	3	F		
P <sub>3</sub>	2	0	0	T		

*NOW...*



# Deadlock detection

- This technique
  - ◆ does not attempt to prevent deadlocks;
  - ◆ instead, it lets them occur.
- The system
  - ◆ detects when this happens, by periodically running an algorithm to detect a **circular wait** condition,
  - ◆ then takes some action to recover after the fact.
- With deadlock detection, requested resources are granted to processes whenever possible.

# Deadlock detection

## *cont.*

- A check for deadlock can be made
  - ◆ as frequently as for each resource request, or
  - ◆ less frequently, depending on how likely it is for a deadlock to occur.
- Checking at each resource request has two advantages:
  - ◆ it leads to early detection, and
  - ◆ the algorithm is relatively simple because it is based on incremental changes to the state of the system.
- On the other hand,
  - ◆ such frequent checks consume considerable processor time.



# Recovering from deadlocks

- Once the deadlock algorithm has successfully detected a deadlock, some strategy is needed for recovery. There are various ways:
  - ◆ Recovery through *Preemption*

In some cases, it may be possible to temporarily take a resource away from its current owner and give it to another.
  - ◆ Recovery through *Rollback*

If it is known that deadlocks are likely, one can arrange to have processes *checkpointed* periodically. For example, can undo transactions, thus free locks on database records.
  - ◆ Recovery through *Termination*

The most trivial way to break a deadlock is to kill one or more processes. One possibility is to kill a process in the cycle. Warning! *Irrecoverable losses may occur, even if this is the least advanced process.*

# Summary of strategies

<i>Principle</i>	<i>Resource Allocation Strategy</i>	<i>Different Schemes</i>	<i>Major Advantages</i>	<i>Major Disadvantages</i>
<i>DETECTION</i>	<ul style="list-style-type: none"> <li>• Very liberal; grant resources as requested.</li> </ul>	<ul style="list-style-type: none"> <li>• Invoke periodically to test for deadlock.</li> </ul>	<ul style="list-style-type: none"> <li>• Never delays process initiation.</li> <li>• Facilitates on-line handling.</li> </ul>	<ul style="list-style-type: none"> <li>• Inherent preemption losses.</li> </ul>
<i>PREVENTION</i>	<ul style="list-style-type: none"> <li>• Conservative; under-commits resources.</li> </ul>	<ul style="list-style-type: none"> <li>• Requesting all resources at once.</li> <li>• Preemption</li> <li>• Resource ordering</li> </ul>	<ul style="list-style-type: none"> <li>• Works well for processes with single burst of activity.</li> <li>• No preemption is needed.</li> <li>• Convenient when applied to resources whose state can be saved and restored easily.</li> <li>• Feasible to enforce via compile-time checks.</li> <li>• Needs no run-time computation.</li> </ul>	<ul style="list-style-type: none"> <li>• Inefficient.</li> <li>• Delays process initiation.</li> <li>• Preempts more often than necessary.</li> <li>• Subject to cyclic restart.</li> <li>• Preempts without immediate use.</li> <li>• Disallows incremental resource requests.</li> </ul>
<i>AVOIDANCE</i>	<ul style="list-style-type: none"> <li>• Selects midway between that of detection and prevention.</li> </ul>	<ul style="list-style-type: none"> <li>• Manipulate to find at least one safe path.</li> </ul>	<ul style="list-style-type: none"> <li>• No preemption necessary.</li> </ul>	<ul style="list-style-type: none"> <li>• Future resource requirements must be known.</li> <li>• Processes can be blocked for long periods.</li> </ul>