

Synchronization Techniques (2)

More on Synchronization

Hardware Support for Synchronization

Very much needed for multi-core processors

Memory barriers

- A system (processor or compiler) can reorder instructions for performance
- Break software solutions to synchronization
- Computer architectures determines what guarantees memory can provide the applications - **memory model**

Hardware instructions

- Instructions in the processor to test and set memory locations **atomically** - one uninterruptible unit
- Micro processors provide a variety of these instructions - accessible in assembly programming

Atomic variables

- This is the way hardware instructions are exposed to the programmer in a programming language
- Use this to build higher level sync primitives
- Atomic variables themselves may not solve all sync issues

Memory Barriers

Also referred to as memory fences

- Relate to shared memory - determines how the loads and stores to the shared memory is visible to the processors that are doing them
- Two type of memory barriers
 - **Strongly ordered** - memory modification on one processor is immediately visible to all other processors
 - **Weakly ordered** - memory modifications to memory in one processor is not immediately visible to another processor

```
boolean flag = false;  
int x = 0;
```

where Thread 1 performs the statements

```
while (!flag)  
    ;  
print x;
```

and Thread 2 performs

```
x = 100;  
flag = true;
```

```
while (!flag)  
    memory_barrier();  
print x;
```

```
x = 100;  
memory_barrier();  
flag = true;
```

Hardware Instructions

- **Test_And_Set (TAS)** - retrieve the current value of a memory and set it to 1 in an uninterruptible instruction
- Implemented by the processor
- Executes a **read-modify-write** atomically

```
boolean test_and_set(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
  
    return rv;  
}
```

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
  
    /* remainder section */  
} while (true);
```

Hardware Instructions

- Compare_And_Swap (CAS) is like TAS, but uses a different mechanism
- It operates on three operands: it sets the memory variable to the new value only if the current value is same as expected - whole operation is atomic - all or none
- Two CAS instructions executed concurrently, they are ordered for execution sequentially in an arbitrary manner

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

```
while (true) {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

Examples using Hardware Instructions

- Implement a mutual exclusion using CAS
- Critical section is protected by a lock - a memory variable
- Threads wanting to enter the CS need to grab the lock
- Only one winning thread - all other threads are waiting to get the lock using busy waiting

```
while (true) {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

Atomic Variables

- Hardware instructions are not used directly
- They are exposed via atomic variables that operate on integers, booleans, etc
- Lets say we want to implement
- `increment()` on an atomic variable
- atomic variables solve part of the synchronization issues in an application such as producer-consumer

```
void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;
    }
    while (temp != compare_and_swap(v, temp, temp+1));
}
```

Mutex Locks

- OS designers provide software tools to solve CS
- Simplest is mutex locks
- We enter the CS by acquiring the lock
- We release the lock at exit
- Calls to `acquire()` and `release()` must be atomic
- If the solution depends on busy waiting, locks are called spin locks

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```

The definition of `acquire()` is as follows:

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

The definition of `release()` is as follows:

```
release() {  
    available = true;  
}
```


Lock Contention & Duration

Lock contention

- Lot of threads can try to grab a lock at the same time
- All except one thread would win and others have to wait
- Overhead at lock acquiring is known as the lock contention - time to grab the lock
- Lock contention rises with increasing number of waiting threads - all wanting to get into the CS at the same time

Lock duration

- Duration for which the lock is going to be held by the acquiring thread
- Short duration locks can use busy waiting with good overall performance

Semaphore

A generalized synchronization primitive

- Mutex locks are used to protect access to the critical sections
- Semaphores provide generalized synchronization for processes (threads) - that can be used to implement mutex locks
- Semaphore was invented by Edsger Dijkstra
- It is a special variable (s) supported by the OS - implemented inside the kernel
- Supports two operations - wait(s) (P(s)) and signal (s) (V(s))
- Two or more processes are cooperate by manipulating the semaphores
- wait() and signal() is considered as message passing using the semaphore variable
- Complex coordination mechanisms can be implemented

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

Semaphore Types

Counting semaphore

- Initial value set to any non negative integer value
- Internal value can be negative

Binary semaphore

- Initial value can be set to either 0 or 1
- Internal value can be negative or restricted to 0 (min) - book vague on this implementation issue
- Strong semaphore - if a strict FIFO order is enforced in serving the blocked processes
- Weak semaphore - no strict FIFO order in serving blocked processes

```
struct semaphore {  
    int count;  
    queueType queue;  
}
```

```
void wait(semaphore s)  
{  
    s.count--;  
    if (s.count < 0)  
    {  
        place this process in  
        s.queue;  
        block this process  
    }  
}
```

```
void signal(semaphore s)  
{  
    s.count++;  
    if (s.count <= 0)  
    {  
        remove a process P  
        from s.queue;  
        place process P on  
        ready list;  
    }  
}
```

Producer Consumer Using Semaphores

- Implement producer-consumer using 3 semaphores
- Protect the critical section that does queue manipulation
- Buffer overflow synchronization
- Buffer underflow synchronization

```
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

producer() {
    int item;

    while(TRUE) {
        item = produce_item();
        wait(&empty);
        wait(&mutex);
        insert_item(item);
        signal(&mutex);
        signal(&full);
    }
}
```

```
// protects the critical section
// counts the empty slots
// counts full buffer slots

consumer() {
    int item;

    while(TRUE) {
        wait(&full);
        wait(&mutex);
        item = remove_item();
        signal(&mutex);
        signal(&empty);
        consume_item(item);
    }
}
```

Monitors

A higher level abstract data type that provides synchronization


- A higher level abstraction that provides a convenient and effective mechanism for synchronization
- Abstract data type that encapsulates synchronization variables so they are manipulatable by code within the ADT

```
monitor monitor-name
{
    // shared variable declarations
    function P1 (...) { ... }

    function P2 (...) { ... }

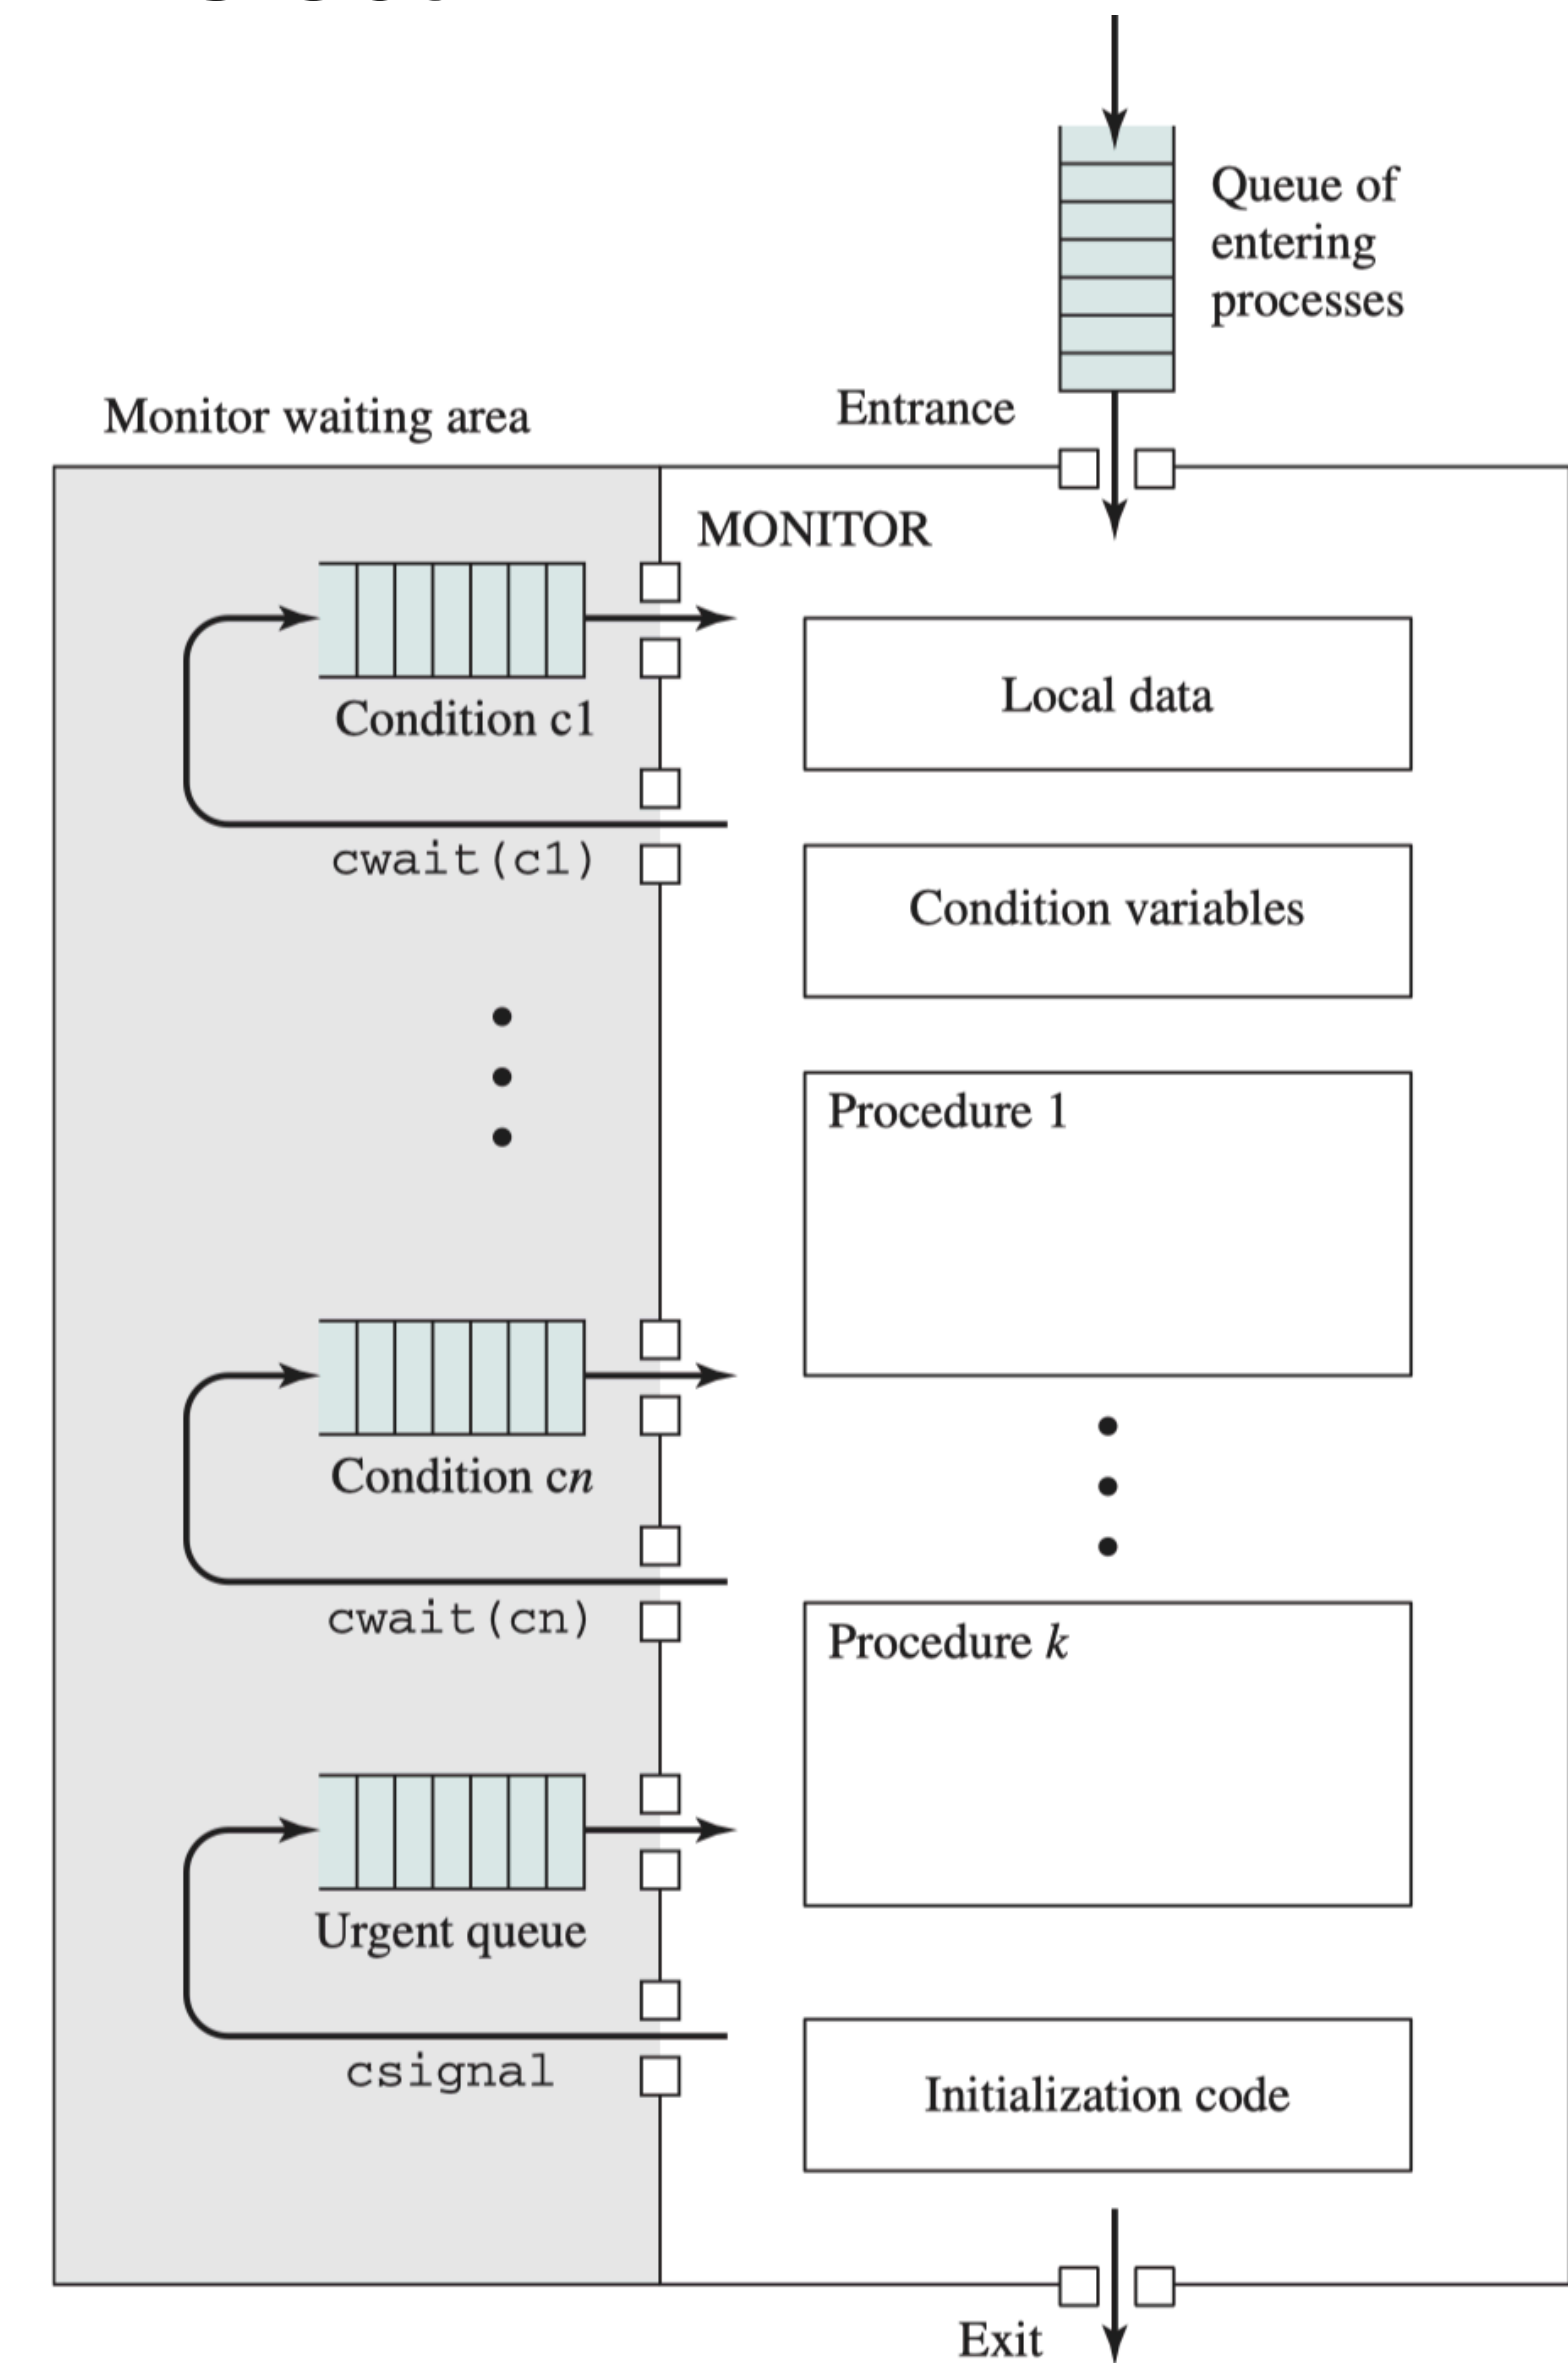
    function Pn (...) {.....}

    initialization code (...) { ... }
}
```



Monitor in More Detail

- Monitor provides a critical section - only one thread can be active inside it
- Monitor provides facilities for threads to sleep inside the monitor - such threads are not active inside the monitor
- You enter the monitor by calling a method exposed by the monitor
- A thread signalling inside the monitor goes to sleep inside the monitor



Monitor Example

- Producer-consumer using monitor
- Has two condition variables
 - Full buffer — notfull
 - Empty buffer — notempty
- No need to protect queue insertion or deletion!

Monitor Definition

```
/* program producerconsumer */
monitor boundedbuffer;
char buffer [N];                      /* space for N items */
int nextin, nextout;                  /* buffer pointers */
int count;                           /* number of items in buffer */
cond notfull, notempty;              /* condition variables for synchronization */

void append (char x)
{
    if (count == N) cwait(notfull);    /* buffer is full; avoid overflow */
    buffer[nextin] = x;
    nextin = (nextin + 1) % N;
    count++;
    /* one more item in buffer */
    csignal (notempty);                /* resume any waiting consumer */
}

void take (char x)
{
    if (count == 0) cwait(notempty);    /* buffer is empty; avoid underflow */
    x = buffer[nextout];
    nextout = (nextout + 1) % N;
    count--;
    /* one fewer item in buffer */
    csignal (notfull);                /* resume any waiting producer */
}

{
    /* monitor body */
    nextin = 0; nextout = 0; count = 0; /* buffer initially empty */
}
```

Using the monitor

```
void producer()
{
    char x;
    while (true) {
        produce(x);
        append(x);
    }
}

void consumer()
{
    char x;
    while (true) {
        take(x);
        consume(x);
    }
}

void main()
{
    parbegin (producer, consumer);
}
```